

# [368] C++ Programming: Welcome!

Tyler Caraza-Harter

# Outline

Welcome

Logistics

Background and Motivation

- Why C/C++: performance
- Why C++ (over C): language features

Demos

# Introductions


Tyler Caraza-Harter

- Long time Badger
- Email: [tharter@wisc.edu](mailto:tharter@wisc.edu)
- Just call me “Tyler” (he/him)



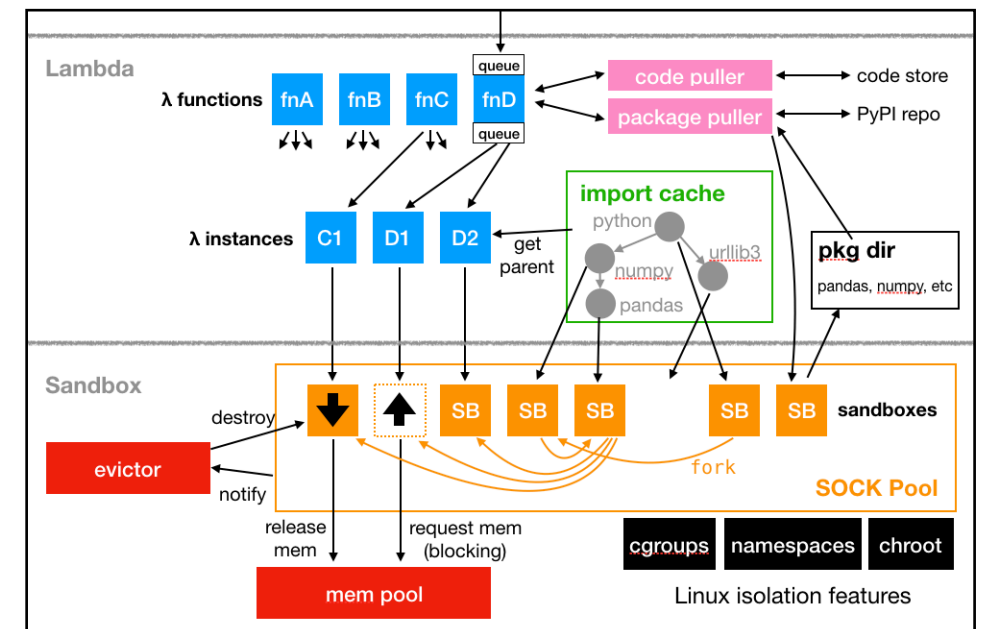
Industry experience

- Worked at Microsoft on SQL Server and Cloud
- Other internships/collaborations: Qualcomm, Google, Facebook, Tintri, Bauplan



Open source

- OpenLambda (serverless cloud platform)
- <https://github.com/open-lambda/open-lambda>



# Who are You?

Year in school? Major?

What CS courses have people taken before?

- 300? 320? 354?

Please fill this form (**due today**):

<https://forms.gle/3BaLREB1upurZDky6>

Why?

- Help me get to know you
- Earn 1 point

# What will you learn in this 368?

## Learning objectives

- Read and understand C++ code
- Write C++ programs making use of the abstractions provided by the language
- Understand the lower level details of memory management like pointers and references
- Organize and build multi-file projects using the make tool
- Solve real world programming problems using C++ as a tool

# What will you learn today?

## Learning objectives

- recall course **logistics and policies**
- describe reasons for using a language like **C/C++**
- describe reasons for using **C++ over C**

# Outline

Welcome

Logistics

Background and Motivation

- Why C/C++: performance
- Why C++ (over C): language features

Demos

# Main Websites

1

<https://tyler.caraza-harter.com/cs368/s24/schedule.html>

- schedule, course content, how to get help
- links to all other resources/tools
- some lecture recordings (review only)

2

<https://github.com/cs368-wisc/s24>

- project specifications
- lecture demo code

3

Canvas

- announcements
- quizzes
- grade summaries



# Other Tools

4

TopHat (me asking you questions during lecture)

- can earn points from this

5

Piazza (asking questions of **general interest**)

- goal: responses < 1 business day
- don't post > 5 lines of project code

6

Email (asking questions of **individual interest**)

- goal: responses < 2 business days
- please keep related issues on the same thread

7

GitHub classroom

- you'll be given a **private** repo for your project

8

Anki Flash Cards

- memory terms, basic ideas using flash cards and spaced repetition

# Lecture

Wednesday:

- **in person** (usually recorded too, barring technical difficulties)
- focus on concepts (lecture, worksheets, etc)
- TopHats

Friday:

- **posted nline**, multiple short videos
- focus on programming demos
- watch before next in-person class!

# Sparrow Project

## Project:

- one big project with six project stages (P1 - P6)
- project name: Sparrow (simple prototype of Arrow)
- Arrow project (<https://arrow.apache.org/>) enables fast in-memory analytics on tables of data; the main implementation is in C++

## Collaboration:

- done individually
- can help each other debug (with citation)
- sharing code is not allowed



## Submission:

- you will push your code to a GitHub classroom repo (keep it private!)
- submit a form when a specific version (commit number) is ready for grading

## Grading:

- autograded using tests I'll release
- I might manually modify grades if anybody tries to "game the tests"

# Grading

This course is credit/no credit:

- so pass/fail, no letter grades
- to pass, you need to earn  $\geq 100$  points
- there will be  $> 150$  points possible to earn, so there are many possible ways to pass
- 100 is a bit of a low bar; 140+ would be a score to really "feel good" about

Scoring:

- projects: 120 points possible (4 per passed test)
- quizzes: 30 points possible (1 per correct answer)
- TopHat: 1 point for correct answer, 0.5 for incorrect
- other: I might offer other opportunities for points as we go

# TODO

books

# Outline

Welcome

Logistics

Background and Motivation

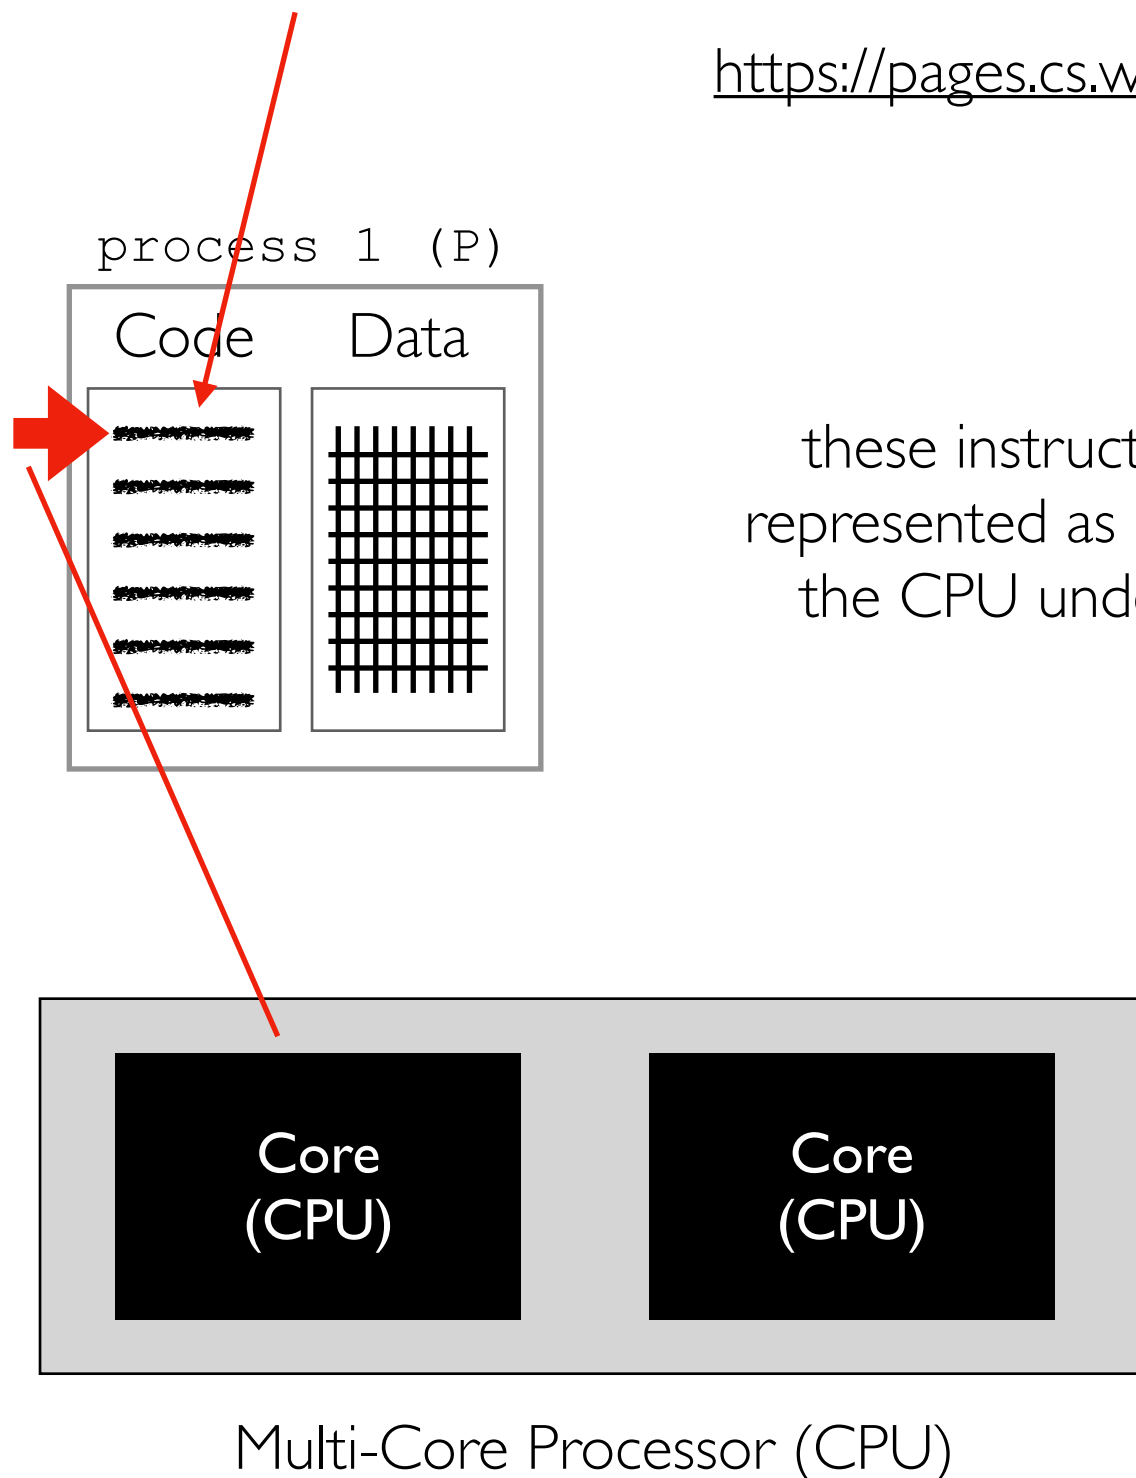
- Why C/C++: performance
  - how code runs
  - cachelines
  - garbage collection
  - safety checks
- Why C++ (over C): language features

Demos

# Background: How to Code Runs on CPUs

these instructions are in "machine code"  
that the CPU can understand

<https://pages.cs.wisc.edu/~deppeler/cs354/reference/x86-cheat-sheet.pdf>



these instructions are  
represented as 1's and 0's  
the CPU understands

## arithmetic

### two operand instructions

```
addl src,dst    dst = dst + src
subl src,dst    dst = dst - src
imull src,dst   dst = dst * src
sall src,dst    dst = dst << src (aka shll)
sarl src,dst    dst = dst >> src (arith)
shrl src,dst    dst = dst >> src (logical)
xorl src,dst    dst = dst ^ src
andl src,dst    dst = dst & src
orl src,dst     dst = dst | src
```

### one operand instructions

```
incl dst       dst = dst + 1
decl dst       dst = dst - 1
negl dst       dst = -dst
notl dst       dst = ~dst
```

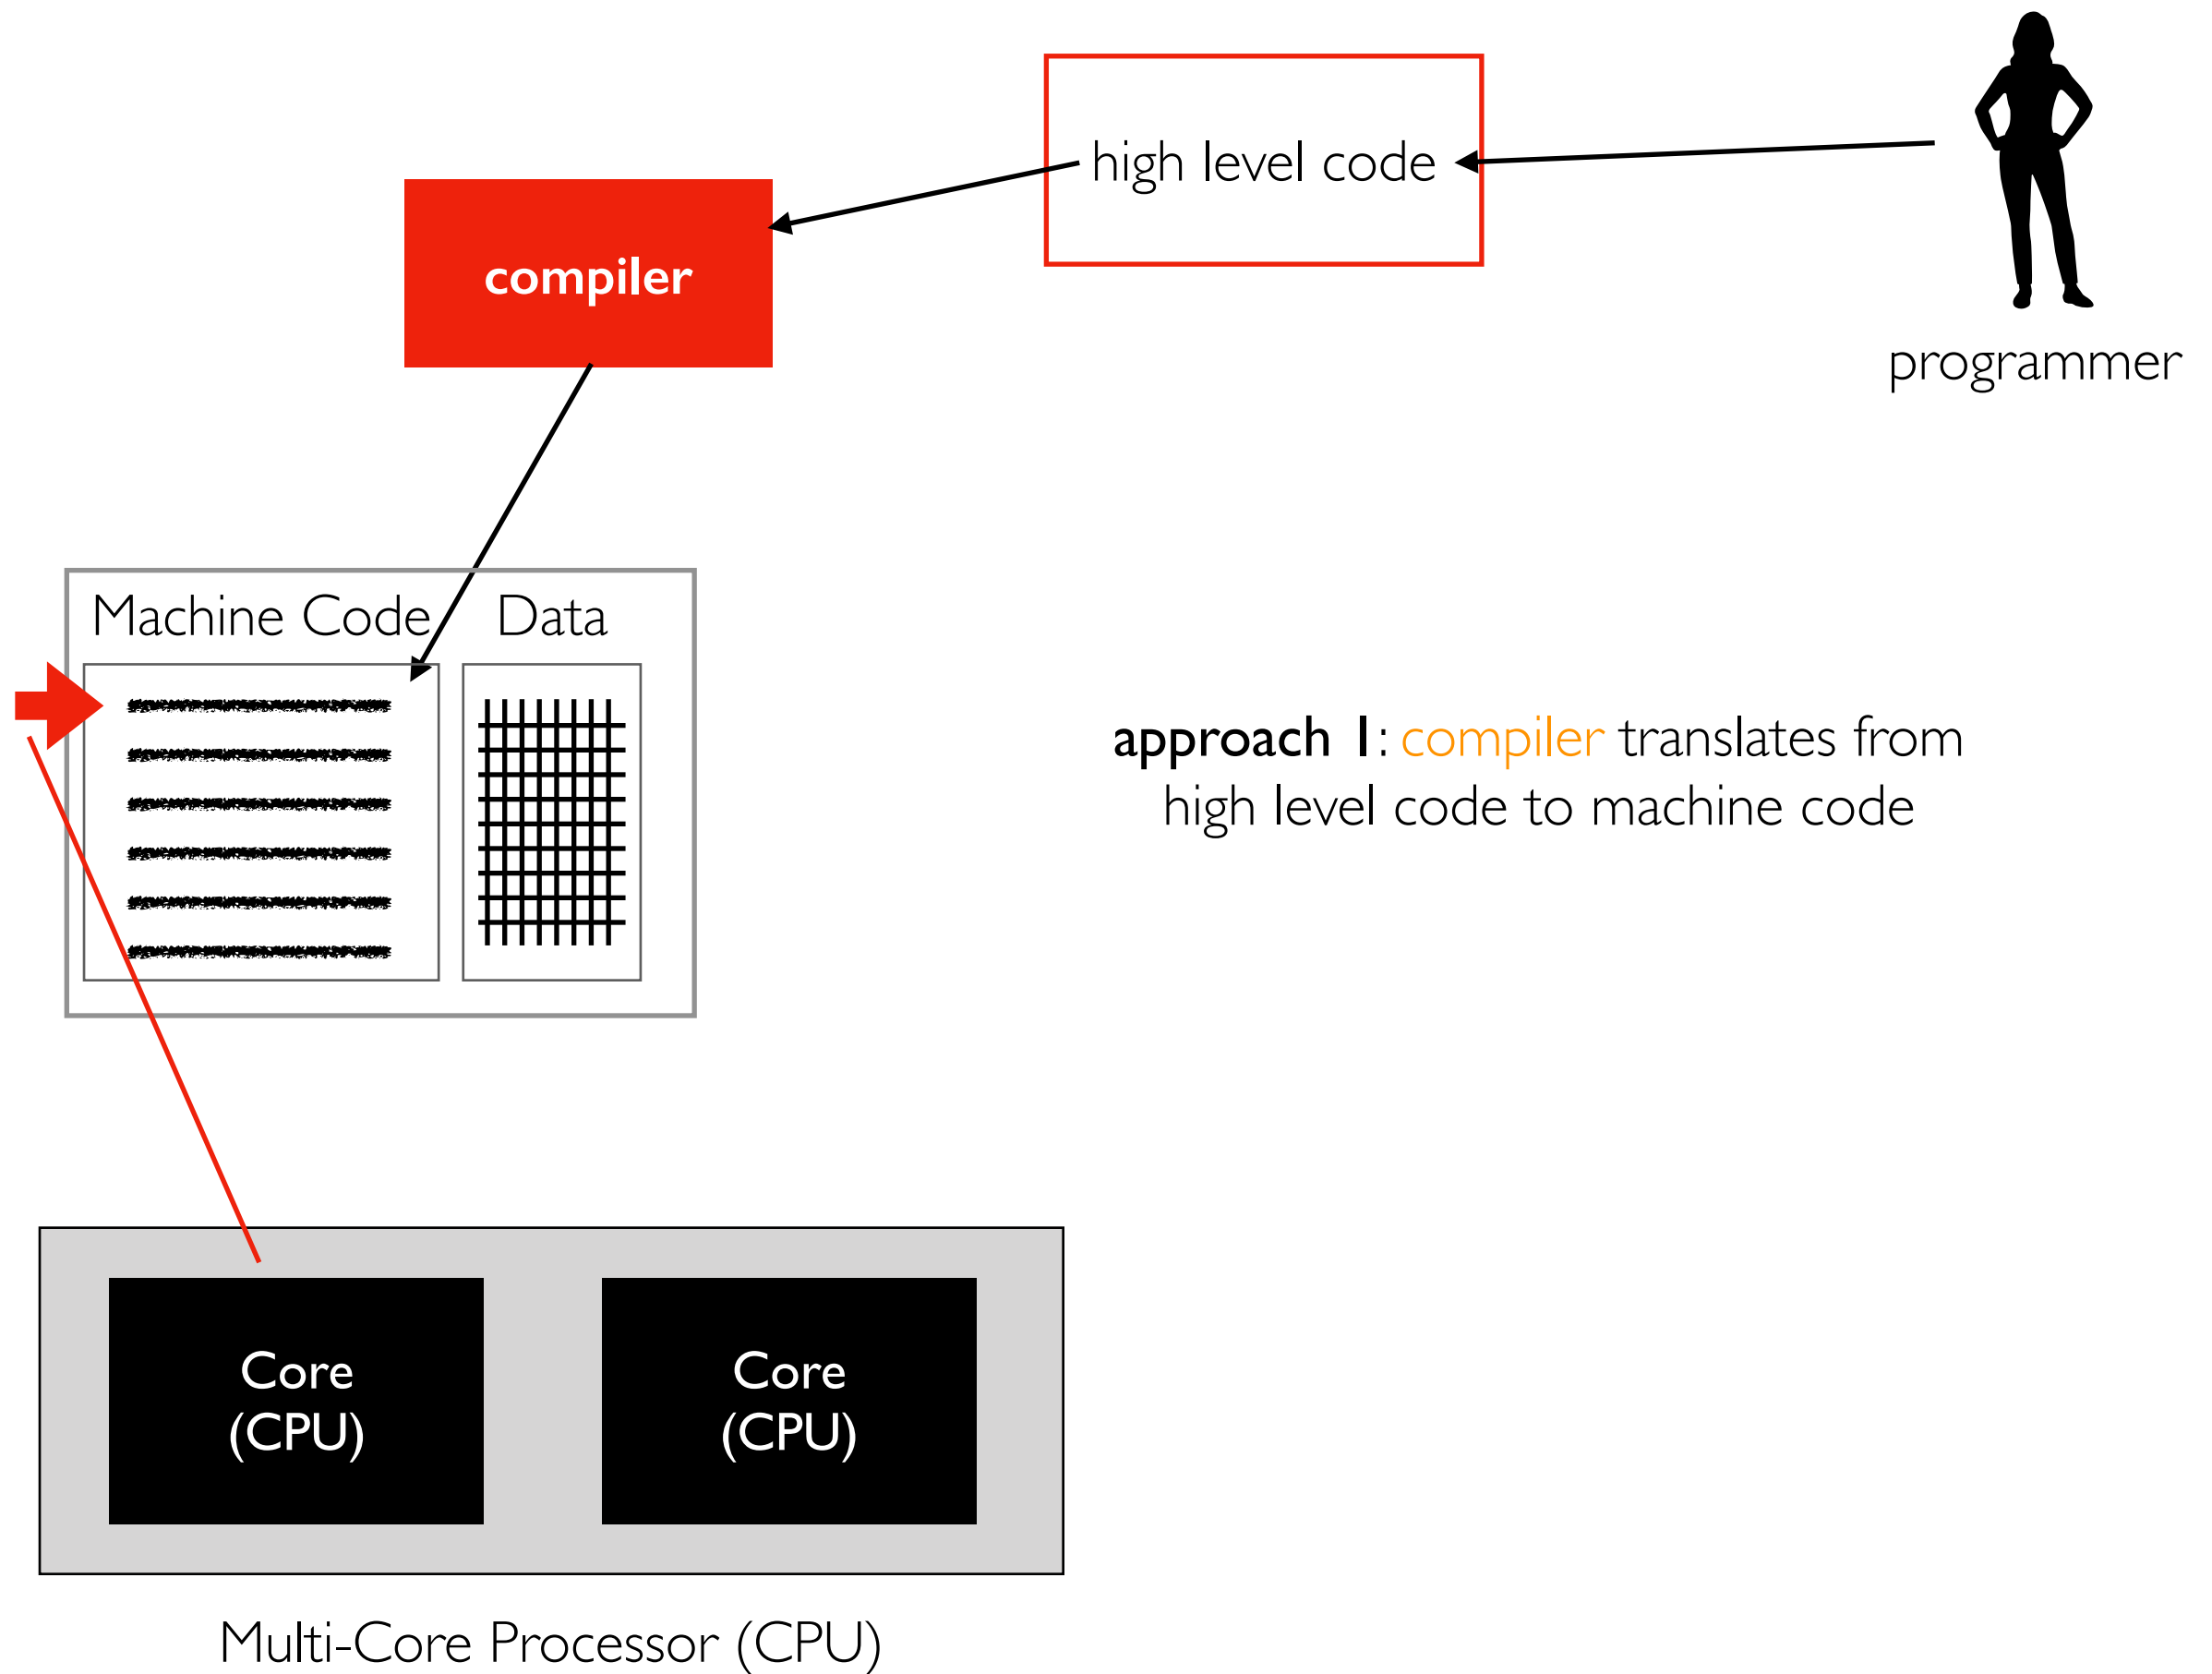
### arithmetic ops set CCs implicitly

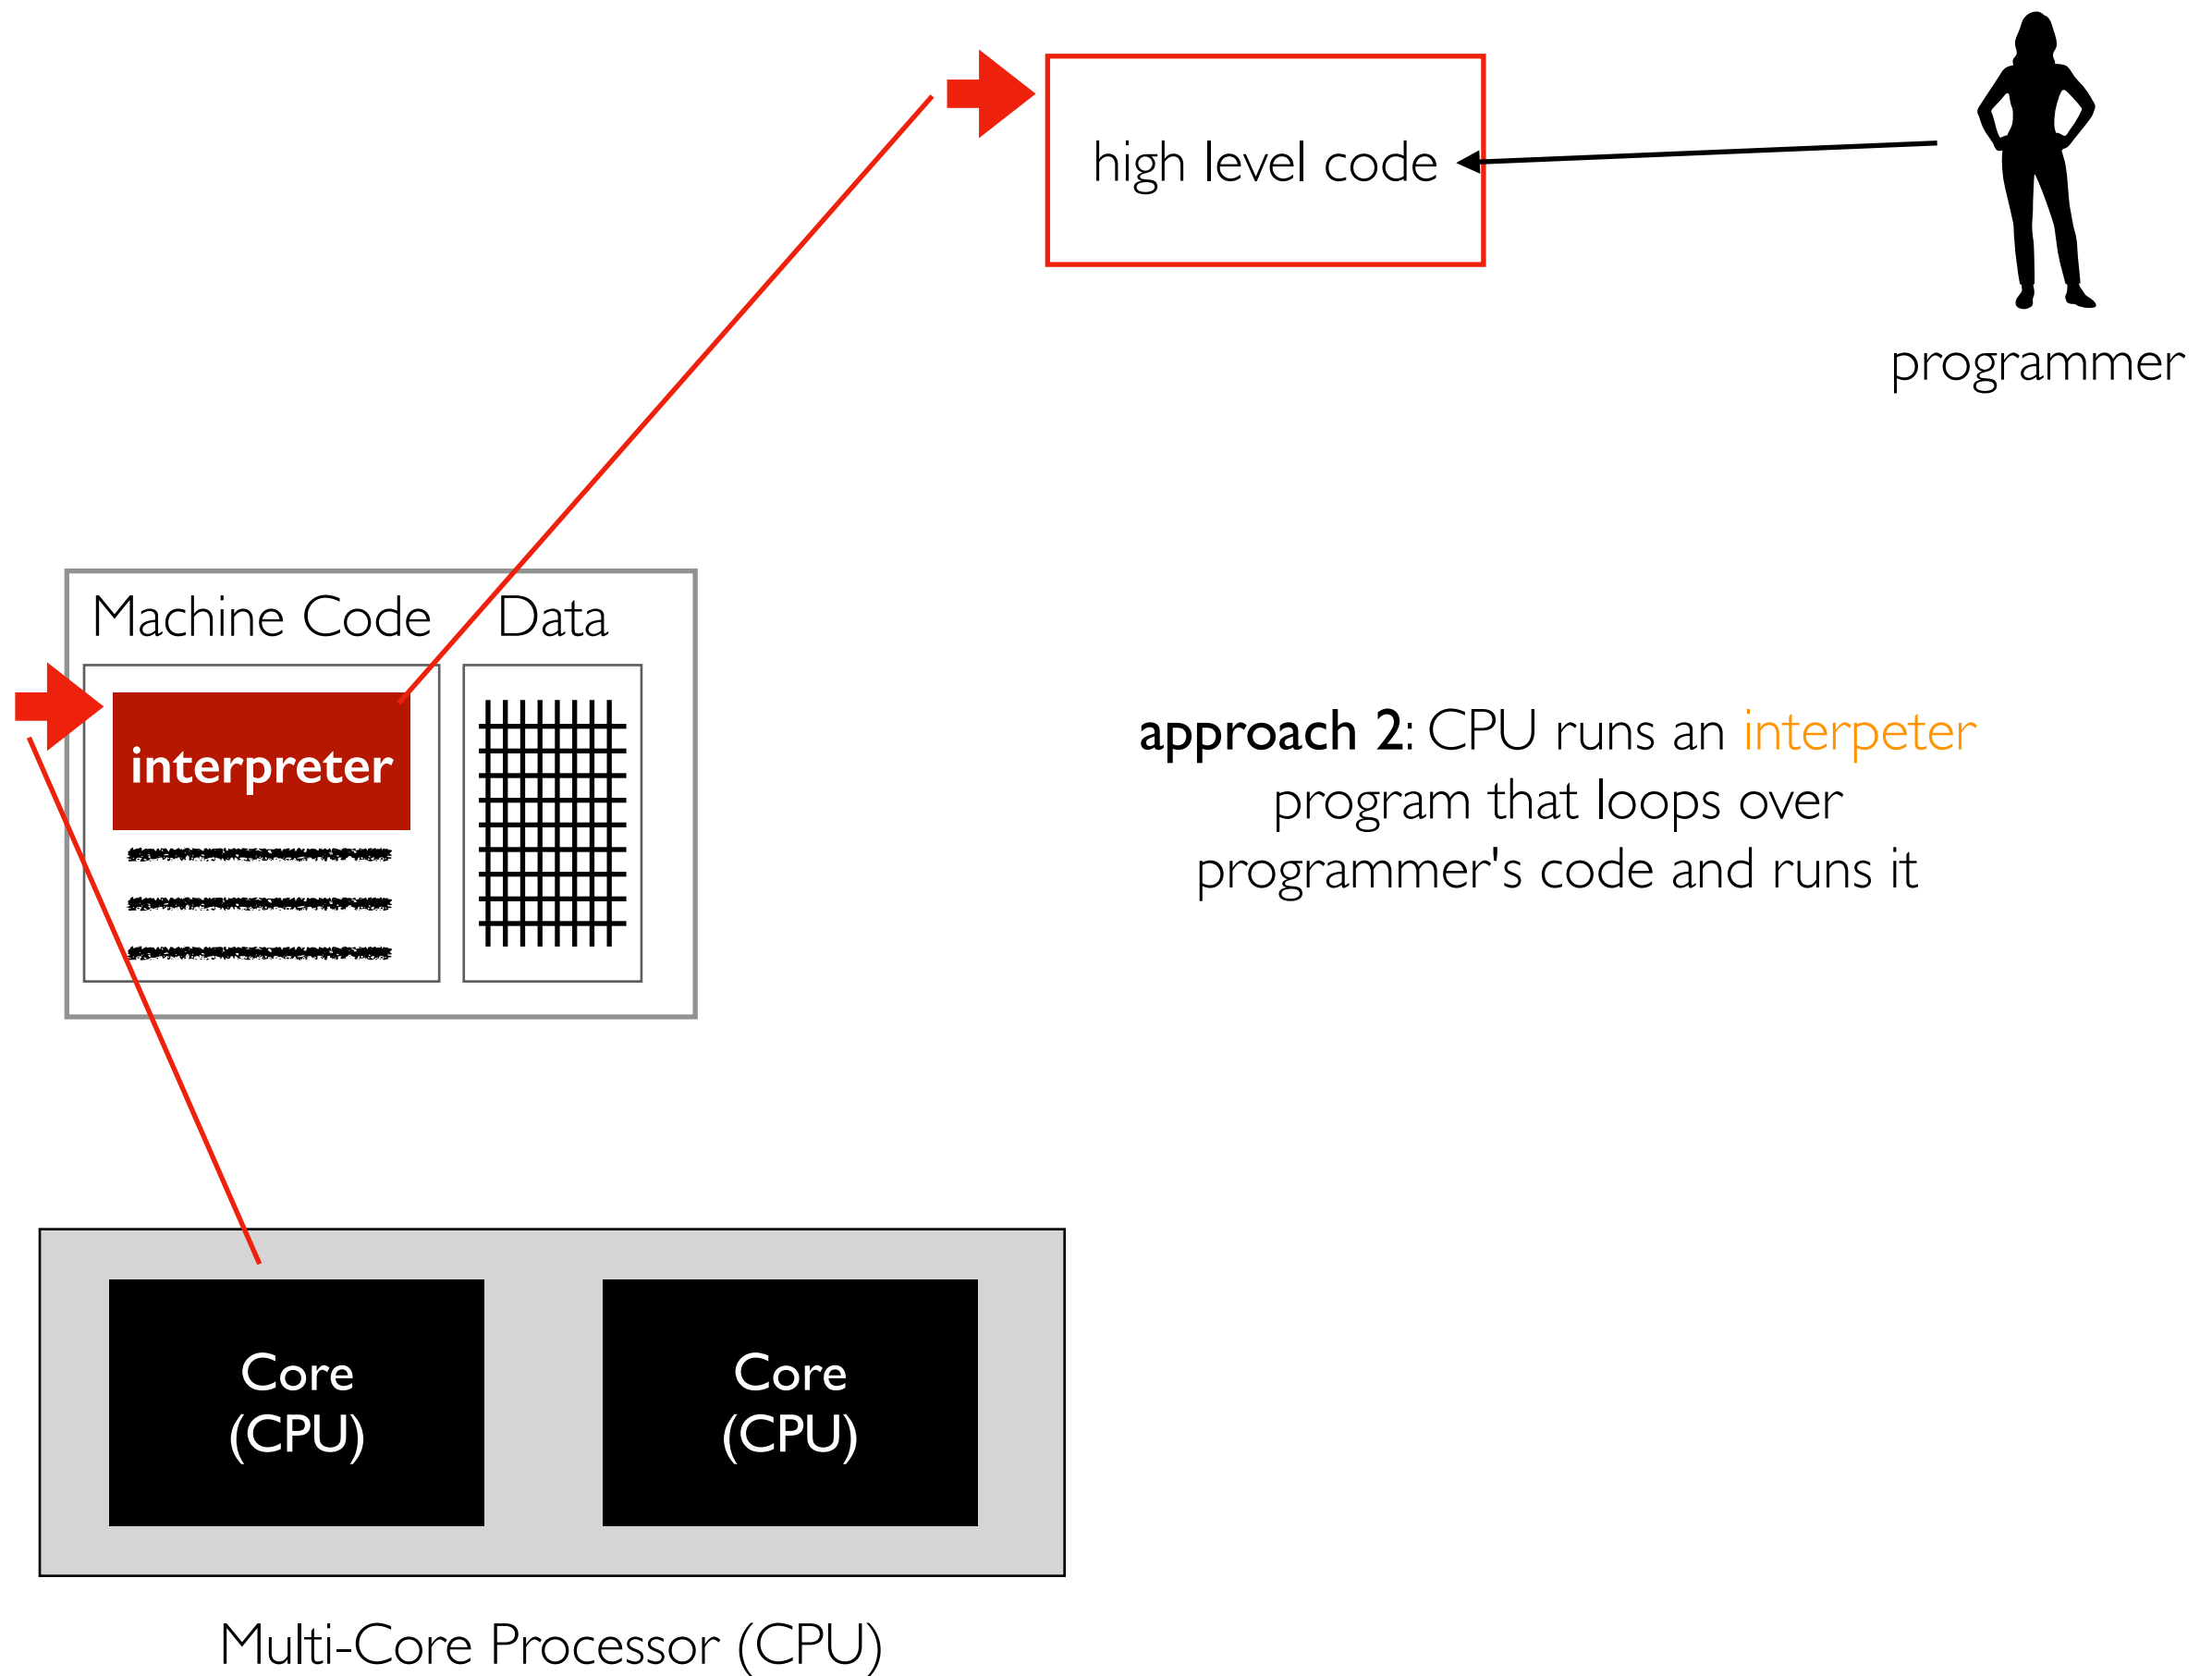
```
cf=1 if carry out from msb
zf=1 if dst==0,
sf=1 if dst < 0 (signed)
of=1 if two's complement
      (signed) under/overflow
```

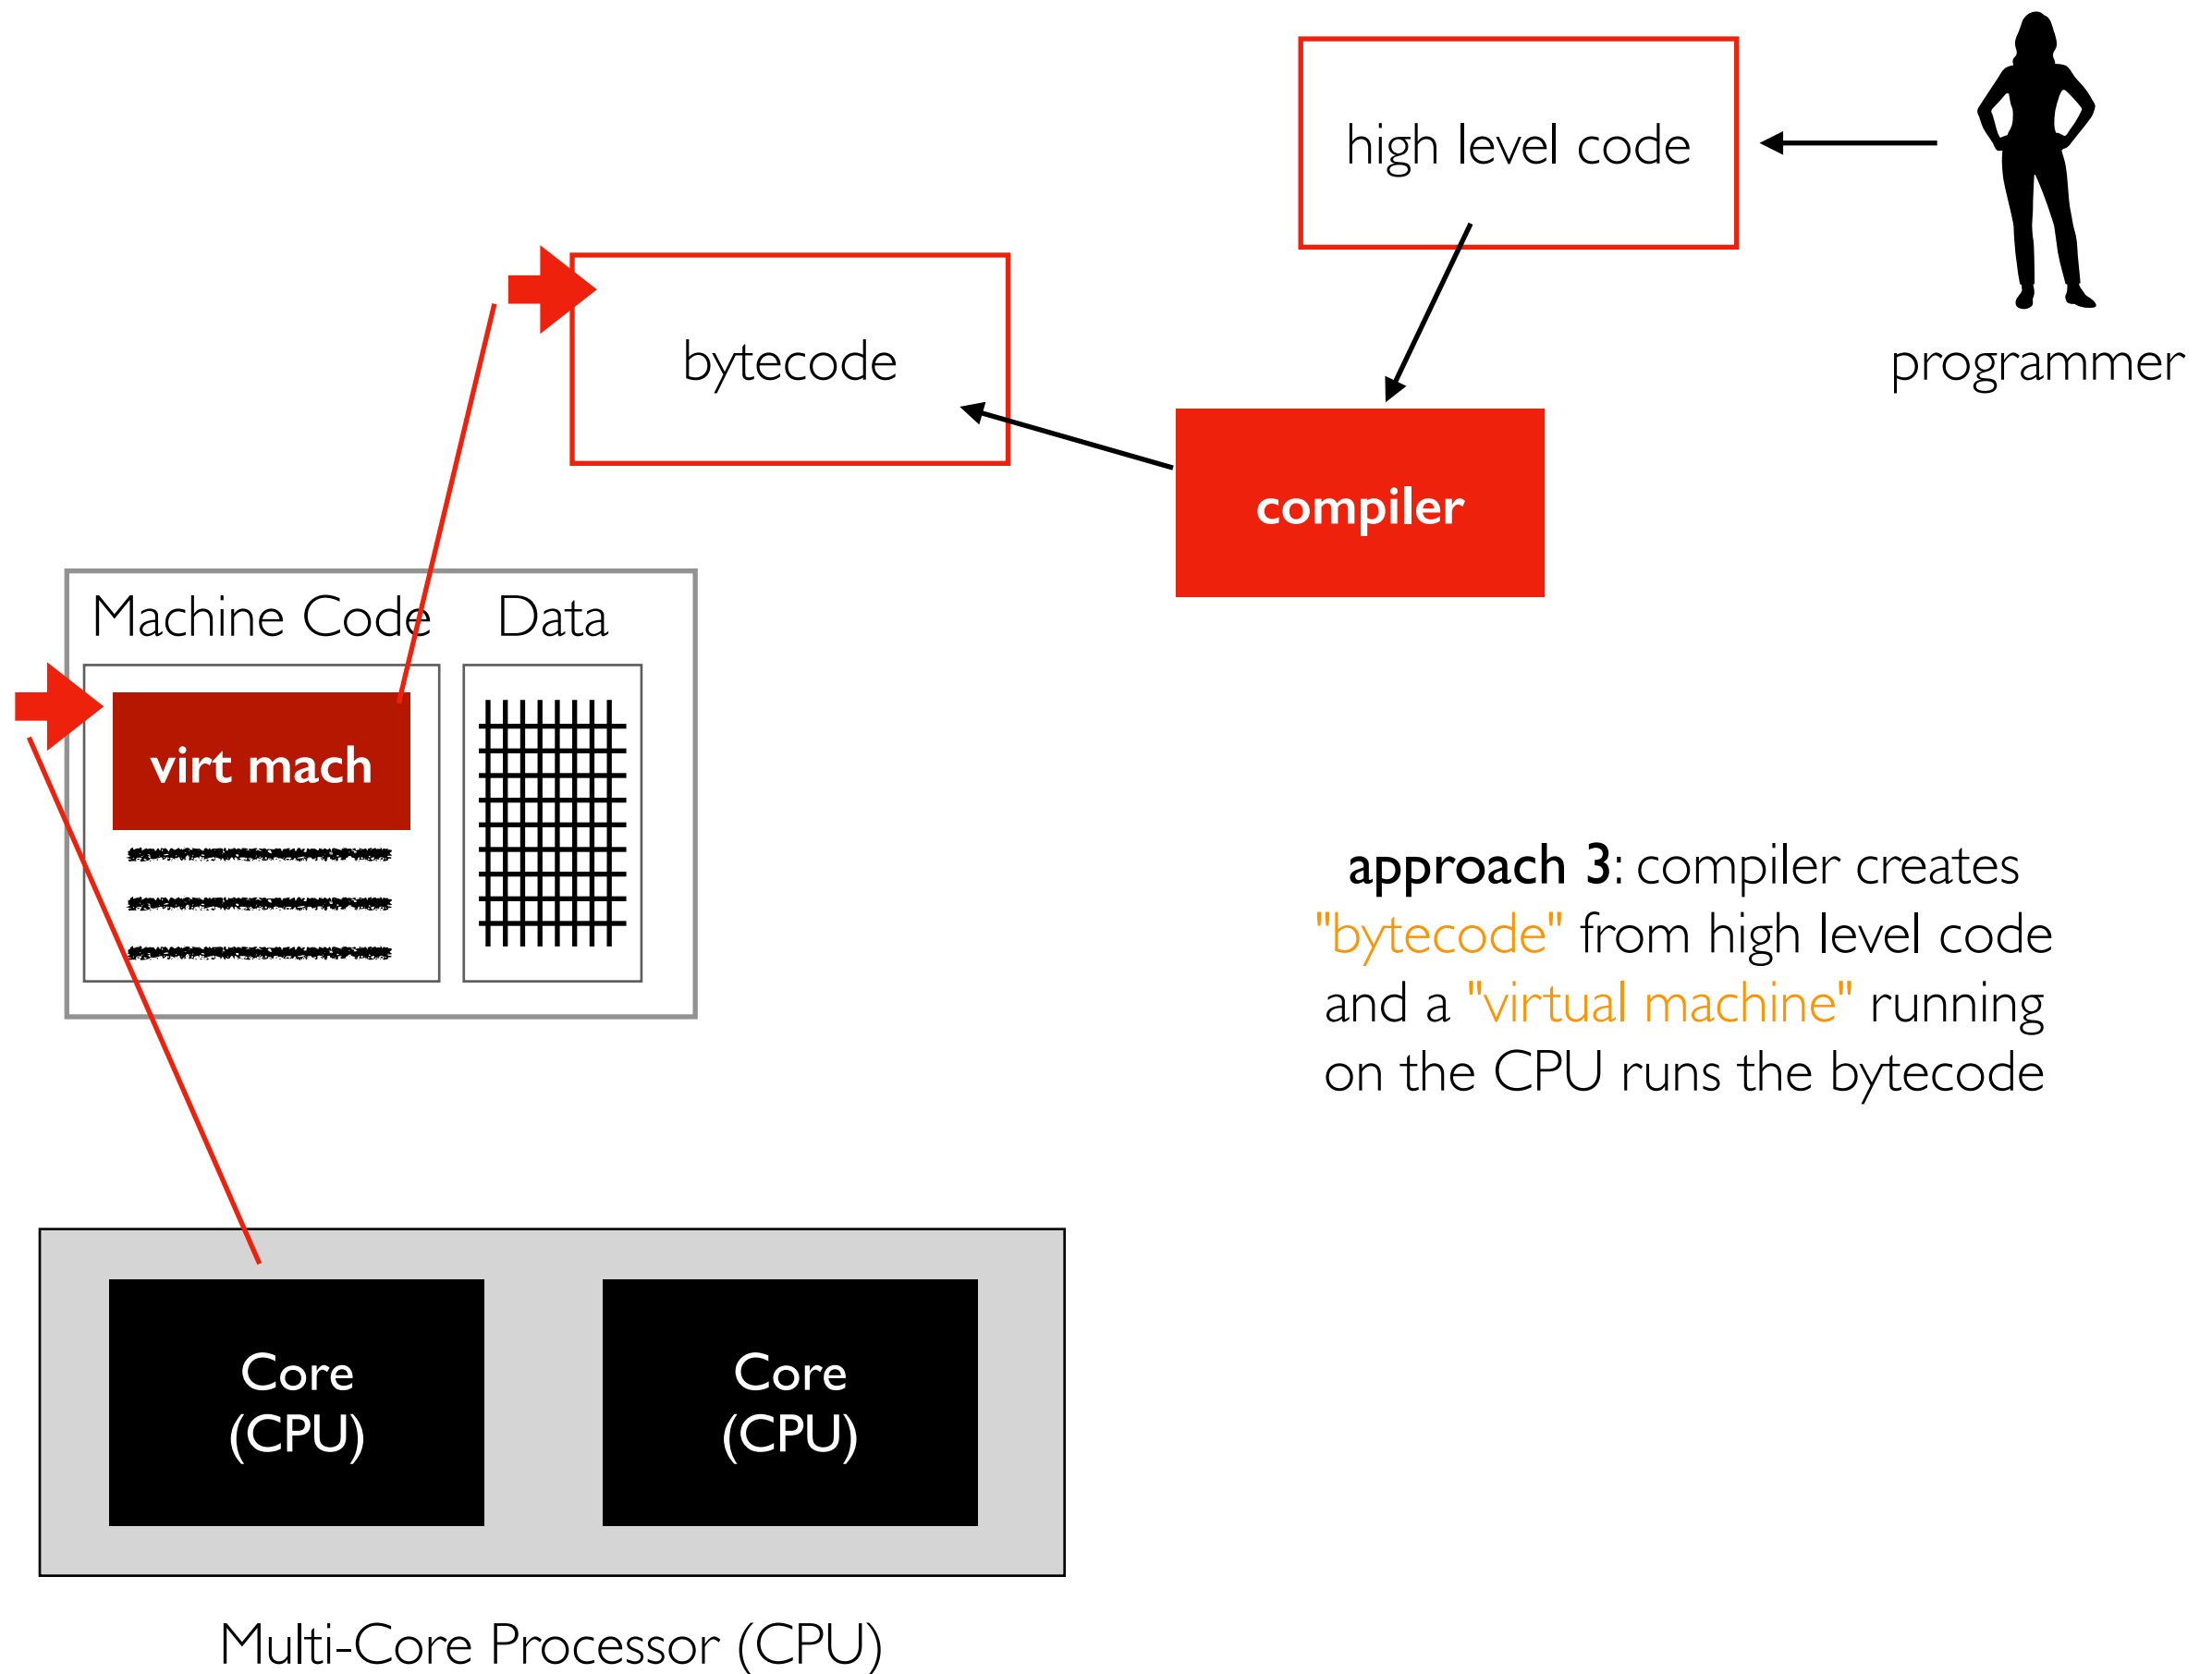
how do we bridge the gap between "high level"  
code (Python/Java/etc) and machine code?

how do we bridge the gap between "high level"  
code (C++/Python/Java/etc) and machine code?

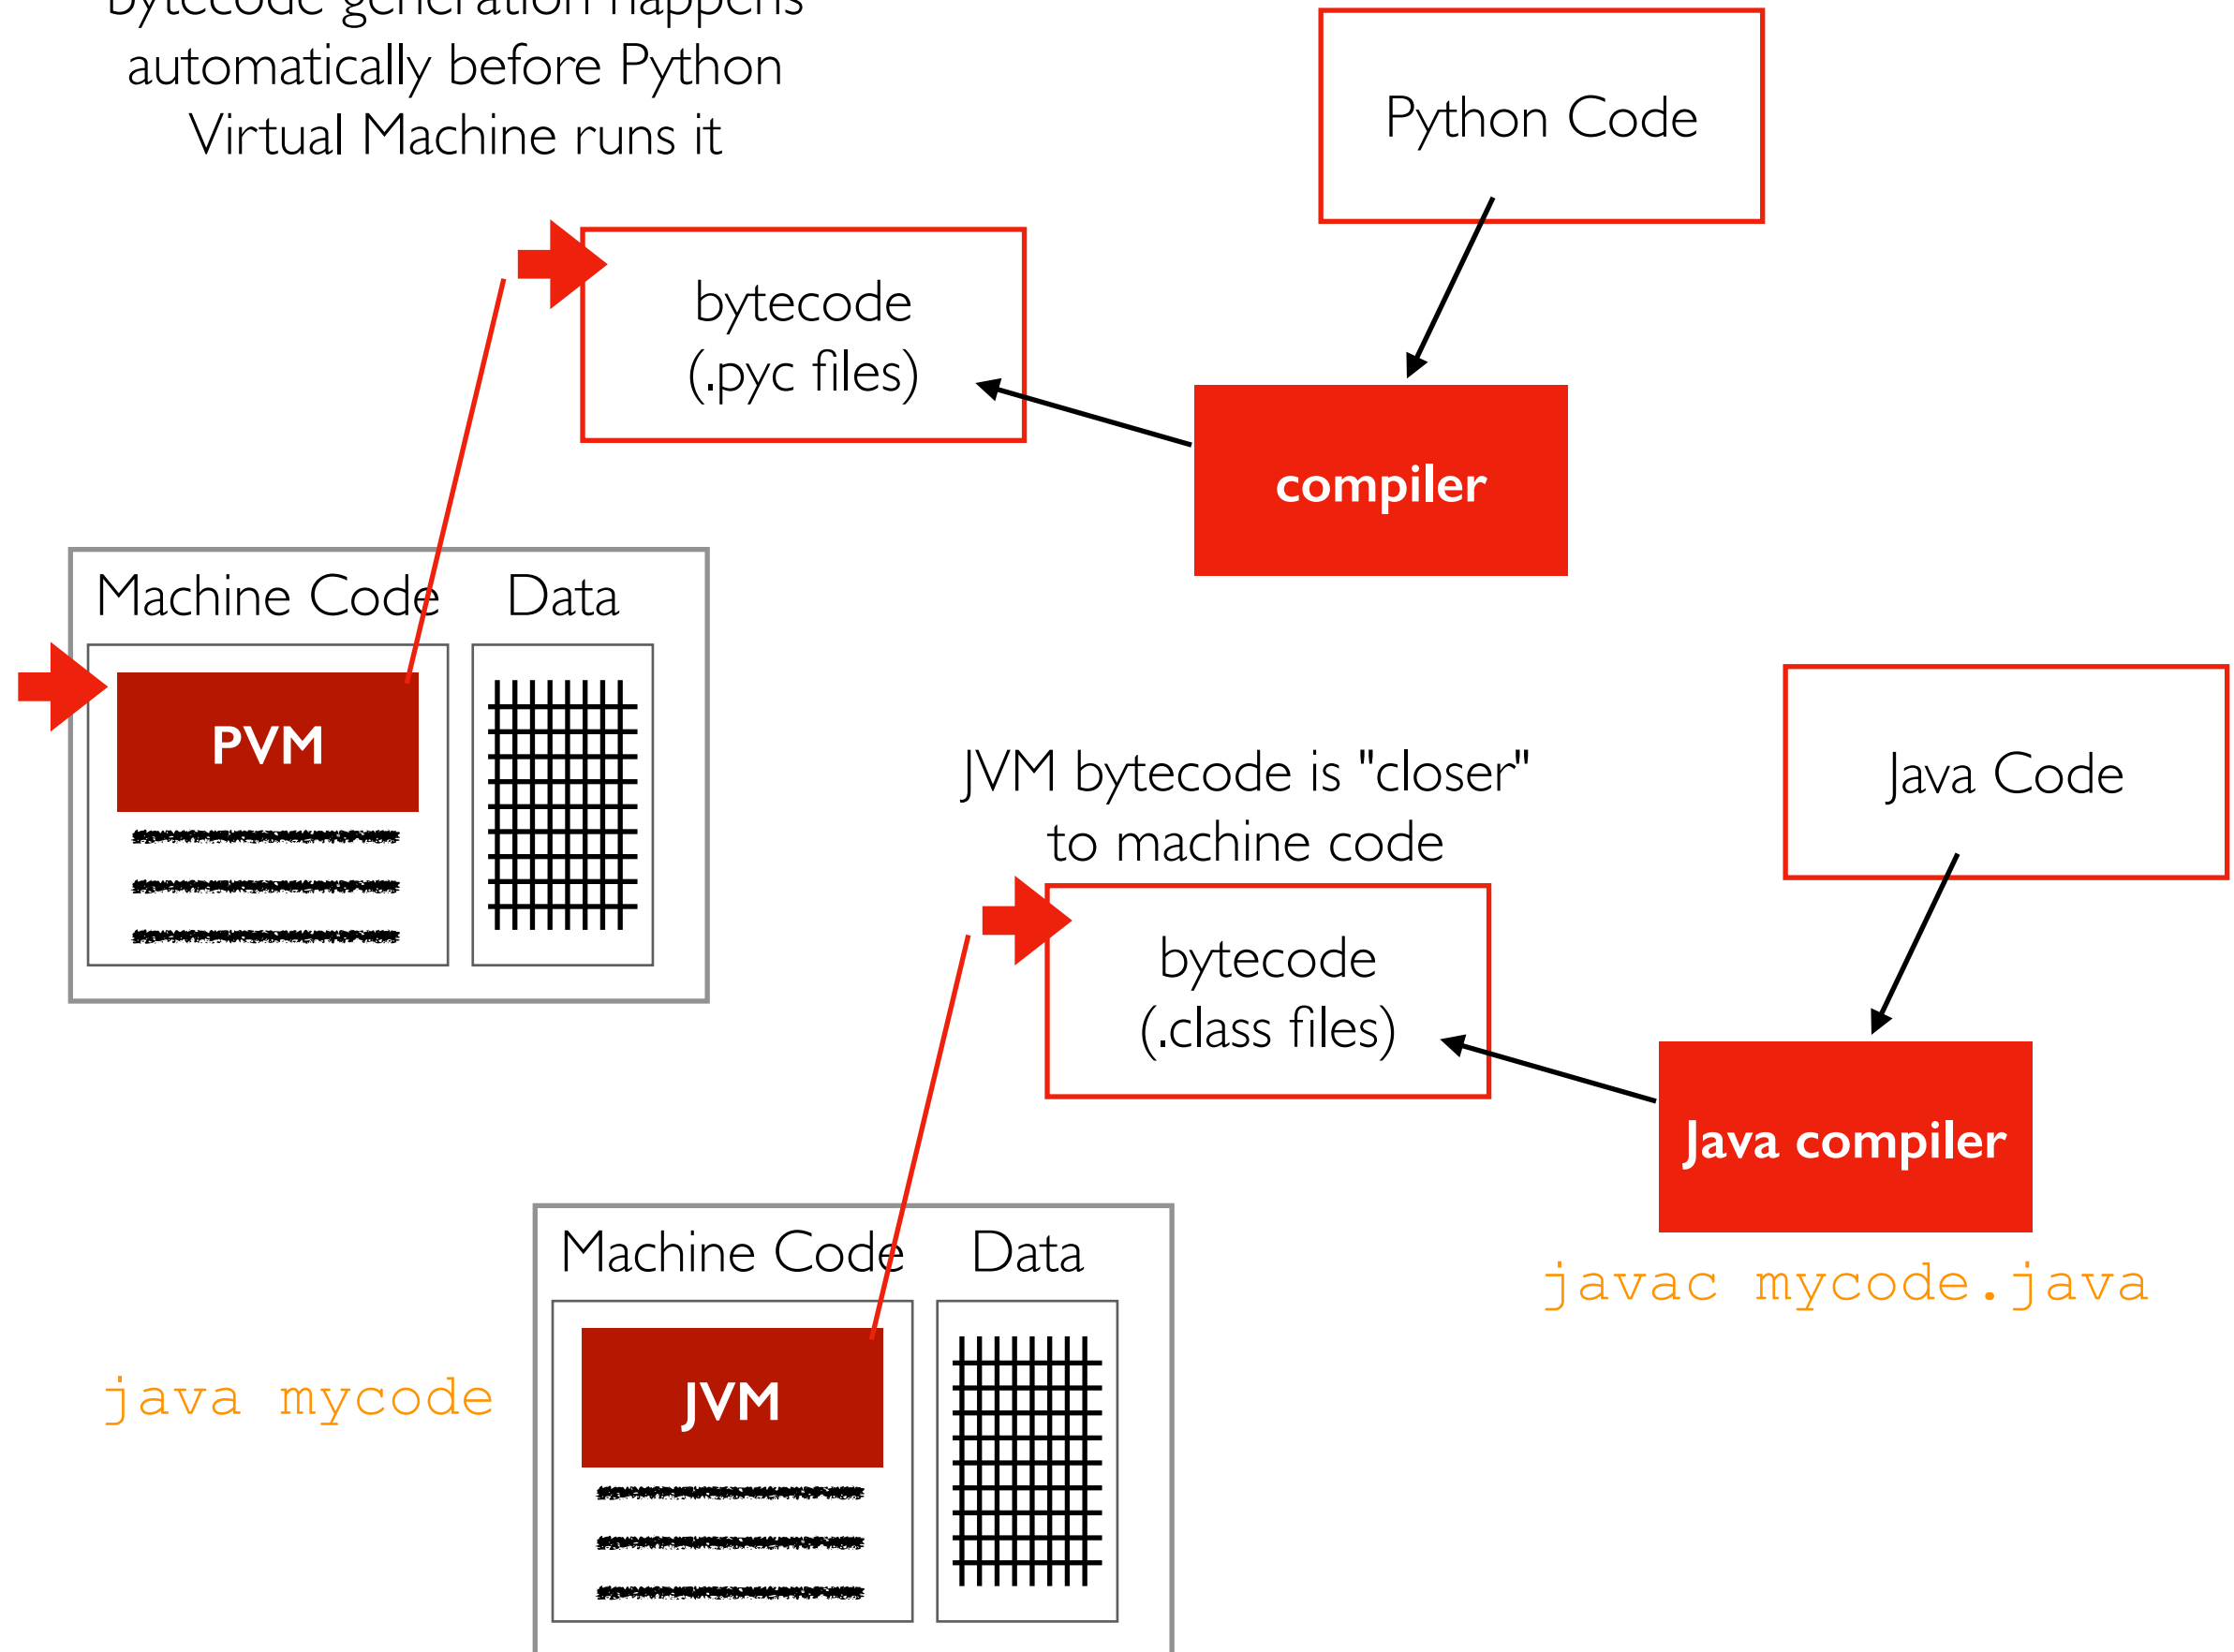








when you run "python3 ..."  
bytecode generation happens  
automatically before Python  
Virtual Machine runs it



# C/C++ Performance

**Advantage 1:** compiled languages are *usually* faster at runtime

- no overhead due to interpreter or language virtual machine
- however, cannot dynamically profile+optimize

# Outline

Welcome


Logistics

Background and Motivation

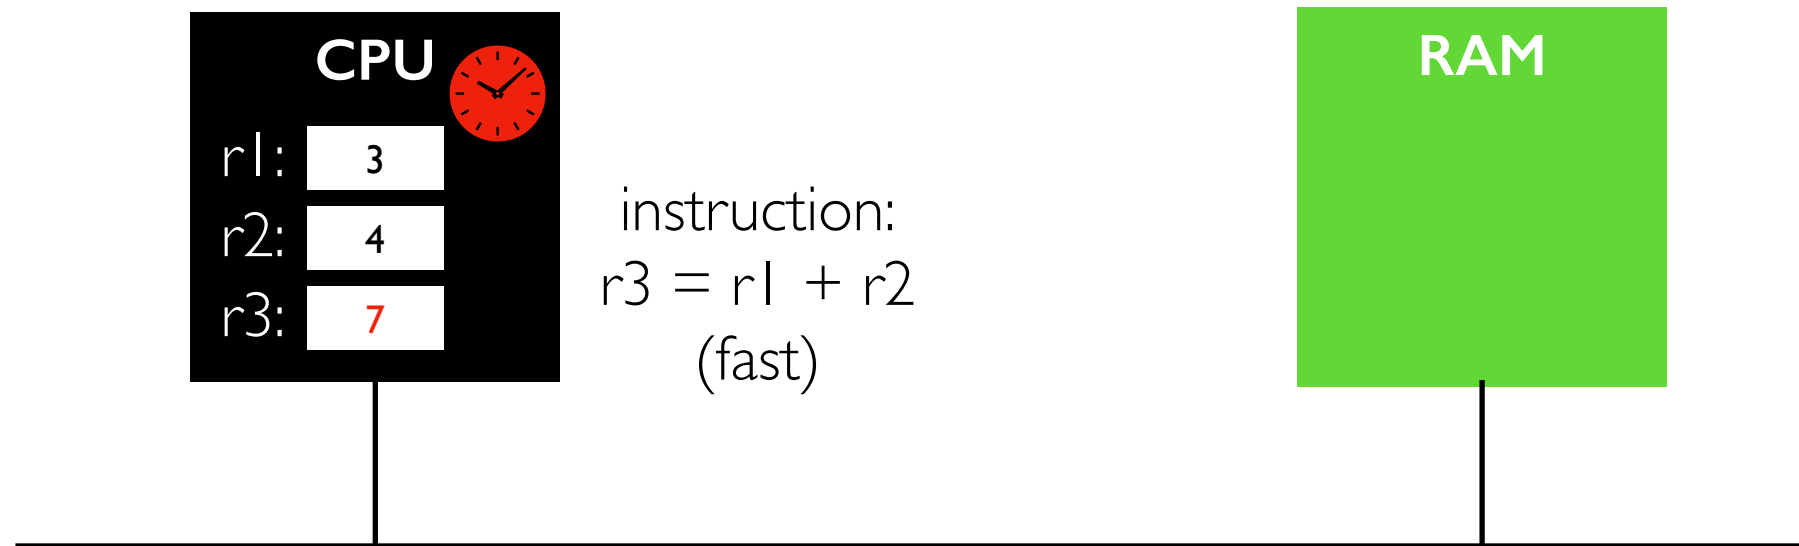
- Why C/C++: performance
  - how code runs
  - cachelines
  - garbage collection
  - safety checks
- Why C++ (over C): language features

Demos

# Background: CPU and RAM



# Background: CPU and RAM






# Load and Store




**challenge:** if we want to add some numbers stored in RAM, we need to **load** before adding and **store** after

# Latency



very slow, but not long enough to  
switch to a different process...



source: visuals, estimates

# Cache



## What happens:

- the value needed (for example, a 4-byte integer) goes to the register
- a whole **cacheline** (often 64 bytes) containing the value goes to the cache
- **future accesses to values in same cacheline will be relatively fast!**


## What matters for performance:

- how many cache misses there are (that is, how many times we need data that is not in the cache)
- how many values we access is less important

# Example 1: Step and Multiply

as K gets bigger, we do fewer multiplications. But does it matter?


```
for (int i = 0; i < arr.Length; i += K) arr[i] *= 3;
```



[Gallery of Processor Cache Effects](http://igoro.com/archive/gallery-of-processor-cache-effects/)

<http://igoro.com/archive/gallery-of-processor-cache-effects/>

# Example 1: Step and Multiply



**performance tip:** think about how many cachelines you're touching, not just about how many values

**k=1 loop:** all the ints, all the cachelines


```
int int int int int int int int int int int int int int int int int ..
```

k=2 loop: **half** the ints, **all** the cachelines


```
int int int int int int int int int int int int int int int int ..
```

# Example 2: Series of Strings

which layout is most cache friendly?



**linked list**



**array of references to strings**



**array of inline strings**

# C/C++ Performance

**Advantage 1:** compiled languages are *usually* faster at runtime

- no overhead due to interpreter or language virtual machine
- cannot dynamically profile+optimize

**Advantage 2:** C/C++ gives us more control over memory layout

- can design cache-friendly data structures

# Outline

Welcome

Logistics

Background and Motivation

- Why C/C++: performance
  - how code runs
  - cachelines
  - garbage collection
  - safety checks
- Why C++ (over C): language features


Demos



# Background: Memory Management


<https://pythontutor.com/>

```
→ 1 nums = [1,2,3]
→ 2 items = nums
  3 nums = [4,5,6]
  4 items = None
```



# Background: Memory Management

```
1  nums = [1,2,3]
→ 2  items = nums
→ 3  nums = [4,5,6]
4  items = None
```




# Background: Memory Management

1	nums = [1,2,3]
2	items = nums
→ 3	nums = [4,5,6]
→ 4	items = None

→ line that has just executed

→ next line to execute


line that has just executed



# Background: Memory Management

- data that can no longer be accessed in any way is "garbage"
- we can release garbage to free up memory
- in simple cases, the garbage objects might be recognizable immediately
- in complicated cases (for example, circular references), a background garbage collection algorithm needs to run to identify garbage
- garbage collection is costly and generally involves pausing execution (perhaps for many seconds!)

```
1  nums = [1,2,3]
2  items = nums
3  nums = [4,5,6]
→ 4  items = None
```



# C/C++ Performance

**Advantage 1:** compiled languages are *usually* faster at runtime

- no overhead due to interpreter or language virtual machine
- cannot dynamically profile+optimize

**Advantage 2:** C/C++ gives us more control over memory layout

- can design cache-friendly data structures

**Advantage 3:** C/C++ lets us manage memory allocation/deallocation manually

- YOU (the programmer) write code to manually delete allocations
- memory is freed up sooner (don't need to wait for garbage collection)
- no overheads for GC; no long pauses during GC

# Outline

Welcome

Logistics

Background and Motivation

- Why C/C++: performance
  - how code runs
  - cachelines
  - garbage collection
  - safety checks
- Why C++ (over C): language features



Demos

# Background: Safety Checks

A = [ 5 , 6 , 7 ]

B = [ 8 , 9 , 1 ]

every process has an address space, which resembles a big array of bytes (indexes are called addresses). All the processes data lives somewhere in that address space.




# Background: Safety Checks

A = [ 5 , 6 , 7 ]

B = [ 8 , 9 , 1 ]

A[ 0 ] → 5



0 unmapped  
range

mapped  
range

...bytes...

N




# Background: Safety Checks

A = [ 5 , 6 , 7 ]

B = [ 8 , 9 , 1 ]

A[ 2 ] → 7




# Background: Safety Checks

A = [5, 6, 7]

B = [8, 9, 1]

A[3] → **IndexError: list index out of range**



Many languages (Python, Java, etc) check bounds for you and raise an exception if you're outside. This checking has a performance cost, but is safer.

# Bounds Checking


```
def print_at(items, idx):  
    if 0 <= idx < len(items):  
        print(items[idx])  
    else:  
        print("bad index")
```

← Python checks that idx is in range, which is  
wasteful because your code already did that!

# C/C++ Approach

Trust programmer to write code that checks bounds.

Generally don't spend time on double checking that!




When programmer makes a mistake, however, are a variety of strange things that could happen...

# C/C++ Approach

Trust programmer to write code that checks bounds.

Generally don't spend time on double checking that!




When programmer makes a mistake, however, are a variety of strange things that could happen...

# C/C++ Approach

Trust programmer to write code that checks bounds.

Generally don't spend time on double checking that!




When programmer makes a mistake, however, are a variety of strange things that could happen...

# C/C++ Approach

Trust programmer to write code that checks bounds.

Generally don't spend time on double checking that!



When programmer makes a mistake, however, are a variety of strange things that could happen...

# C/C++ Performance

**Advantage 1:** compiled languages are *usually* faster at runtime

- no overhead due to interpreter or language virtual machine
- cannot dynamically profile+optimize

**Advantage 2:** C/C++ gives us more control over memory layout

- can design cache-friendly data structures

**Advantage 3:** C/C++ lets us manage memory allocation/deallocation manually

- YOU (the programmer) write code to manually delete allocations
- memory is freed up sooner (don't need to wait for garbage collection)
- no overheads for GC; no long pauses during GC

**Advantage 4:** C/C++ doesn't spend much compute time to catch programming mistakes

- avoids duplicated checking effort
- runs a little faster



# C/C++ Performance

**Advantage 1:** compiled languages are *usually* faster at runtime

- no overhead due to interpreter or language virtual machine
- cannot dynamically profile+optimize

**Advantage 2:** C/C++ gives us more control over memory layout

- can design cache-friendly data structures

**Advantage 3:** C/C++ lets us manage memory allocation/deallocation manually

- YOU (the programmer) manage memory
- memory is more difficult to manage
- no overhead of garbage collection

**Observation:** almost all these performance features make programming more difficult and introduce new occasions for bugs.

**Note:** there are many tools for calling from one language to another (Python to C, Java to C++, etc).

**Advantage 4:** C/C++

- avoids duplication of code
- runs a little faster

**Suggestion:** if 80% of execution time is spent on 20% of your code, consider writing the critical 20% in a fast language (like C++) and the rest in an "easy" language (like Python)

# Outline

Welcome

Logistics

Background and Motivation

- Why C/C++: performance
- Why C++ (over C): language features

Demos

# A Few Language Features in C++ but not C

## Function overloading

- multiple functions with the same name that accept different types

## Type deduction

- use "auto" type (or other features) to let C++ decide what the type should be
- templating, so you don't need many different similar functions to handle different types

## Alternatives to pointers

- references, smart pointers (for example, unique and shared)

## OOP (Object Oriented Programming)

- classes, inheritance (multiple!), public/private

## Resource management with RAII (Resource Acquisition is Initialization)

- use destructors to make sure resources are freed when necessary
- differentiate copy/move, manager ownership of objects over resources

## Rich STL (Standard Library)

- containers, iterators, algorithms

## Functional programming

- anonymous lambda functions
- many standard library functions that take function references

# Outline

Welcome

Logistics

Background and Motivation

- Why C/C++: performance
- Why C++ (over C): language features

Demos