

[544] Networking

Tyler Caraza-Harter

Outline

Networks

Internets and "The Internet"

Transport Protocols

Network Interface Controllers and MAC Addresses




NICs can connect a computer to different physical mediums, such as:

- Ethernet (wired)
- Wi-Fi (wireless)

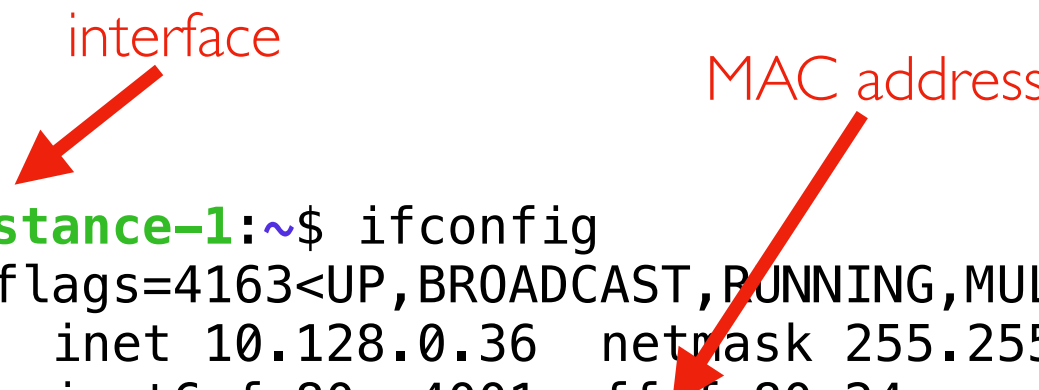
Every NIC in the world has a unique MAC (media access control) address

- 28 trillion possible addrs
- some devices randomly change their MAC addr for privacy

MAC address:



ifconfig (Interface Config)



```
trh@instance-1:~$ ifconfig
ens4: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1460
    inet 10.128.0.36 netmask 255.255.255.255 broadcast 0.0.0.0
    inet6 fe80::4001:aff:fe80:24 prefixlen 64 scopeid 0x20<link>
    ether 42:01:0a:80:00:24 txqueuelen 1000 (Ethernet)
    RX packets 2332795 bytes 6456770667 (6.4 GB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1994116 bytes 718670305 (718.6 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 407321 bytes 417582056 (417.5 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 407321 bytes 417582056 (417.5 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

`sudo apt install net-tools`

Virtual Interfaces

```
trh@instance-1:~$ ifconfig
ens4: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1460
    inet 10.128.0.36 netmask 255.255.255.255 broadcast 0.0.0.0
    inet6 fe80::4001:aff:fe80:24 prefixlen 64 scopeid 0x20<link>
    ether 42:01:0a:80:00:24 txqueuelen 1000 (Ethernet)
    RX packets 2332795 bytes 6456770667 (6.4 GB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1994116 bytes 718670305 (718.6 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 407321 bytes 417582056 (417.5 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 407321 bytes 417582056 (417.5 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

loopback (lo) device a virtual interface (not actual hardware)
connecting to a mini network containing just your computer

Google Console: Adding Interfaces (NICs)


Create Instance > Advanced Options > Networking

Network interfaces

Network interface is permanent

| | |
|---------------------------------------|---|
| default default (10.128.0.0/20) | ▼ |
| other-net subnet (10.0.0.0/24) | ▼ |
| ADD NETWORK INTERFACE | |

Virtual Machine Summary

| | | | | | |
|--------------------------|---|----------------------------|---------------|---|---|
| <input type="checkbox"/> |  | instance-2 | us-central1-a | 10.128.0.37 (nic0) | 34.29.220.248 (nic0) |
| | | | | 10.0.0.2 (nic1) | 35.202.74.234 (nic1) |

Google Console: Adding Interfaces

Create Instance > Advanced Options > Networking

Network interfaces 

```
trh@instance-2:~$ ifconfig
```

```
ens4: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1460
    inet 10.128.0.37 netmask 255.255.255.255 broadcast 0.0.0.0
    inet6 fe80::4001:aff:fe80:25 prefixlen 64 scopeid 0x20<link>
    ether 42:01:0a:80:00:25 txqueuelen 1000 (Ethernet)
    RX packets 637 bytes 546000 (546.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 612 bytes 97265 (97.2 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

```
ens5: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1460
    inet 10.0.0.2 netmask 255.255.255.255 broadcast 0.0.0.0
    inet6 fe80::4001:aff:fe00:2 prefixlen 64 scopeid 0x20<link>
    ether 42:01:0a:00:00:02 txqueuelen 1000 (Ethernet)
    RX packets 51 bytes 9955 (9.9 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 61 bytes 6834 (6.8 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```


```
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 120 bytes 13534 (13.5 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 120 bytes 13534 (13.5 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

| | |
|---|--|
| 7 | 34.29.220.248 (nic0) 35.202.74.234 (nic1) |
|---|--|

Networks

A **network** has **nodes** that send bytes to other nodes **by MAC address**

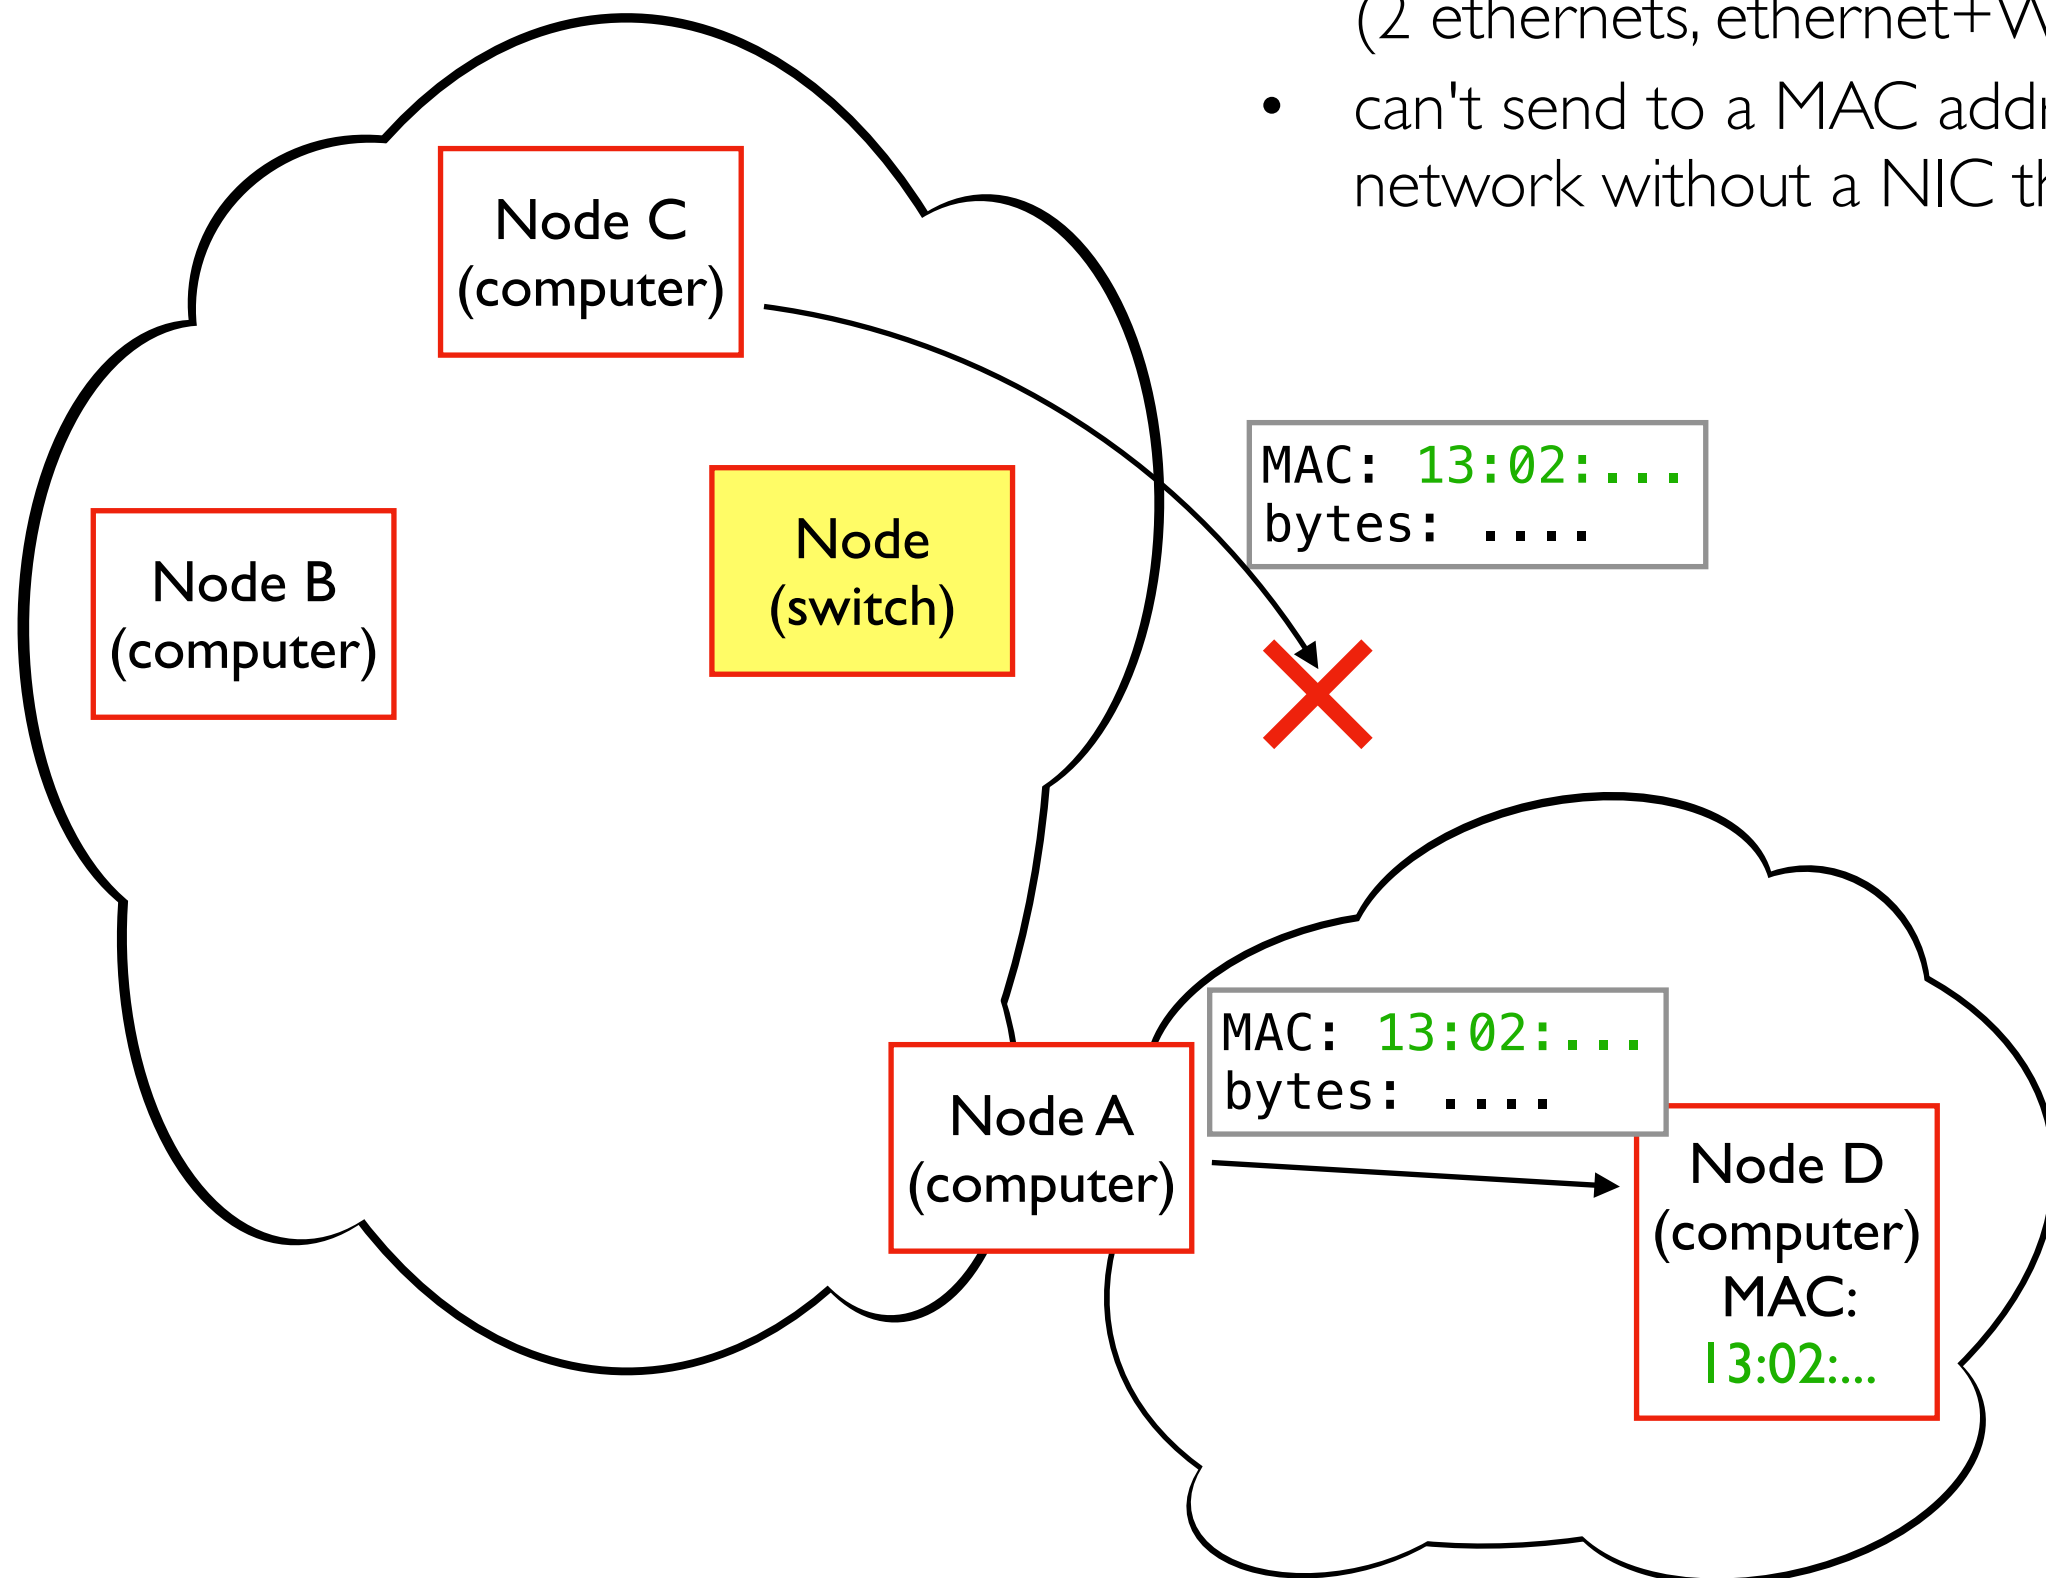
- **nodes**: computer, switch, etc
- direct, or **forwarded by switches**
- whole network uses same physical tech (Wi-Fi, Ethernet, etc)



Networks

Computers can have multiple NICs

- can be on multiple networks (2 ethernets, ethernet+Wi-Fi, etc)
- can't send to a MAC addr in another network without a NIC there



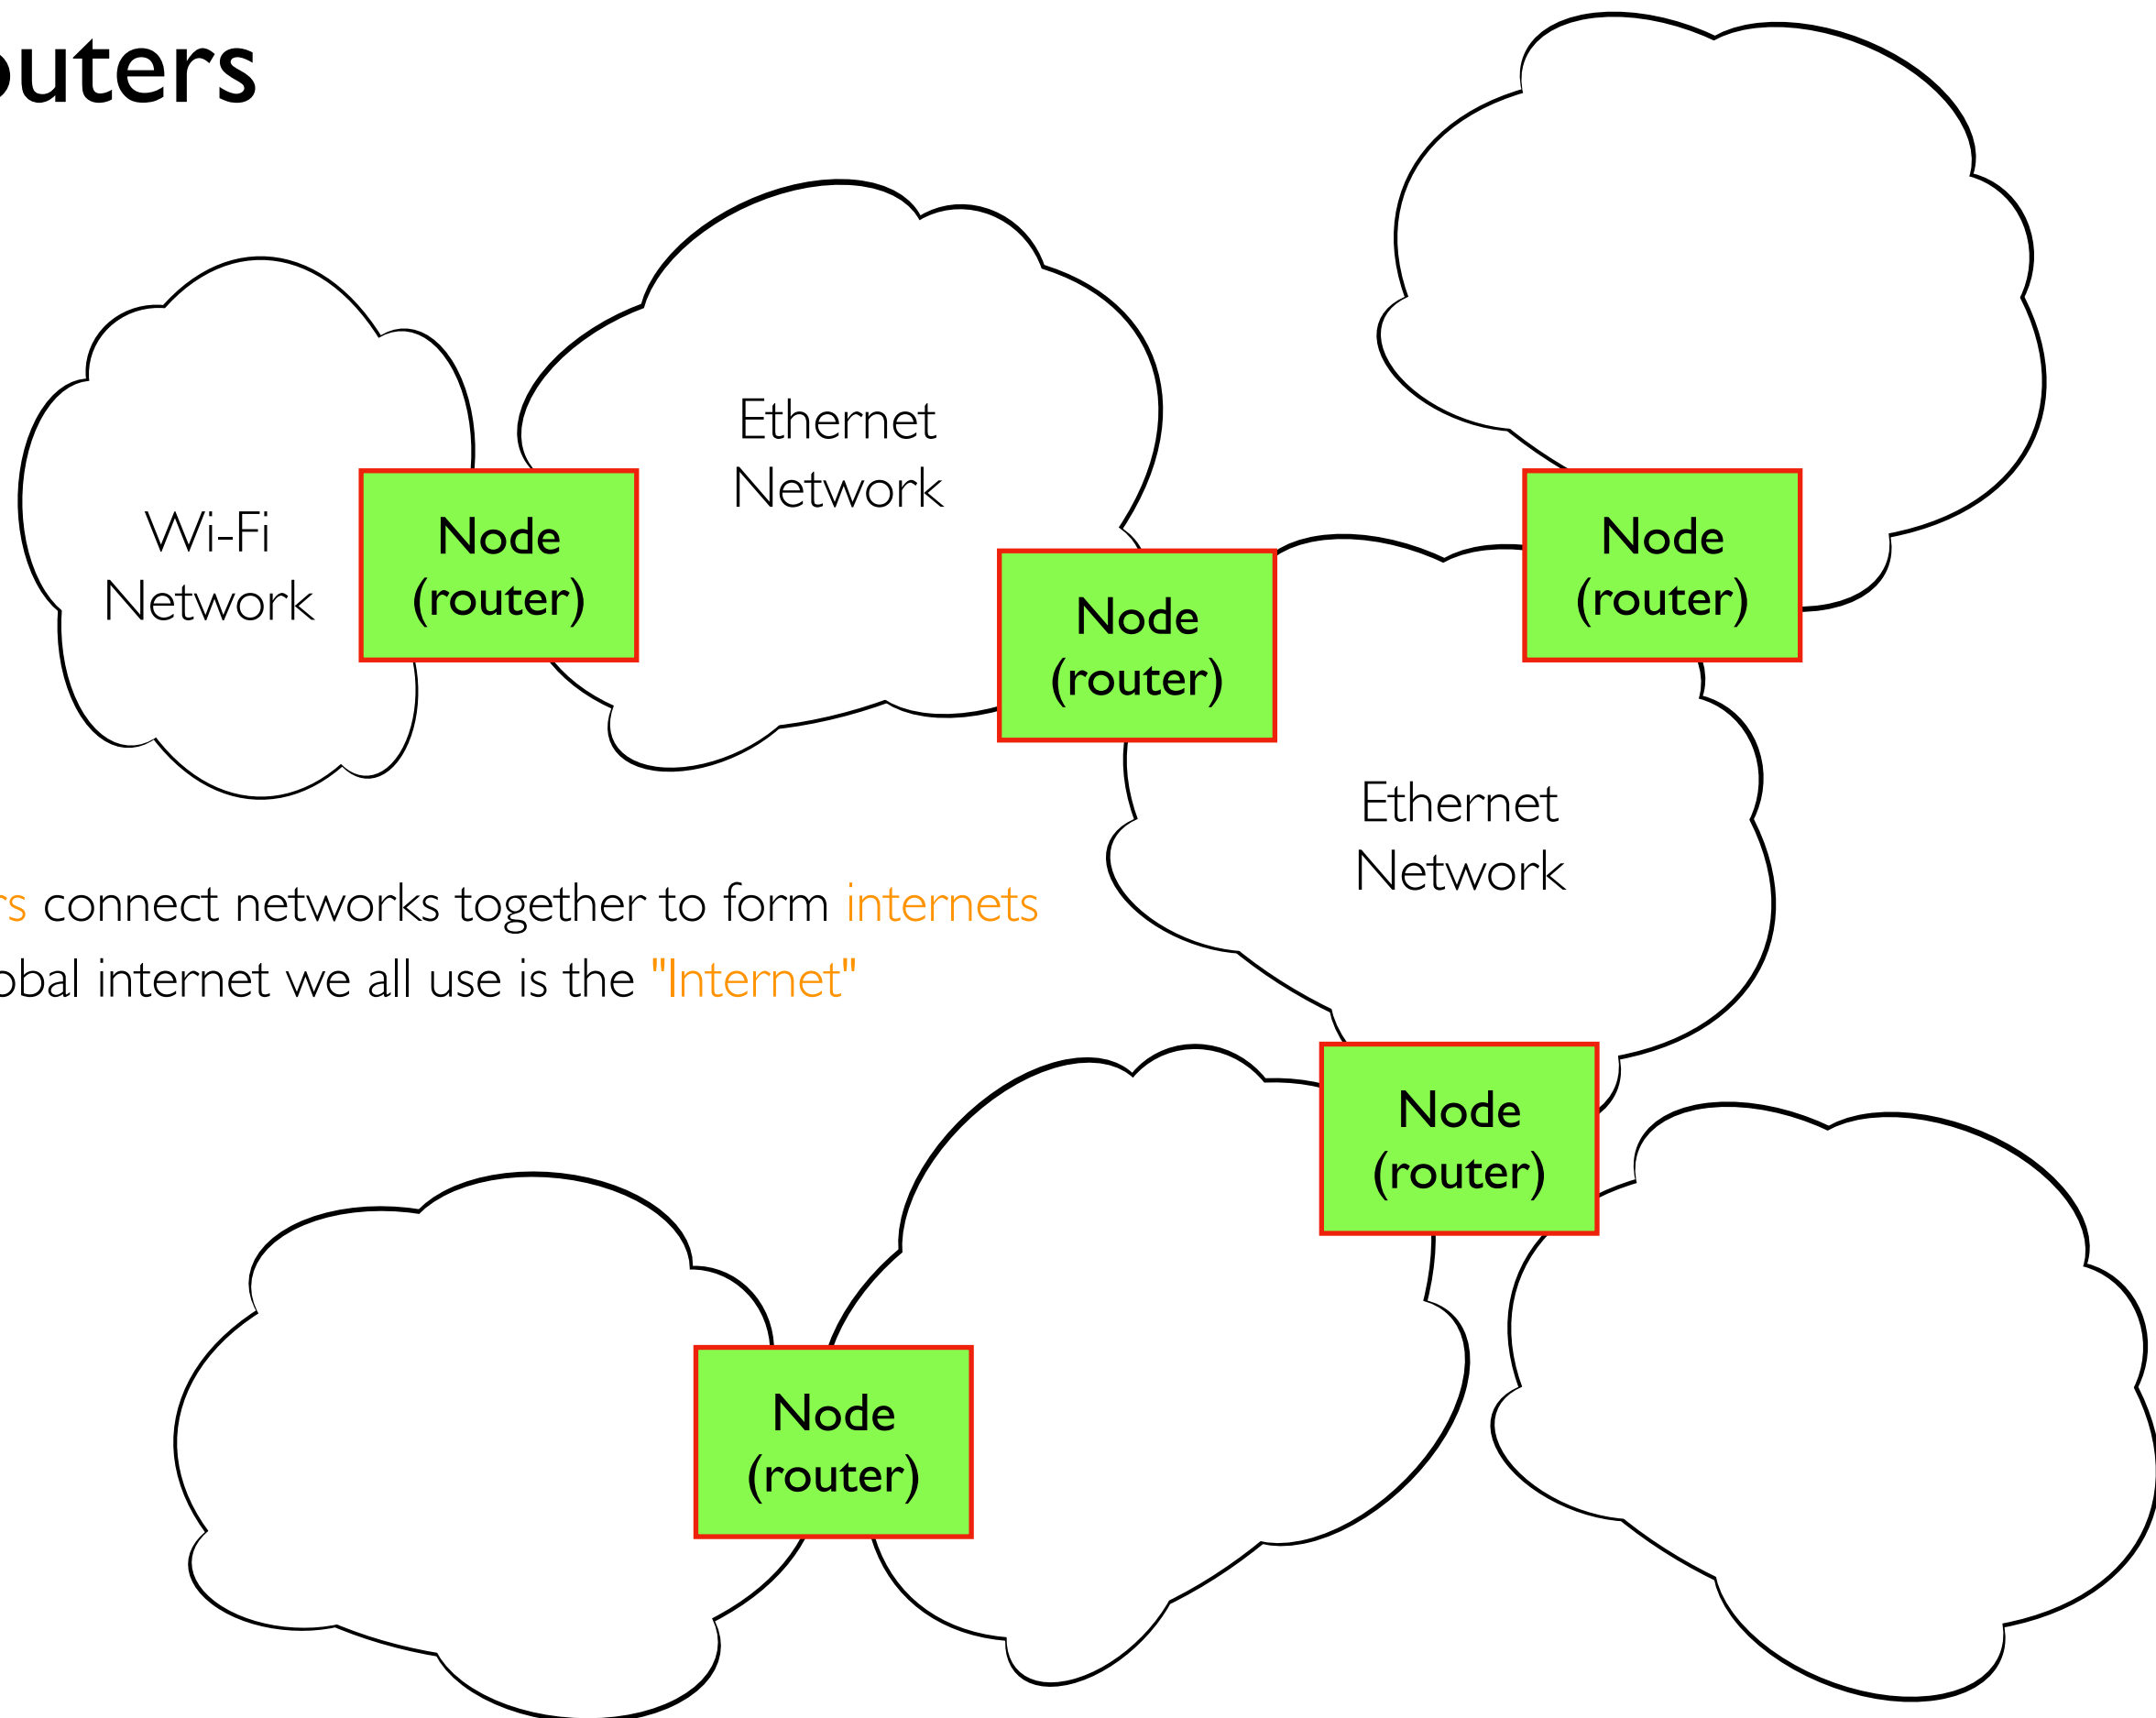
Outline

Networks

Internets and "The Internet"

Transport Protocols

Routers

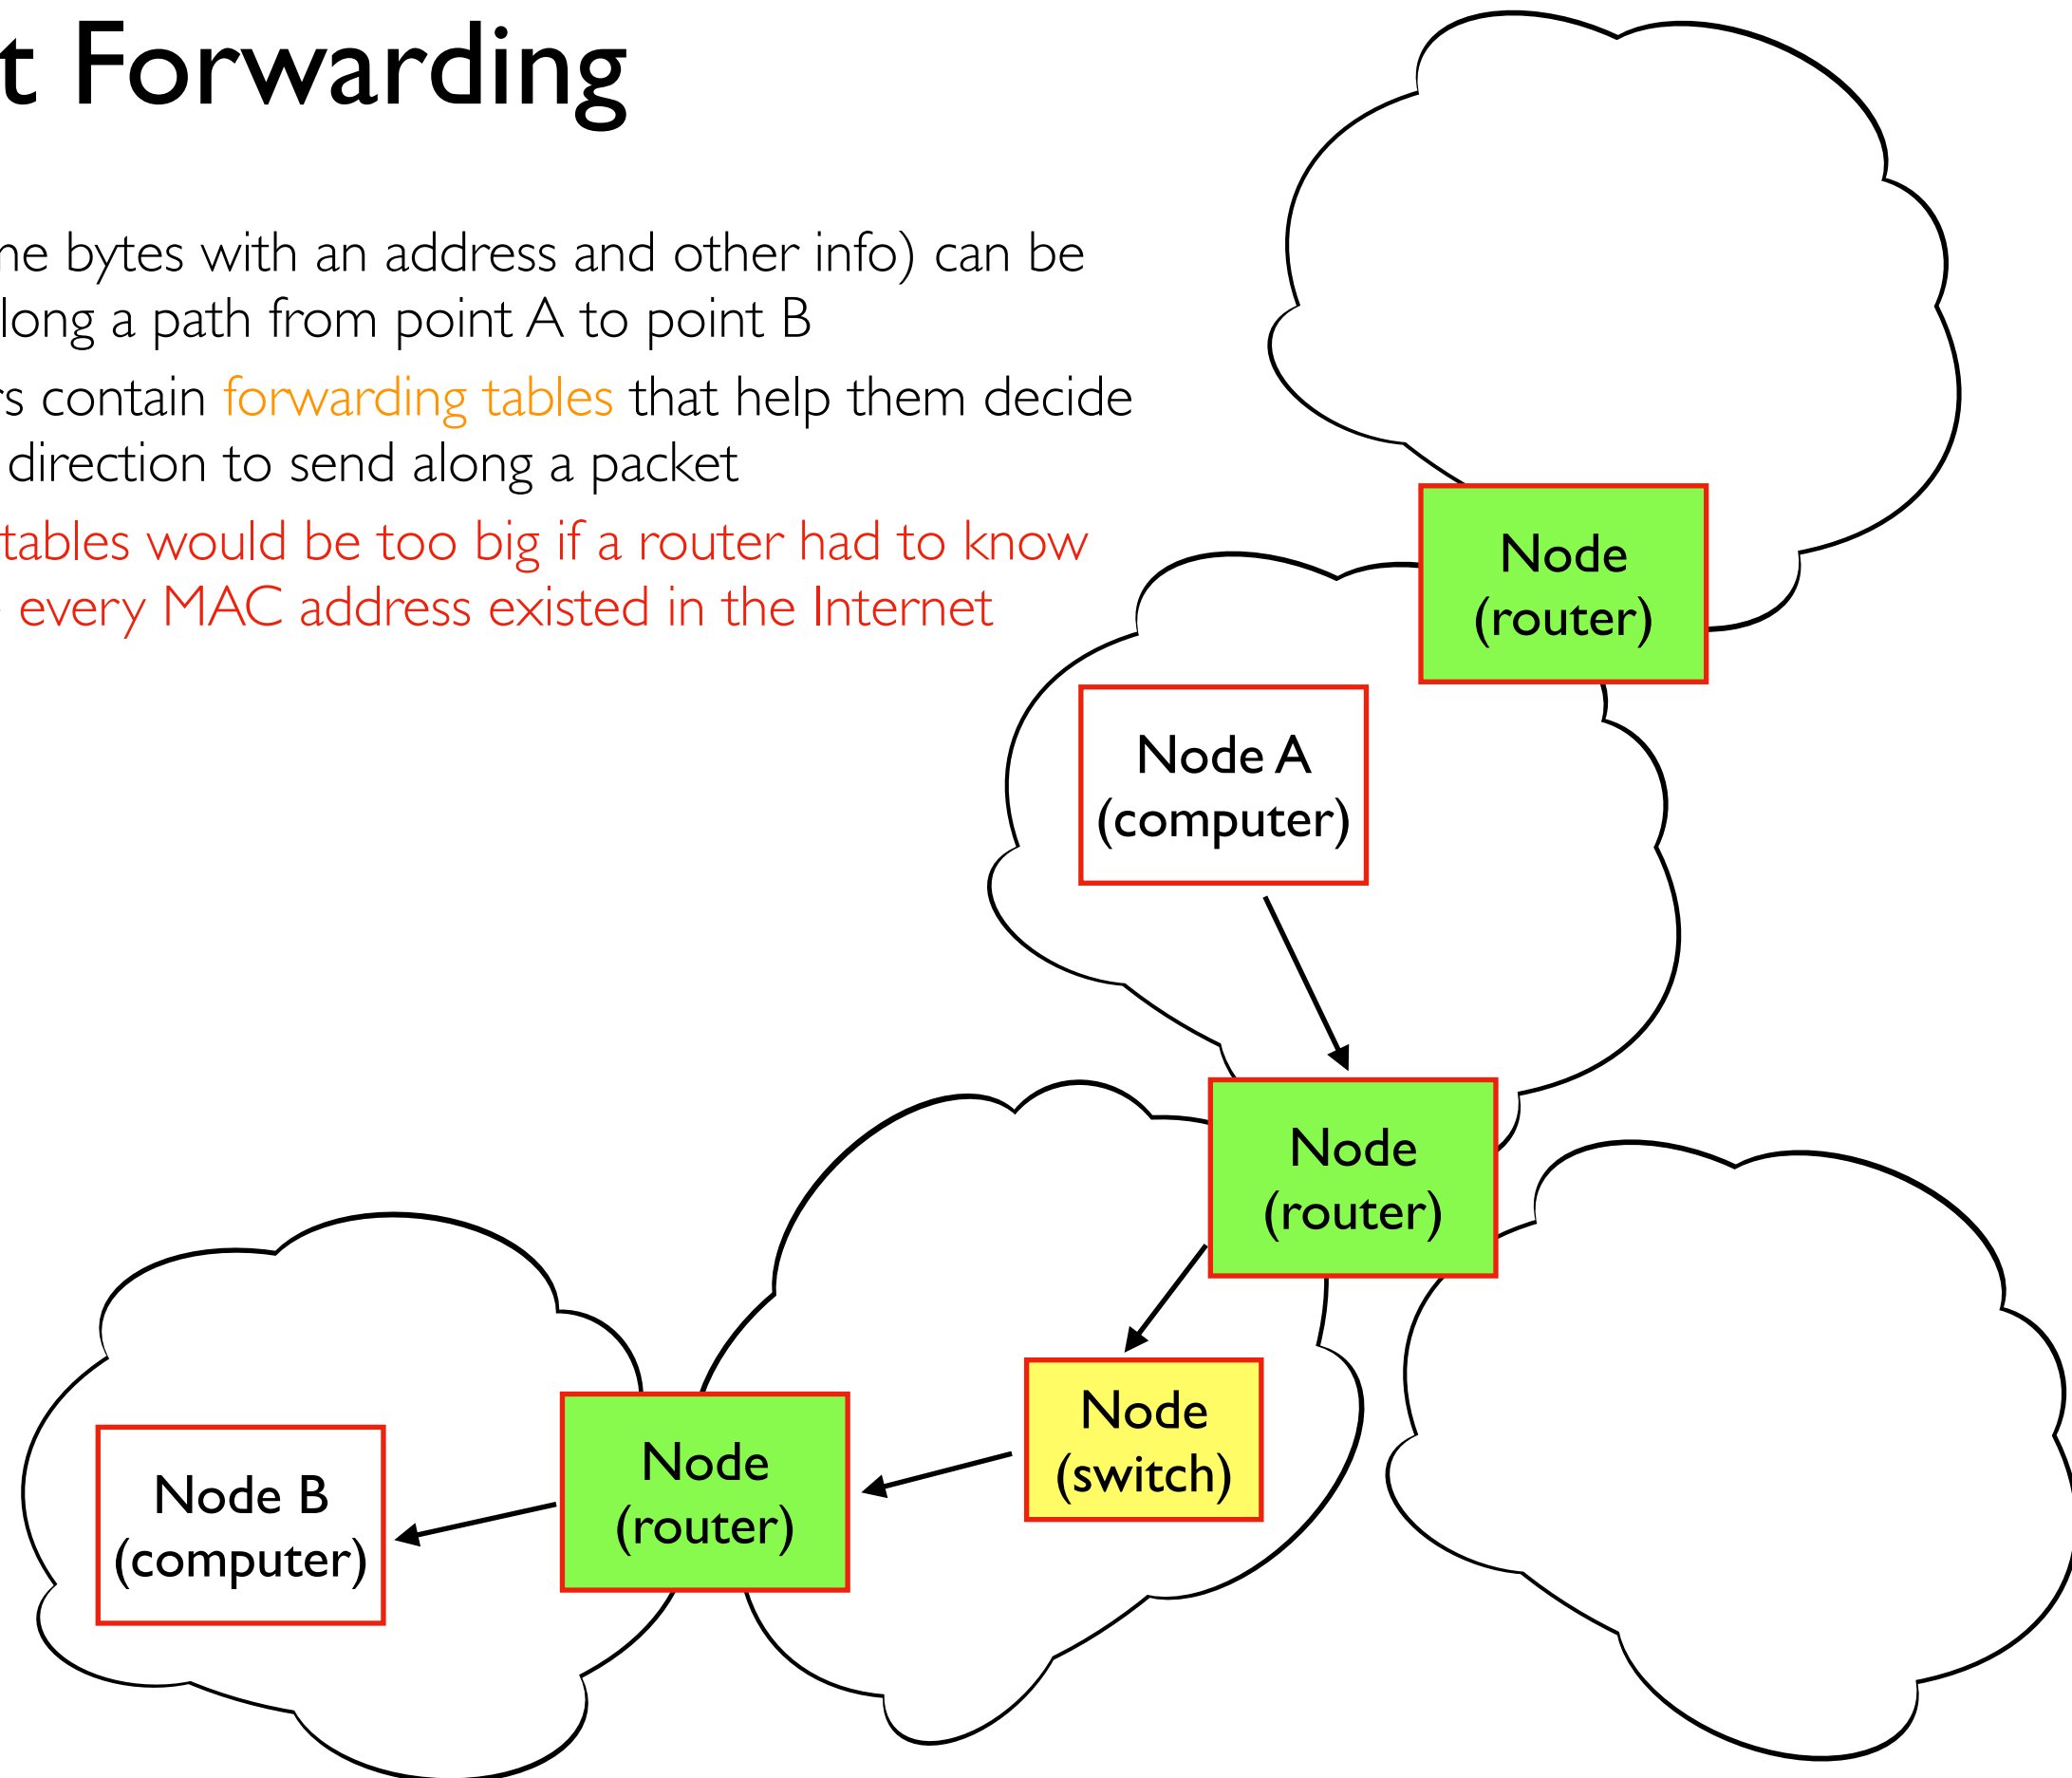


- routers connect networks together to form internets
- the global internet we all use is the "Internet"

Packet Forwarding

Packets (some bytes with an address and other info) can be forwarded along a path from point A to point B



- routers contain **forwarding tables** that help them decide which direction to send along a packet
- those tables would be too big if a router had to know where every MAC address existed in the Internet



Internet Protocol

IP addresses are used to send packets across an internet

- example: 34 . 29 . 237 . 29 (domains can map to IP addrs)
- there are about 4 billion possible IP addresses (IPv4)
- IPv6 (less used) are 4x longer
- forwarding tables only need to know which way to send for a given network number



Listening on an Interface

```
trh@instance-2:~$ ifconfig
ens4: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1460
    inet 10.0.1.2 netmask 255.255.255.255 broadcast 0.0.0.0
    inet6 fe80::214:0000:0000:0000%ens4<link>
    ether 42:00:00:00:00:00
    ...

ens5: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1430
    inet 10.0.3.2 netmask 255.255.255.255 broadcast 10.0.3.2
    inet6 fe80::214:0000:0000:0000%ens5<link>
    ether 42:00:00:00:00:00
    ...

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1
    loop txqueuelen 1000 (Ethernet)
    ...
```

all of them: `python3 -m http.server --bind 0.0.0.0`

Private Networks

Challenges


- we don't have enough IPv4 addresses
- we don't want every machine to be able to receive packets from anywhere

Private ranges:

- 192.168.0.0 to 192.168.255.255
- 172.16.0.0 to 172.31.255.255
- 10.0.0.0 to 10.255.255.255

these can be divided into "sub networks"
(subnets) to create different networks in
a bigger org

Private networks allow duplicates and unreachable machines



Private Networks

```
trh@instance-2:~$ ifconfig
ens4: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1460
    inet 10.0.1.2
    inet6 fe80::40
    ether 42:01:0a:00:01:02  txqueuelen 1000  (Ethernet)
    ...
```


```
python3 -m http.server --bind 10.0.1.2
```

Private ranges:

- 192.168.0.0 to 192.168.255.255
- 172.16.0.0 to 172.31.255.255
- 10.0.0.0 to 10.255.255.255

http://10.0.1.2:....
won't work in web browser!


Private networks allow duplicates and unreachable machines



Network Address Translation

Google Console (view NAT config)

| <input type="checkbox"/> | Status | Name ↑ | Internal IP | External IP |
|--------------------------|--------|----------------------------|--|--|
| <input type="checkbox"/> | ✓ | instance-1 | 10.128.0.36 (nic0) | 34.29.237.29 (nic0) |
| <input type="checkbox"/> | ✓ | instance-2 | 10.0.1.2 (nic0) 10.0.3.2 (nic1) | 35.202.74.234 (nic0) 34.29.220.248 (nic1) |




Network Address Translation

Google Console (view NAT config)

| <input type="checkbox"/> | Status | Name ↑ | Internal IP | External IP |
|--------------------------|--------|----------------------------|--|--|
| <input type="checkbox"/> | ✓ | instance-1 | 10.128.0.36 (nic0) | 34.29.237.29 (nic0) |
| <input type="checkbox"/> | ✓ | instance-2 | 10.0.1.2 (nic0) 10.0.3.2 (nic1) | 35.202.74.234 (nic0) 34.29.220.248 (nic1) |

By default, the external IPs are "ephemeral" (change upon reboot). You can rent "static" IPs that don't change.

Browser



Server

```
trh@instance-2:~/temp$ python3 -m http.server --bind 10.0.1.2
Serving HTTP on 10.0.1.2 port 8000 (http://10.0.1.2:8000/) ...
72.33.0.184 - - [10/Feb/2023 21:12:53] "GET / HTTP/1.1" 200 -
...
```

Outline

Networks

Internets and "The Internet"

Transport Protocols


Port Numbers

Computers might be running multiple processes using the network

- IP address => which NIC?
- Port number => which process?

```
trh@instance-2:~$ python3 -m http.server --directory=A --bind 10.0.1.2 8000 &  
[1] 13502  
Serving HTTP on 10.0.1.2 port 8000 (http://10.0.1.2:8000/) ...
```

```
trh@instance-2:~$ python3 -m http.server --directory=B --bind 10.0.1.2 9000 &  
[2] 13503  
Serving HTTP on 10.0.1.2 port 9000 (http://10.0.1.2:9000/) ...
```



TopHat

42:01:0a:80:00:25 is an example of what?

Transport Protocols

Most common

- **UDP** (User Datagram Protocol)
- **TCP** (Transmission Control Protocol)

BOTH build on IP networking and BOTH provide **port numbers**

```
trh@instance-2:~/temp$ sudo lsof -i tcp -P
```

```
COMMAND      PID  NODE  NAME
```

```
...
```

```
sshd          863  TCP  *:22 (LISTEN)
```

```
sshd          863  TCP  *:22 (LISTEN)
```

```
...
```

```
python3      13607 TCP  instance-2...internal:8000 (LISTEN)
```

```
python3      13608 TCP  instance-2...internal:9000 (LISTEN)
```

Reliability: UDP vs. TCP


Packets may be

- dropped
- reordered
- split

TCP saves+reassembles **packets** in order to provide original **message** (when possible).
For packet drops, it retries. We'll mostly use TCP.

UDP doesn't do this extra work. Why ever use UDP?

Network Stack: Common Implementations



Network applications (like most complex systems) are not built as one single system. Layers are built upon other layers to provide additional functionality.


Outline

Networks

Internets and "The Internet"

Transport Protocols

HTTP (Hypertext Transfer Protocol)



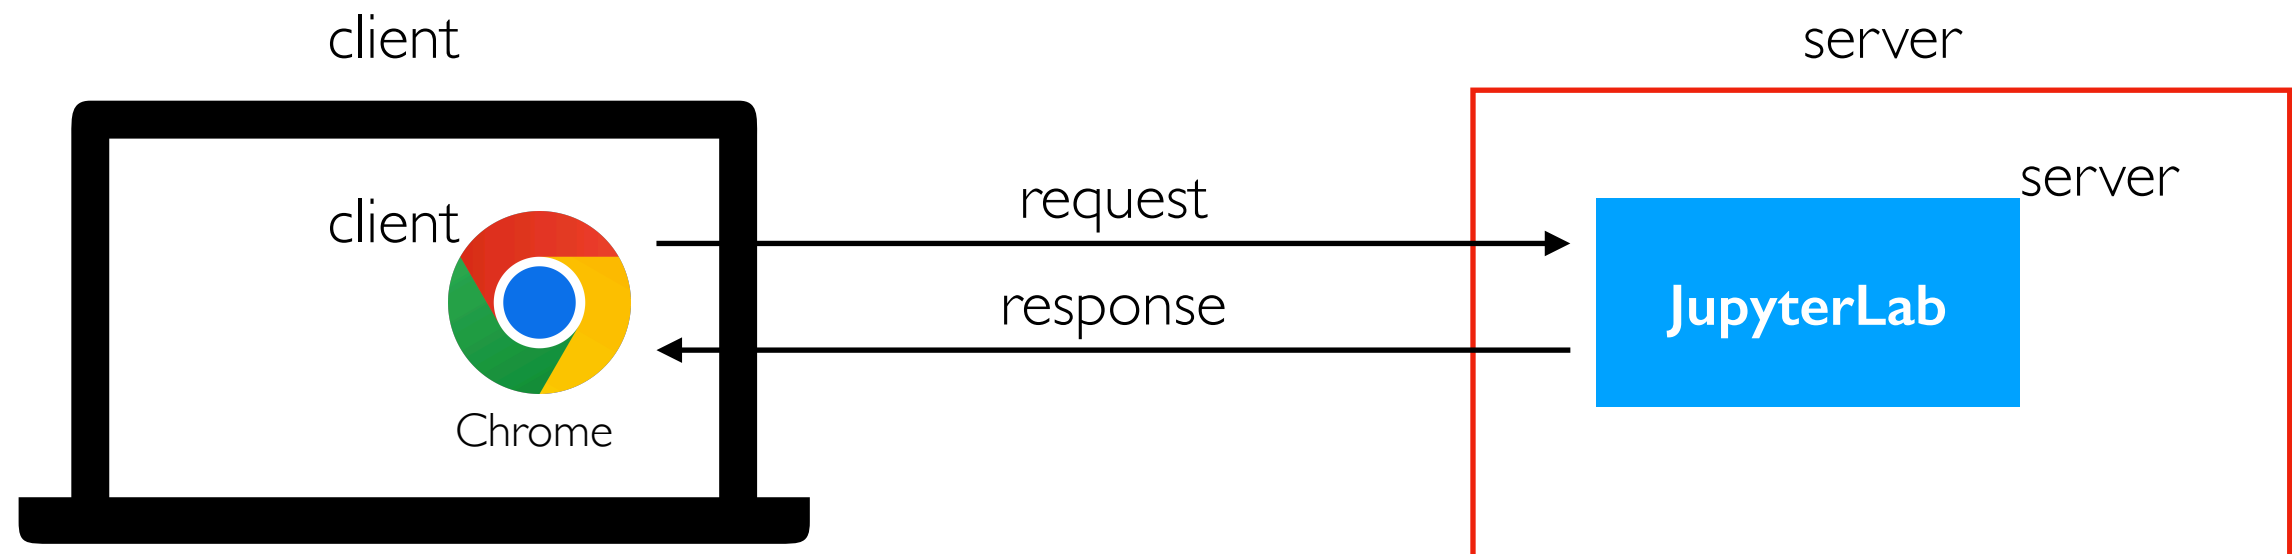
<https://tyler.caraza-harter.com:443/cs544/s23/schedule.html>

domain name
(mapped to an IP)


port
(443 is default
for https)

resource

HTTP Messages Between Clients and Servers



Parts: method, resource, status code, headers, body



<https://developer.mozilla.org/en-US/docs/Web/HTTP/Messages>

HTTP Methods (types of messages)

Types of request

- **POST**: create a new resource (request+response have body)
- **PUT**: update a resource (request+response have body, usually)
- **GET**: fetch a resource (response has body)
- **DELETE**: delete a resource ()
- others...

Canvas API example:

GET <https://canvas.wisc.edu/api/v1/conversations>
(see all Canvas conversations in JSON format)

POST <https://canvas.wisc.edu/api/v1/conversations>
(create new Canvas conversation)

<https://canvas.instructure.com/doc/api/conversations.html>


Outline

Networks

Internets and "The Internet"

Transport Protocols

Remote Procedure Calls (RPCs)



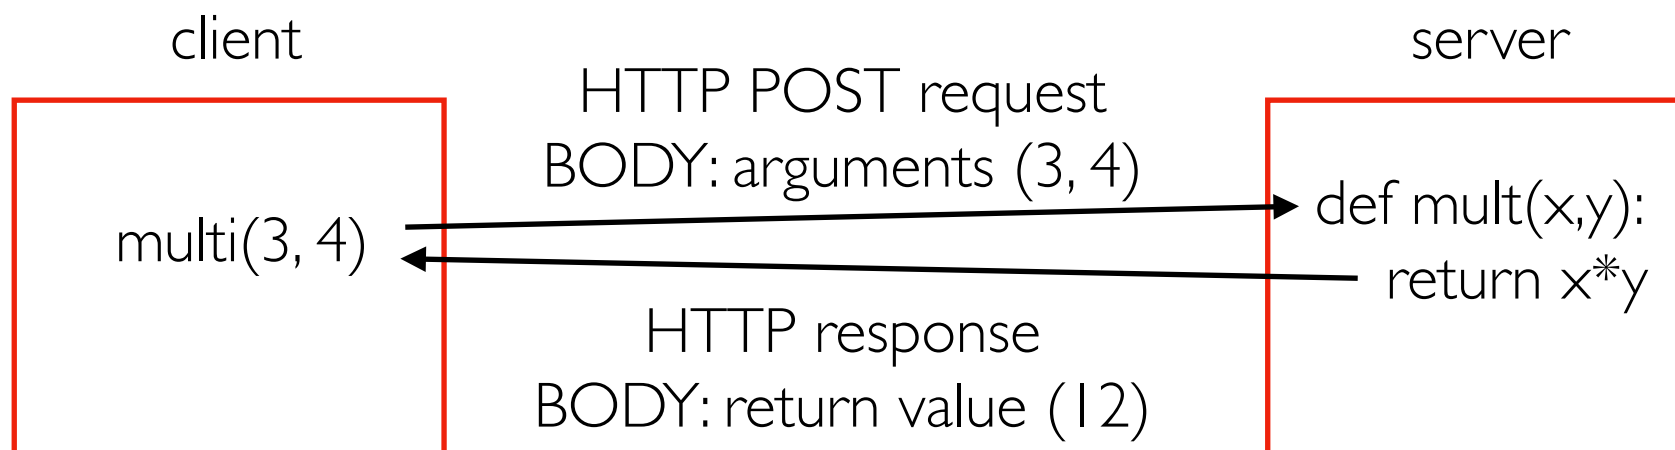

procedure = function

- **main** calling **add** is a regular procedure call
- **main** call **mult** is a remote procedure call

There are MANY tools to do RPCs

- Thrift (developed at Meta)
- gRPC (developed at Google) -- this semester

gRPC builds on HTTP



Serialization/deserialization (Protobufs)

How do we represent arguments and return values as bytes in a request/response body?

Serialization: various types (ints, strs, lists, etc) to **bytes** ("wire format")

Deserialization: **bytes** to various types

Challenge 1: every language has different types and we want cross-languages calls

gRPC uses Google's **Protocol Buffers** provide a uniform type system.

Challenge 2: different CPUs order bytes differently

cpu A int32:

| | | | |
|--------|--------|--------|--------|
| byte 1 | byte 2 | byte 3 | byte 4 |
|--------|--------|--------|--------|


cpu B int32:

| | | | |
|--------|--------|--------|--------|
| byte 4 | byte 3 | byte 2 | byte 1 |
|--------|--------|--------|--------|

| .proto | C++ | Java | Pytho |
|--------|--------|------------|-------|
| double | double | double | float |
| float | float | float | float |
| int32 | int32 | int | int |
| int64 | int64 | long | int |
| uint32 | uint32 | int | int |
| uint64 | uint64 | long | int |
| sint32 | int32 | int | int |
| sint64 | int64 | long | int |
| bool | bool | boolean | bool |
| string | string | String | str |
| bytes | string | ByteString | bytes |

<https://protobuf.dev/programming-guides/proto/>


Variable-Length Encoding



For **computational efficiency**, int32's use 4 bytes during computation. Also helps w/ offsets.

For **space efficiency**, smaller numbers in int32s use fewer bytes (4 bytes is max). This reduces network traffic.

Example nums in a protobuf:



Demos...