

[544] Spark Streaming

Tyler Caraza-Harter

Outline: Spark Streaming

DStreams

Grouped Aggregates

Watermarks

Pivoting

Joining

Exactly-Once Semantics


Review RDD Data Lineage: Transformations and Actions

```
data = [  
    ("A", 1),  
    ("B", 2),  
    ("A", 3),  
    ("B", 4)  
]
```


```
def mult2(row):  
    return (row[0], row[1] * 2)  
  
def onlyA(row):  
    return row[0] == "A"
```

goal: get 2 times the second column wherever the first column is "A"

```
table = sc.parallelize(data)  
double = table.map(mult2)  
doubleA = double.filter(onlyA)  
doubleA.collect()
```



```
[ ('A', 2),  
  ('A', 6) ]
```



Handling Data Changes: Re-Calculate Everything


```
data = [  
    ("A", 1),  
    ("B", 2),  
    ("A", 3),  
    ("B", 4),  
    new data ("A", 5),  
    ("C", 6)  
]
```

```
def mult2(row):  
    return (row[0], row[1] * 2)  
  
def onlyA(row):  
    return row[0] == "A"
```

Round 1



Round 2



re-doing work is wasteful!

Handling Data Changes: Incremental Computation

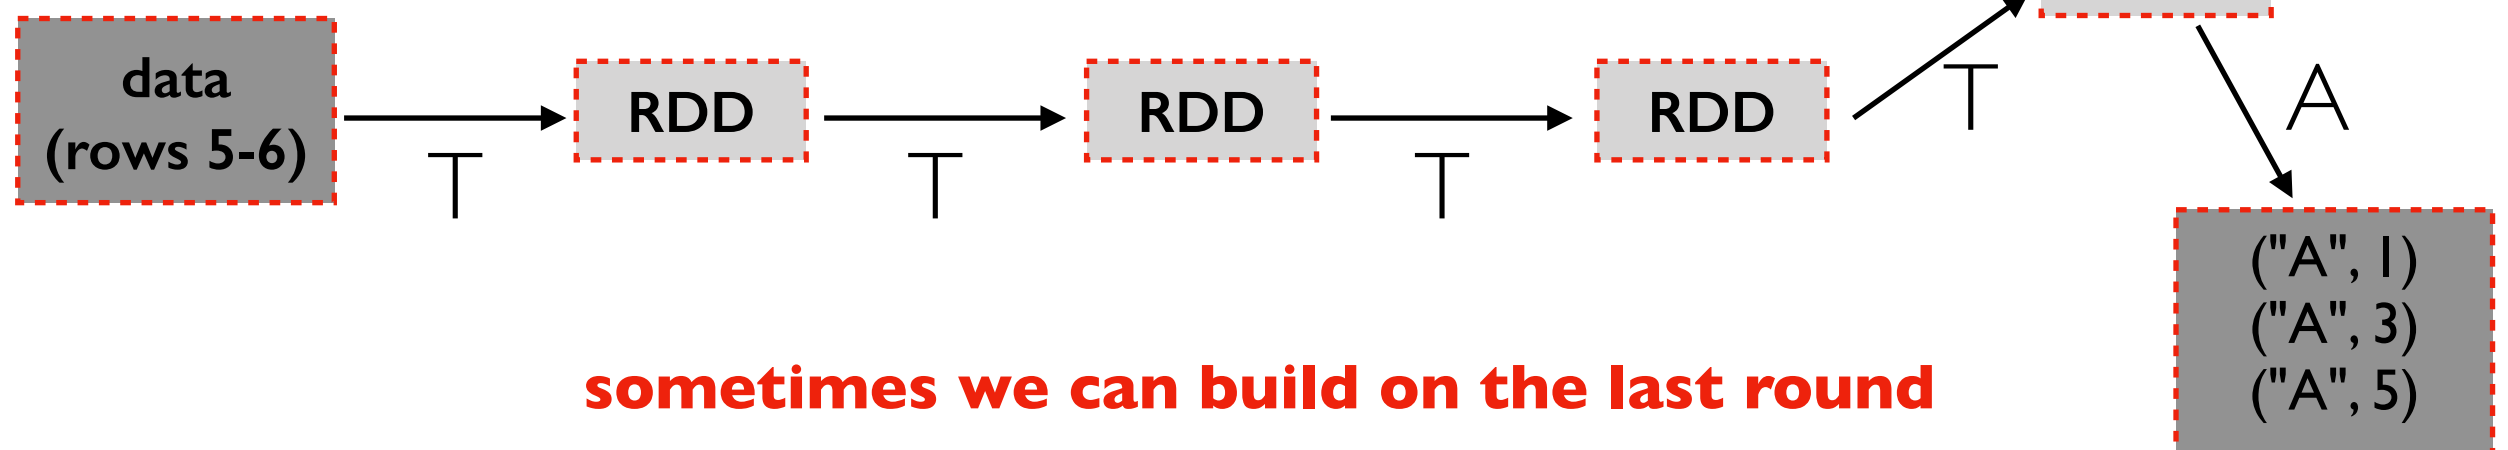
```
data = [  
    ("A", 1),  
    ("B", 2),  
    ("A", 3),  
    ("B", 4),  
    new data ("A", 5),  
    ("C", 6)  
]
```

```
def mult2(row):  
    return (row[0], row[1] * 2)  
  
def onlyA(row):  
    return row[0] == "A"
```

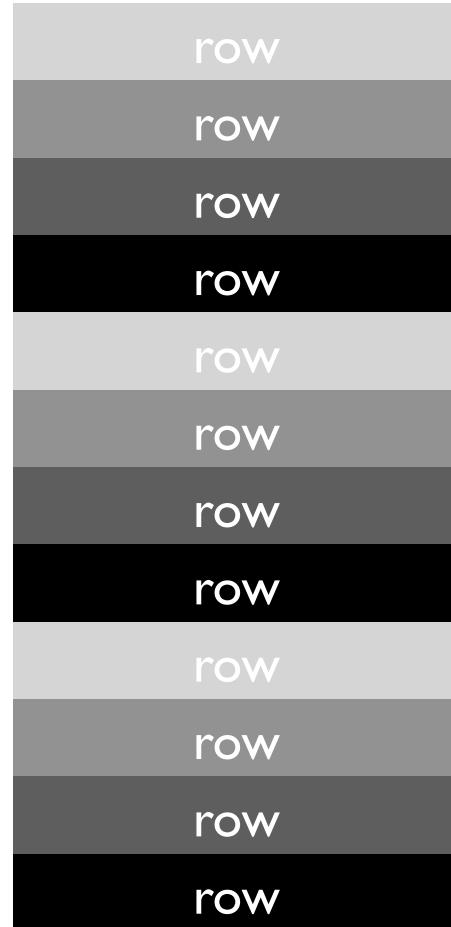
Round 1



Round 2



Some DataFrames constantly grow



row
row
row
row
row
row
row
row
row
row
row
row
row
row
row
row

continuously growing table

Mini Batches

row
row
row
row
row
row

mini batch


row
row
row
row
row
row

mini batch

continuously growing table

```
trigger (processingTime="12 seconds")
```

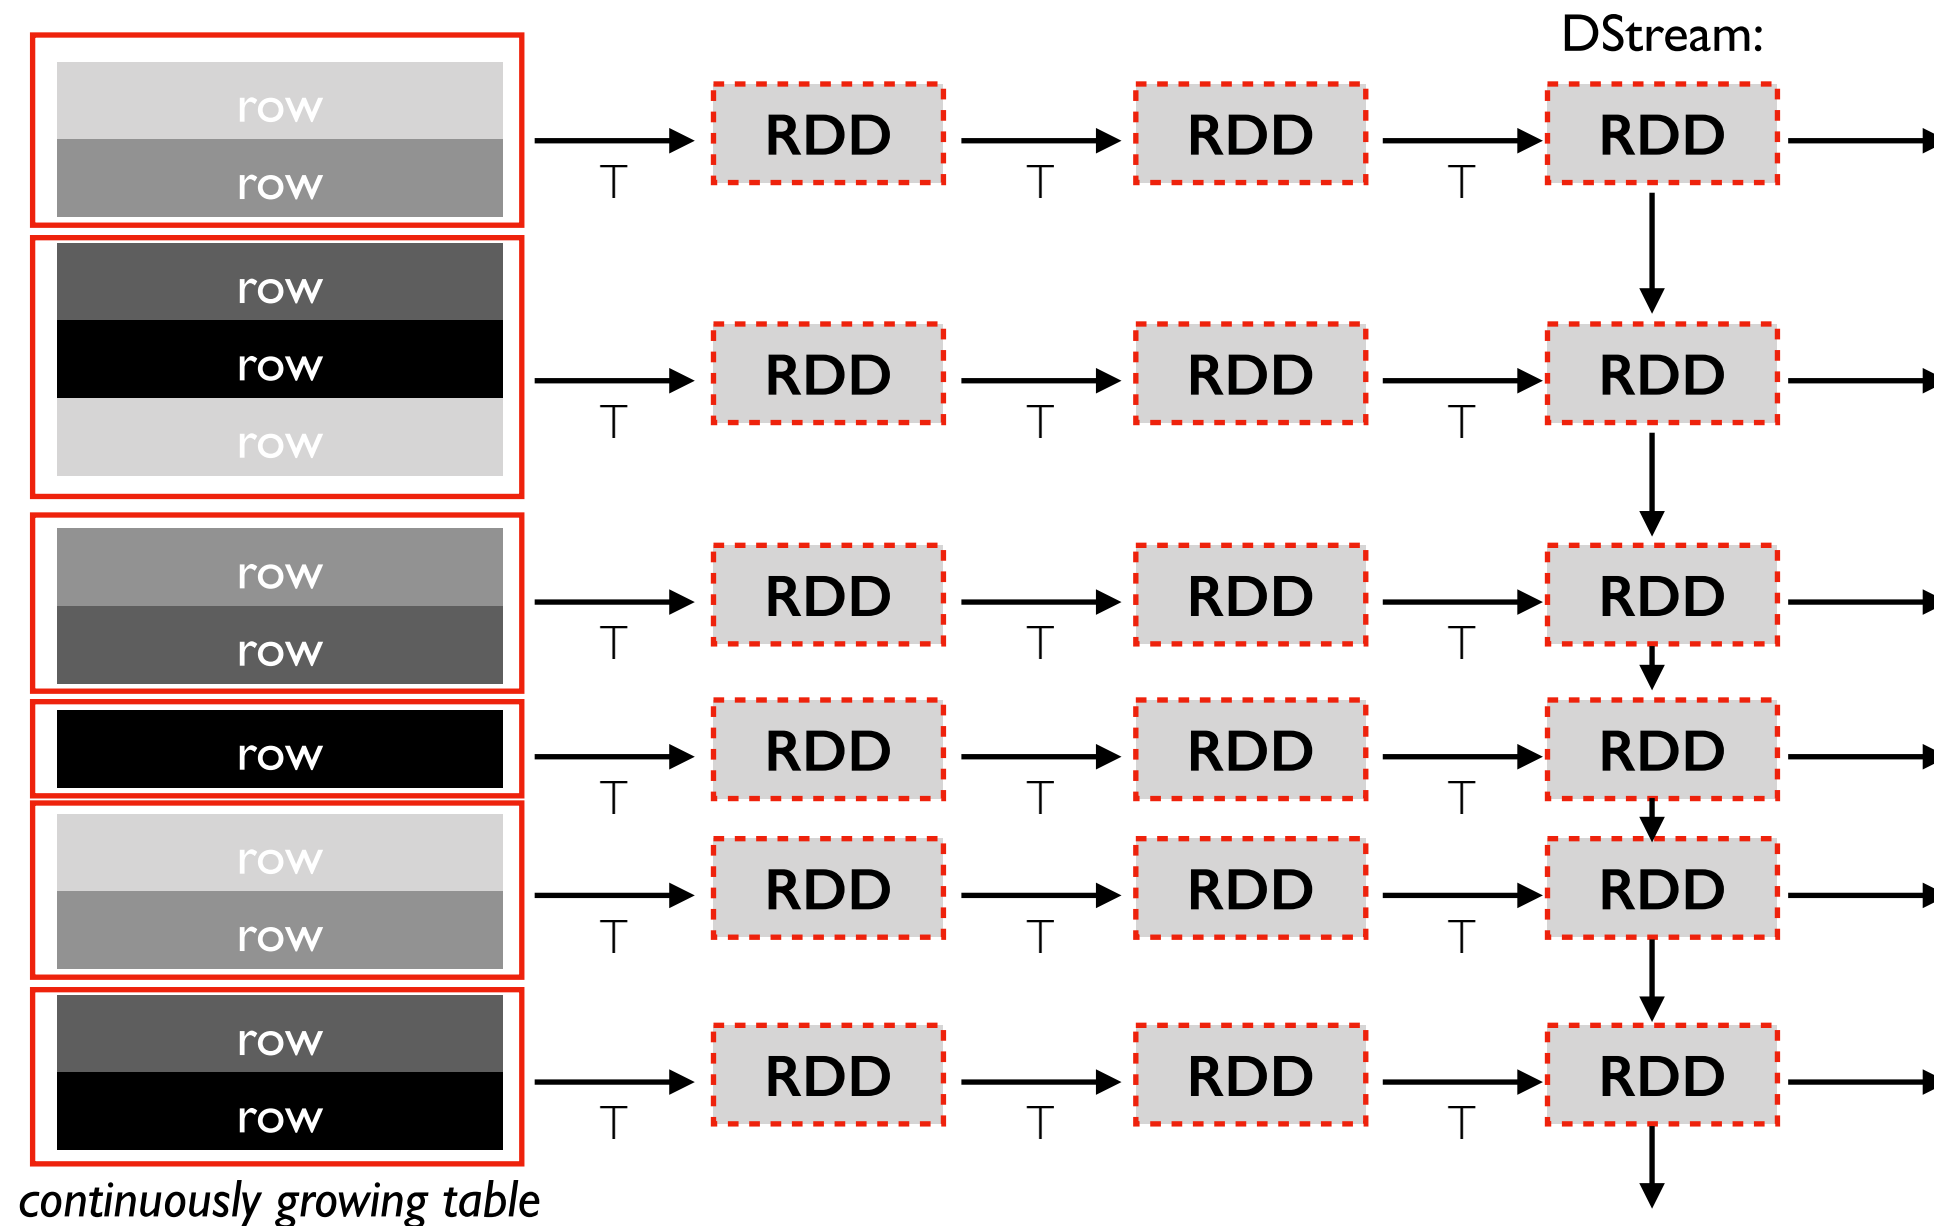
Trigger Frequency



continuously growing table

```
trigger (processingTime="4 seconds")
```

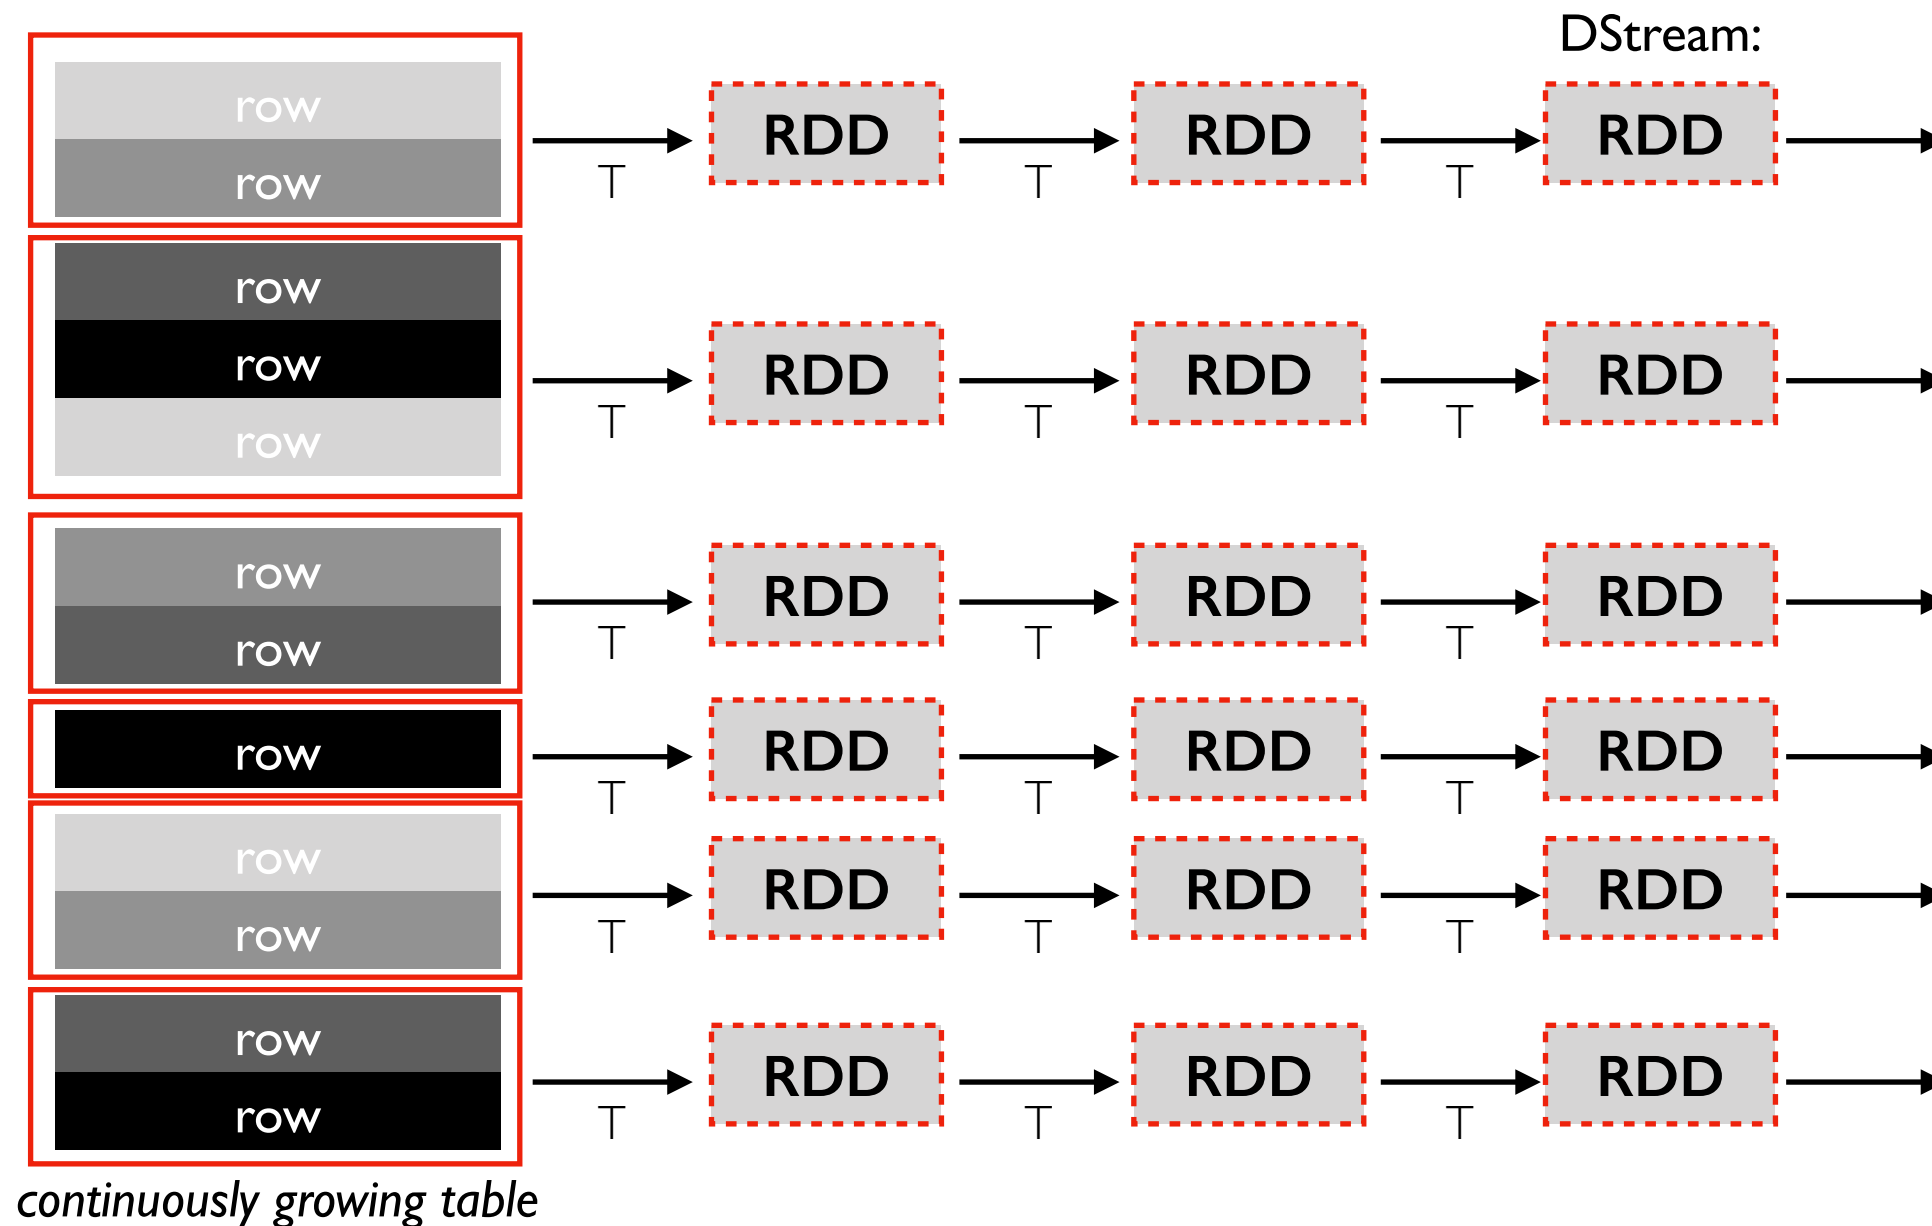

DStream (Stateful)



A Spark **DStream** is a series of RDDs

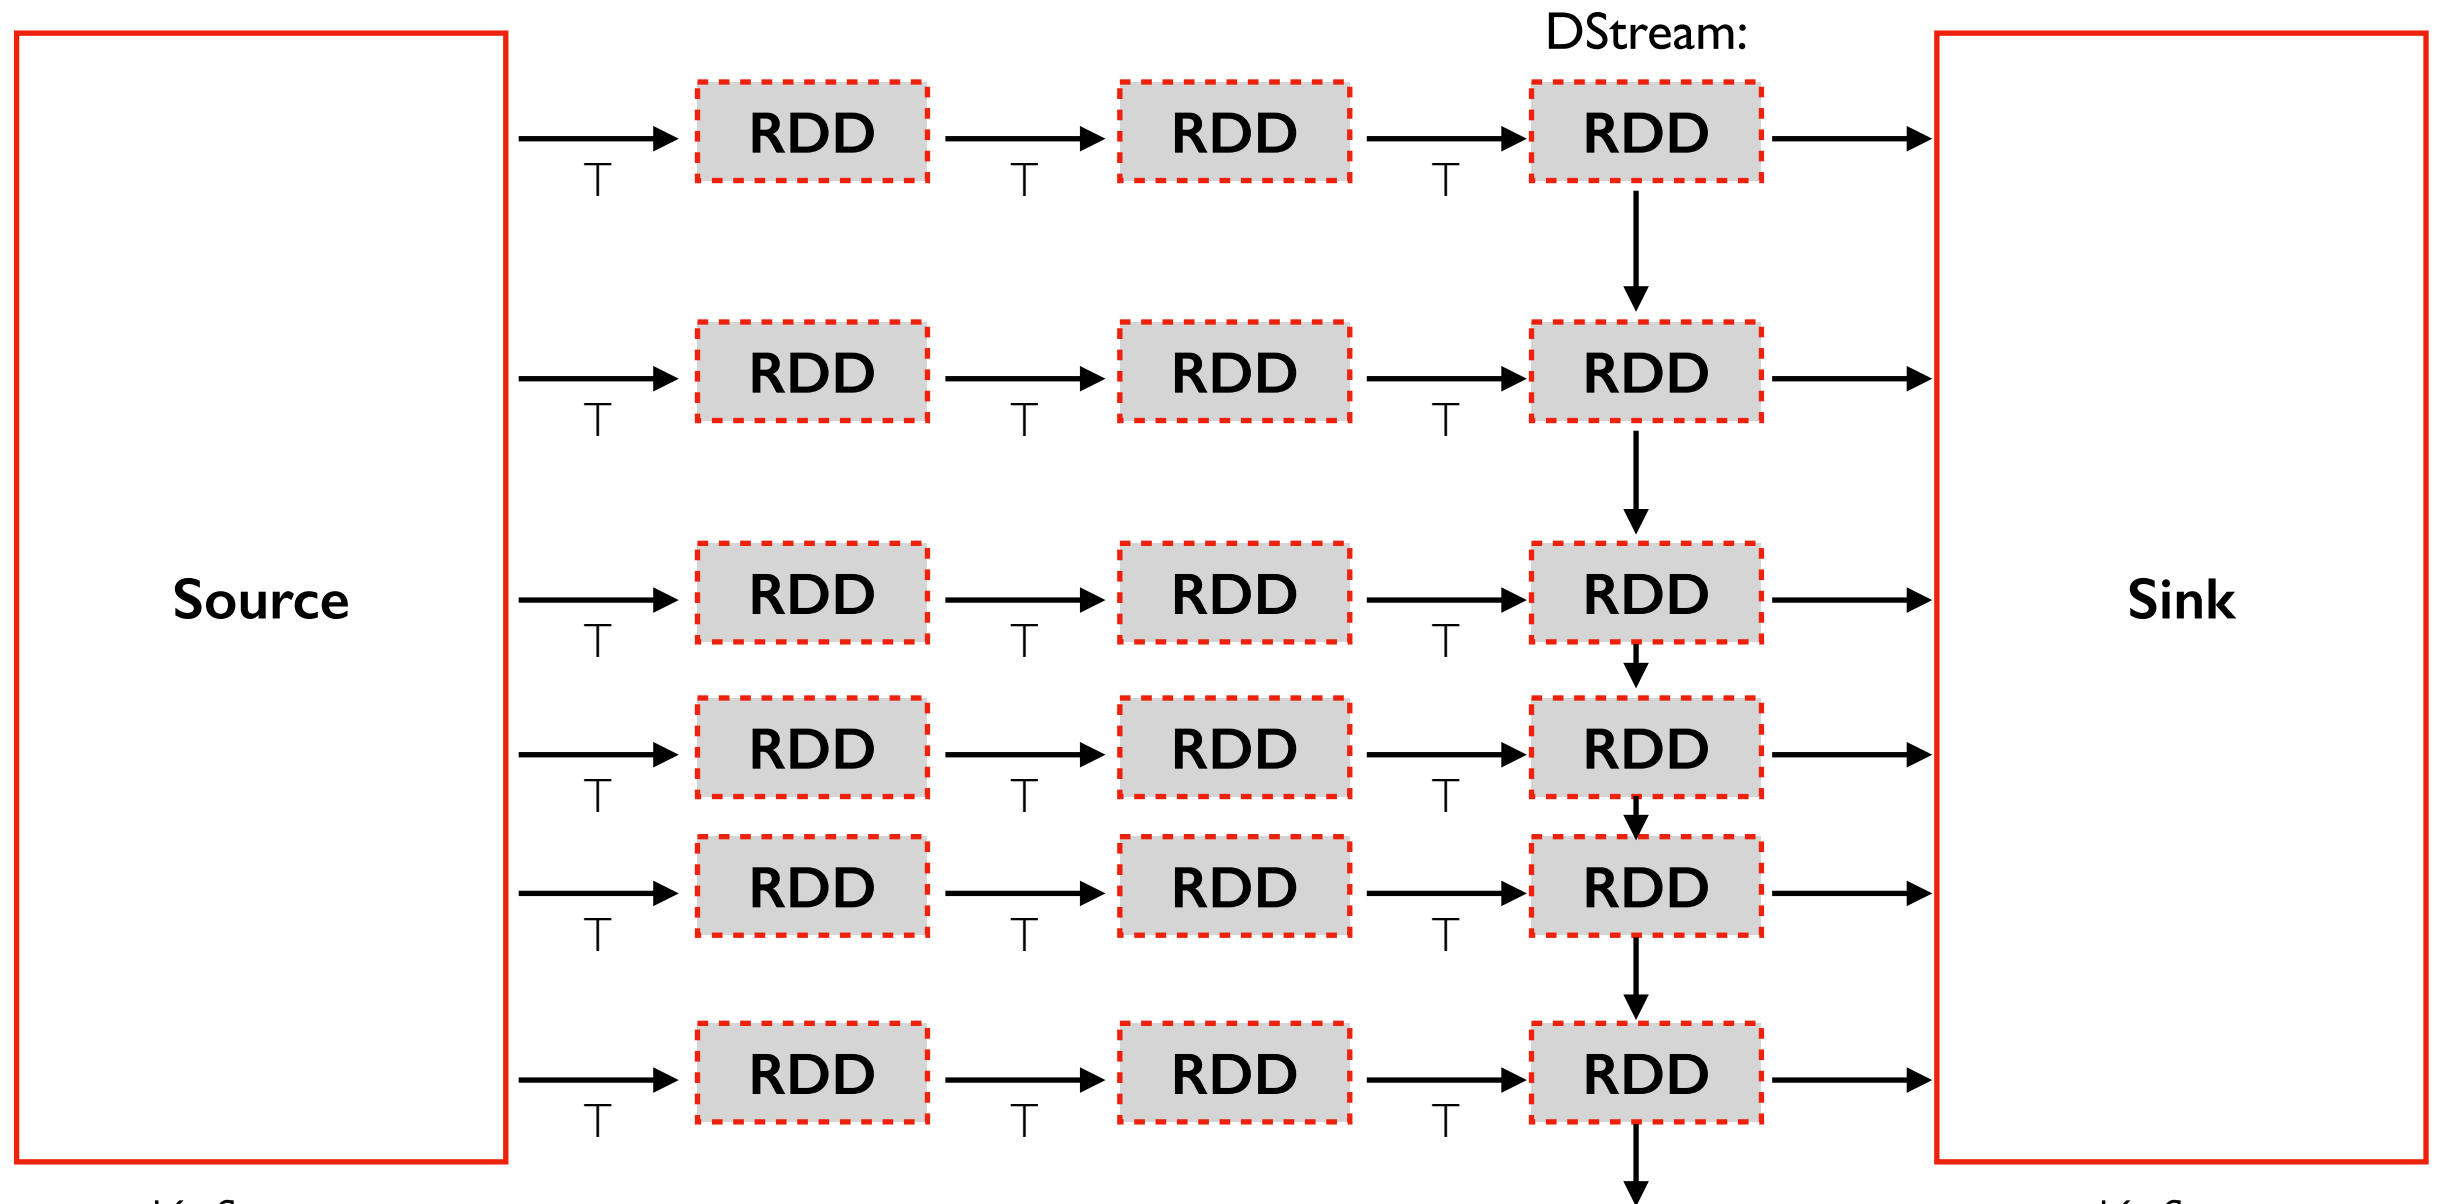
DStream (Stateless)

TopHat



If we can compute on each batch without using state from previous computations, it is **stateless**.

Source => DStream => Sink




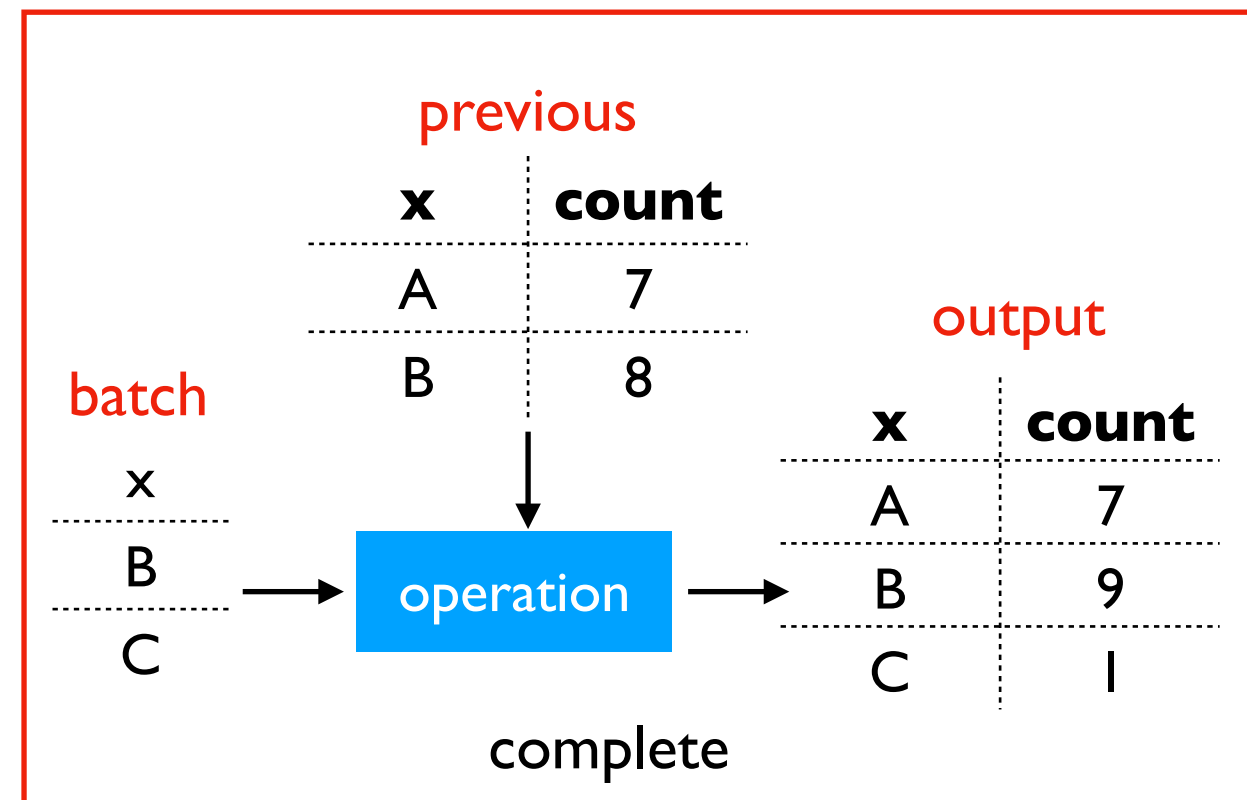
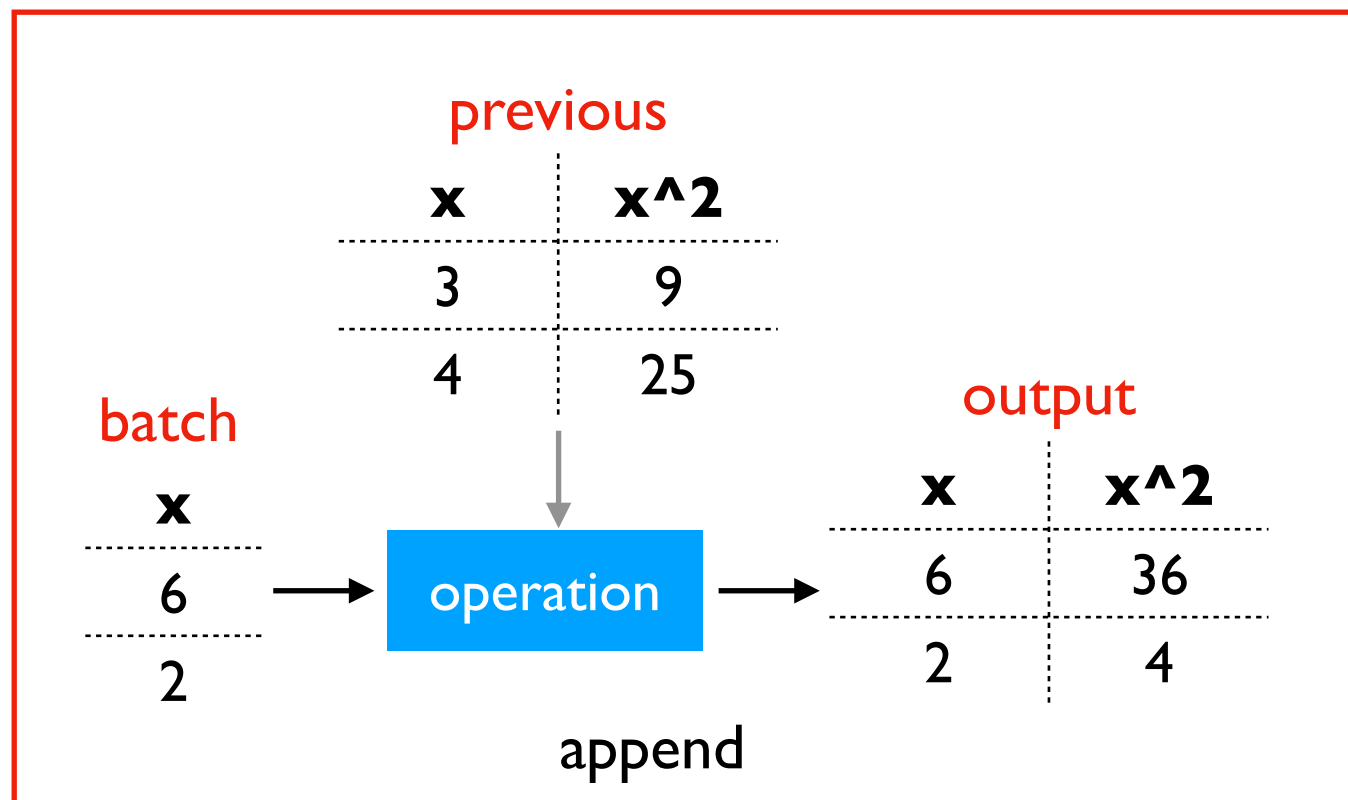
- Kafka
- HDFS files
- Cassandra
- etc.

A DStream continuously pulls data from a source, transforms it, and sends it to a sink

many possible source/sink formats

- Kafka
- HDFS files
- console
- etc.

Output Modes: Update, Complete, Append

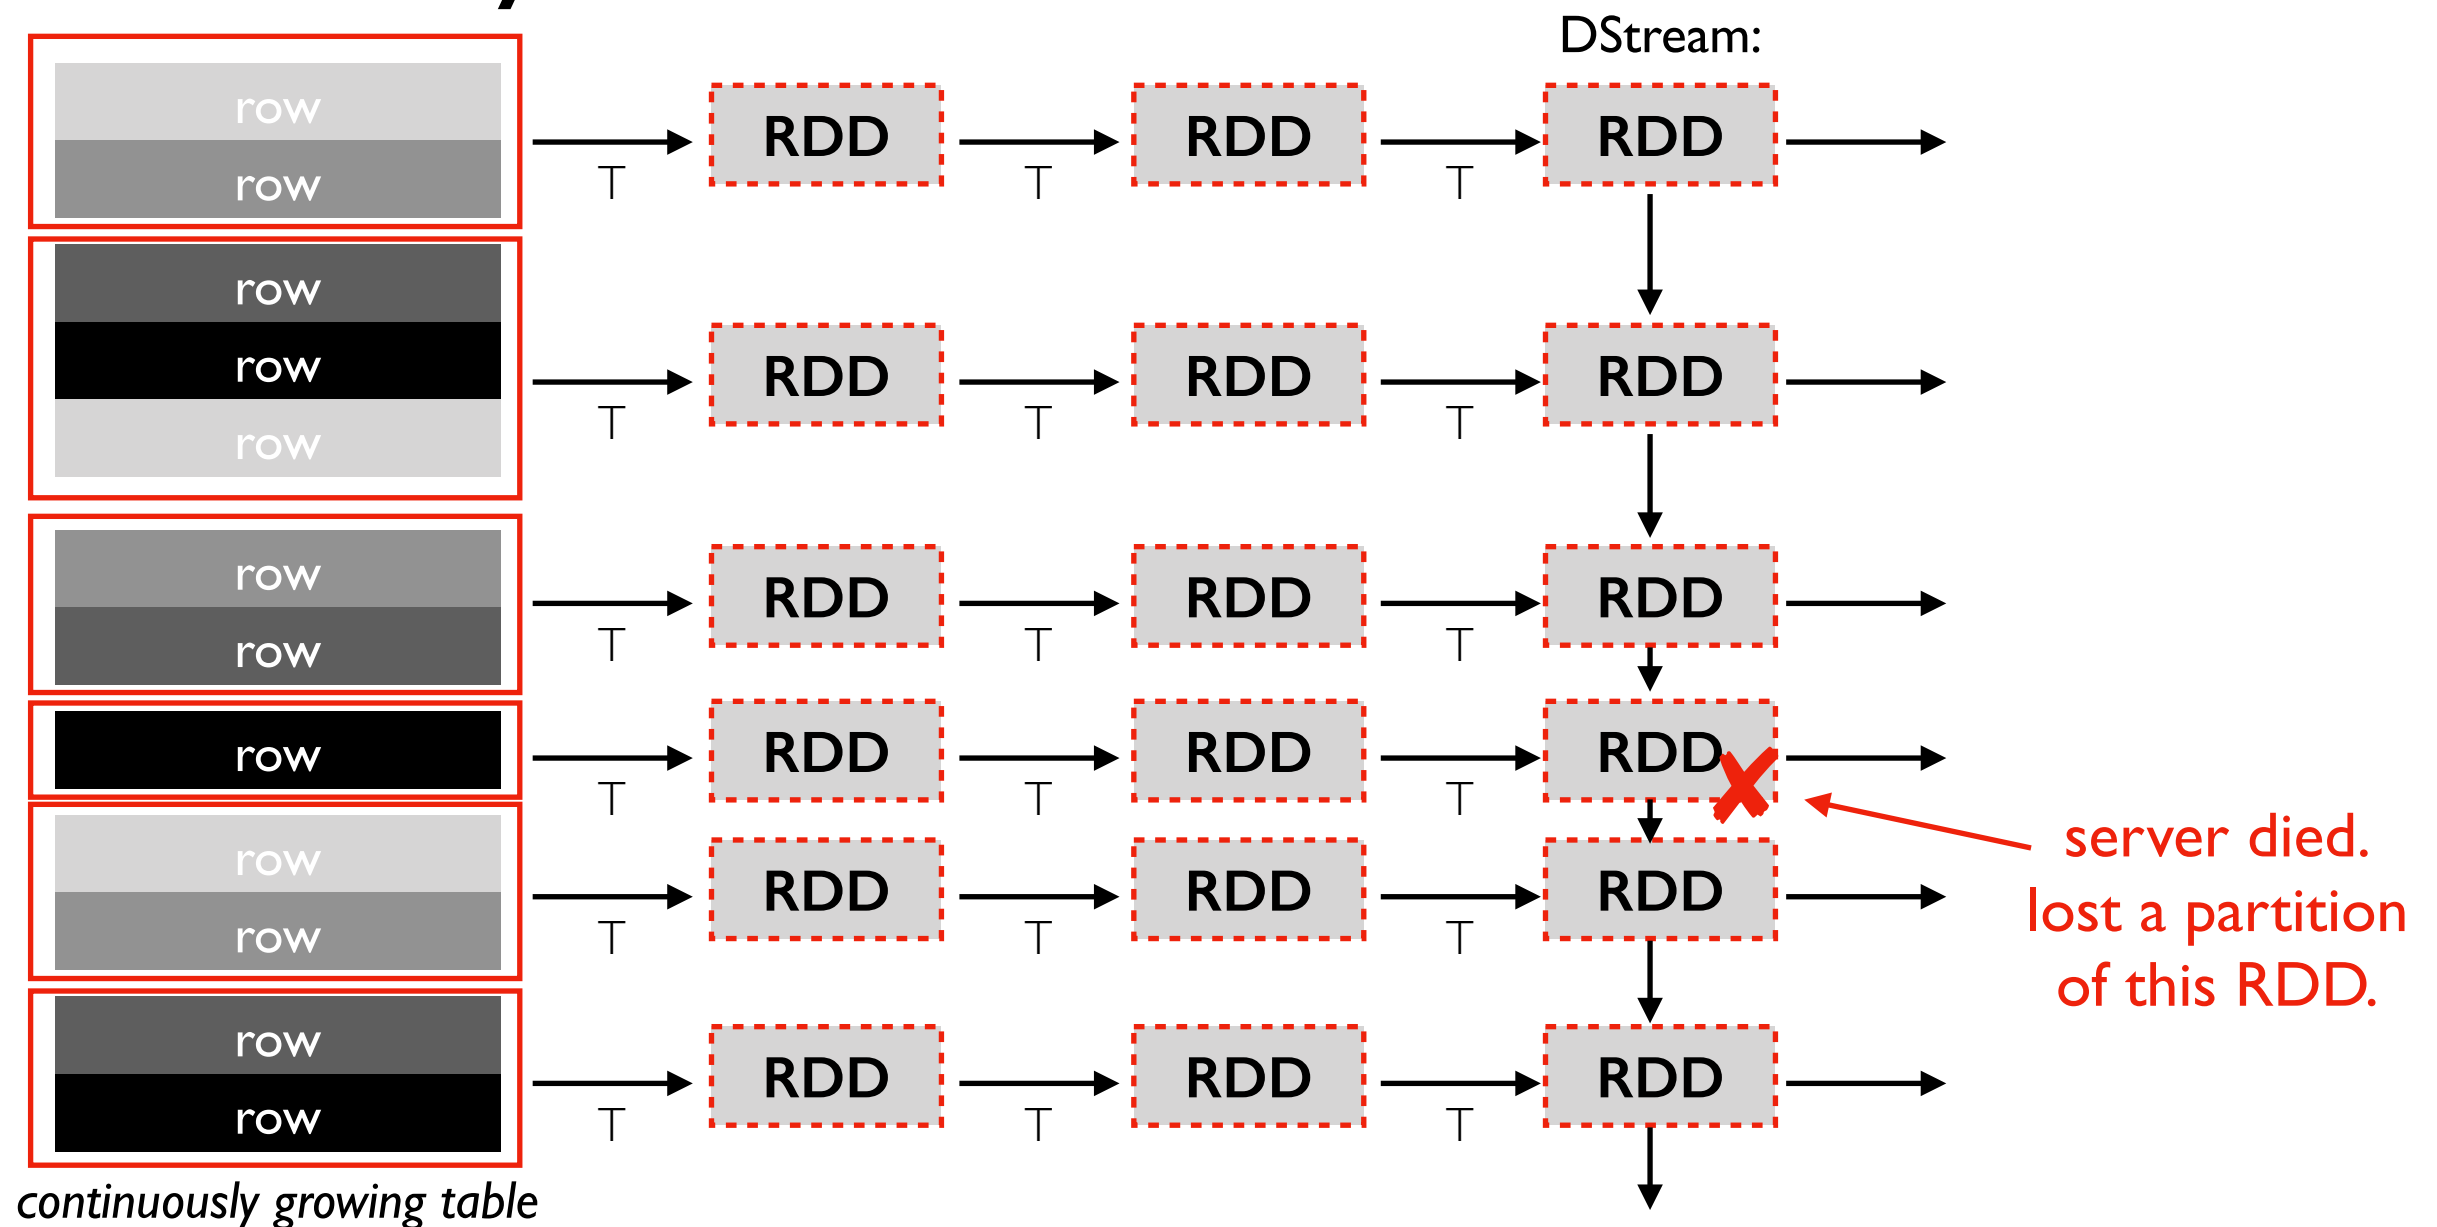


Different modes are available depending on transformation and output format.

Examples:

- **update**: output is usually a DB
- **append**: generally narrow transformations (previous output rows cannot change)
- **complete**: often for aggregates (otherwise too expensive so not allowed)

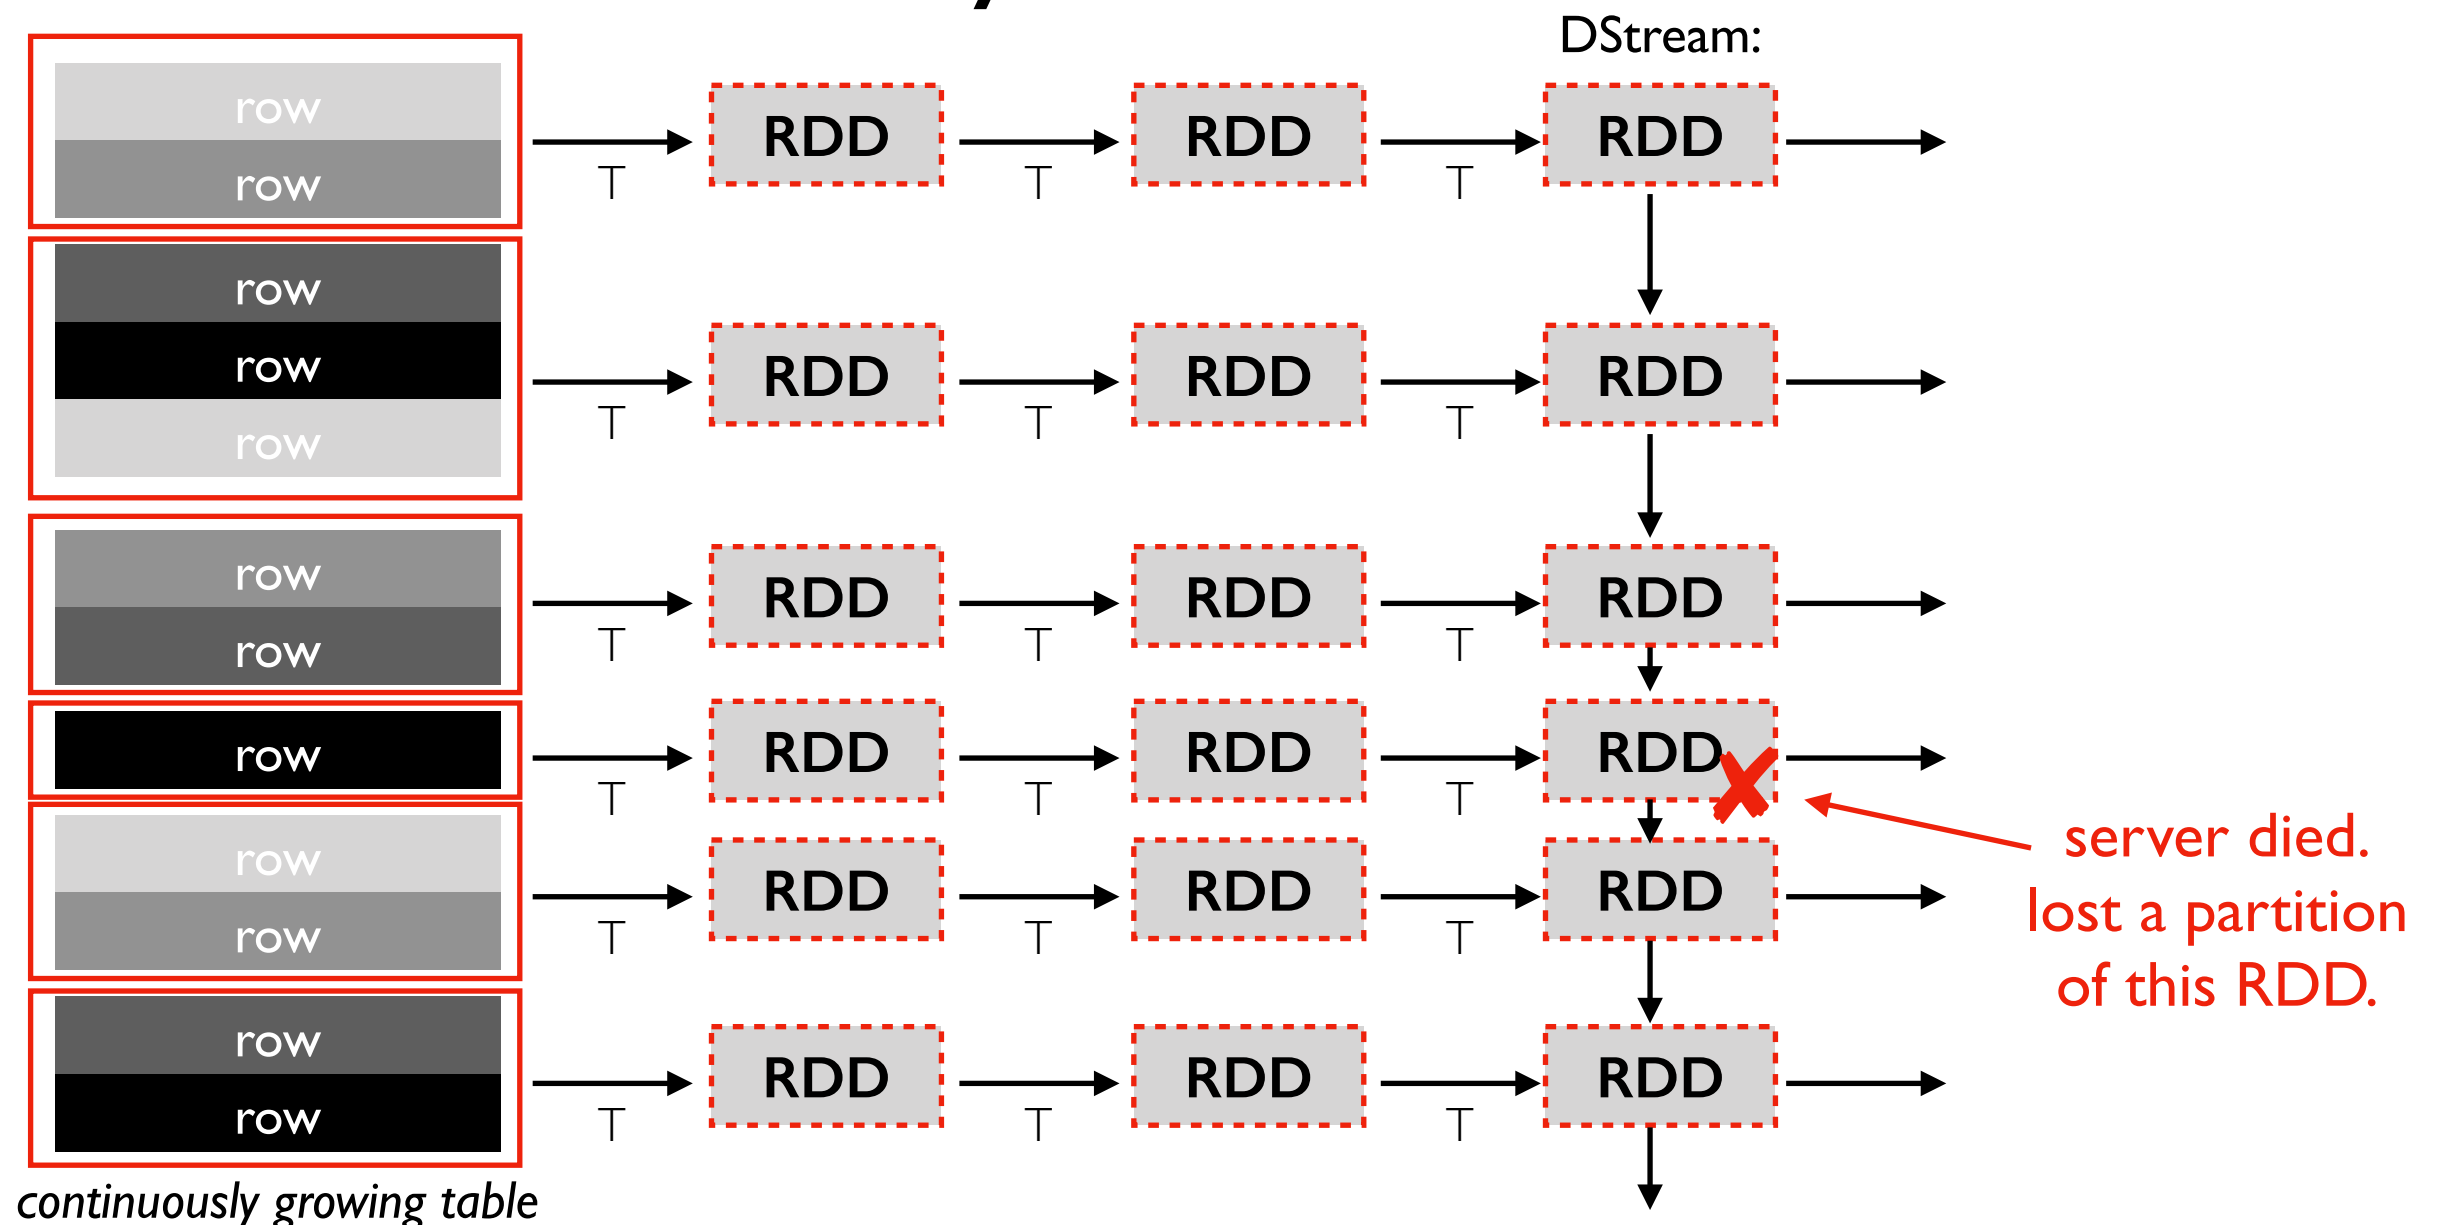
Recovery



Recovery:

- Spark usually doesn't replicate data because RDDs tell us how to recompute lost data
- What if source data is no longer available? (e.g., beyond Kafka retention time)
- What if it takes too long to recover?

Efficient Recovery



Recovery:

- Spark usually doesn't replicate data because RDDs tell us how to recompute lost data
- What if source data is no longer available? (e.g., beyond Kafka retention time)
- What if it takes too long to recover?

Spark Optimizations:

- Often, every worker can help with recovery work (i.e., recomputing data for an RDD)
- Checkpoint DStream once every 10 batches.

Outline: Spark Streaming

DStreams

Grouped Aggregates

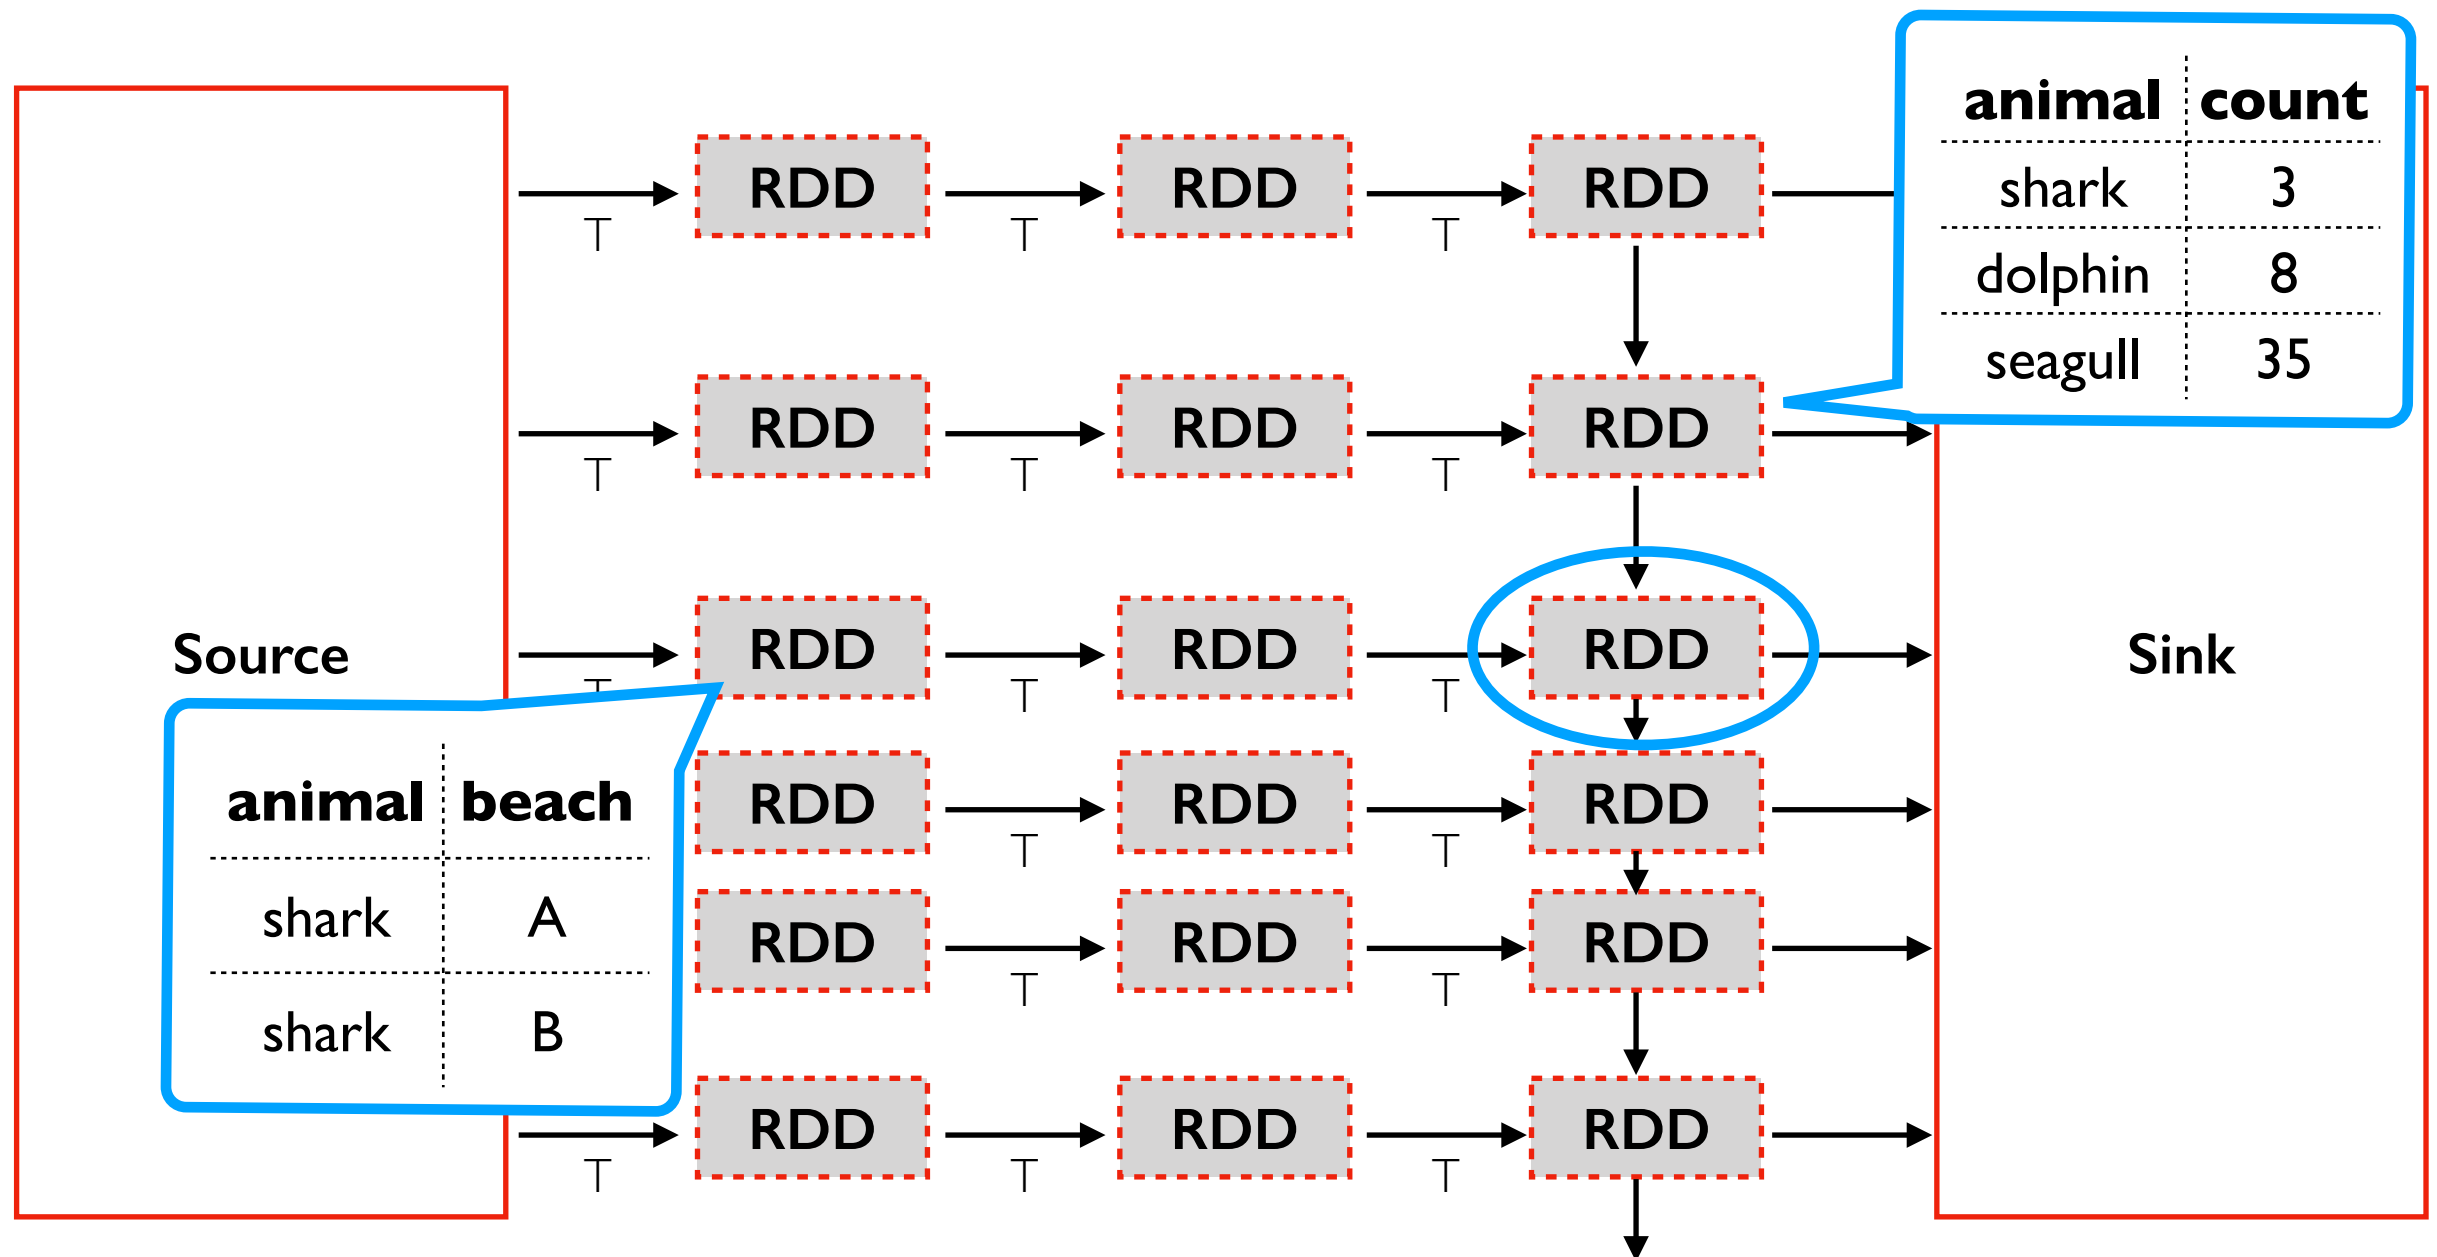
Watermarks

Pivoting

Joining

Exactly-Once Semantics


Incremental Aggregations



```
SELECT animal, COUNT(*)  
FROM sightings  
GROUP BY animal
```


- many aggregations are easy to compute incrementally
- mode: update or complete (append usually not valid because previous rows change)
- space for state proportional to unique categories

Grouped Aggregate Internals: Shuffle Partitions



How many partitions will we have?

- `spark.sql.shuffle.partitions` (default 200) sets this -- fixed for whole application
- Often need to reduce for streaming jobs
- Batch jobs can automatically coalesce small partitions into bigger ones?
- Why not optimized for streaming? One challenge: coalescing based on data so far probably isn't good for future data. Avoid re-shuffling existing counts.



see Epilogue:
Apache Spark 3.0

Outline: Spark Streaming

DStreams

Grouped Aggregates

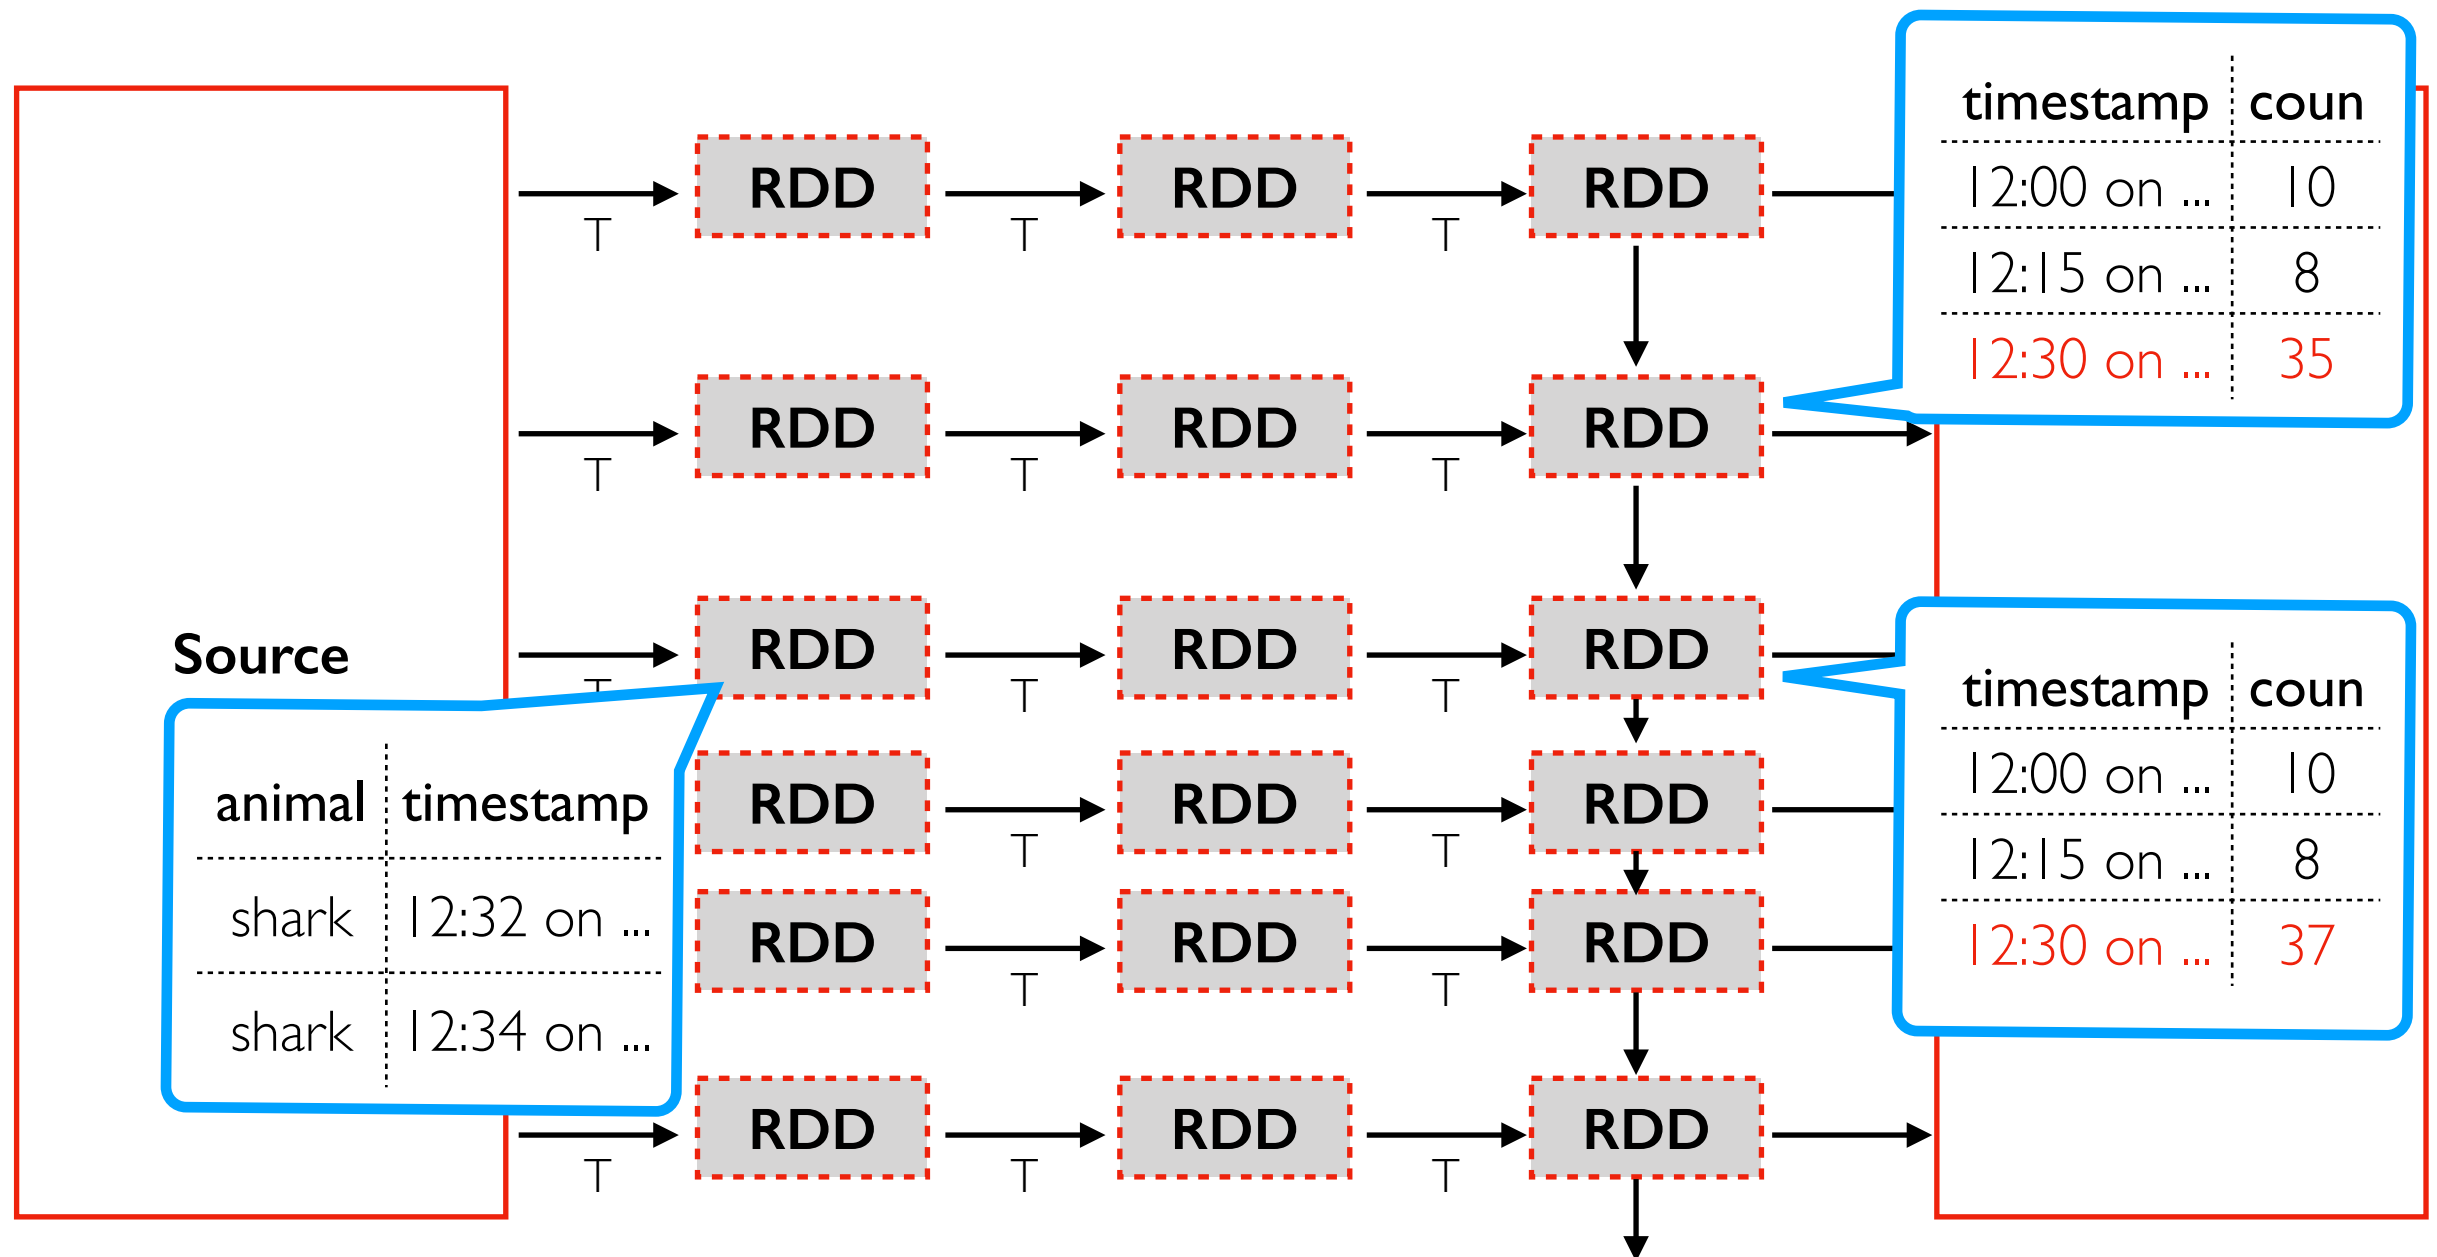
Watermarks

Pivoting

Joining

Exactly-Once Semantics

Grouping By Time Intervals



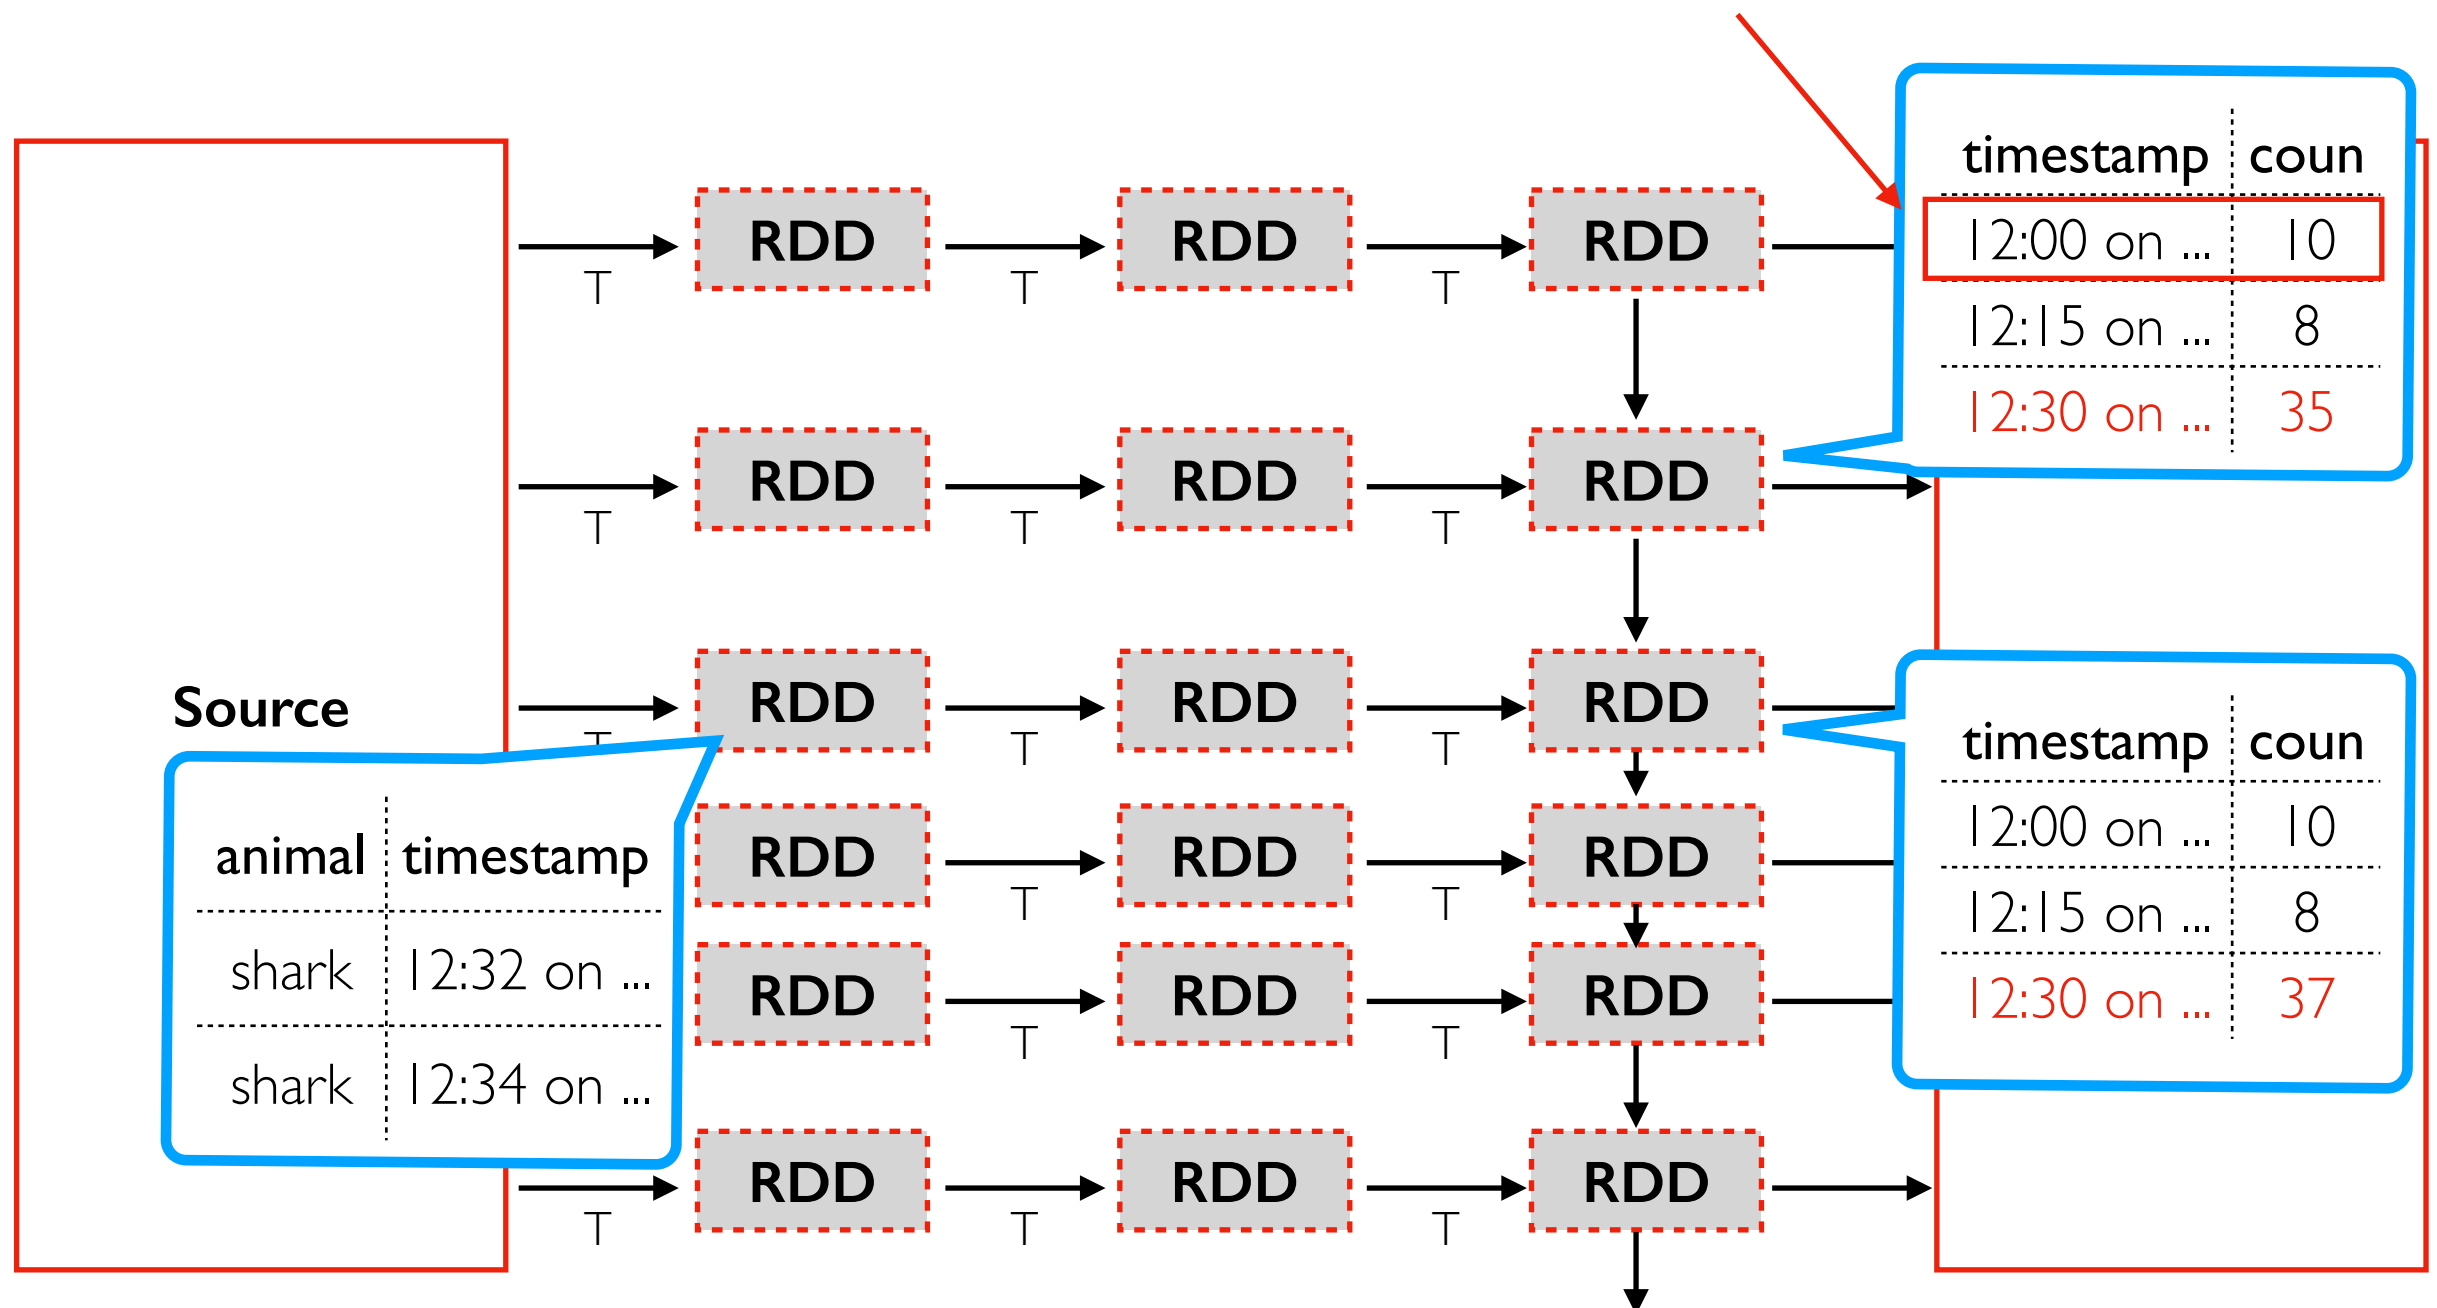
Observations:

- number of groups (and RAM needed) grows indefinitely with time
- new batches contain recent times
- old times might occasionally pop up (Kafka delays)

```
(animals
  .groupBy(window("timestamp",
                  "15 minute"))
  .count())
```

Watermarks

Spark can discard this running count after 8:15pm because it is unlikely the pipeline will fall 8 hours behind



```
(animals
  .withWatermark("timestamp",
    "8 hours")
  .groupBy(window("timestamp",
    "15 minute"))
  .count())
```

Behavior:

- never throw away rows/aggregates newer than watermark time
- might throw away older data to save space

TopHat

Outline: Spark Streaming

DStreams

Grouped Aggregates

Watermarks


Pivoting

Joining

Exactly-Once Semantics

Pivots

beach	animal
A	seagull
B	seagull
B	dolphin
C	seagull
A	seagull
A	dolphin
B	dolphin




beach	seagull	dolphin
A	2	1
B	1	2
C	1	0

what if we add a row with previously unseen values?

Pivots

beach	animal
A	seagull
B	seagull
B	dolphin
C	seagull
A	seagull
A	dolphin
B	dolphin
D	shark



beach	seagull	dolphin	shark
A	2	1	0
B	1	2	0
C	1	0	0
D	0	0	1

- **new row:** OK for batching and streaming
- **new col:** only OK for batching
- with streaming, it would cause confusion if columns were added mid query (how would somebody even query from our results?)
- some operations like pivot are supported for batching but not streaming

Outline: Spark Streaming

DStreams

Grouped Aggregates

Watermarks

Pivoting

Joining

Exactly-Once Semantics

JOIN Scenarios

static-static
(previously covered)




fixed size



fixed size

stream-static



growing



fixed size

stream-stream



growing



growing

static-static review:

- shuffle sort merge join
- broadcast hash join

- Spark has at least some support for each scenario
- stream-stream can use an ever increasing amount of memory if we're not careful (need watermarking)

JOIN Scenarios

static-static
(previously covered)




fixed size



fixed size

stream-static




growing




fixed size

stream-stream



growing



growing

static-static review:

- shuffle sort merge join
- broadcast hash join

- Spark has at least some support for each scenario
- stream-stream can use an ever increasing amount of memory if we're not careful (need watermarking)

Stream-Static INNER JOIN

animals

id	name
1	dolphin
2	shark
3	seagull

fixed

what *known* animals do we see?

```
SELECT beach, name
FROM sightings
INNER JOIN animals
ON sightings.animal_id=animals.id
```

sightings

beach	animal_id
A	3
B	3
A	2
C	4

growing



results

beach	name
A	seagull
B	seagull
A	shark

growing



is the JOIN stateless?

Stream-Static LEFT JOIN

animals

id	name
1	dolphin
2	shark
3	seagull

fixed

are there any sightings of unknown animals?

```
SELECT beach, animal_id
FROM sightings
LEFT JOIN animals
ON sightings.animal_id=animals.id
WHERE name IS NULL
```

sightings

beach	animal_id
A	3
B	3
A	2
C	4

growing



results

beach	name
C	4

growing



Stream-Static RIGHT JOIN

animals

id	name
1	dolphin
2	shark
3	seagull

fixed

sightings

beach	animal_id
A	3
B	3
A	2
C	4

growing



are there any animals that are never seen?

```
SELECT name, beach
FROM sightings
RIGHT JOIN animals
ON sightings.animal_id=animals.id
WHERE beach IS NULL
```

results

name	beach
dolphin	NULL

fixed

why is it impossible to compute the results, even though it would be easy for static-static?

Cannot RIGHT JOIN if right is static; Cannot LEFT JOIN if left is static

animals

id	name
1	dolphin
2	shark
3	seagull

fixed

are there any animals that are never seen?

```
SELECT name, beach
FROM sightings
RIGHT JOIN animals
ON sightings.animal_id=animals.id
WHERE beach IS NULL
```

sightings

beach	animal_id
A	3
B	3
A	2
C	4

growing



results

name	beach
dolphin	NULL

fixed

we can never say an animal is never seen if we keep seeing animals forever, so this query is illogical (and unsupported by Spark)

JOIN Scenarios

when possible, cache this.
It JOINS against every micro batch.
Don't want to re-read every time!

static-static
(previously covered)




fixed size



fixed size

stream-static



growing




fixed size

stream-stream



growing



growing


static-static review:

- shuffle sort merge join
- broadcast hash join


- Spark has at least some support for each scenario
- stream-stream can use an ever increasing amount of memory if we're not careful (need watermarking)

JOIN Scenarios


static-static
(previously covered)



stream-static



stream-stream



static-static review:

- shuffle sort merge join
- broadcast hash join

- Spark has at least some support for each scenario
- stream-stream can use an ever increasing amount of memory if we're not careful (need watermarking)

Stream-Stream

closures

date	type
4/10/23	"all day"
4/15/23	"part day"
4/20/23	"all day"

growing



sightings

date	animal
4/13/23	seagull
4/14/23	seagull
4/14/23	shark
4/15/23	dolphin

growing



how many sharks are seen on
days when the beach is closed?

```
SELECT COUNT(*)  
FROM sightings  
INNER JOIN closures  
ON sightings.date=closures.date  
WHERE animal = 'shark'
```

challenge: we can't "forget" about this row if we might later learn about a beach closure on the 14th (for example, from a lagging Kafka stream)

solution: use watermarks (like for grouped aggregates)

note: Spark works without watermarks; it just keeps using more memory indefinitely

Outline: Spark Streaming

DStreams

Grouped Aggregates

Watermarks

Pivoting

Joining

Exactly-Once Semantics


Exactly-Once Semantics

If a task crashes, we can restart a new one, but we don't want to:

- double count any row
- miss any row

Spark can achieve exactly-once semantics given 3 features

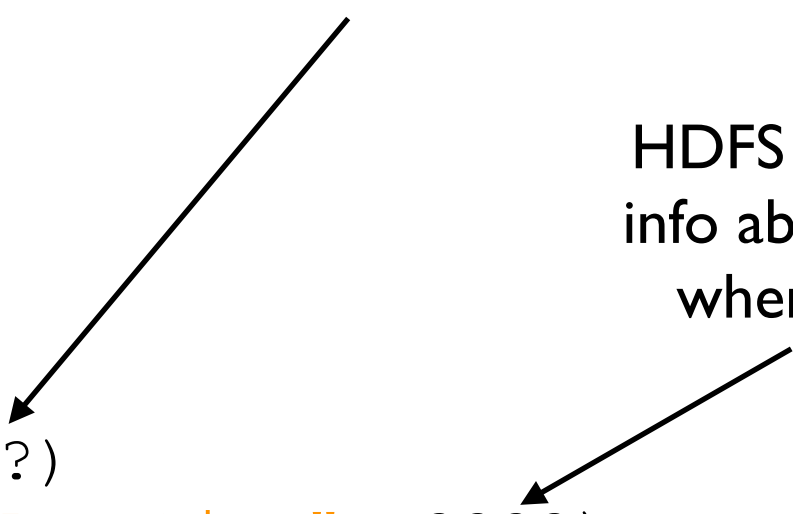
- your code is "deterministic" (does same thing each time given same inputs)
- **source:** it's possible to go back and re-read older inputs that the previous task was processing when it crashed (Kafka makes this easy, within the retention period)
- **sink:** it is "idempotent" (can suppress duplicates)



file sink (parquet files on HDFS) supports this --
Spark writes checkpoint files that identify which
output files correspond to which input messages

Parquet on HDFS

```
query = (df
  .writeStream
  .format("parquet")
  .option("path", ????)
  .option("checkpointLocation", ????)
  .start())
```



The diagram consists of two arrows. The first arrow originates from the text 'HDFS directory that will accumulate parquet files' and points to the 'path' option in the code. The second arrow originates from the text 'HDFS directory where Spark stores info about how to suppress duplicates when reading those parquet files' and points to the 'checkpointLocation' option in the code.

HDFS directory that will accumulate parquet files

HDFS directory where Spark stores info about how to suppress duplicates when reading those parquet files

When Spark reads a directory of parquet files, it automatically suppresses duplicates. But be careful reading individual parquet files in a directory yourself, because then you might see those duplicates.

Conclusion

Spark streaming is frequent batch computing

- DStream is series of RDDs
- Most things we can do with regular DataFrames can be done with streams
- Not quite realtime, but fast crash recovery

Performance

- choose shuffle partition count carefully
- apply watermarks to limit memory consumption
- in stream-static JOIN, try to cache the static table