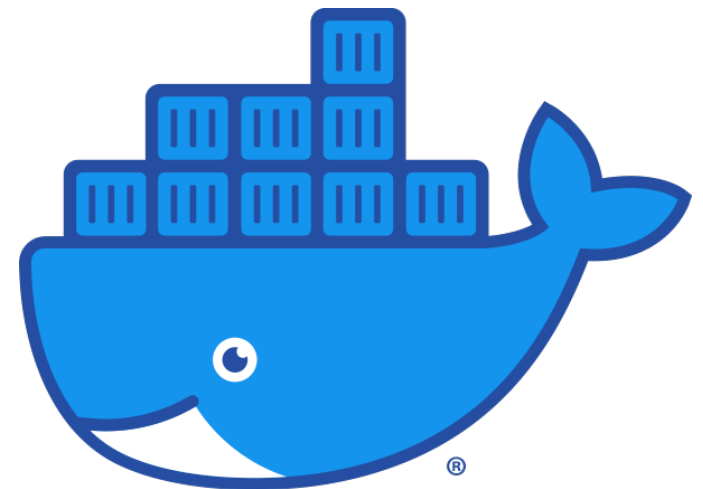


[544] Docker Deployment

Tyler Caraza-Harter



Outline

Virtualization

Images, Containers, and Dockerfiles

Ports

Docker Compose

What is virtualization?

Definition: the illusion of **private** resources, provided by software

Contexts this semester

- Virtual Memory
- Virtual Machines (hardware)
- Virtual Machines (languages)
- Virtual Operating System (container) *new today*


What is virtualization?

Definition: the illusion of **private** resources, provided by software

Contexts this semester

- Virtual Memory
- Virtual Machines (hardware)
- Virtual Machines (languages)
- Virtual Operating System (container)

each process using a virtual address space
isn't aware of other processes using memory
(illusion of private memory)



What is virtualization?

Definition: the illusion of **private** resources, provided by software

Contexts this semester

- Virtual Memory
- Virtual Machines (hardware)
- Virtual Machines (languages)
- Virtual Operating System (container)

virtualized resources include CPU, RAM, disks, network devices, etc

VMs rarely use all their allocated resources, so overbooking is possible

VM: 8 GB of RAM
and 4 cores

VM: 6 GB of RAM
and 3 cores

VM: 8 GB of RAM
and 6 cores

virtual machines
for rent (by you)

Physical Machine: 16 GB of RAM and 8 CPU cores

actual hardware bought by Google for their cloud services (GCP)

What is virtualization?

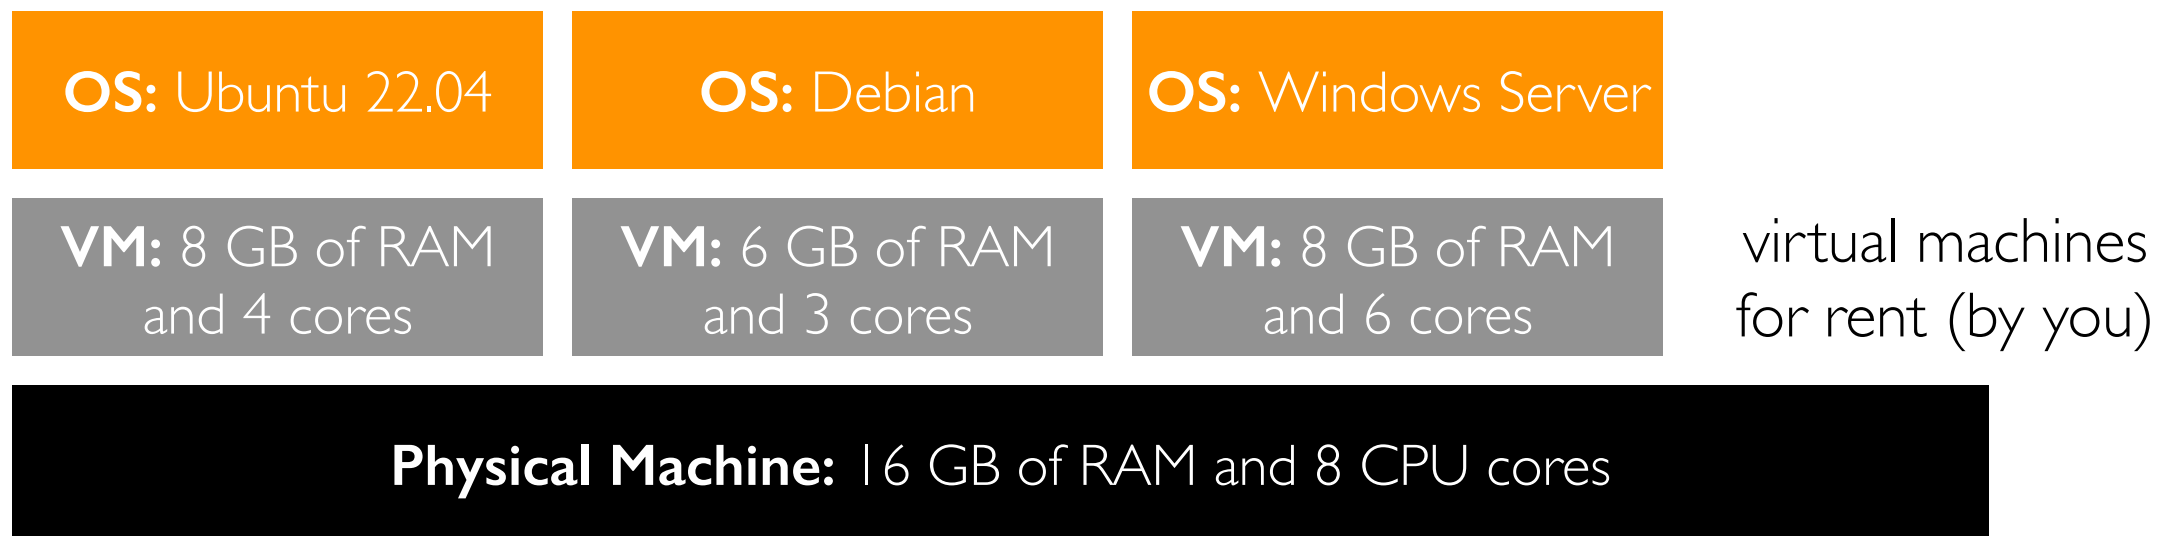
Definition: the illusion of **private** resources, provided by software

Contexts this semester

- Virtual Memory
- Virtual Machines (hardware)
- Virtual Machines (languages)
- Virtual Operating System (container)

problem: if each program is deployed to a different VM, operating system overheads will dominate

these operating systems are mostly unaware that their on VMs instead of physical hardware



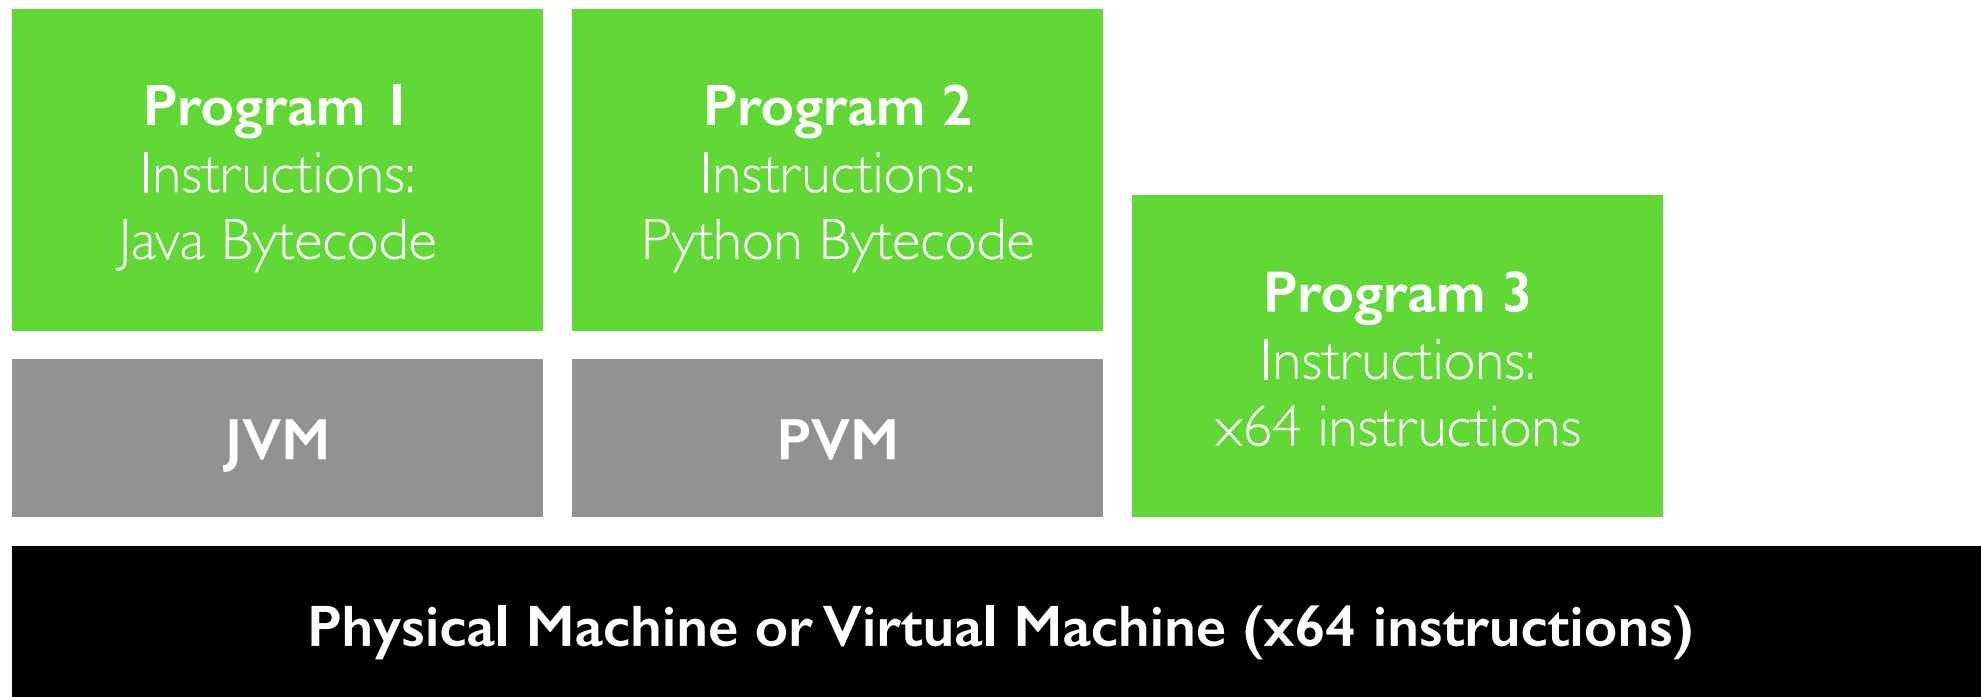
actual hardware bought by Google for their cloud services (GCP)

What is virtualization?

Definition: the illusion of **private** resources, provided by software

Contexts this semester

- Virtual Memory
- Virtual Machines (hardware)
- Virtual Machines (languages)
- Virtual Operating System (container)



actual hardware bought by Google for their cloud services (GCP)

What is virtualization?


Definition: the illusion of **private** resources, provided by software

Contexts this semester

- Virtual Memory
- Virtual Machines (hardware)
- Virtual Machines (languages)
- Virtual Operating System (container)

Linux containers

- Docker makes creation easy
- The "physical" OS is shared, which is very efficient
- Programs in different containers can use different flavors of Linux
- Cannot have a Windows container on Linux



actual hardware bought by Google for their cloud services (GCP)

Containers and Virtual Machines are "Sandboxes"

A Computer

student submission:

```
import subprocess
subprocess.check_output([
    "rm", "-rf", "/"
])
```

Sandbox



process


Sandbox



process

python3.9

Sandbox



process

Sandbox



process

python3.10

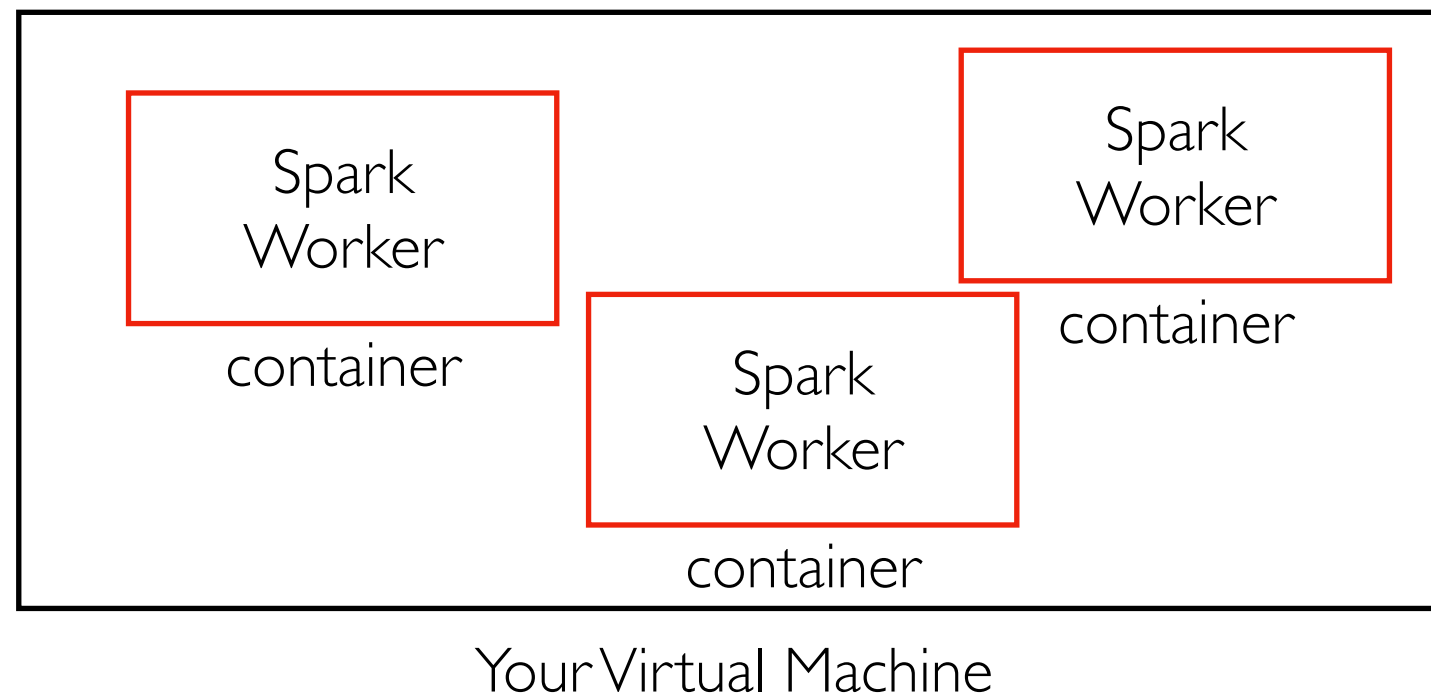


process

Docker containers

Containers are a lightweight alternative to virtual machines.

You'll run Docker containers this semester to have your own "mini cluster"



Resources of the "cluster" are limited to those of a single VM, so we'll scale projects accordingly. But the techniques will apply to large clusters and datasets.

Outline

Virtualization

Images, Containers, and Dockerfiles

Ports

Docker Compose

TIP: make notes of docker commands

docker **SOME-COMMAND** arg1, arg2, ...

Docker Install

For Ubuntu 22.04:

```
sudo snap install docker
```

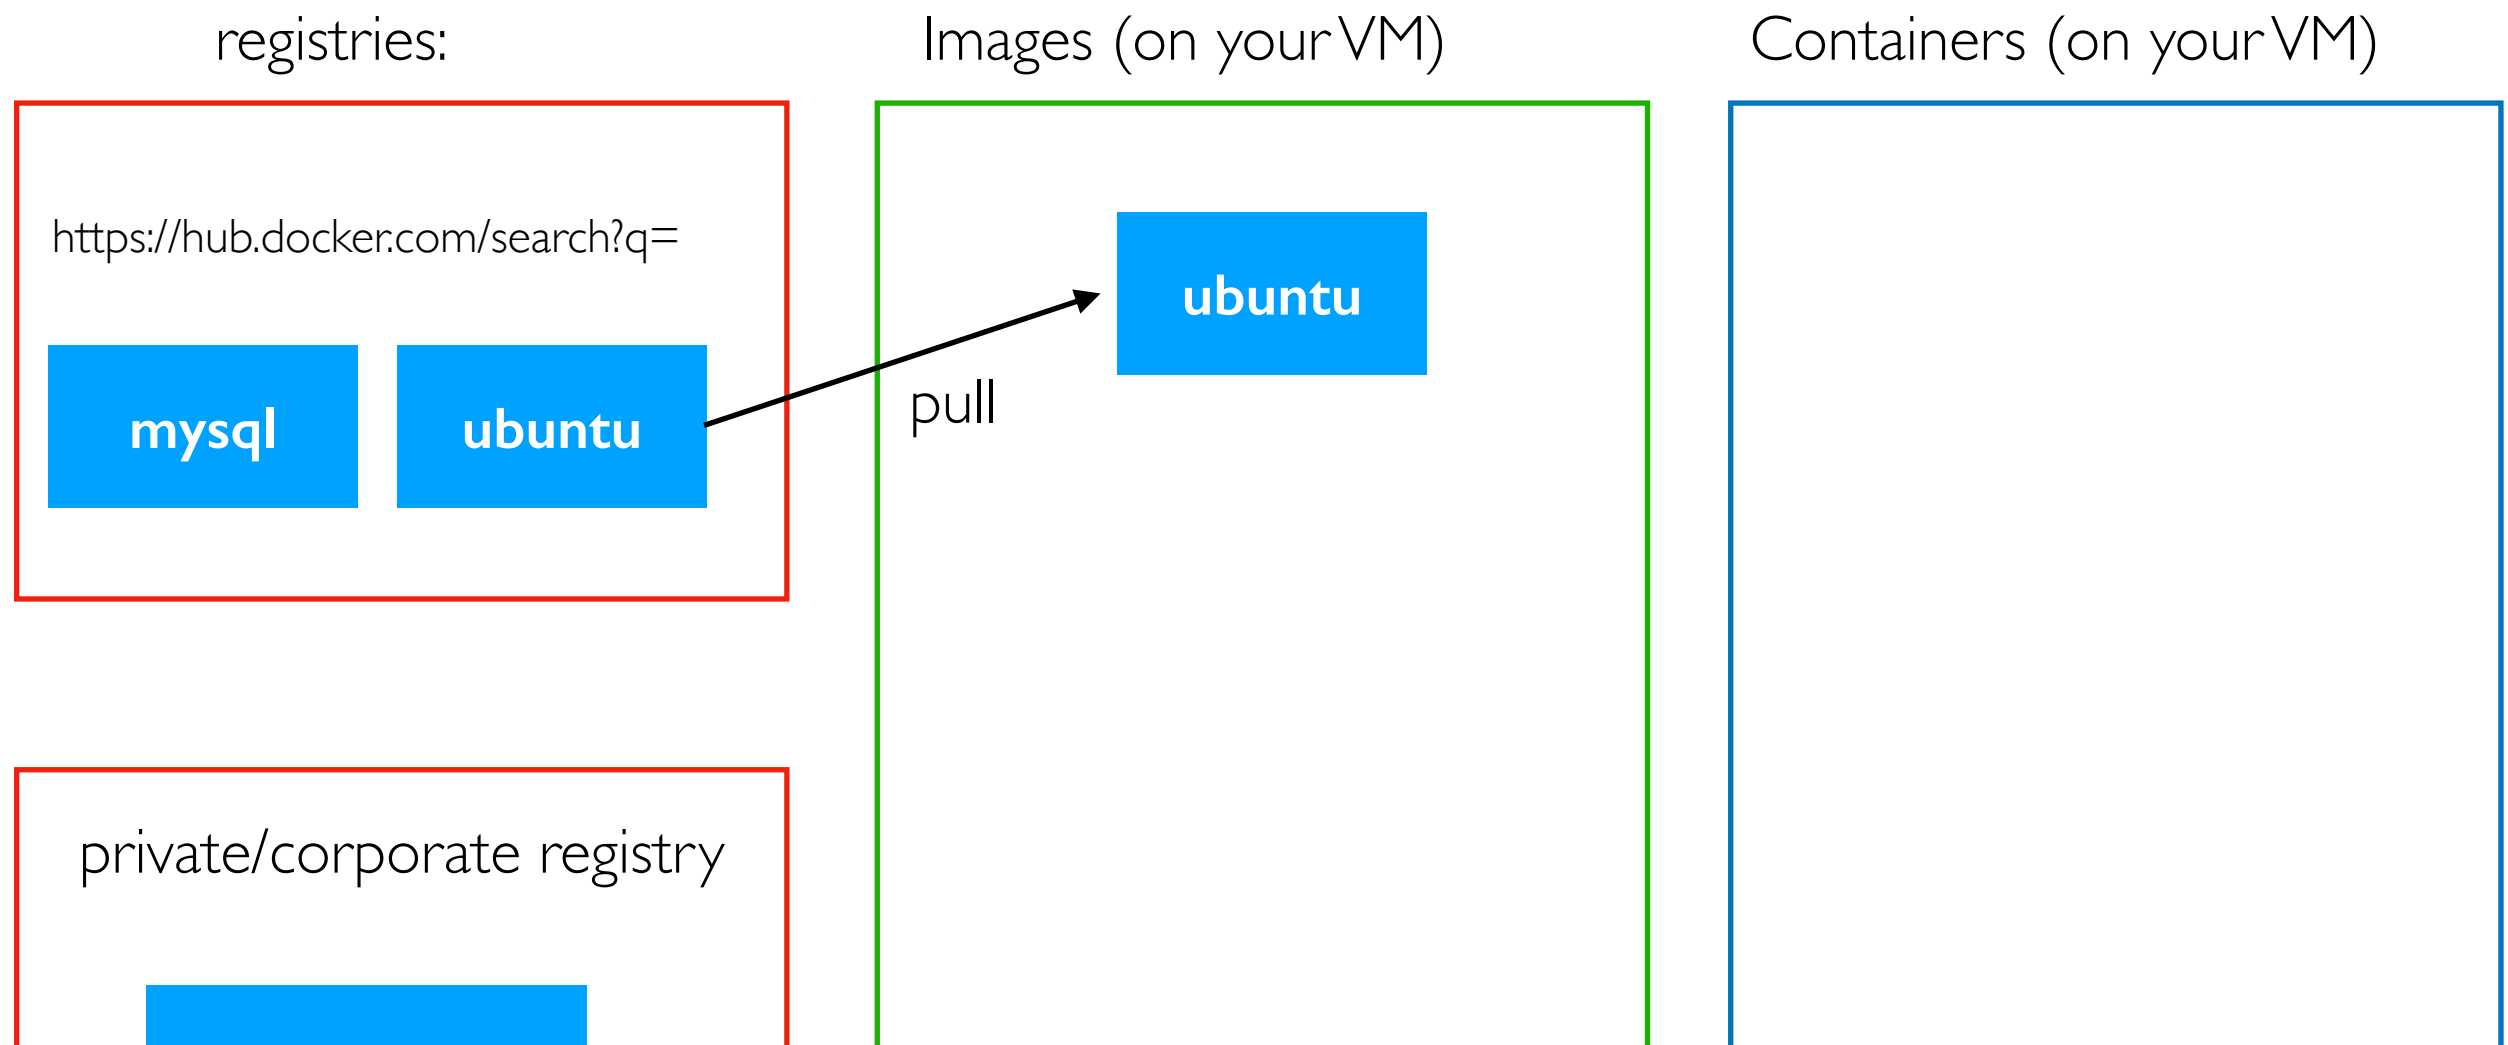
```
# allow regular user to use docker
```

```
sudo groupadd docker
```

```
sudo usermod -aG docker $USER
```

```
sudo chgrp docker /var/run/docker.sock
```

Registries, Images, Containers, and Dockerfiles

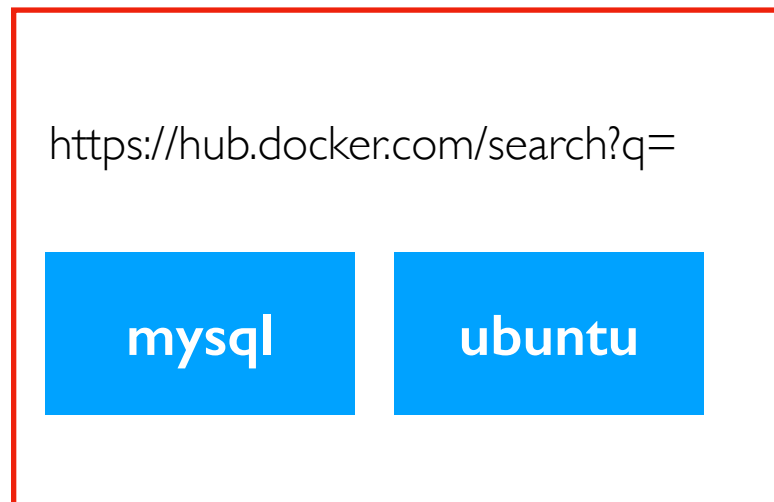


Images

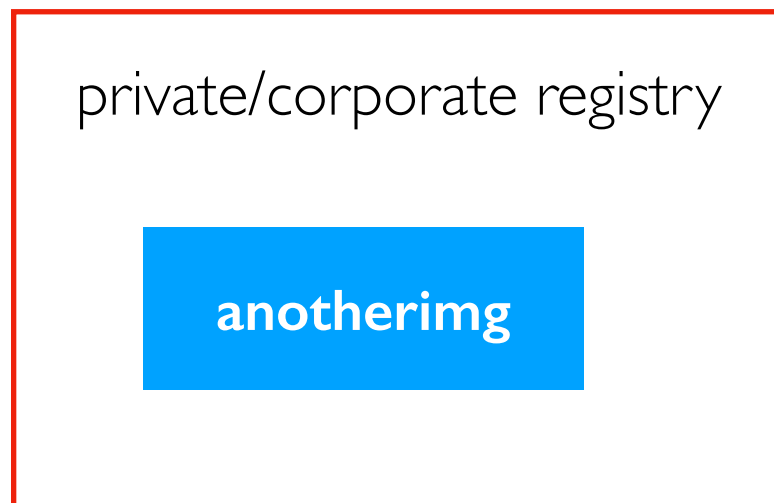
- snapshot of installed software from which to create a container
- `docker pull ubuntu:22.04`

Registries, Images, Containers, and Dockerfiles

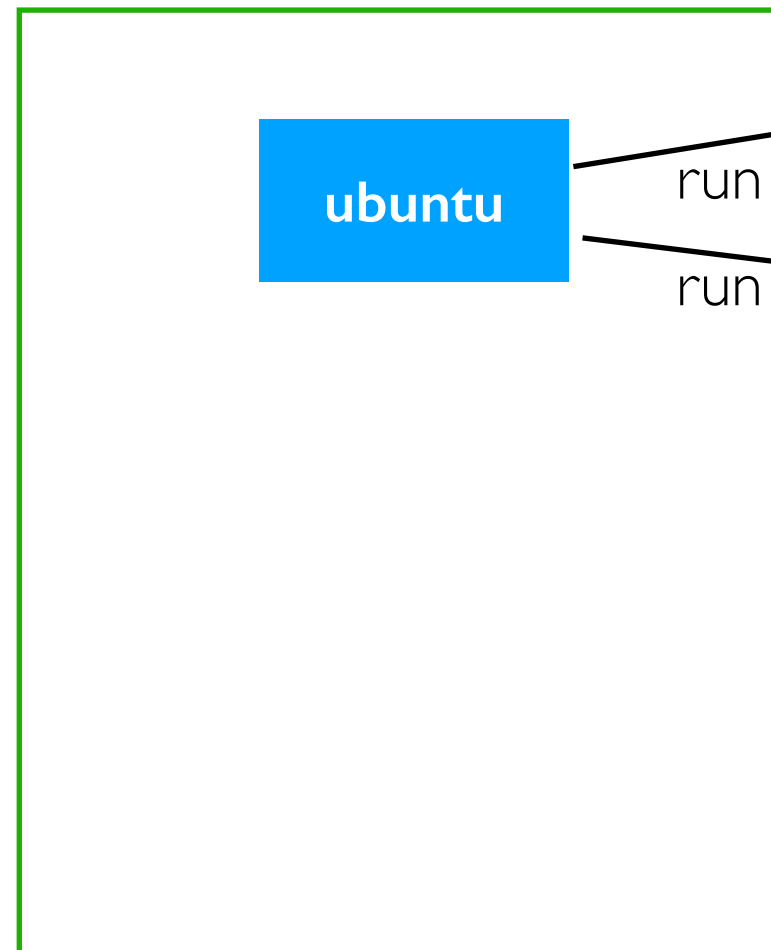
registries:



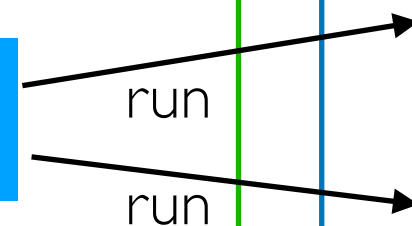
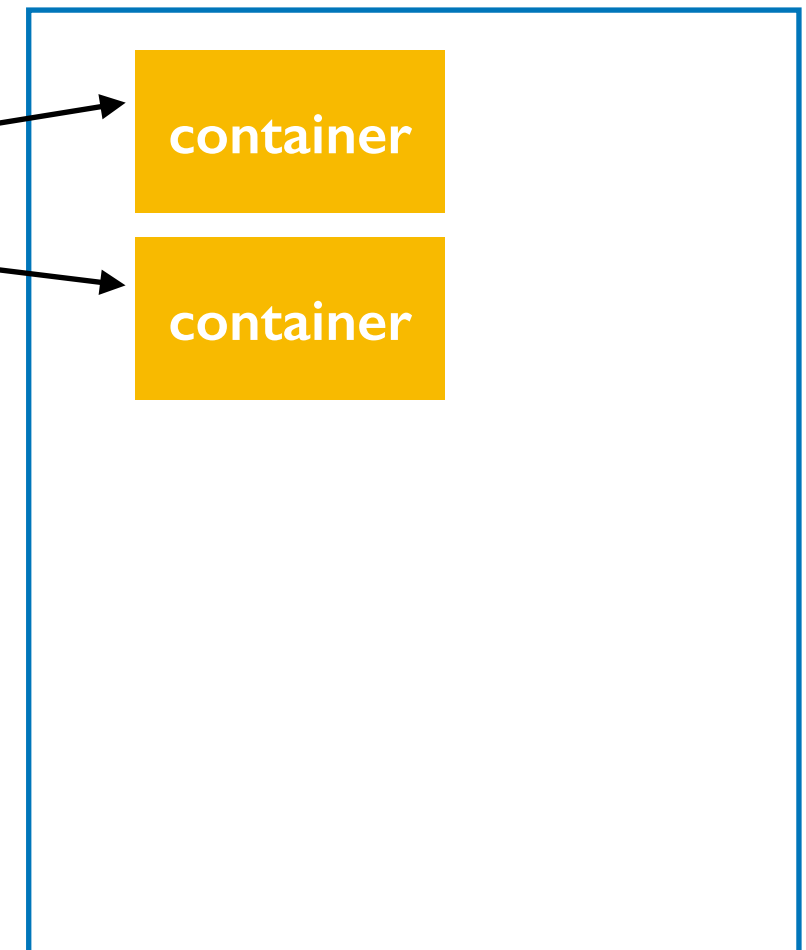
private/corporate registry



Images (on your VM)



Containers (on your VM)

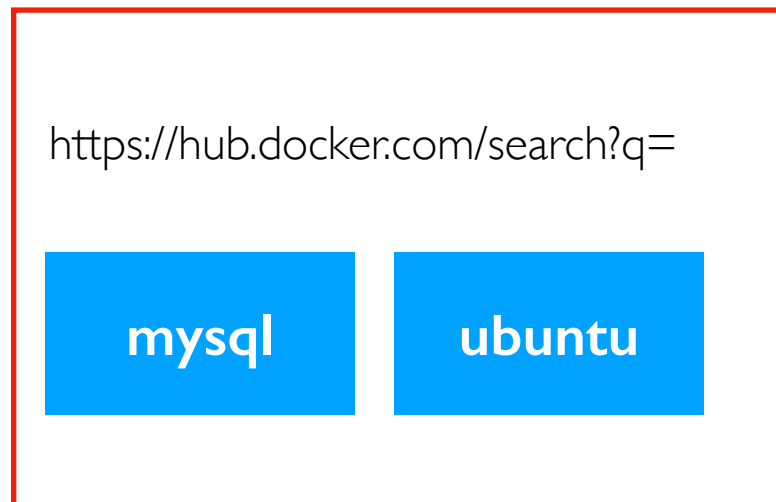


Containers


- Linux sandbox in which to run processes
- docker **run** ubuntu

Registries, Images, Containers, and Dockerfiles


registries:



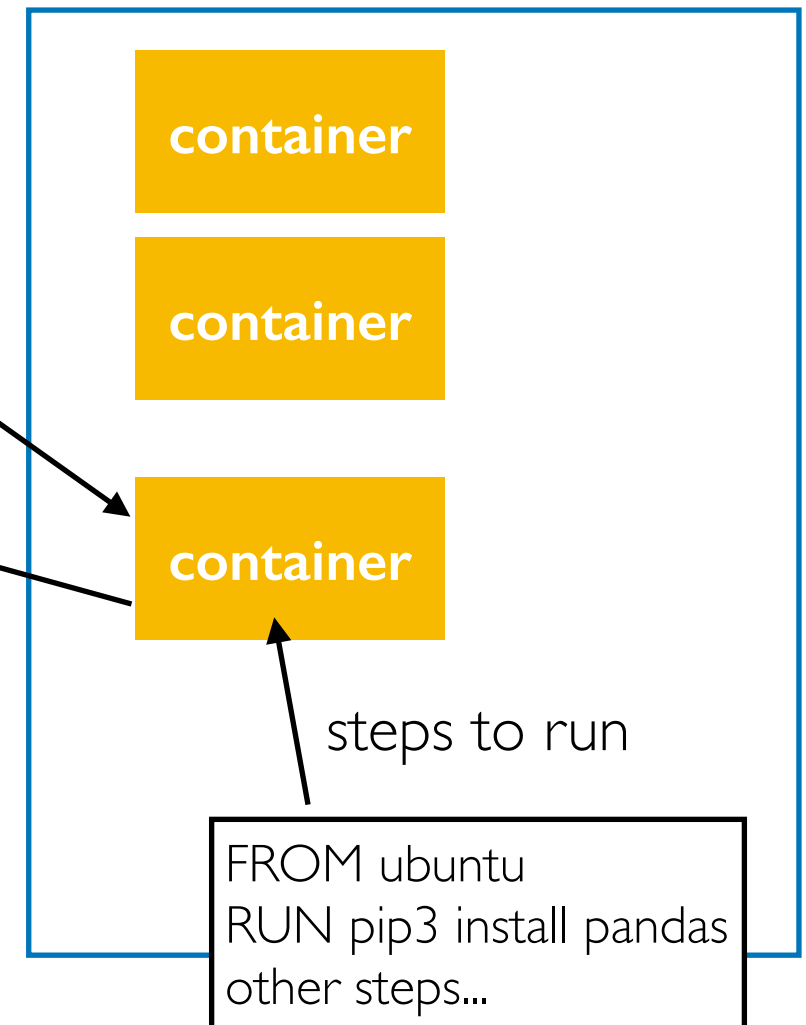
private/corporate registry



Images (on your VM)



Containers (on your VM)



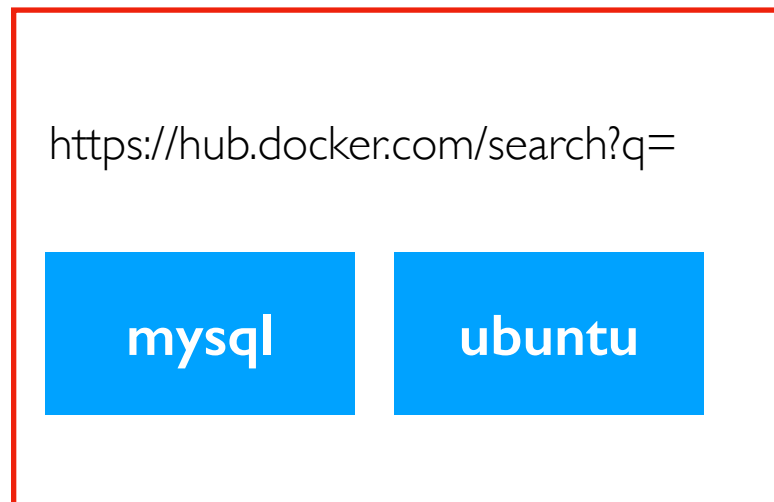
Dockerfiles

Dockerfile

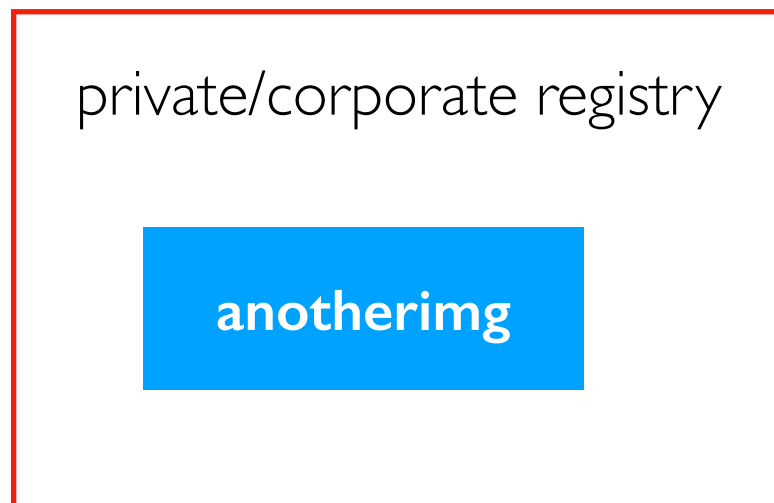
- steps to run in a container (like installs)
- creates a new image
- docker **build** myimg -t pandas

Registries, Images, Containers, and Dockerfiles

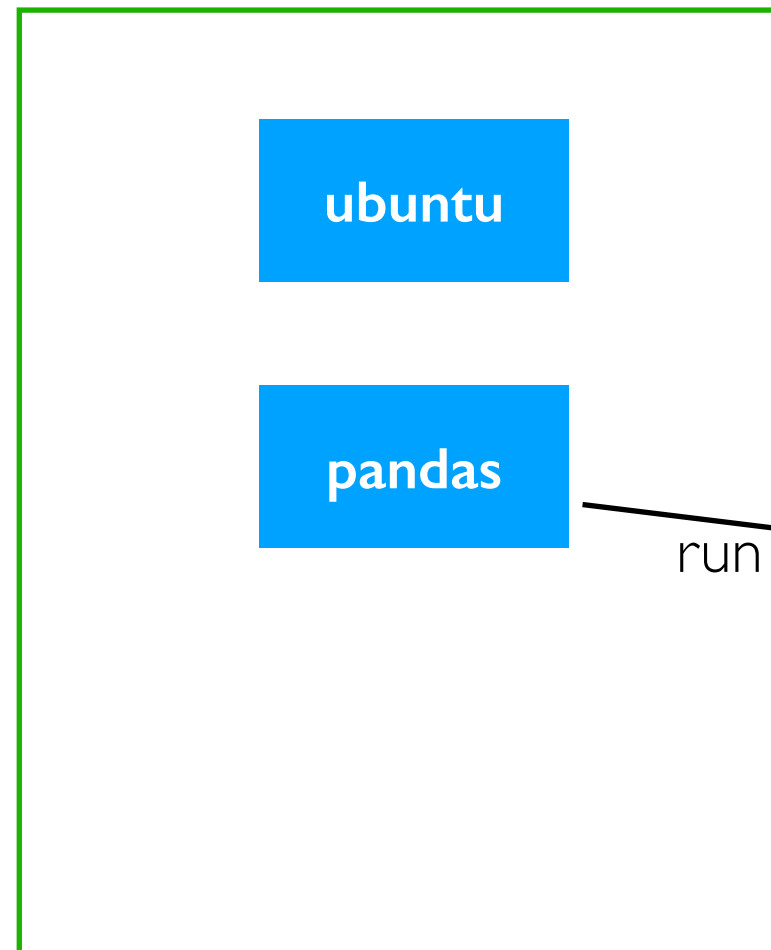
registries:



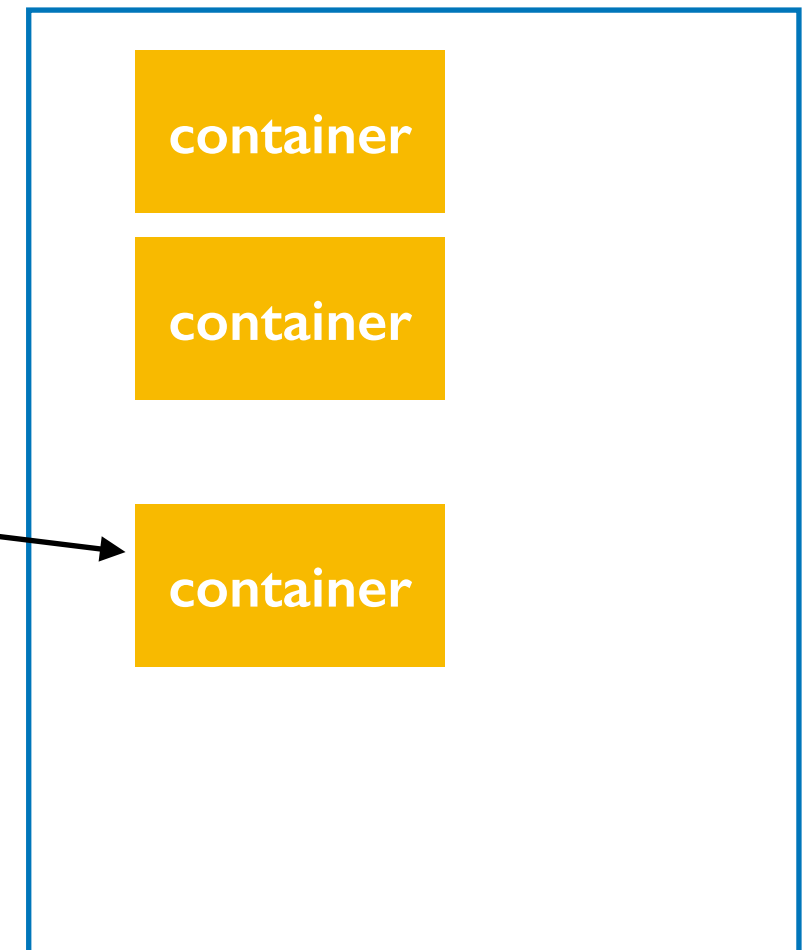
private/corporate registry



Images (on your VM)



Containers (on your VM)



run

Reproducibility

- Docker files unambiguously describe the setup
- Others can get all the same version numbers

Demos...

Outline

Virtualization

Images, Containers, and Dockerfiles


Ports

Docker Compose

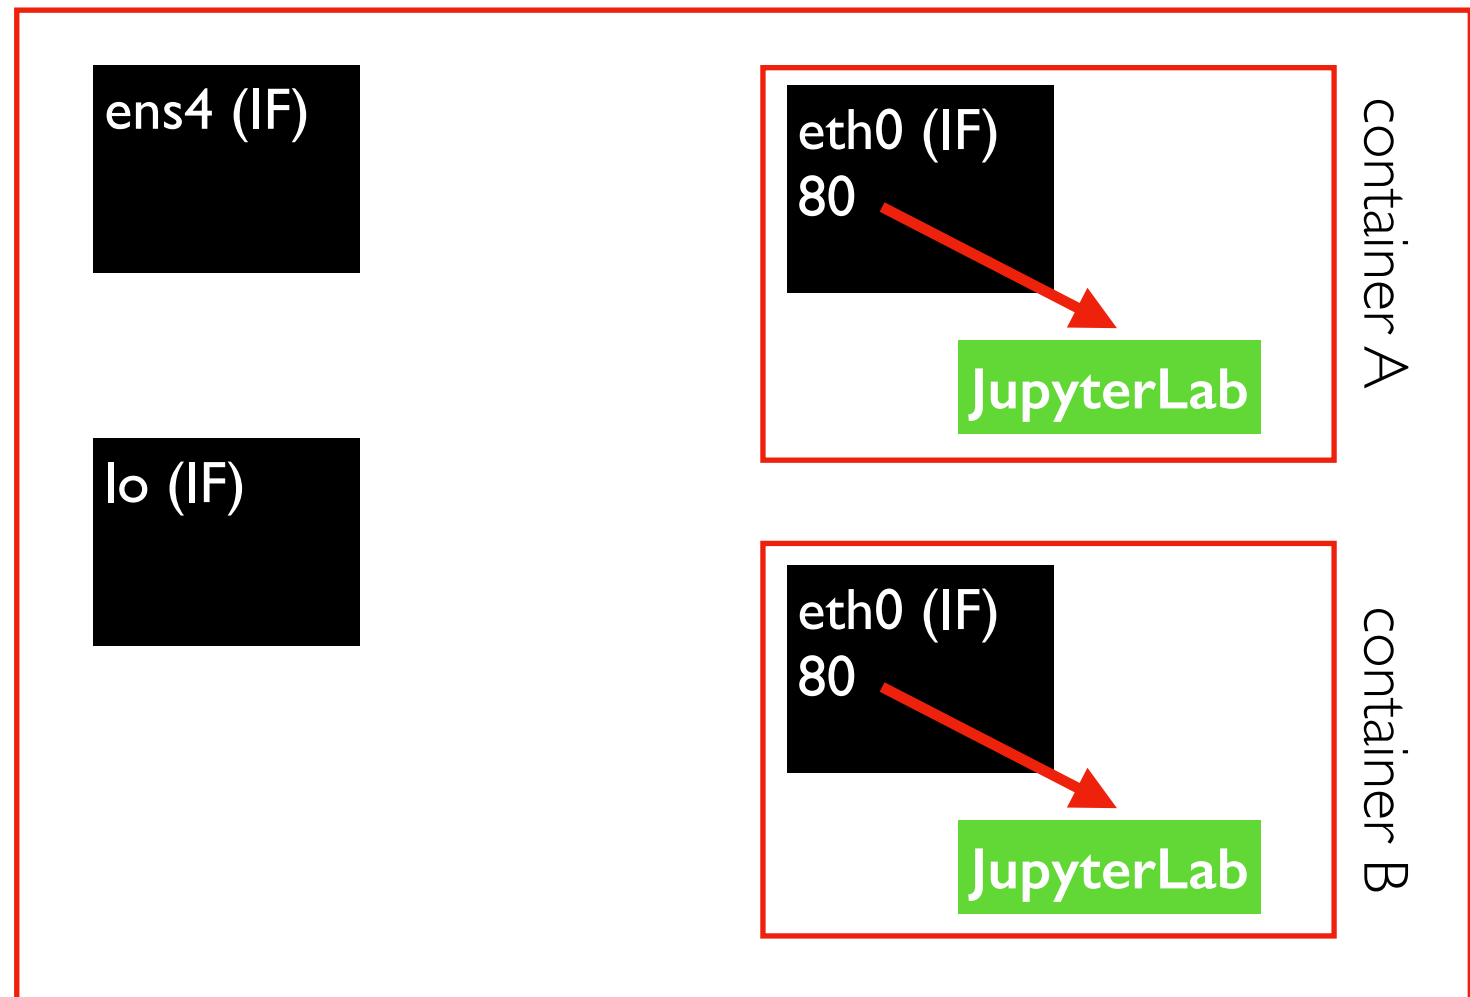
Interfaces (IF) and Ports

both containers have
a virtual port 80

laptop

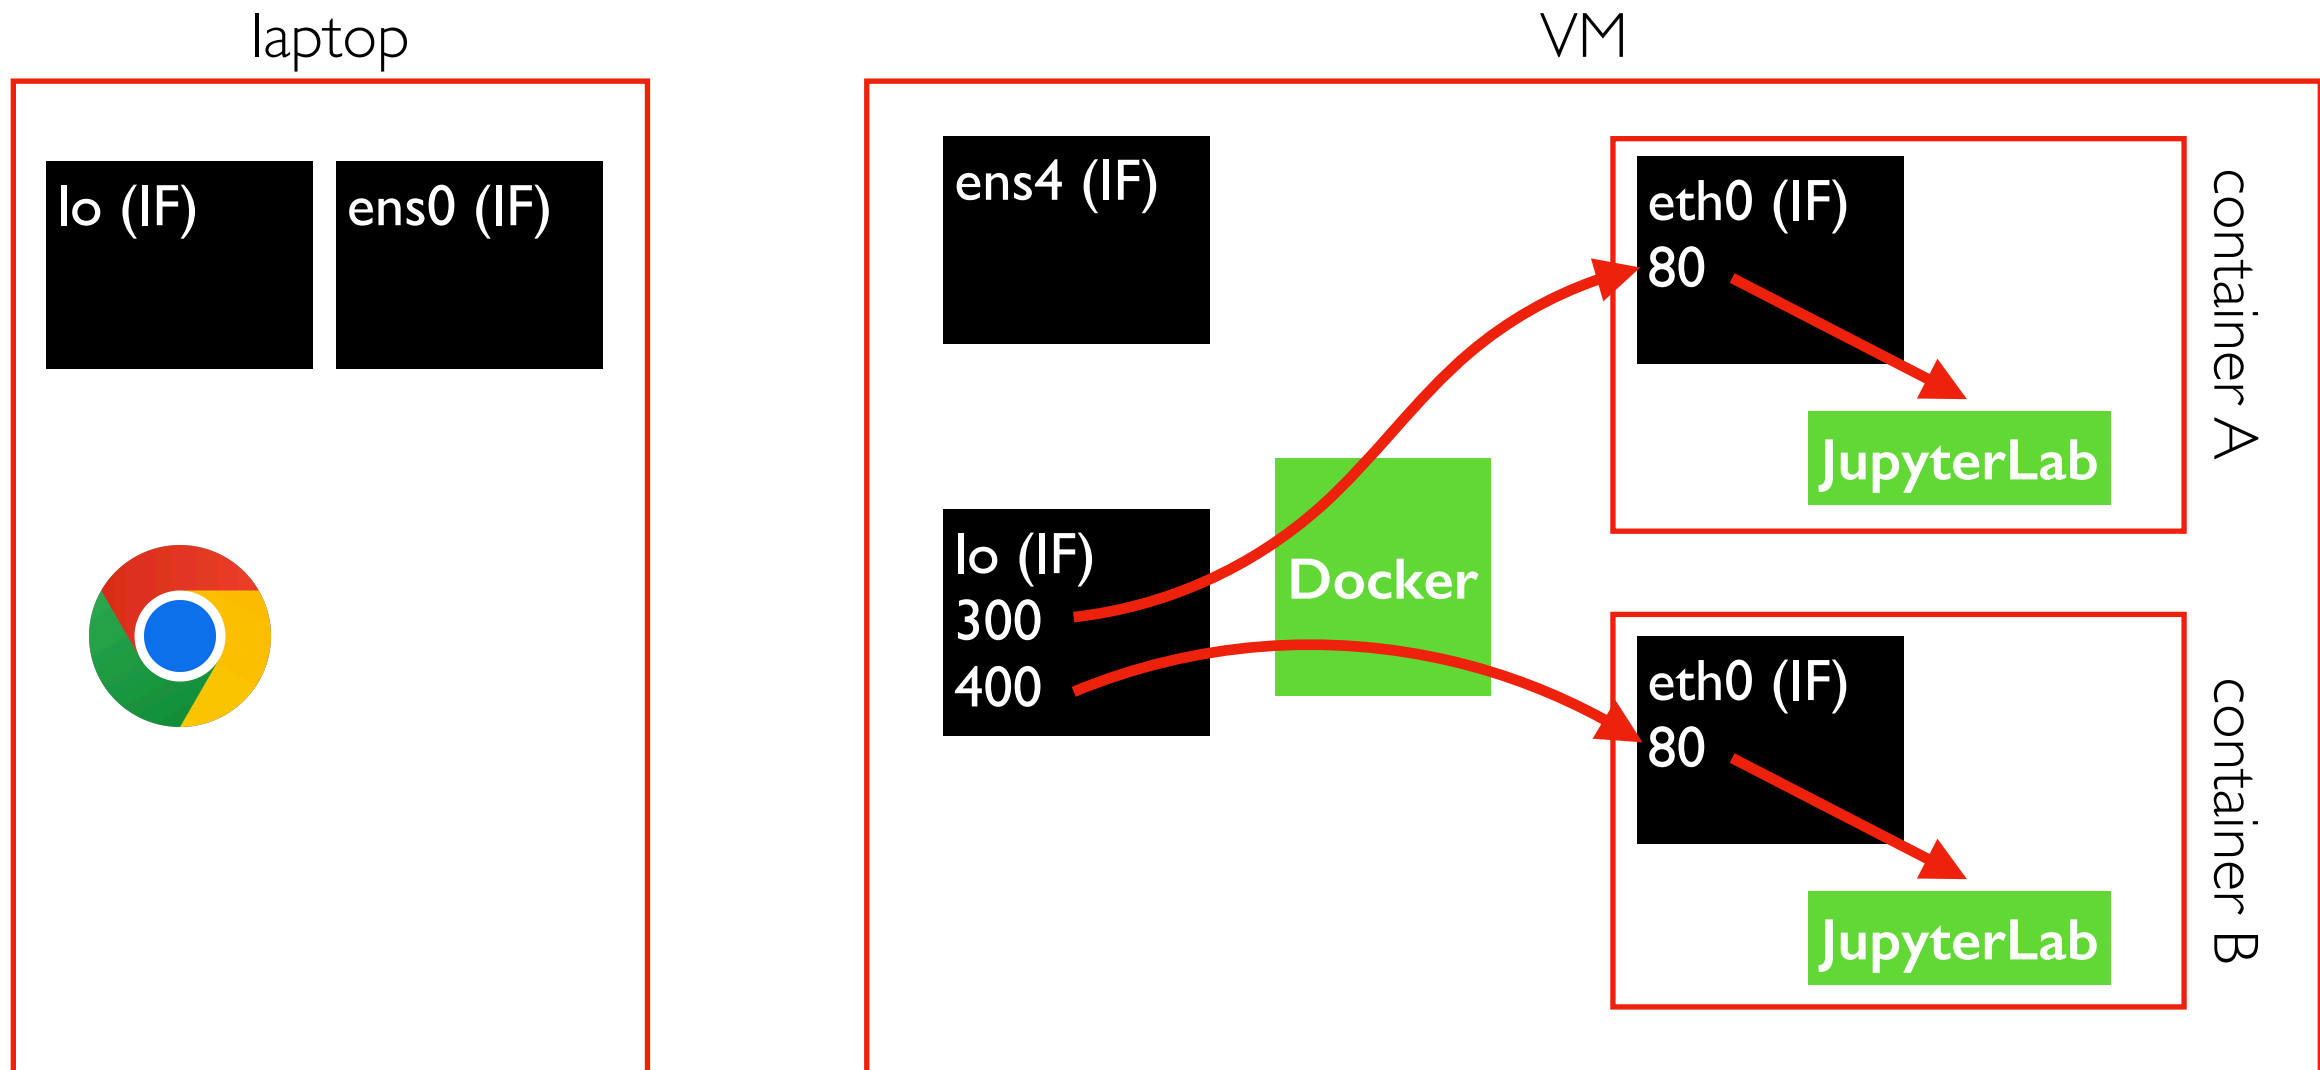


VM



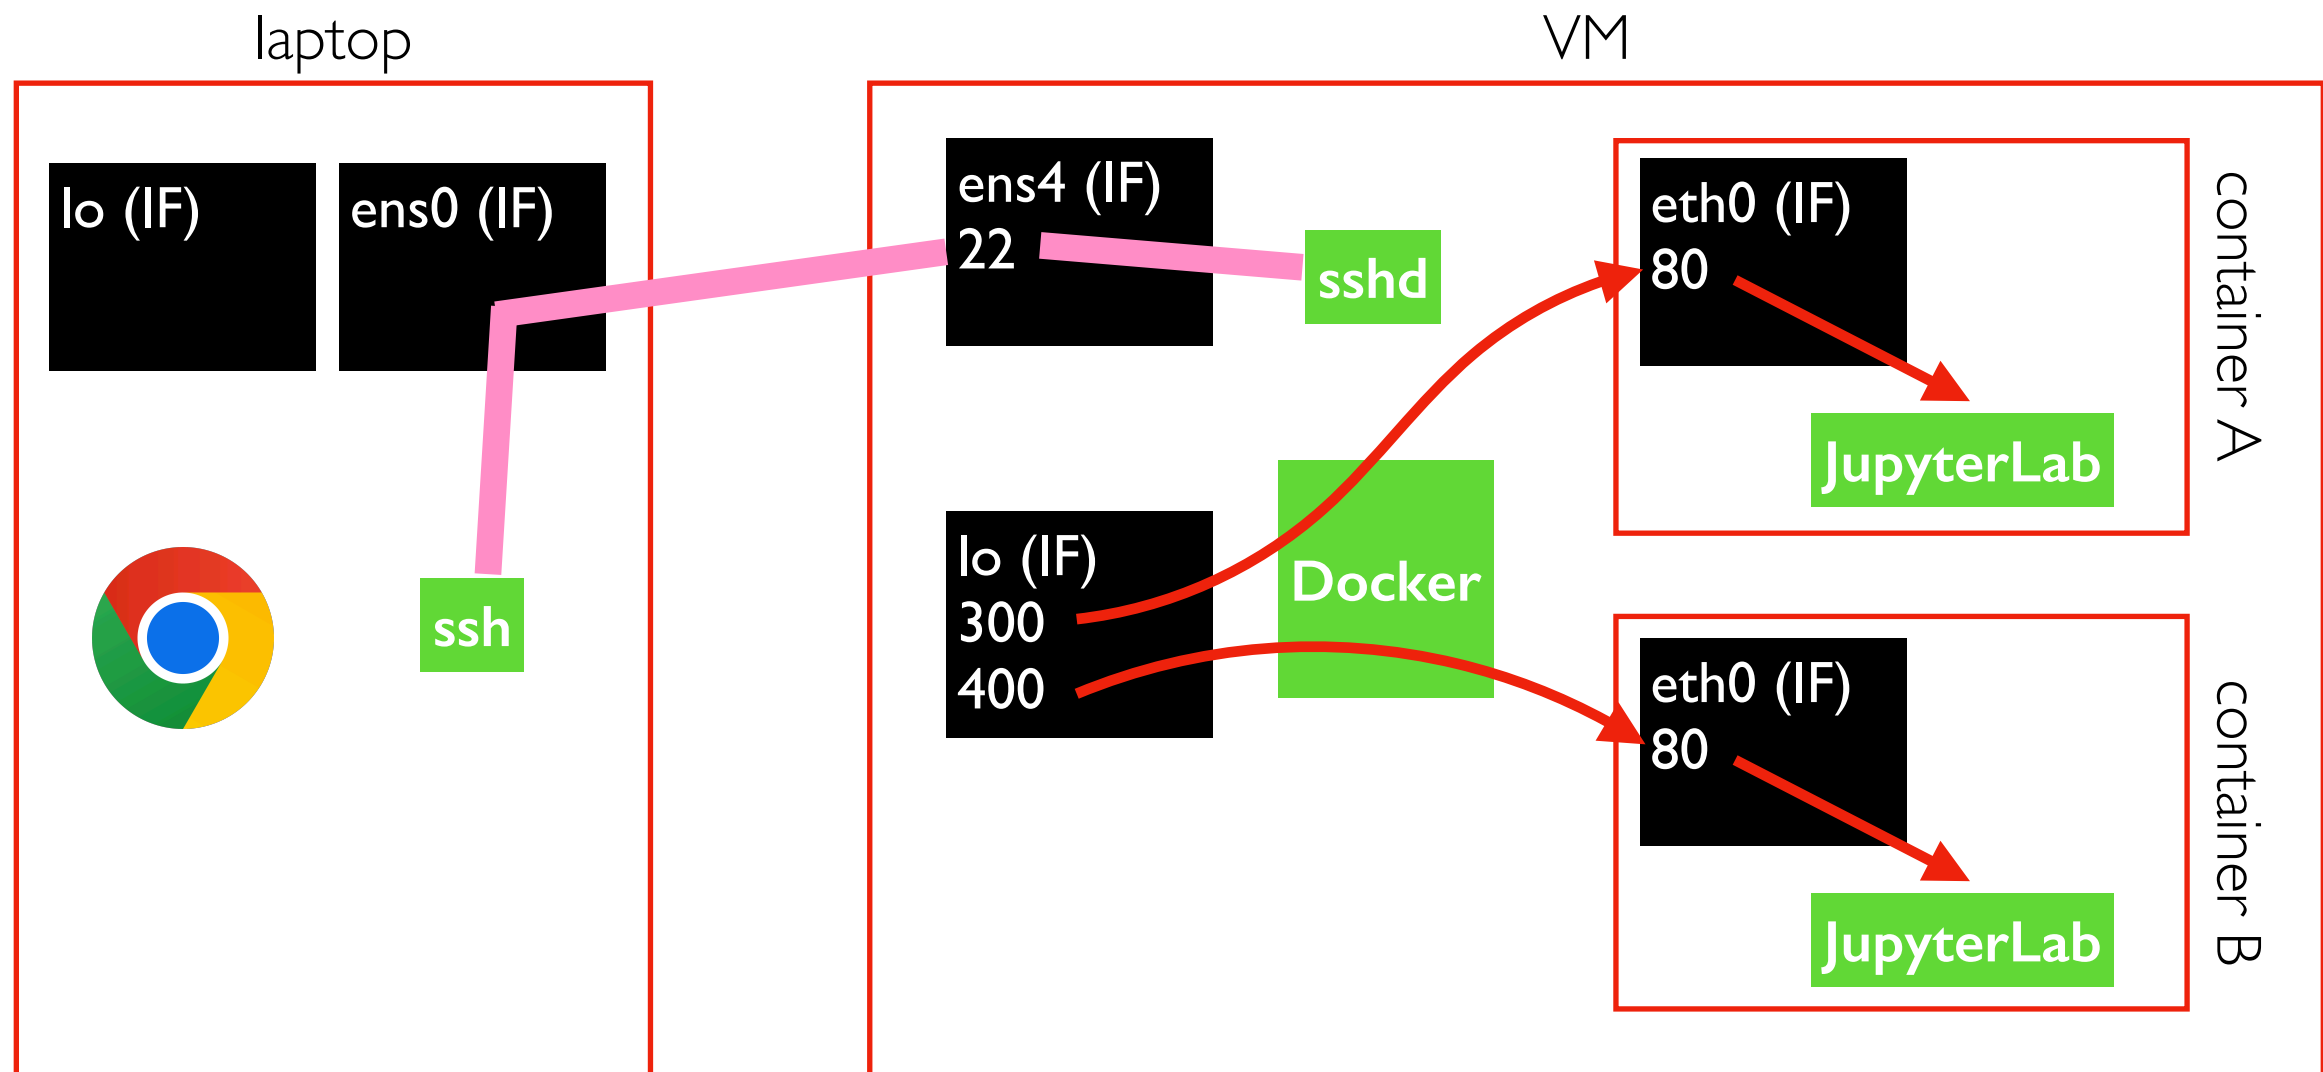
```
docker run -d myimg  
docker run -d myimg
```

Interfaces (IF) and Ports



```
docker run -d -p 127.0.0.1:300:80 myimg  
docker run -d -p 127.0.0.1:400:80 myimg
```

Interfaces (IF) and Ports

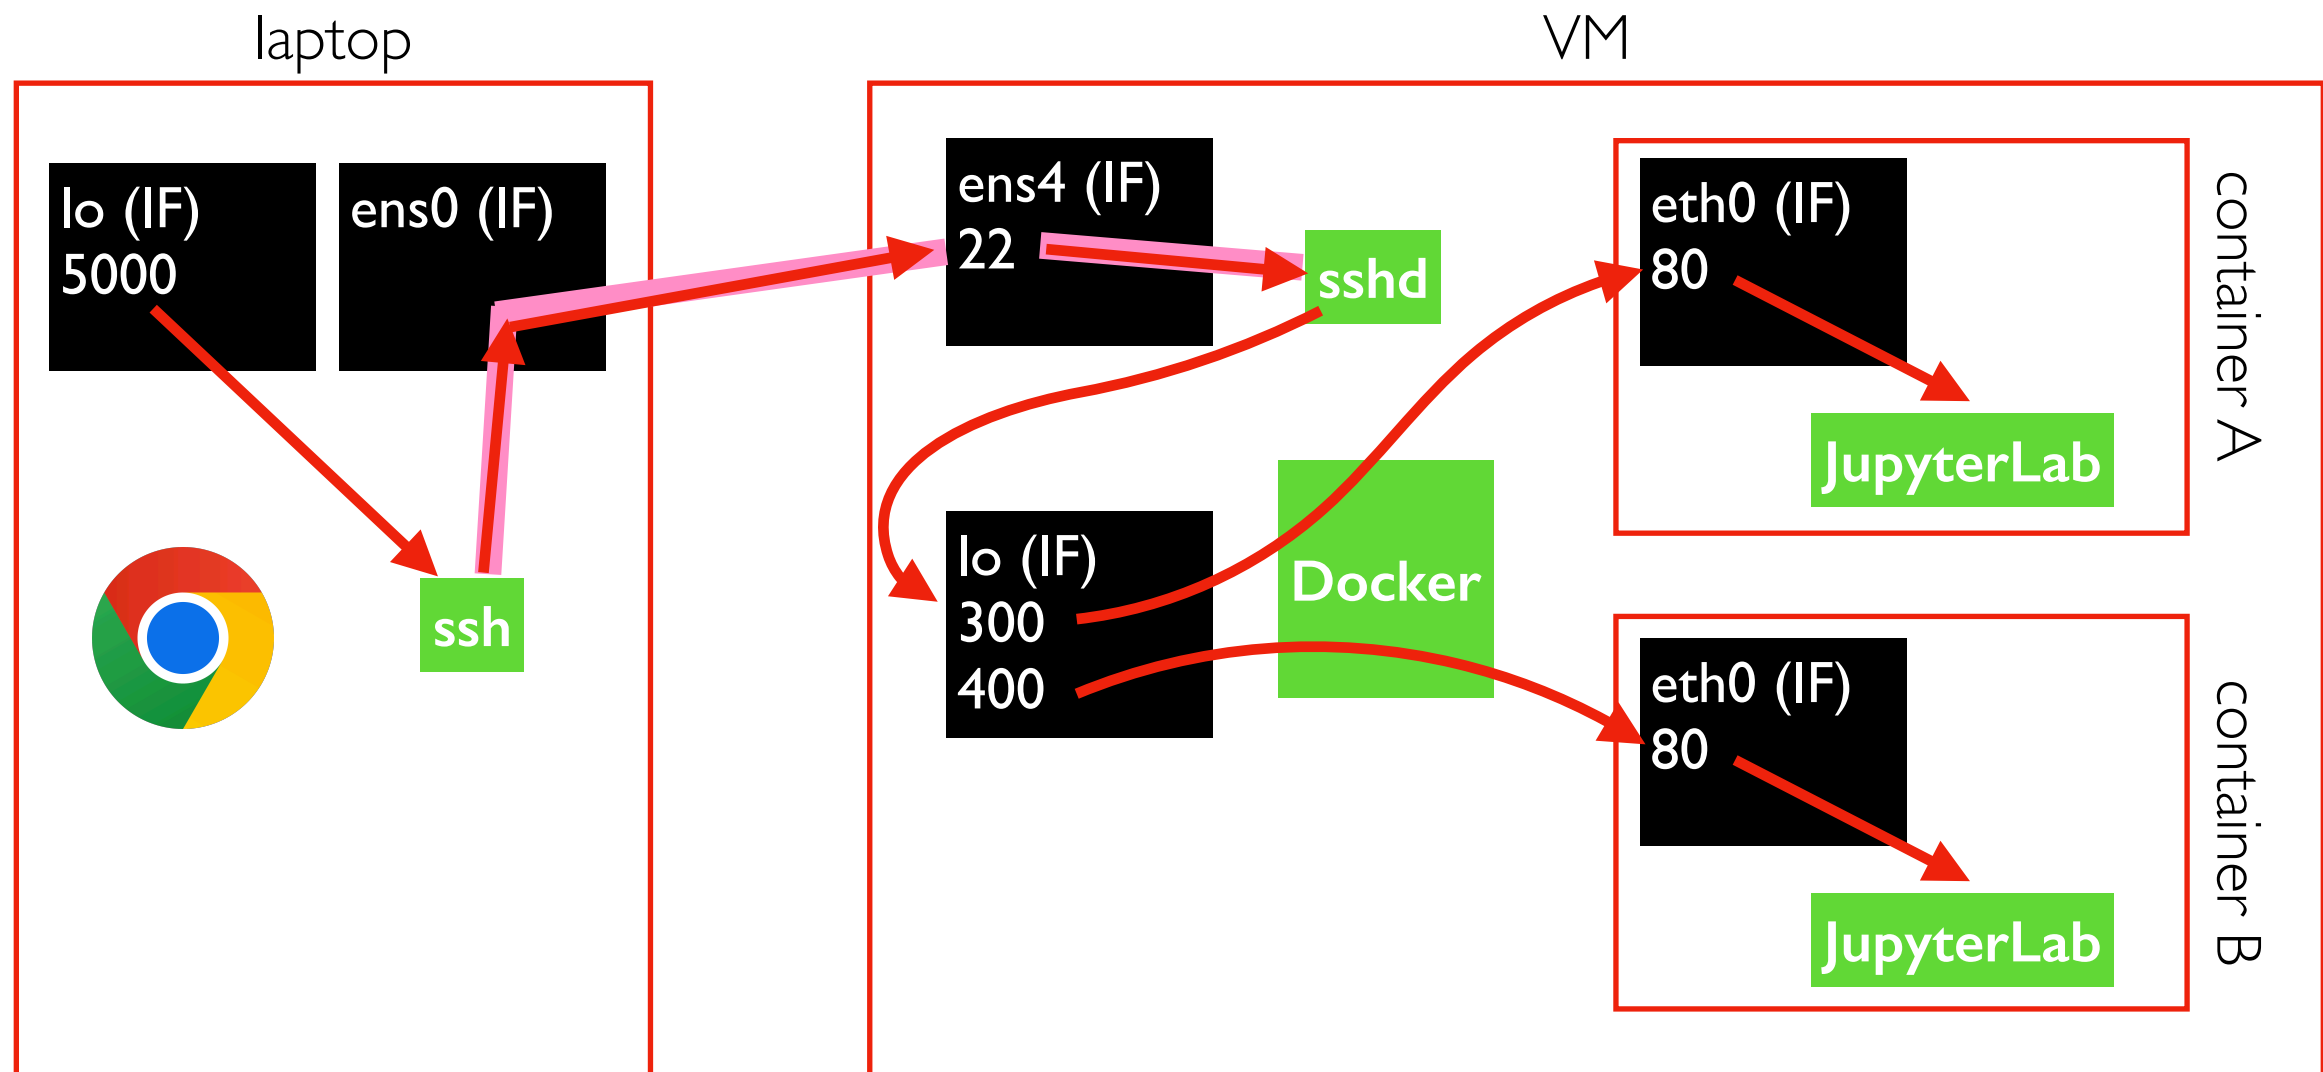


```
ssh USER@VM
```

```
docker run -d -p 127.0.0.1:300:80 myimg  
docker run -d -p 127.0.0.1:400:80 myimg
```

the SSH connection can be used to send commands and/or forward network traffic

Interfaces (IF) and Ports




```
ssh USER@VM -L localhost:5000:localhost:300
```

```
docker run -d -p 127.0.0.1:300:80 myimg  
docker run -d -p 127.0.0.1:400:80 myimg
```

the SSH connection can be used to send
comands and/or forward network traffic

Interfaces (IF) and Ports

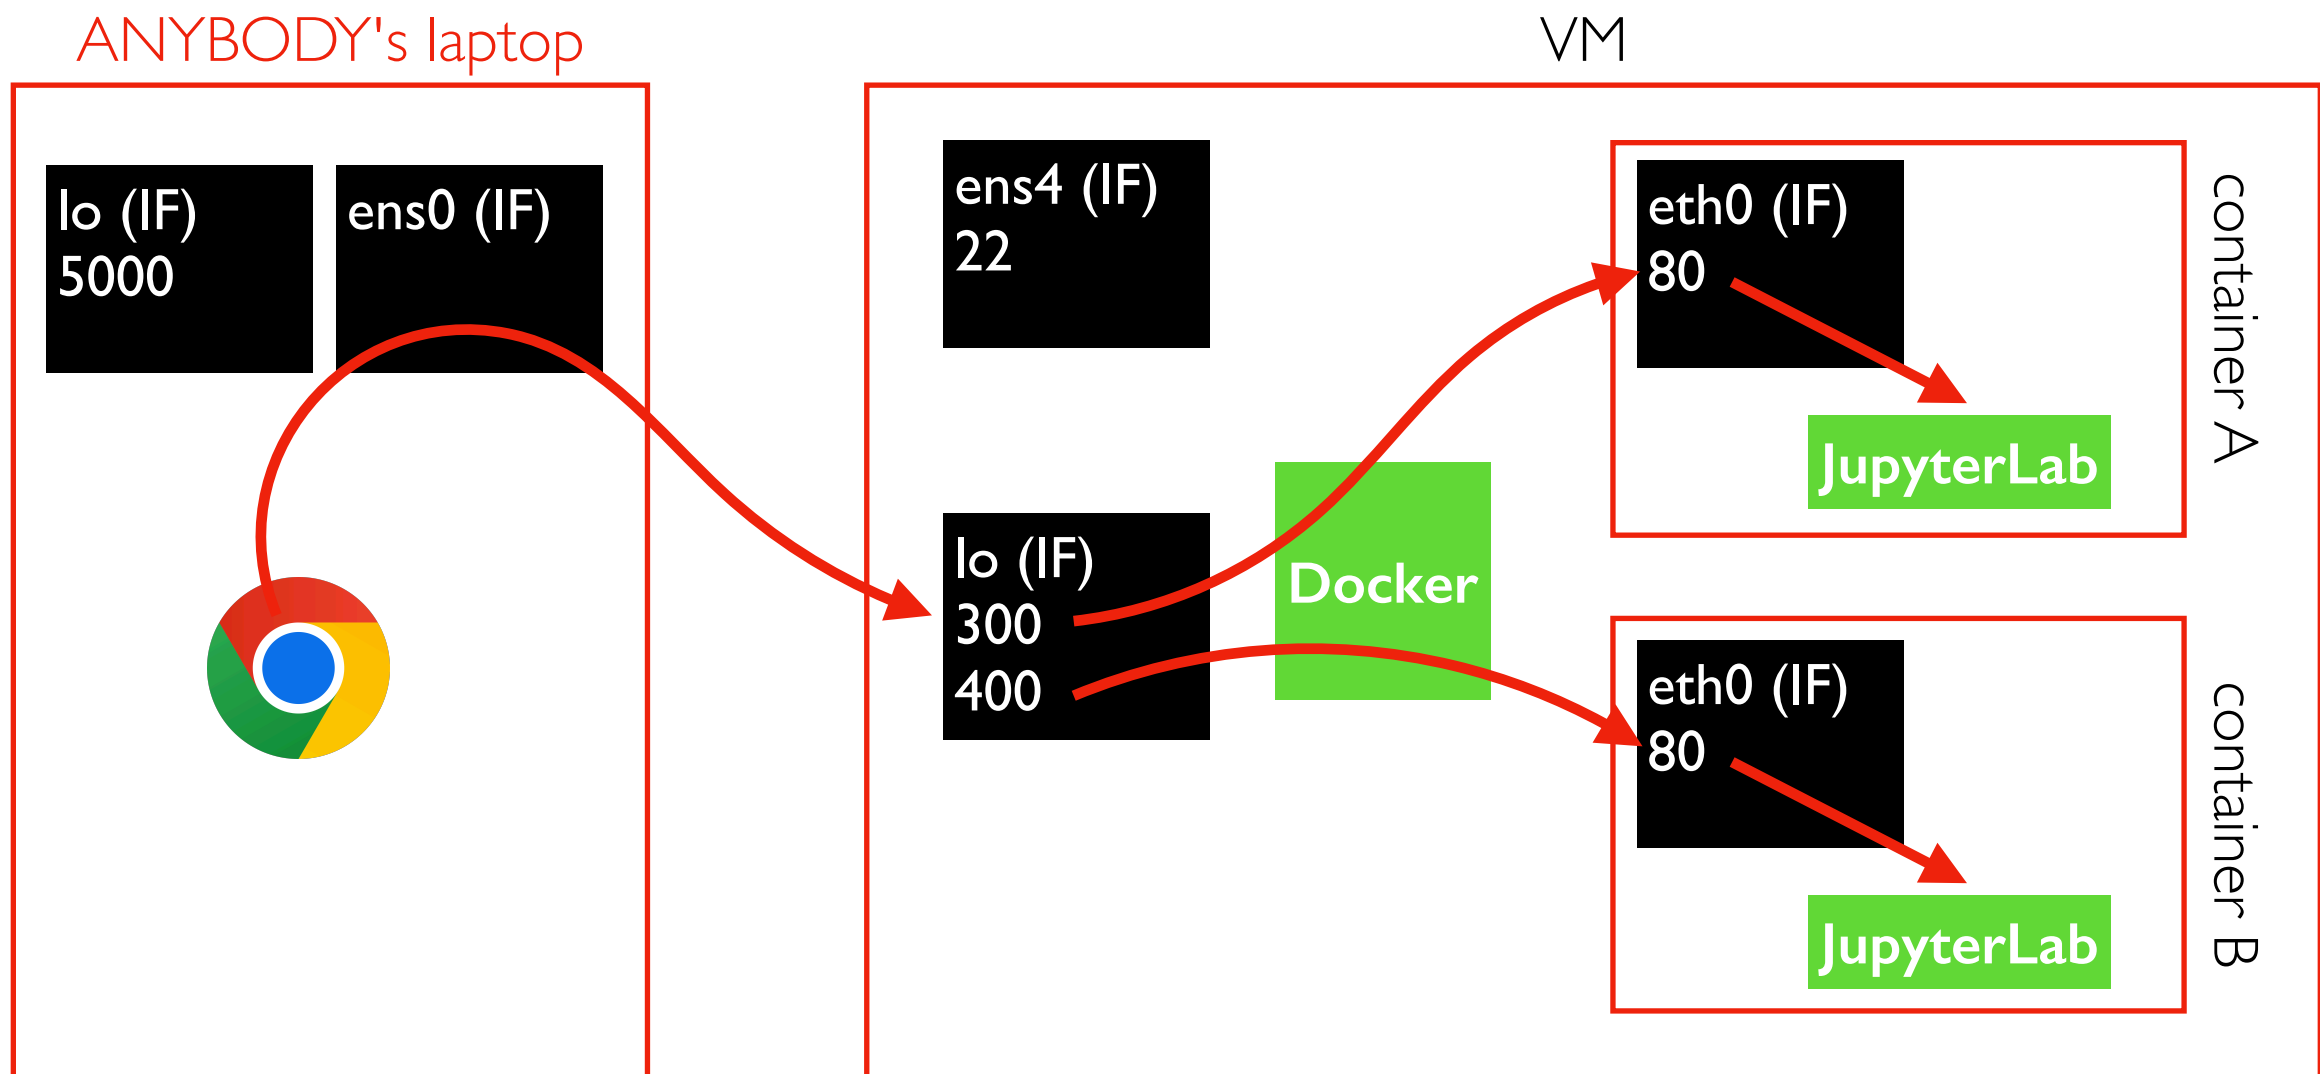


`ssh USER@VM -L localhost:5000:localhost:300` `docker run -d -p 127.0.0.1:300:80 myimg`
`docker run -d -p 127.0.0.1:400:80 myimg`

<http://localhost:5000/lab> (in browser)

yay! You can connect to JupyterLab
inside a container running on your VM

Interfaces (IF) and Ports



`docker run -d -p 300:80 myimg`



`docker run -d -p 0.0.0.0:300:80 myimg`

Careful, default is to listen on all ports!

Other security:

- firewall (block port 300)
- password (in JupyterLab)

Demos...

Outline

Virtualization

Images, Containers, and Dockerfiles

Ports


Docker Compose

Container Orchestration

Orchestration lets you deploy many cooperating containers across a cluster of Docker workers.

Kubernetes is the most well known.

Docker **compose** is a simpler tool that lets you deploy cooperating containers to a single worker.



Demos...