

[544] Networking

Tyler Caraza-Harter

Learning Objectives

- explain how MAC addresses, IP addresses, and port numbers provide addressing, used to facilitate communication between processes on different machines
- select IP addresses correctly for binding and using in a browser to achieve connection in the context of a server running behind a NAT
- identify the port number being used by a process
- select between transport methods (TCP and UDP) based on the functionality needed (on top of IP functionality)

Outline

Networks

Internets and "The Internet"

Transport Protocols

Network Interface Controllers and MAC Addresses




NICs can connect a computer to different physical mediums, such as:

- Ethernet (wired)
 - Wi-Fi (wireless)

Every NIC in the world has a unique MAC (media access control) address

- 281 trillion possible addrs
 - some devices randomly change their MAC addr for privacy

MAC address:



ip address

interface

```
trh@instance-20240903-151711:~$ ip address
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group
default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host noprefixroute
        valid_lft forever preferred_lft forever
2: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc mq state UP group
default qlen 1000
    link/ether 42:01:0a:80:00:04 brd ff:ff:ff:ff:ff:ff
    altname enp0s4
    inet 10.128.0.4/32 metric 100 scope global dynamic ens4
        valid_lft 1870sec preferred_lft 1870sec
    inet6 fe80::4001:aff:fe80:4/64 scope link
        valid_lft forever preferred_lft forever
```

Virtual Interfaces

```
trh@instance-20240903-151711:~$ ip address
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group
default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host noprefixroute
        valid_lft forever preferred_lft forever
2: ens4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1460 qdisc mq state UP group
default qlen 1000
    link/ether 42:01:0a:80:00:04 brd ff:ff:ff:ff:ff:ff
    altname enp0s4
    inet 10.128.0.4/32 metric 100 scope global dynamic ens4
        valid_lft 1870sec preferred_lft 1870sec
    inet6 fe80::4001:aff:fe80:4/64 scope link
        valid_lft forever preferred_lft forever
```

loopback (lo) device a virtual interface (not actual hardware)
connecting to a mini network containing just your computer

Google Console: Adding Interfaces (NICs)

Create Instance > Advanced Options > Networking

Network interfaces ?

Network interface is permanent

default default (10.128.0.0/20)	▼
other-net subnet (10.0.0.0/24)	▼
ADD NETWORK INTERFACE	

Virtual Machine Summary

<input type="checkbox"/>	<input checked="" type="checkbox"/>	instance-2	us-central1-a	10.128.0.37 (nic0) 10.0.0.2 (nic1)	34.29.220.248 (nic0) 35.202.74.234 (nic1)
--------------------------	-------------------------------------	----------------------------	---------------	---	--

Google Console: Adding Interfaces (NICs)

Create Instance > Advanced Options > Networking

Network interfaces ?

```
trh@instance-2:~$ ip address
ens4: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1460
    inet 10.128.0.37 netmask 255.255.255.255 broadcast 0.0.0.0
    inet6 fe80::4001:aff:fe80:25 prefixlen 64 scopeid 0x20<link>
        ether 42:01:0a:80:00:25 txqueuelen 1000 (Ethernet)
        RX packets 637 bytes 546000 (546.0 KB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 612 bytes 97265 (97.2 KB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens5: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1460
    inet 10.0.0.2 netmask 255.255.255.255 broadcast 0.0.0.0
    inet6 fe80::4001:aff:fe00:2 prefixlen 64 scopeid 0x20<link>
        ether 42:01:0a:00:00:02 txqueuelen 1000 (Ethernet)
        RX packets 51 bytes 9955 (9.9 KB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 61 bytes 6834 (6.8 KB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
        loop txqueuelen 1000 (Local Loopback)
        RX packets 120 bytes 13534 (13.5 KB)
        RX errors 0 dropped 0 overruns 0 frame 0
        TX packets 120 bytes 13534 (13.5 KB)
        TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

7


34.29.220.248

[\(nic0\)](#)

35.202.74.234

[\(nic1\)](#)


Networks



A **network** has **nodes** that send bytes (**packets**) to other nodes **by MAC address**

- **nodes**: computer, switch, etc
- direct, or **forwarded by switches**
- whole network uses same physical tech (Wi-Fi, Ethernet, etc)

Networks



Computers can have multiple NICs

- can be on multiple networks (2 ethernets, ethernet+Wi-Fi, etc)
- can't send to a MAC addr in another network without a NIC there


Outline

Networks

Internets and "The Internet"

Transport Protocols


Routers



Packet Forwarding

Packets (some bytes with an address and other info) can be forwarded along a path from point A to point B



- routers contain **forwarding tables** that help them decide which direction to send along a packet
- those tables would be too big if a router had to know where every MAC address existed in the Internet



Internet Protocol

IP addresses are used to send packets across an internet

- example: 34.29.237.29 (domains can map to IP addrs)
- there are about 4 billion possible IP addresses (IPv4)
- IPv6 (less used) are 4x longer
- forwarding tables only need to know which way to send for a given network number



Listening on an Interface

```
trh@instance-2:~$ ip address
ens4: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1460
  inet 10.0.1.2 netmask 255.255.255.255 broadcast 0.0.0.0
  inet6 fe80::c228:3ff:fe01:2%ens4 brd fe80::ff:fe01:2 mtu 1280
    python3 -m http.server --bind 10.0.1.2
  ether 42:01:00:00:00:02 brd ff:ff:ff:ff:ff:ff link-layer
...
ens5: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1430
  inet 10.0.3.2 netmask 255.255.255.255 broadcast 10.0.3.2
  inet6 fe80::c228:3ff:fe01:2%ens5 brd fe80::ff:fe01:2 mtu 1280
    python3 -m http.server --bind 10.0.3.2
  ether 42:01:00:00:00:02 brd ff:ff:ff:ff:ff:ff link-layer
...
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
  inet 127.0.0.1 netmask 255.0.0.0
  inet6 ::1 brd :: link-layer
    python3 -m http.server --bind 127.0.0.1
...
all of them: python3 -m http.server --bind 0.0.0.0
```

Private Networks

Challenges


- we don't have enough IPv4 addresses
- we don't want every machine to be able to receive packets from anywhere

Private ranges:

- 192.168.0.0 to 192.168.255.255
- 172.16.0.0 to 172.31.255.255
- 10.0.0.0 to 10.255.255.255

these can be divided into "sub networks" (subnets) to create different networks in a bigger org

Private networks allow duplicates and unreachable machines



Private Networks


```
trh@instance-2:~$ ip address
ens4: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1460
  inet 10.0.1.2  brd 10.0.1.2  netmask 255.255.255.0  broadcast 10.0.1.2
    python3 -m http.server --bind 10.0.1.2
  inet6 fe80::42:1ff:fe01:02%ens4  brd fe80::ff:1ff:fe01:02  scope link
    ether 42:01:0a:00:01:02  txqueuelen 1000  (Ethernet)
...
...
```

Private ranges:

- 192.168.0.0 to 192.168.255.255
- 172.16.0.0 to 172.31.255.255
- 10.0.0.0 to 10.255.255.255

<http://10.0.1.2:...>
won't work in web browser!


Private networks allow duplicates and unreachable machines



Network Address Translation

Google Console (view NAT config)

<input type="checkbox"/> Status	Name 	Internal IP	External IP
<input type="checkbox"/>	 instance-1	10.128.0.36 (nic0)	34.29.237.29 (nic0)
<input type="checkbox"/>	 instance-2	10.0.1.2 (nic0) 10.0.3.2 (nic1)	35.202.74.234 (nic0) 34.29.220.248 (nic1)




Network Address Translation

Google Console (view NAT config)

<input type="checkbox"/> Status	Name 	Internal IP	External IP
<input type="checkbox"/>	 instance-1	10.128.0.36 (nic0)	34.29.237.29 (nic0)
<input type="checkbox"/>	 instance-2	10.0.1.2 (nic0) 10.0.3.2 (nic1)	35.202.74.234 (nic0) 34.29.220.248 (nic1)

a firewall can limit what traffic is allowed




Network Address Translation

Google Console (view NAT config)

<input type="checkbox"/> Status	Name 	Internal IP	External IP
<input type="checkbox"/>	 instance-1	10.128.0.36 (nic0)	34.29.237.29 (nic0)
<input type="checkbox"/>	 instance-2	10.0.1.2 (nic0) 10.0.3.2 (nic1)	35.202.74.234 (nic0) 34.29.220.248 (nic1)

By default, the external IPs are "ephemeral" (change upon reboot). You can rent "static" IPs that don't change.



Directory listing for /

• [temp/](#)

Server

```
trh@instance-2:~/temp$ python3 -m http.server --bind 10.0.1.2
Serving HTTP on 10.0.1.2 port 8000 (http://10.0.1.2:8000/) ...
72.33.0.184 - - [10/Feb/2023 21:12:53] "GET / HTTP/1.1" 200 -
...
```

Outline

Networks

Internets and "The Internet"

Transport Protocols


Port Numbers

Computers might be running multiple processes using the network

- IP address => which NIC?
- Port number => which process?

```
trh@instance-2:~$ python3 -m http.server --directory=A --bind 10.0.1.2 8000 &
[1] 13502
Serving HTTP on 10.0.1.2 port 8000 (http://10.0.1.2:8000/) ...
```

```
trh@instance-2:~$ python3 -m http.server --directory=B --bind 10.0.1.2 9000 &
[2] 13503
Serving HTTP on 10.0.1.2 port 9000 (http://10.0.1.2:9000/) ...
```



TopHat...

Transport Protocols

Most common

- **UDP** (User Datagram Protocol)
- **TCP** (Transmission Control Protocol)

BOTH build on IP networking and BOTH provide **port numbers**

-t: tcp, -u: udp



```
trh@instance-2:~/temp$ sudo ss -tlpn
State      Local Address:Port      Peer Address:Port      Process
LISTEN    10.128.0.4:8000          0.0.0.0:*
LISTEN    10.128.0.4:9000          0.0.0.0:*
LISTEN    *:22                      *:*
```

Reliability: UDP vs. TCP

Packets may be


- dropped
- reordered
- split

TCP saves+reassembles **packets** in order to provide original **message** (when possible).

For packet drops, it retries. We'll mostly use TCP.

UDP doesn't do this extra work. Why ever use UDP?

Network Stack: Common Implementations



Network applications (like most complex systems) are not built as one single system. Layers are built upon other layers to provide additional functionality.