

[544] Docker Networking

Tyler Caraza-Harter

Learning Objectives

- configure SSH tunneling and Docker port forwarding to communicate with an app in a container on a different machine
- deploy multi-container apps with Docker compose
- identify situations where replication and/or some variant of partitioning is useful

Outline

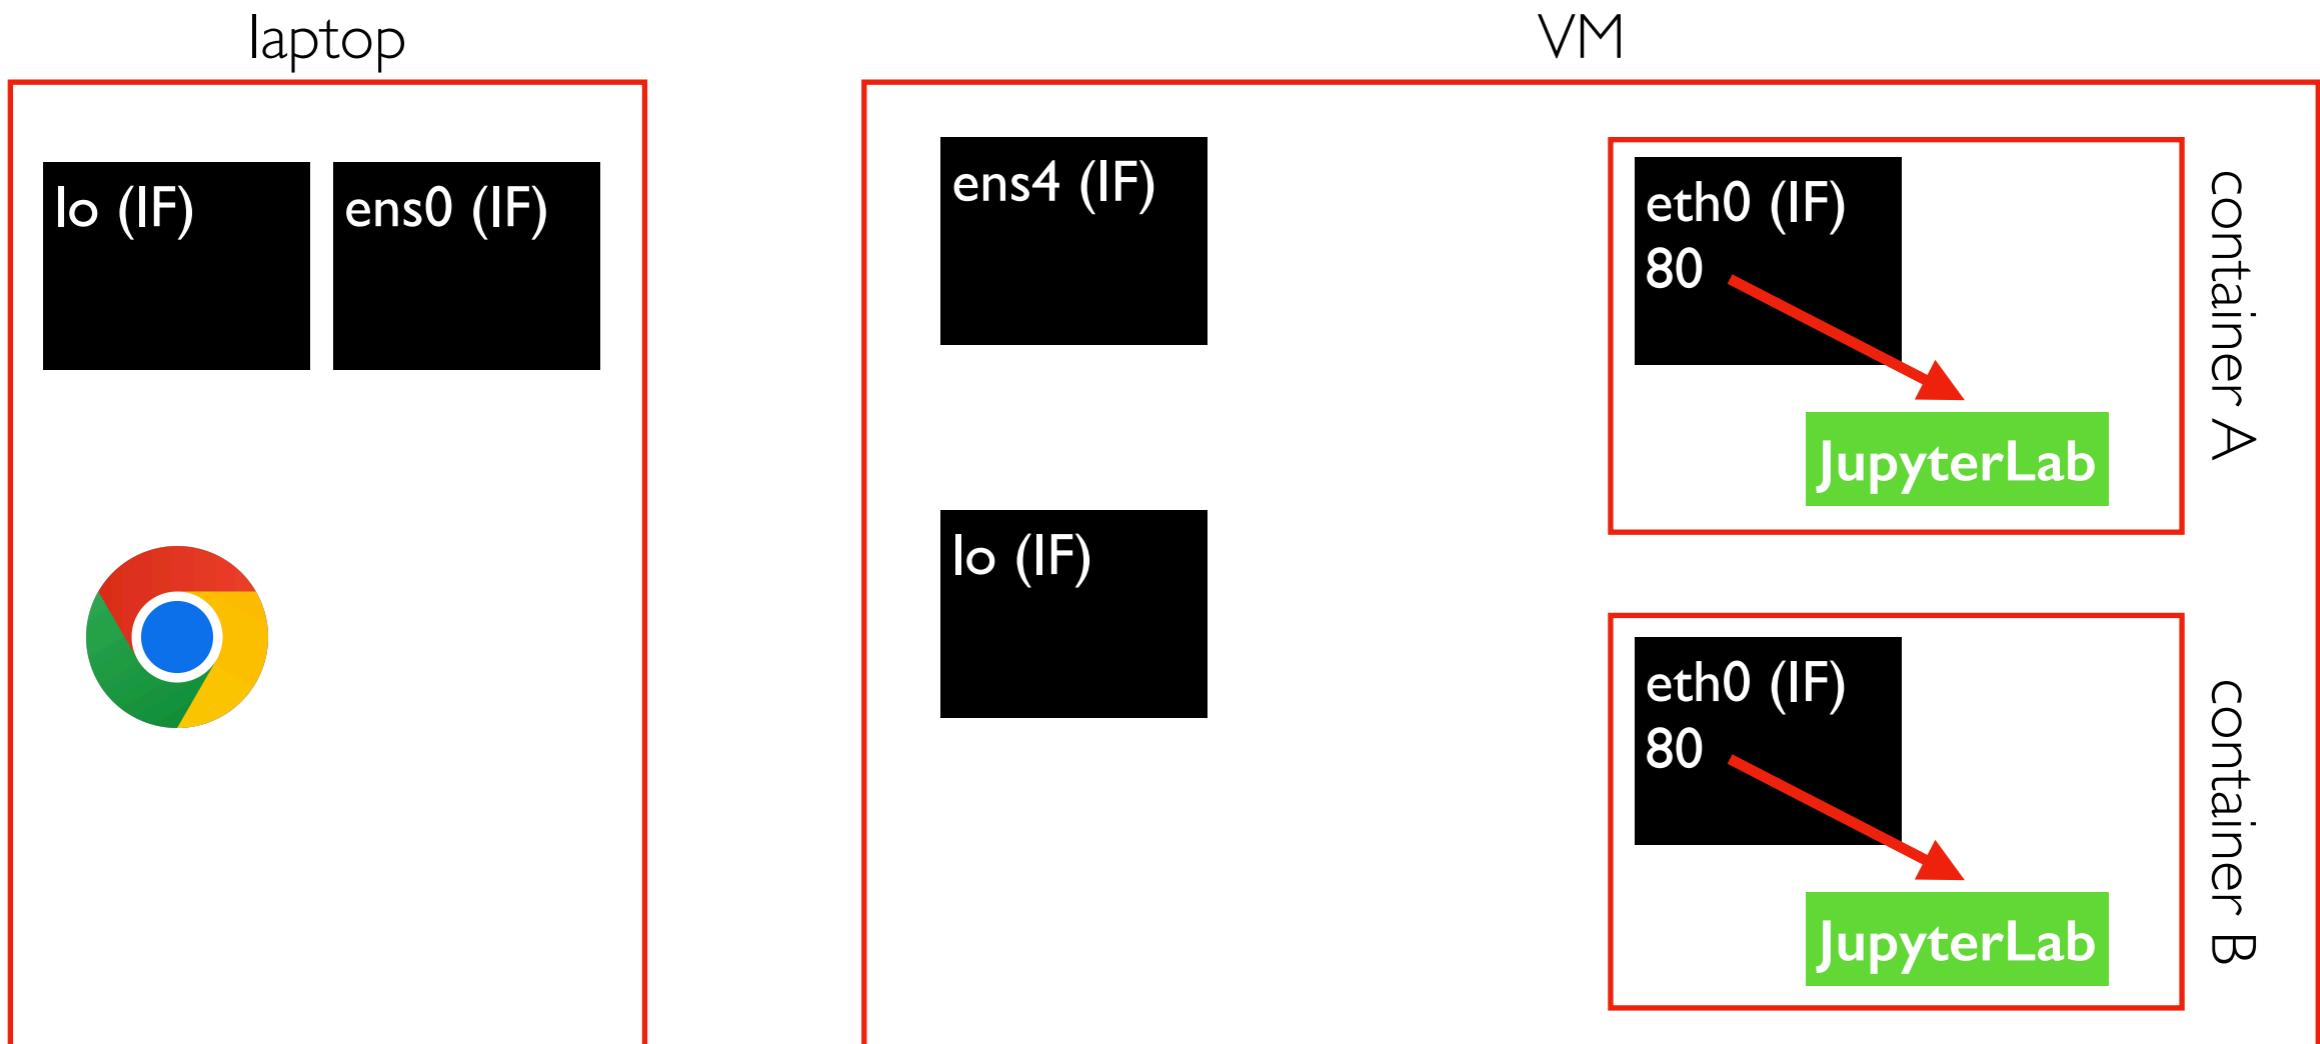
Docker Port Forwarding

Docker Compose

Partitioning and Replication

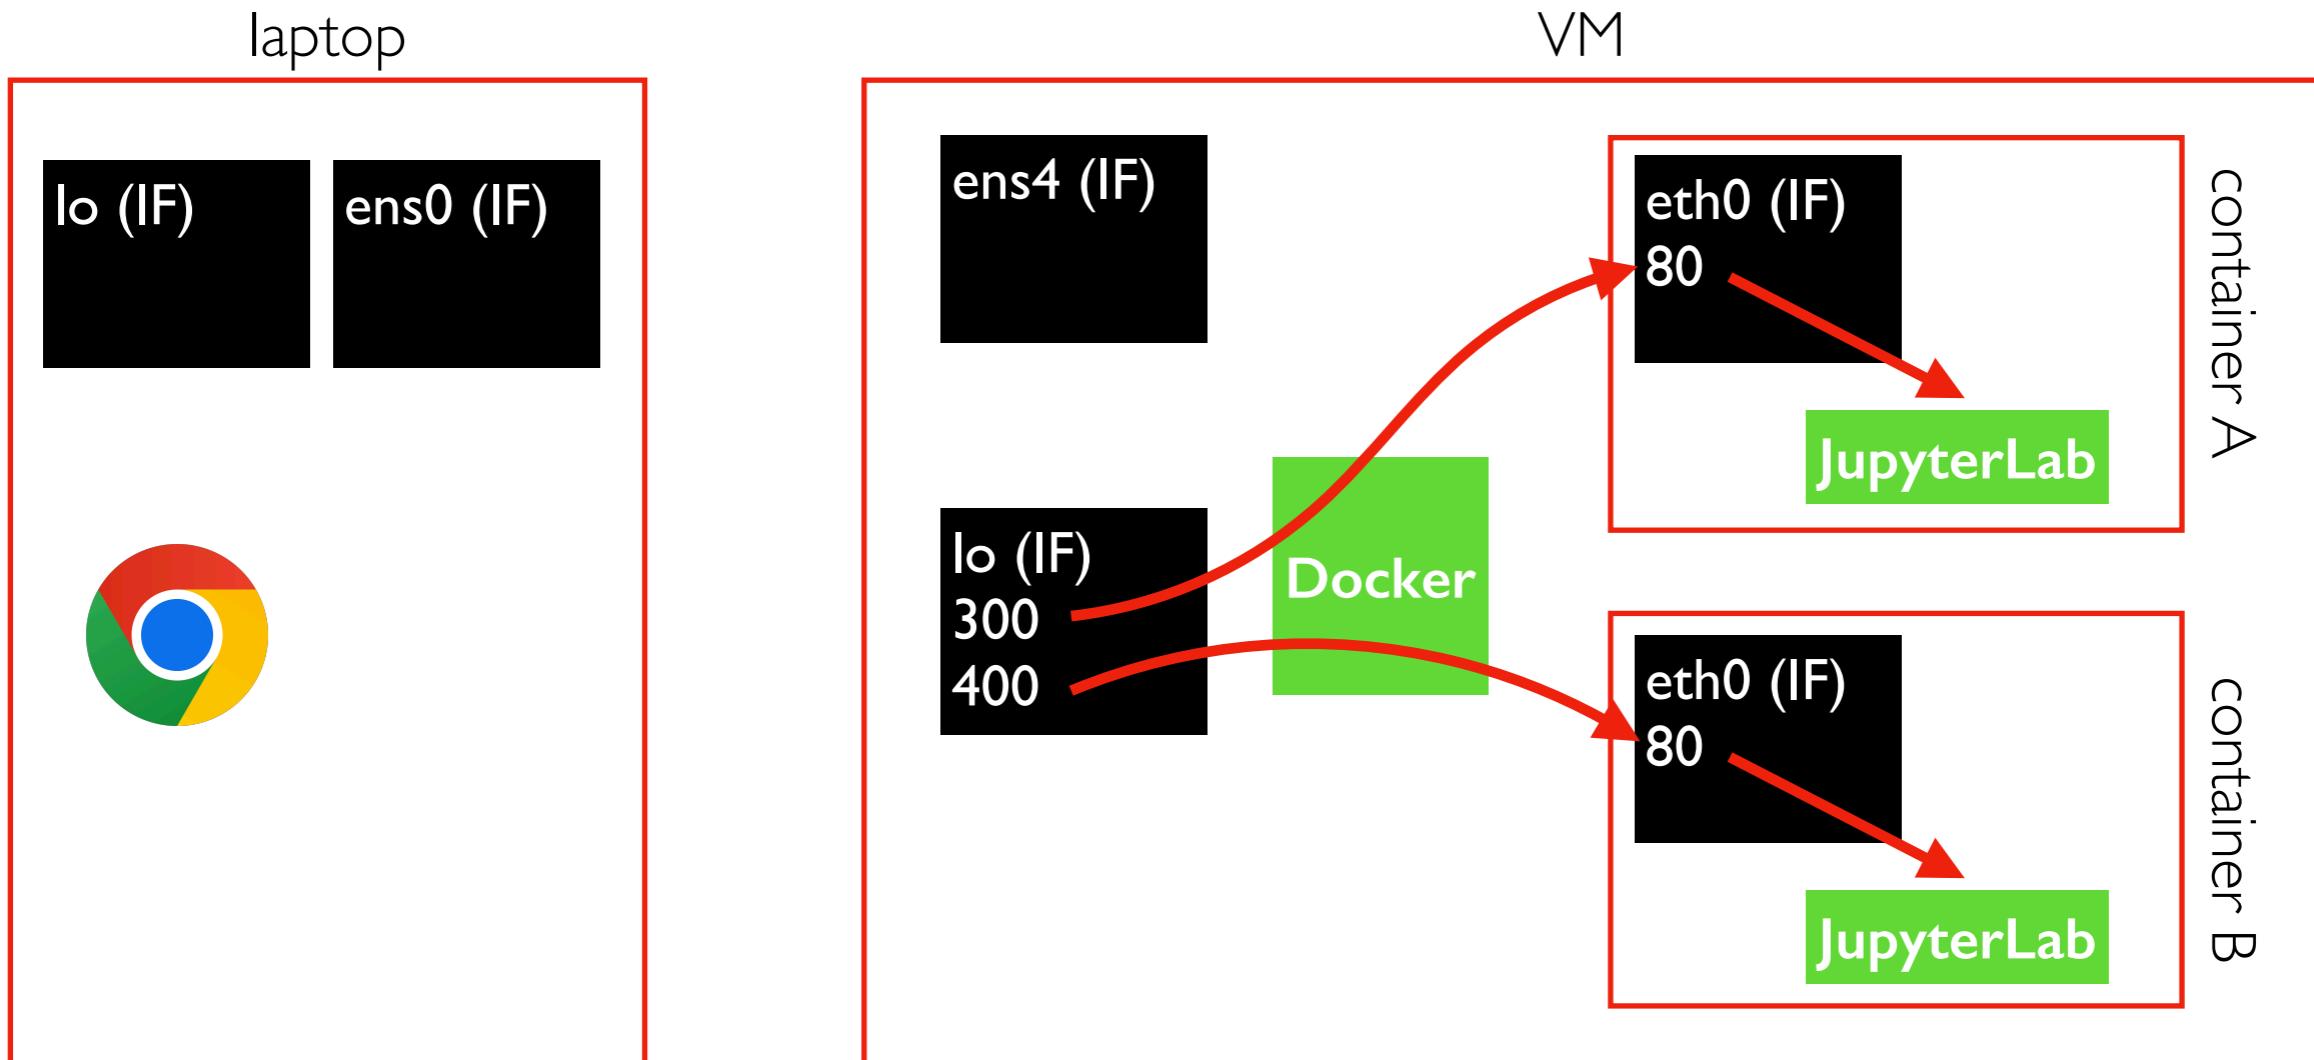
Interfaces (IF) and Ports

both containers have
a virtual port 80



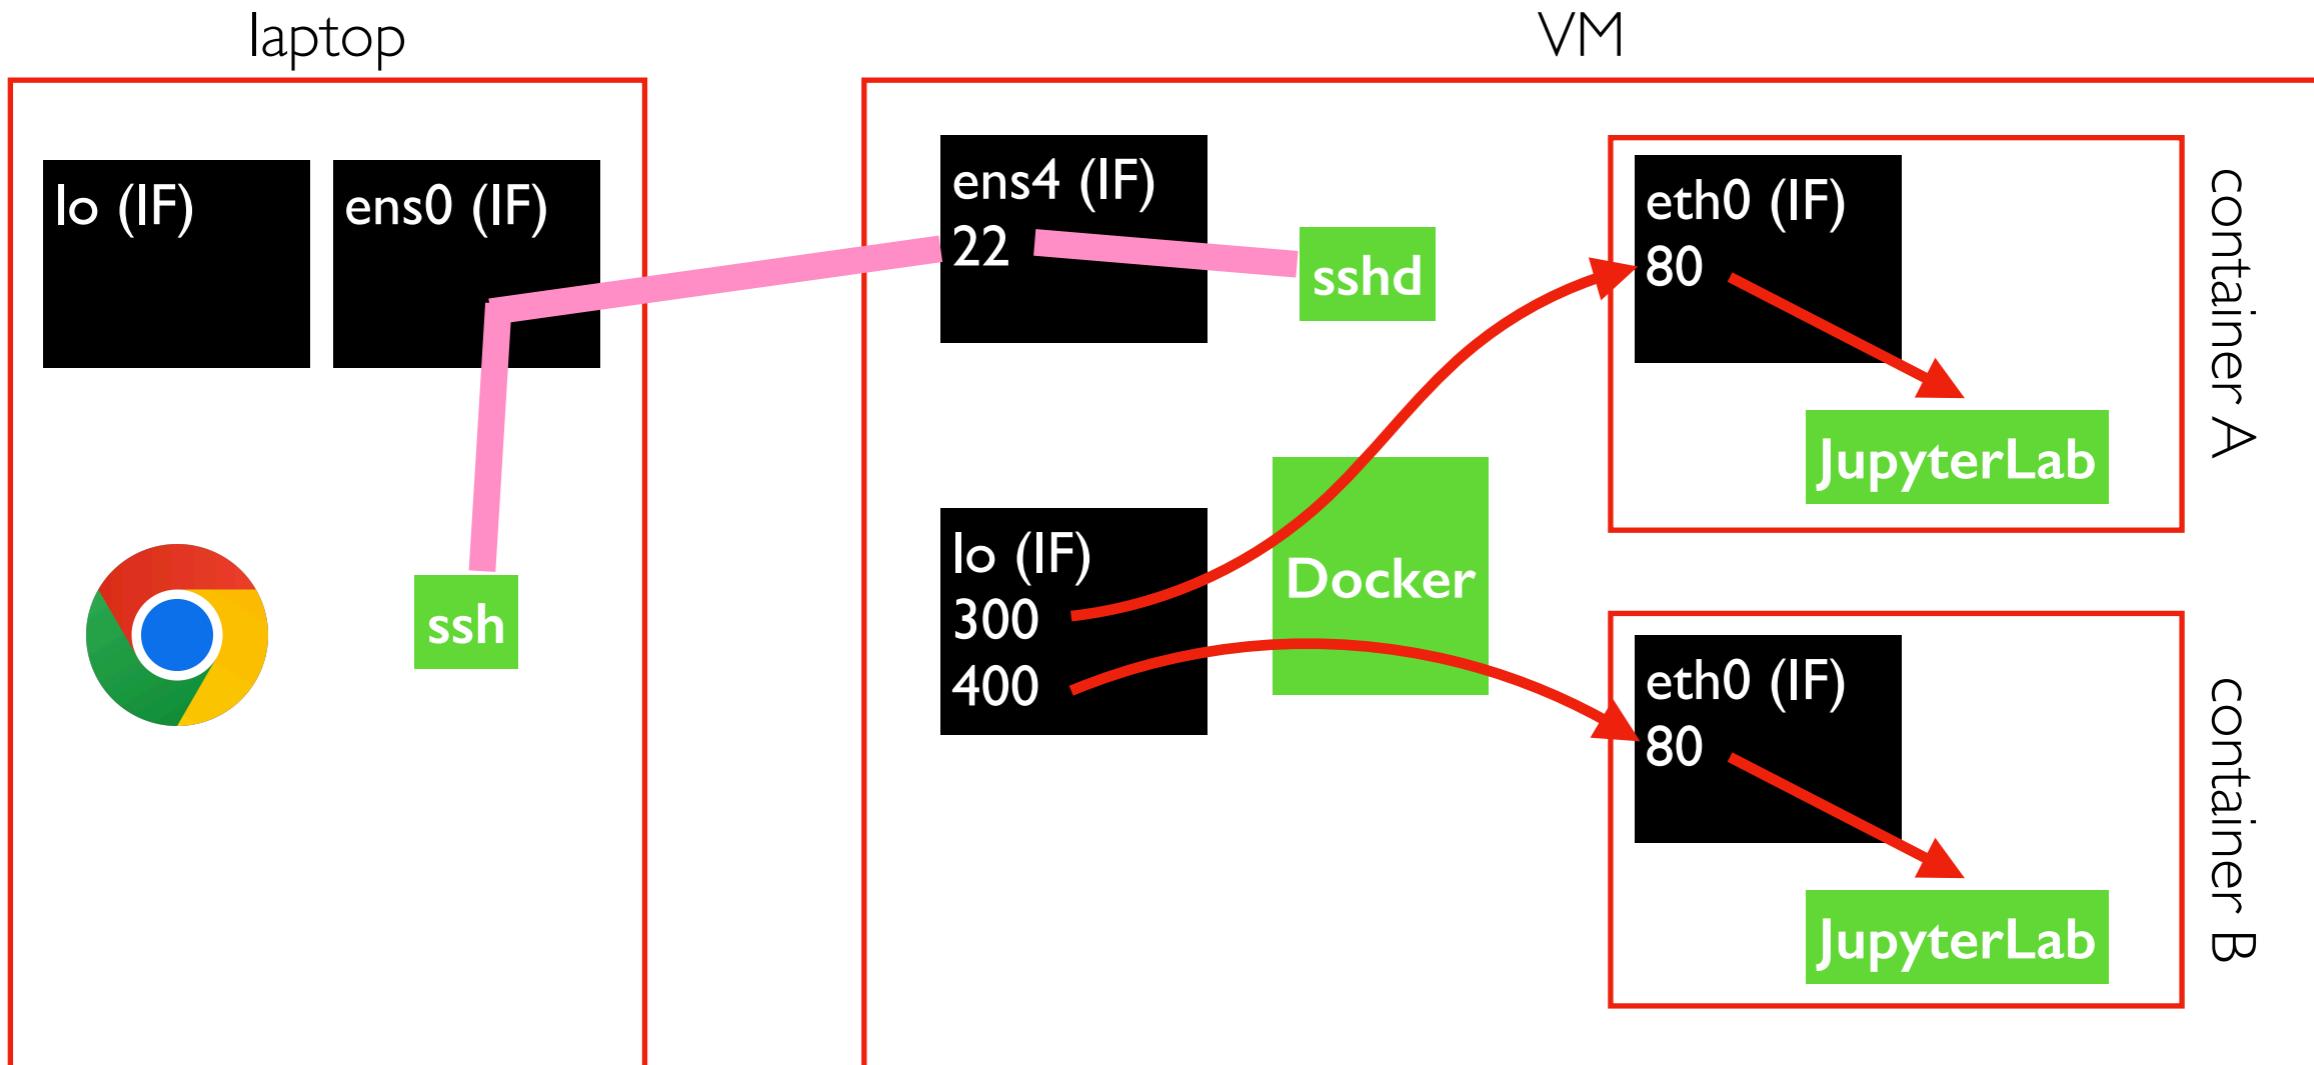
```
docker run -d myimg  
docker run -d myimg
```

Interfaces (IF) and Ports



```
docker run -d -p 127.0.0.1:300:80 myimg  
docker run -d -p 127.0.0.1:400:80 myimg
```

Interfaces (IF) and Ports

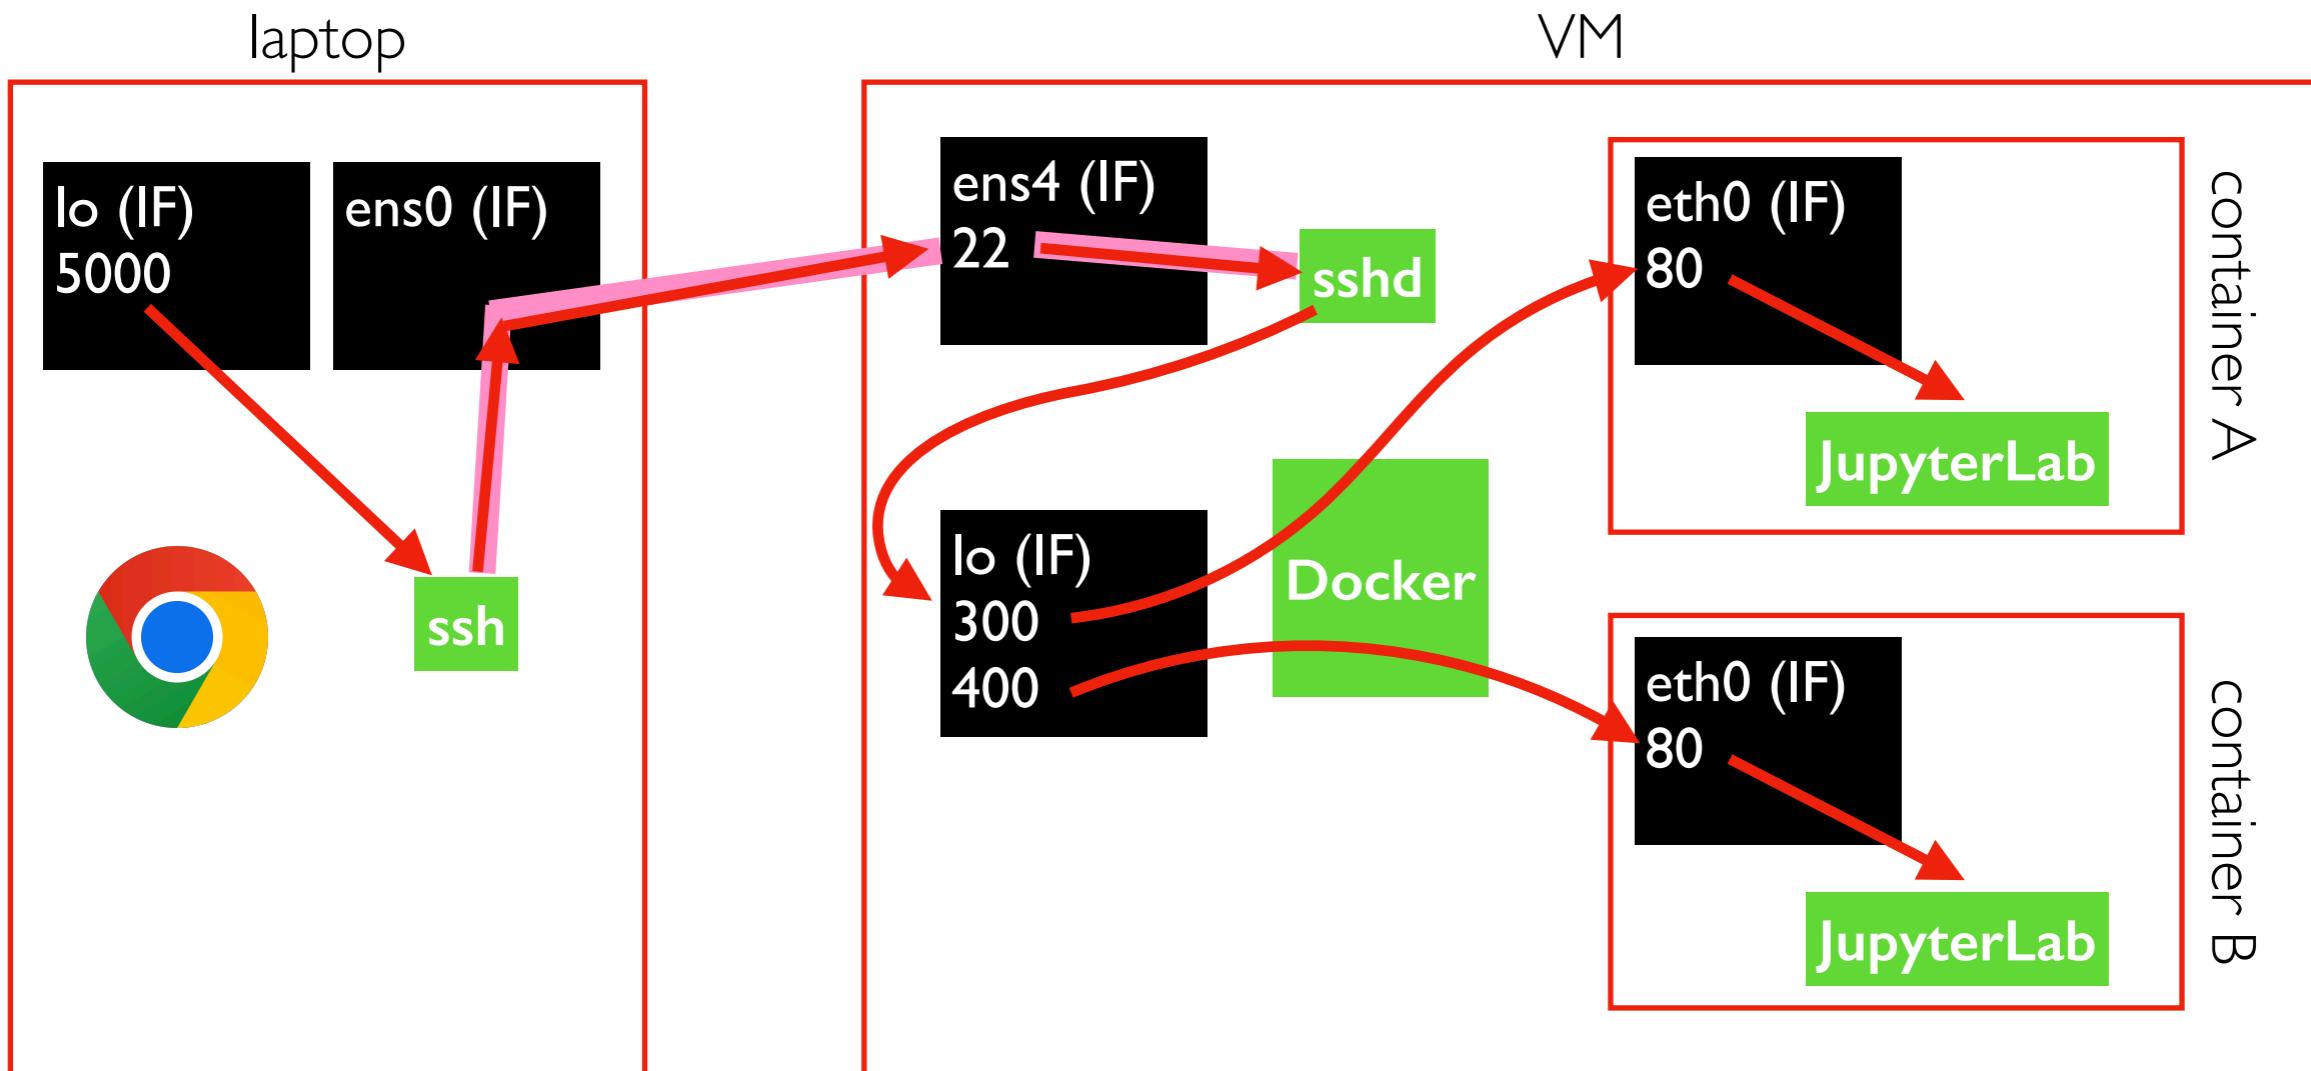


ssh USER@VM

```
docker run -d -p 127.0.0.1:300:80 myimg  
docker run -d -p 127.0.0.1:400:80 myimg
```

the SSH connection can be used to send commands and/or forward network traffic

Interfaces (IF) and Ports

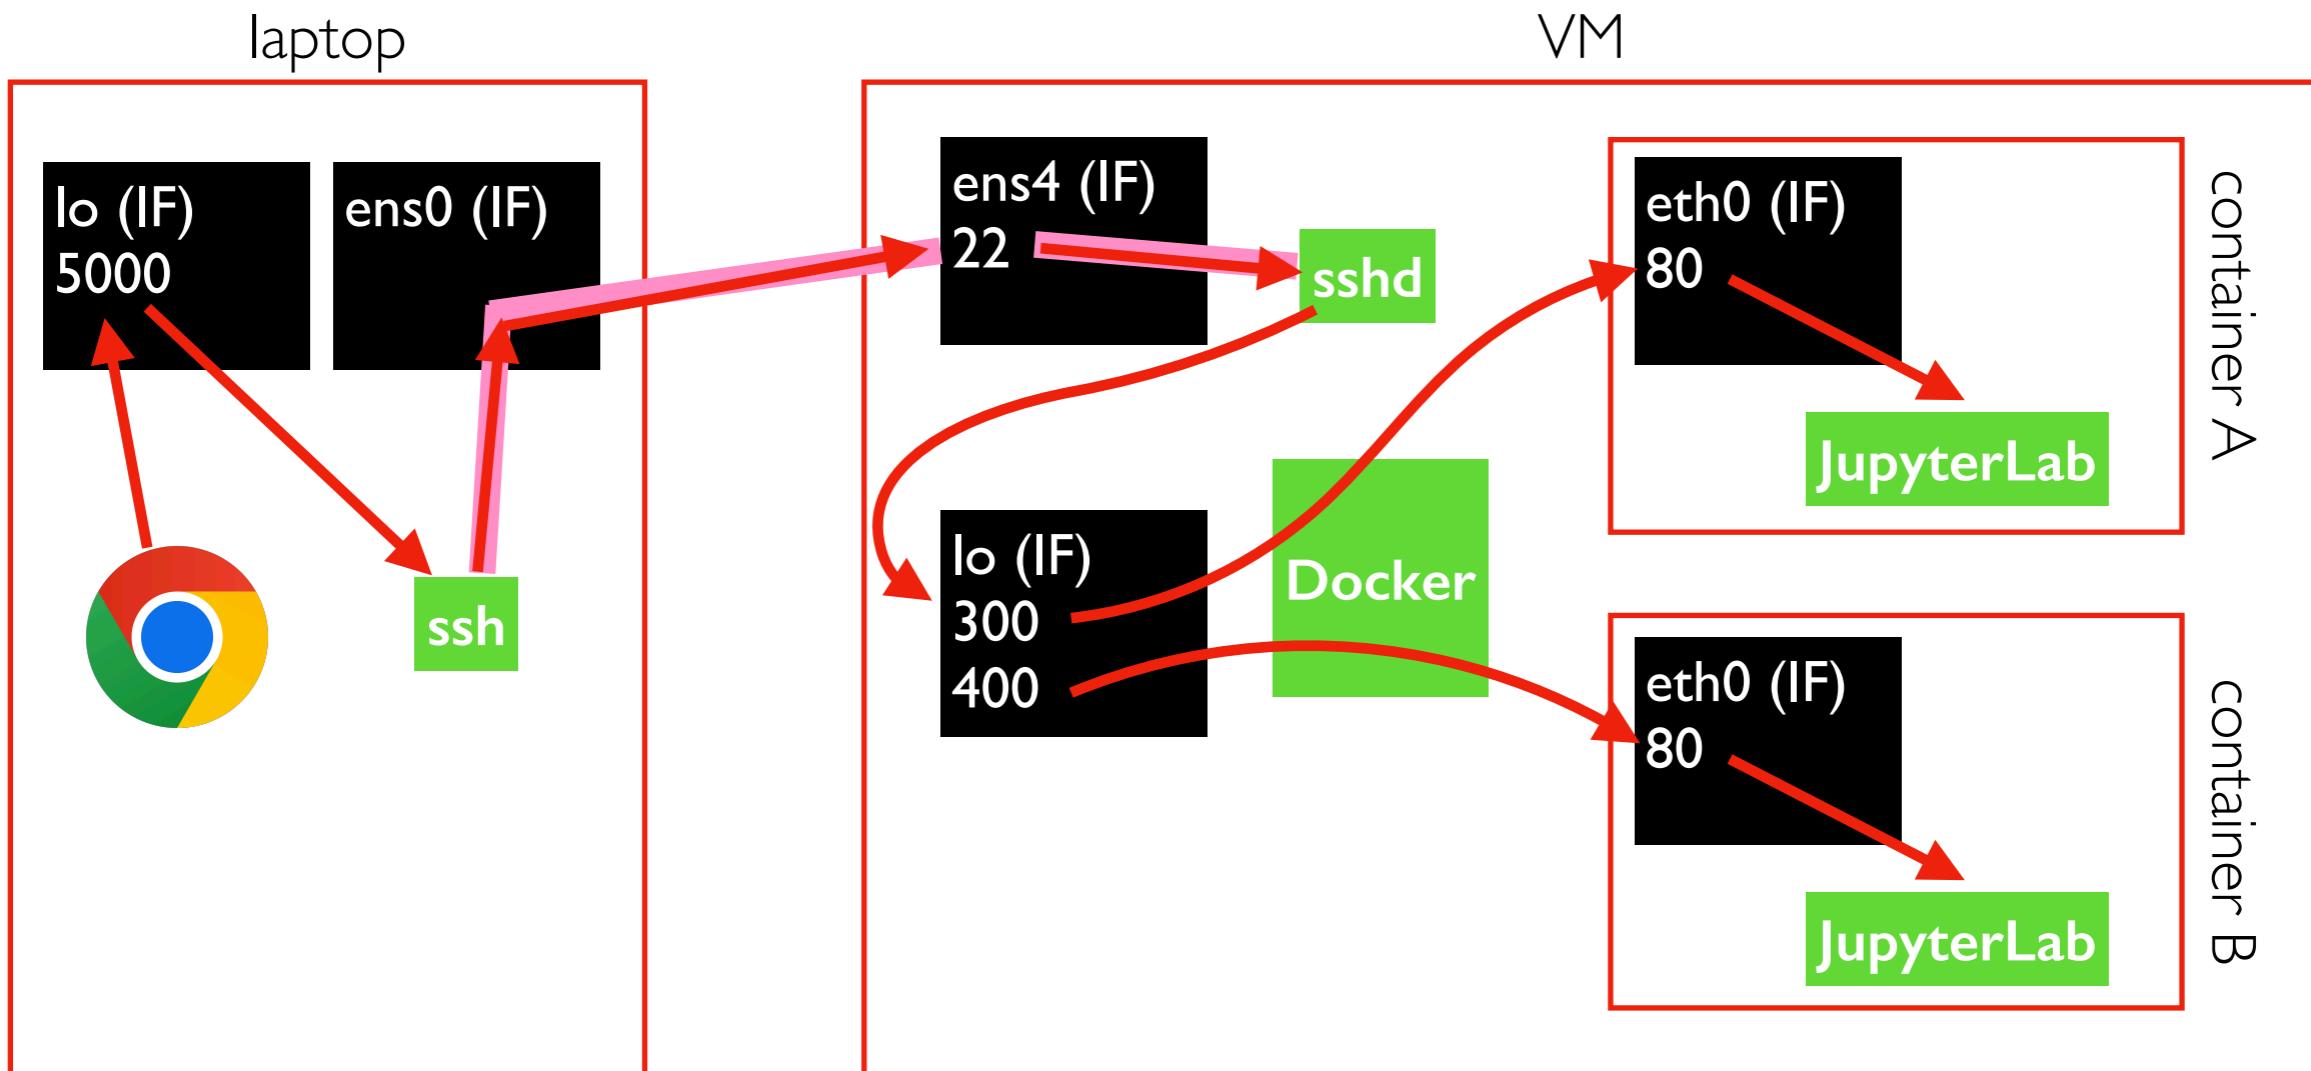


```
ssh USER@VM -L localhost:5000:localhost:300
```

```
docker run -d -p 127.0.0.1:300:80 myimg  
docker run -d -p 127.0.0.1:400:80 myimg
```

the SSH connection can be used to send commands and/or forward network traffic

Interfaces (IF) and Ports



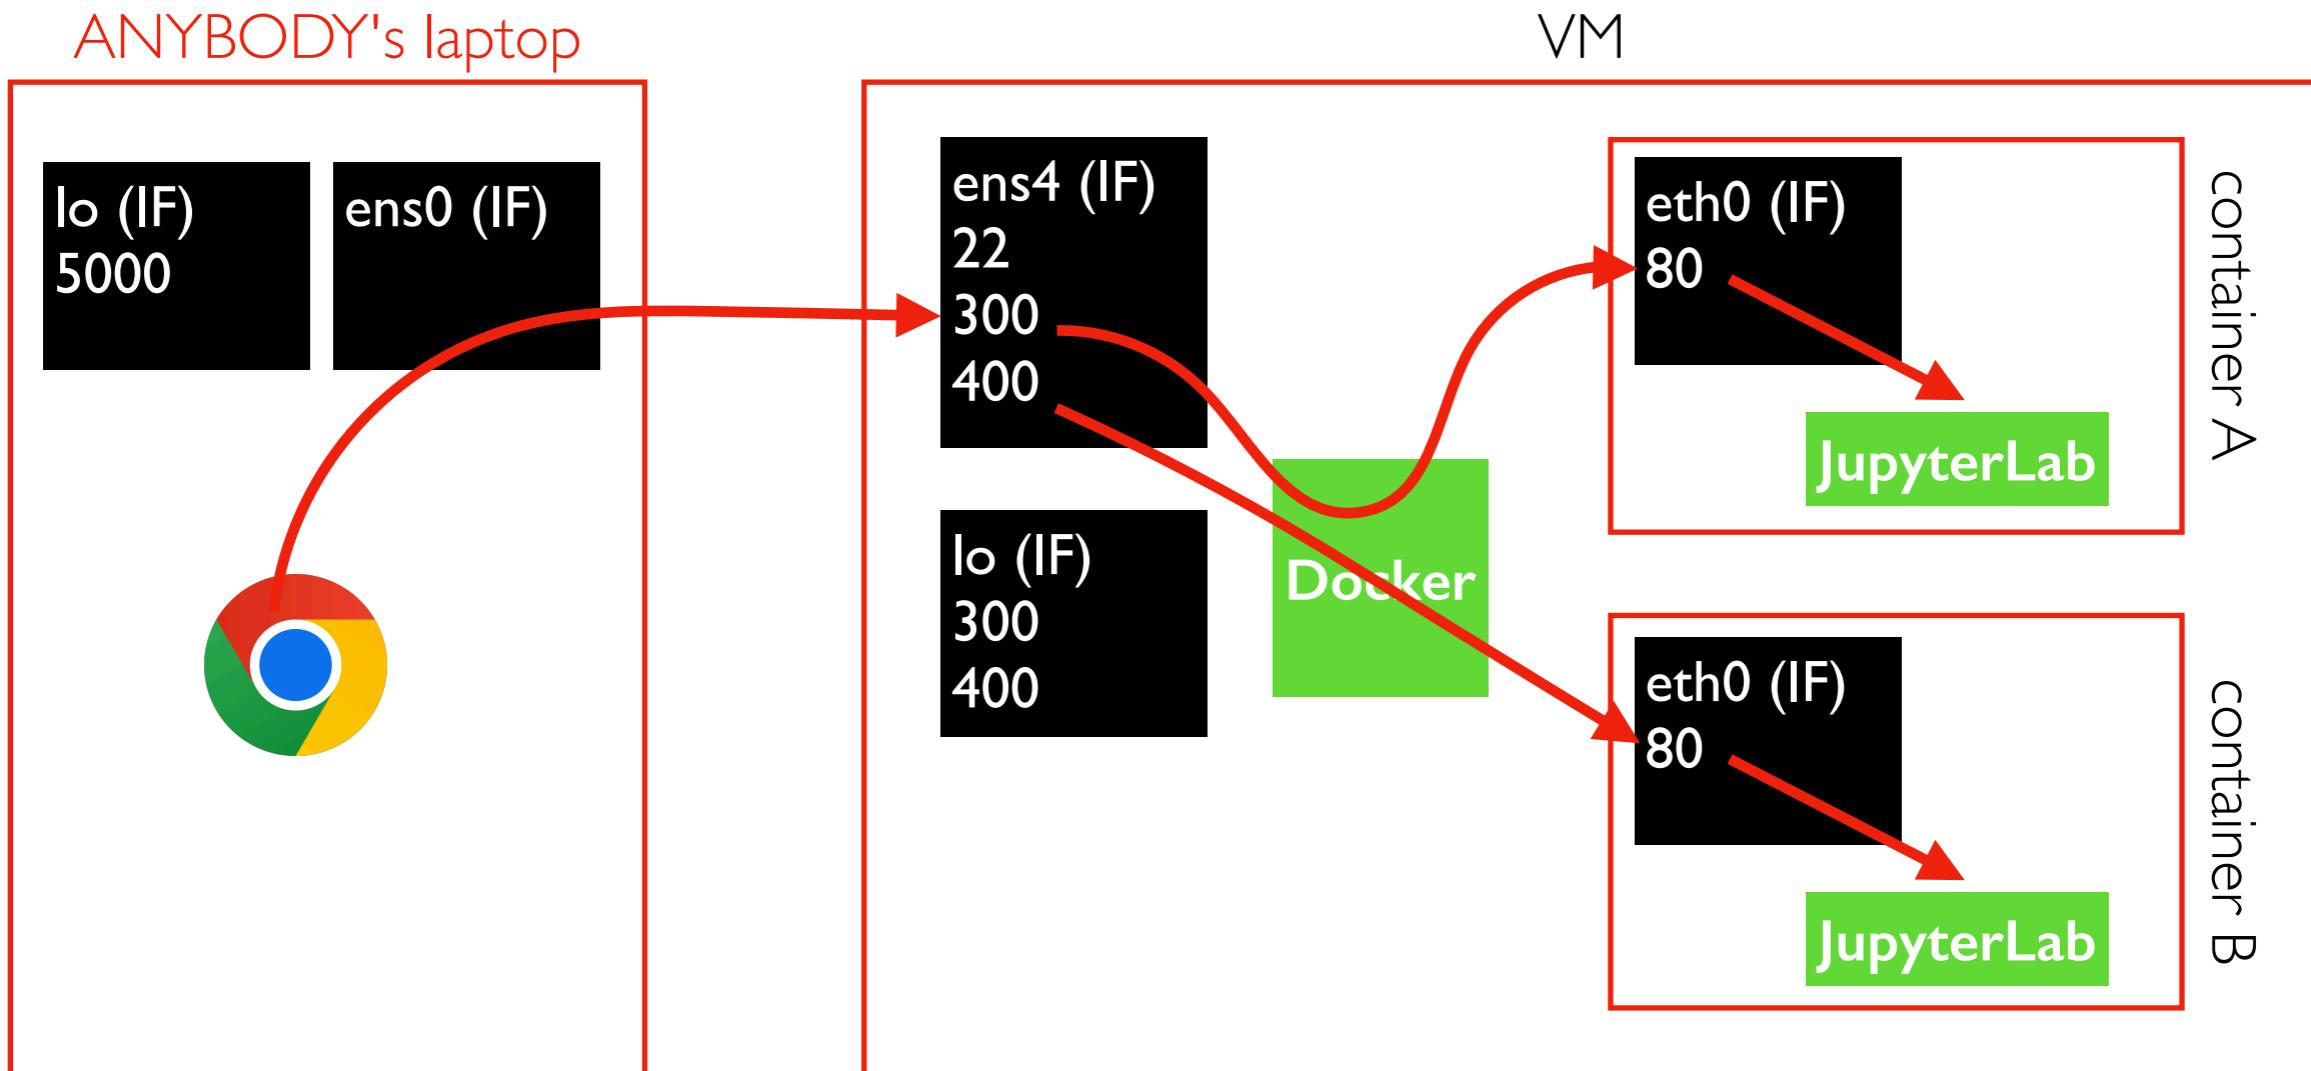
```
ssh USER@VM -L localhost:5000:localhost:300
```

<http://localhost:5000/lab> (in browser)

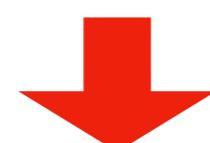
```
docker run -d -p 127.0.0.1:300:80 myimg  
docker run -d -p 127.0.0.1:400:80 myimg
```

yay! You can connect to JupyterLab
inside a container running on your VM

Interfaces (IF) and Ports



```
docker run -d -p 300:80 myimg
```



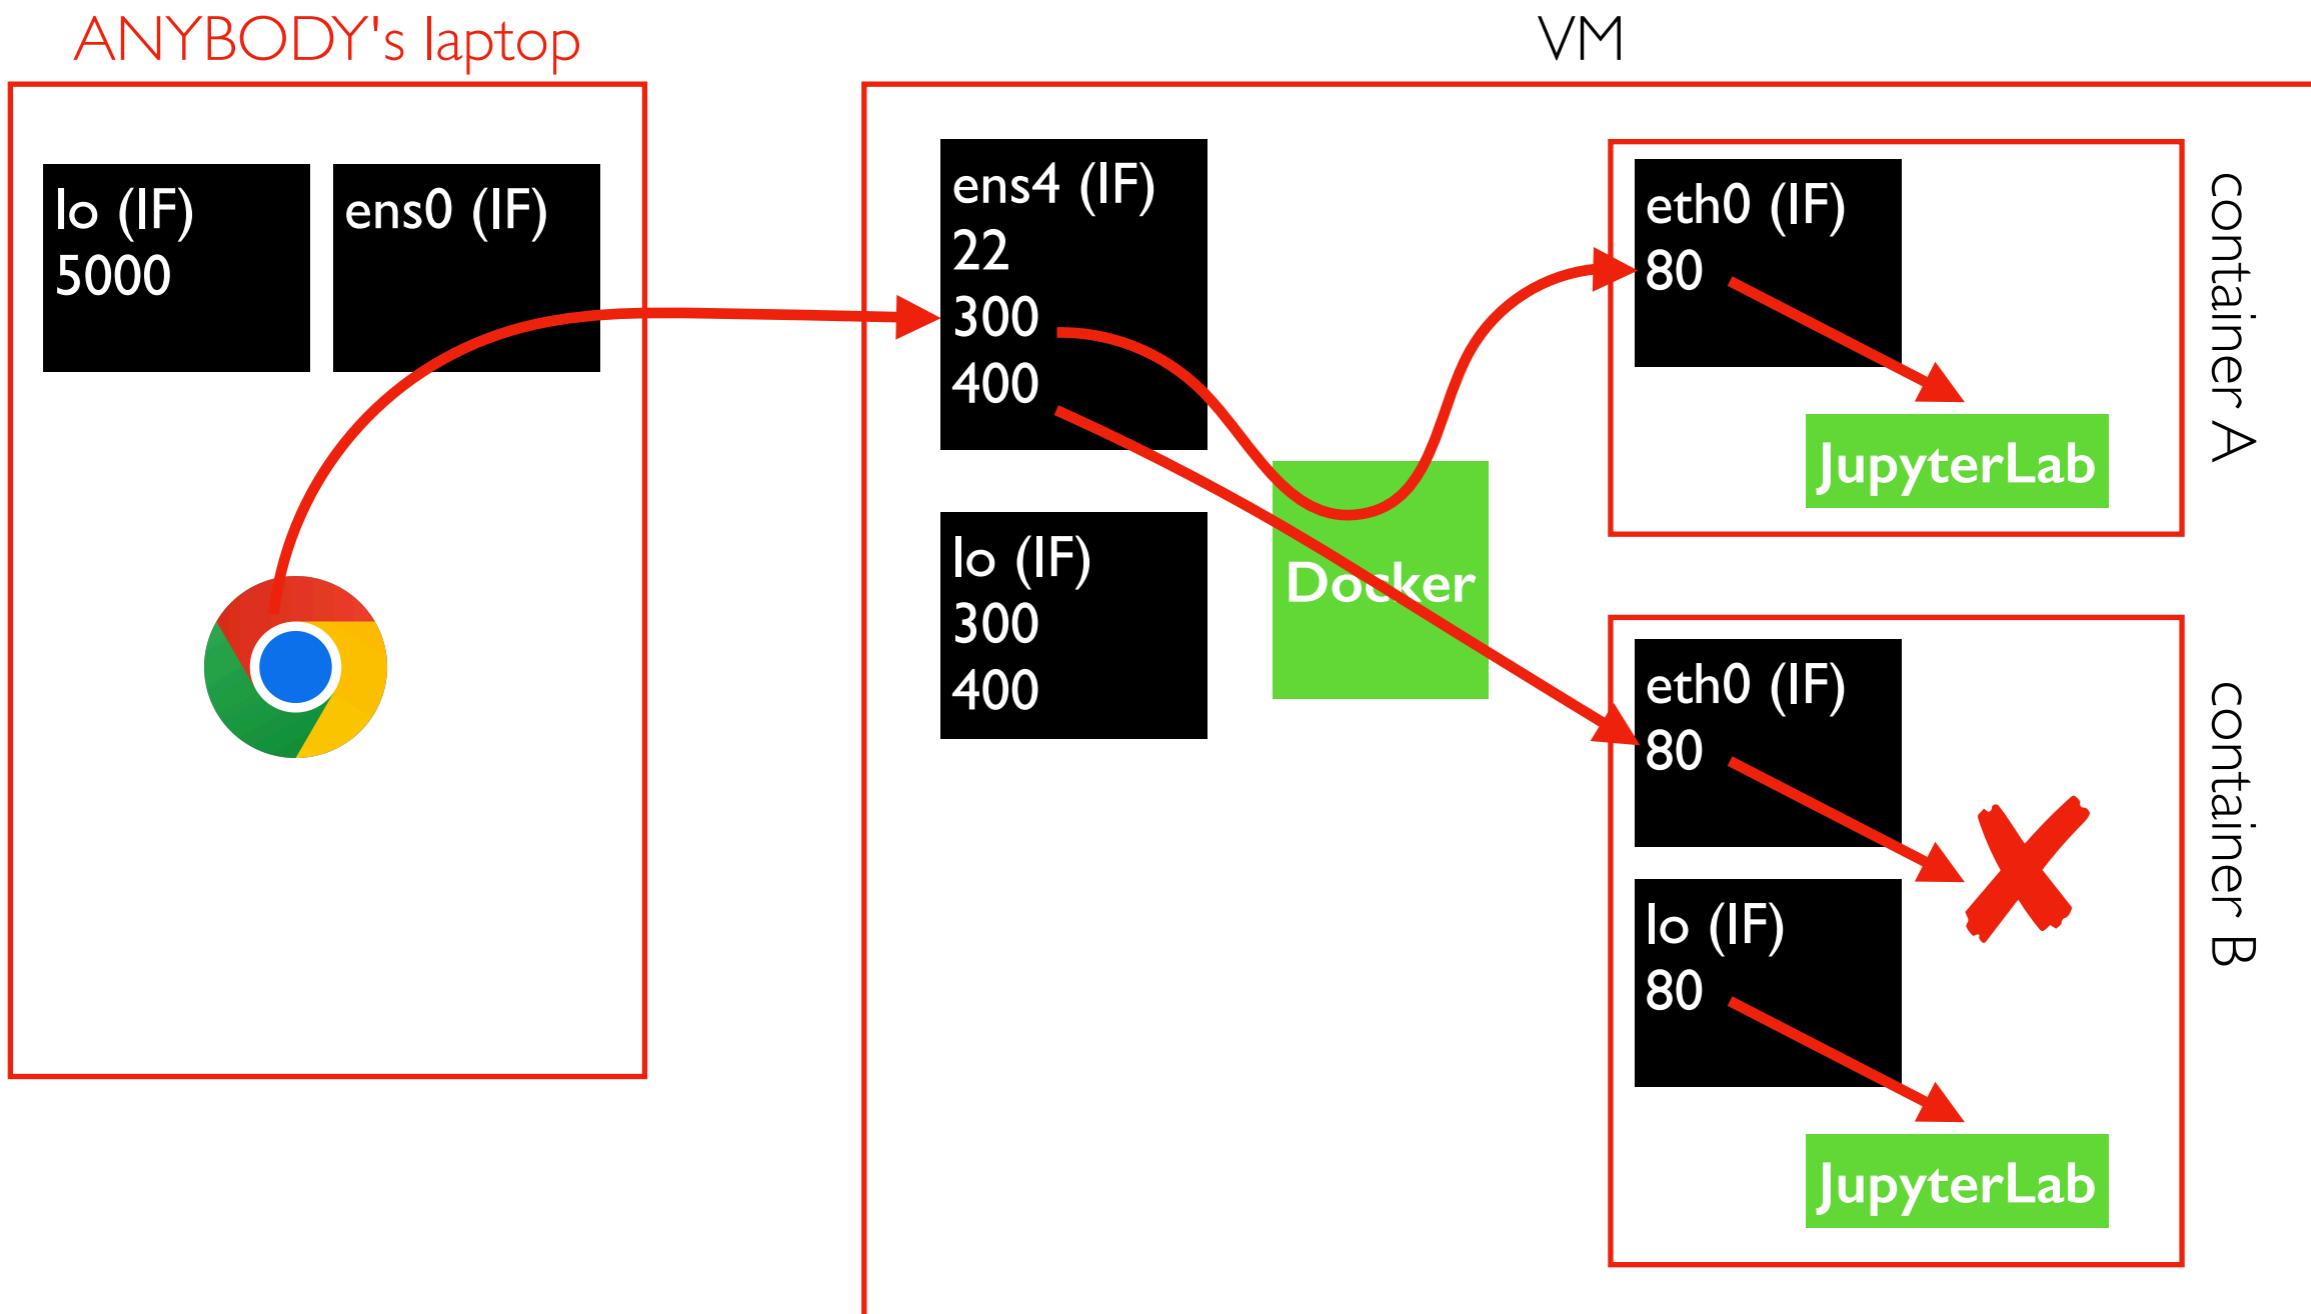
```
docker run -d -p 0.0.0.0:300:80 myimg
```

Careful, default is to listen on all NICs!

Other security options:

- firewall (block port 300)
- password (in JupyterLab)

Interfaces (IF) and Ports



Port forwarding never goes to loopback inside container

- don't use localhost or 127.0.0.1 inside container!
- easiest: use 0.0.0.0 inside container (for all) to port-forwarded traffic

Demo and TopHat...

Outline

Docker Port Forwarding

Docker Compose

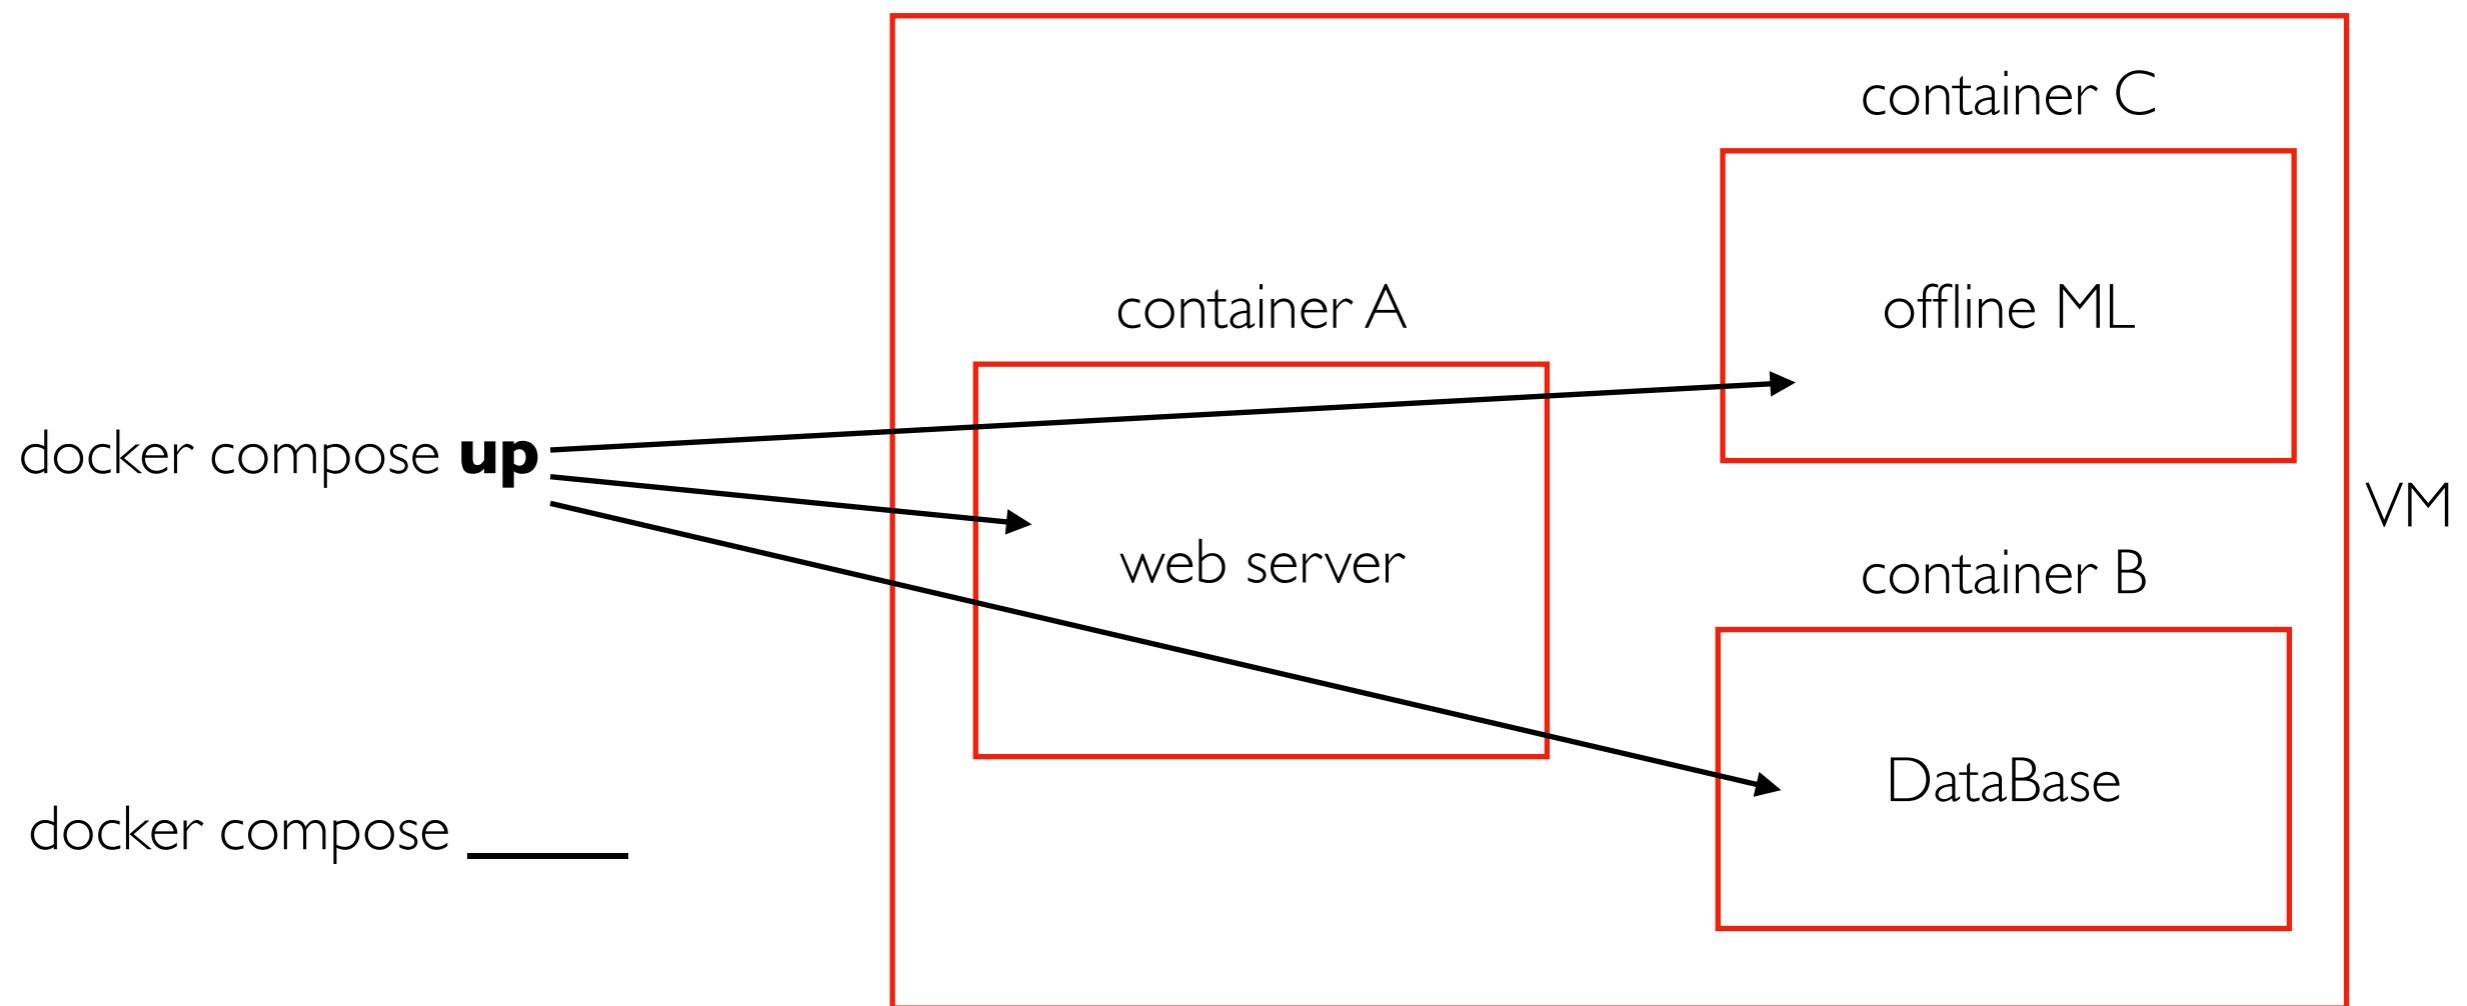
Partitioning and Replication

Container Orchestration

Orchestration lets you deploy many cooperating containers across a cluster of Docker workers.

Kubernetes (K8s) is the most well known.

Docker compose is a simpler tool that lets you deploy cooperating containers to a single worker.



Demos...

Outline

Docker Port Forwarding

Docker Compose

Partitioning and Replication

Data Placement

Say we have large dataset, and many machines.

- can we breakup the dataset (**partitioning**) so different machines can each help with part of it?
- should we have multiple copies (**replication**) of the same data so that we don't lose information if a machine fails?

Partitioning

Scenario: we have two computers, and want an app that lets instructors lookup student IDs by name.

dataset

| Name | Student ID |
|-----------------|------------|
| Aarav Patel | 9031231234 |
| Chen Wei | 8123456789 |
| Fatima Al-Farsi | 7234567890 |
| Hiroshi Tanaka | 6345678901 |
| Isabella Rossi | 5456789012 |
| John Smith | 4567890123 |
| Liam O'Connor | 3678901234 |
| Maria Garcia | 2789012345 |
| Nia Kofi | 1890123456 |
| Yuki Nakamura | 1001234567 |

computer 1



computer 2



Simple Partitioning

Scenario: we have two computers, and want an app that lets instructors lookup student IDs by name.

dataset

| Name | Student ID |
|-----------------|------------|
| Aarav Patel | 9031231234 |
| Chen Wei | 8123456789 |
| Fatima Al-Farsi | 7234567890 |
| Hiroshi Tanaka | 6345678901 |
| Isabella Rossi | 5456789012 |
| John Smith | 4567890123 |
| Liam O'Connor | 3678901234 |
| Maria Garcia | 2789012345 |
| Nia Kofi | 1890123456 |
| Yuki Nakamura | 1001234567 |

computer 1

| first half | |
|-----------------|------------|
| Name | Student ID |
| Aarav Patel | 9031231234 |
| Chen Wei | 8123456789 |
| Fatima Al-Farsi | 7234567890 |
| Hiroshi Tanaka | 6345678901 |
| Isabella Rossi | 5456789012 |

computer 2

| second half | |
|---------------|------------|
| Name | Student ID |
| John Smith | 4567890123 |
| Liam O'Connor | 3678901234 |
| Maria Garcia | 2789012345 |
| Nia Kofi | 1890123456 |
| Yuki Nakamura | 1001234567 |

Challenge: might not easily know which computer to "ask" for a given name (less efficient to ask both each time)

Range Partitioning

If we partition by range, we definitely know which compute to ask for a given name.

dataset

| Name | Student ID |
|-----------------|------------|
| Aarav Patel | 9031231234 |
| Chen Wei | 8123456789 |
| Fatima Al-Farsi | 7234567890 |
| Hiroshi Tanaka | 6345678901 |
| Isabella Rossi | 5456789012 |
| John Smith | 4567890123 |
| Liam O'Connor | 3678901234 |
| Maria Garcia | 2789012345 |
| Nia Kofi | 1890123456 |
| Yuki Nakamura | 1001234567 |

computer 1

| A-M | |
|-----------------|------------|
| Name | Student ID |
| Aarav Patel | 9031231234 |
| Chen Wei | 8123456789 |
| Fatima Al-Farsi | 7234567890 |
| Hiroshi Tanaka | 6345678901 |
| Isabella Rossi | 5456789012 |
| John Smith | 4567890123 |
| Liam O'Connor | 3678901234 |
| Maria Garcia | 2789012345 |

computer 2

| N-Z | |
|---------------|------------|
| Name | Student ID |
| Nia Kofi | 1890123456 |
| Yuki Nakamura | 1001234567 |

Challenge: it might be hard to find good split points, especially if the dataset is changing.

Hash Partitioning

First, choose key column, then hash it. A hash function returns a seemingly arbitrary number for any input, but the same input always produces the same number.

dataset

key

| Name | Student ID | hash(Name) |
|----------------|------------|------------|
| Aarav Patel | 9031231234 | 360993 |
| Chen Wei | 8123456789 | 70525 |
| Fatima Al- | 7234567890 | 913591 |
| Hiroshi | 6345678901 | 121696 |
| Isabella Rossi | 5456789012 | 258452 |
| John Smith | 4567890123 | 438815 |
| Liam | 3678901234 | 588279 |
| Maria Garcia | 2789012345 | 388236 |
| Nia Kofi | 1890123456 | 679776 |
| Yuki | 1001234567 | 160849 |

Challenge: not good if you want to do lookup
for all names in an alphabetic range

computer 1

even hash

| Name | Student ID | hash(Name) |
|----------|------------|------------|
| Hiroshi | 6345678901 | 121696 |
| Isabella | 5456789012 | 258452 |
| Maria | 2789012345 | 388236 |
| Nia Kofi | 1890123456 | 679776 |

computer 2

odd hash

| Name | Student ID | hash(Name) |
|------------|------------|------------|
| Aarav | 9031231234 | 360993 |
| Chen Wei | 8123456789 | 70525 |
| Fatima Al- | 7234567890 | 913591 |
| John | 4567890123 | 438815 |
| Liam | 3678901234 | 588279 |
| Yuki | 1001234567 | 160849 |

Partitioning Vocabulary

TERMINOLOGICAL CONFUSION

What we call a *partition* here is called a *shard* in MongoDB, Elasticsearch, and SolrCloud; it's known as a *region* in HBase, a *tablet* in Bigtable, a *vnode* in Cassandra and Riak, and a *vBucket* in Couchbase. However, *partitioning* is the most established term, so we'll stick with that.

Chapter 6. Partitioning



Replication

Sometimes, we want multiple copies (called replicas) on different computers so we don't lose data when a machine dies.

Replication and partitioning can be used together or independently.

Example of replicated files in HDFS (later lecture...)

F1: "ABCD"

F2: "EFGHIJKL"

3x replication

2x replication

Computers

"ABCD" (F1.1)

"EFGH" (F2.1)

"ABCD" (F1.1)

"EFGH" (F2.1)

"IJKL" (F2.2)

"IJKL" (F2.2)

"ABCD" (F1.1)