

[544] Caching

Tyler Caraza-Harter

Learning Objectives

- describe the cache hierarchy
- trace through access patterns with LRU and FIFO policies
- calculate cache performance metrics

Outline

Challenge: Latency

Cache Hierarchy

- CPU, RAM, SSD, Disk, Network
- Tradeoffs

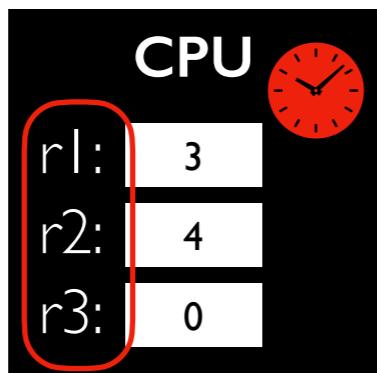
Policy: what data should be cached?

- manual
- expiration
- eviction policies: random, FIFO, LRU

Practice

CPU and RAM

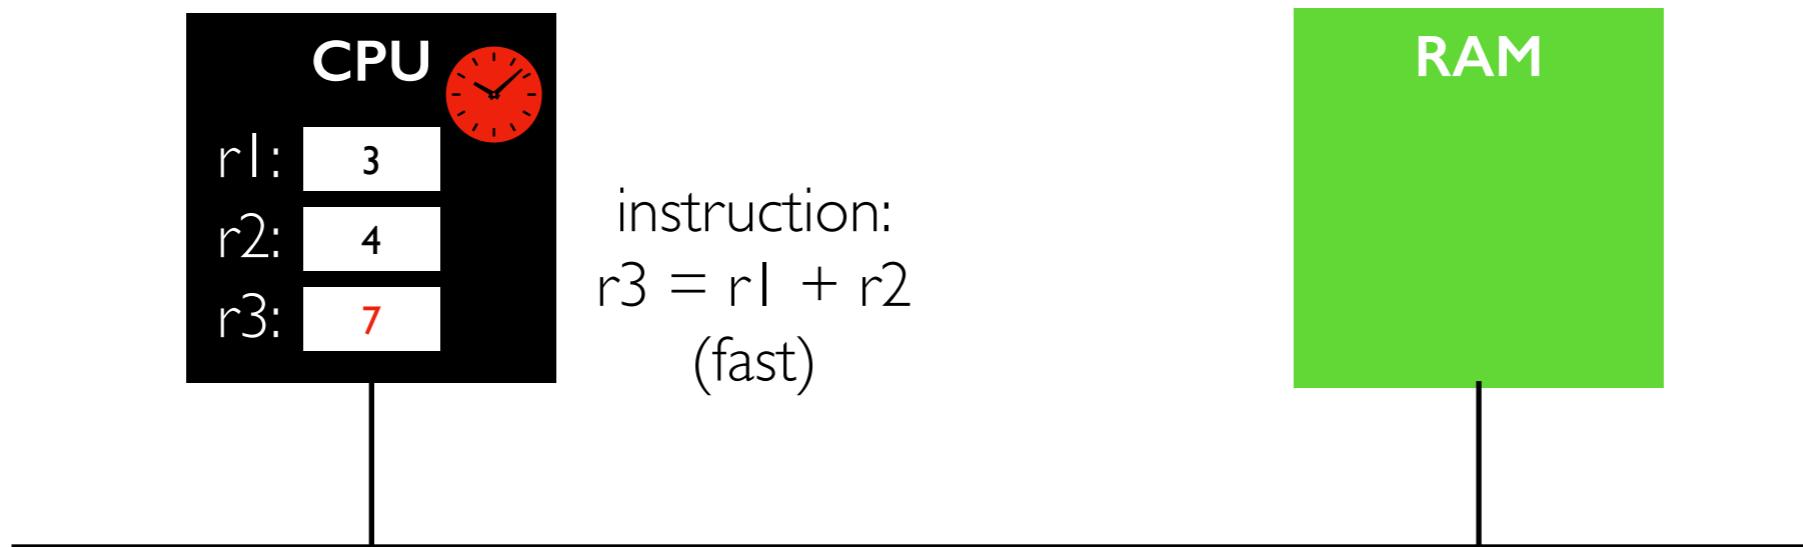
registers are
like variables
built into the CPU



CPU clock has
billions of cycles
per second



CPU and RAM



Load and Store

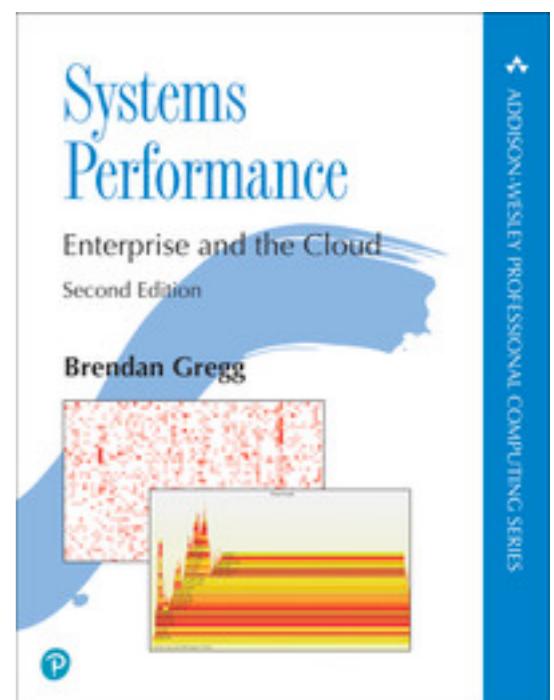


challenge: if we want to add some numbers stored in RAM, we need to **load** before adding and **store** after

Latency

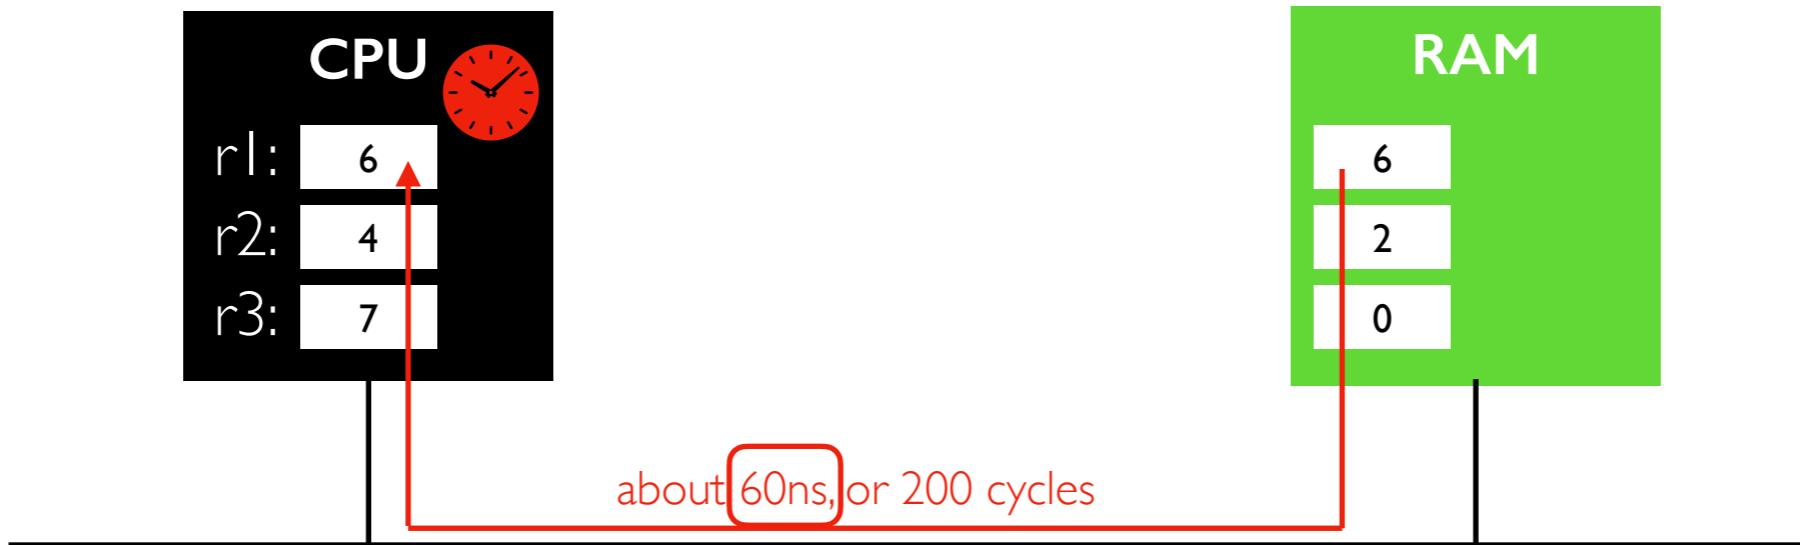


very slow, but not long enough to switch to a different process...

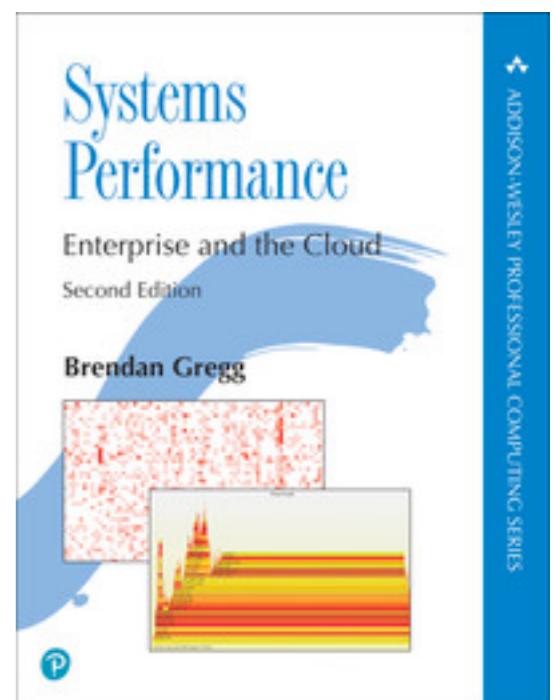


source: visuals, estimates

Latency



"how much time" is a **latency** measure.
Throughput (bytes/second) would depend on how many loads like these we can do simultaneously.

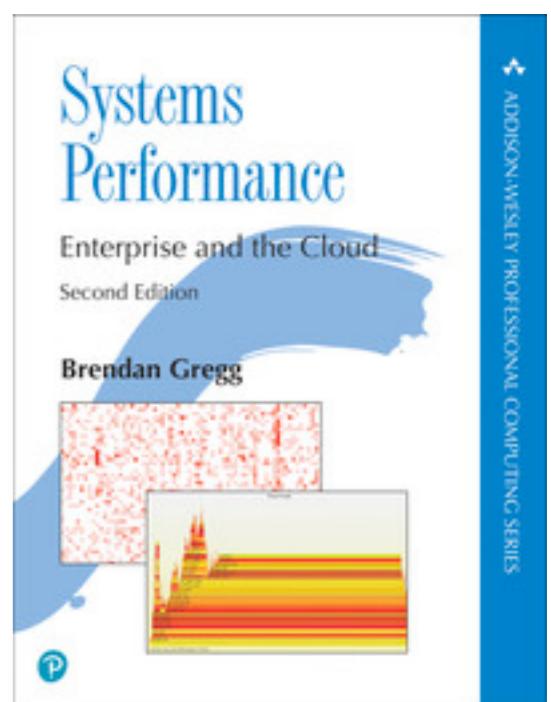


source: visuals, estimates

Cache



Idea: CPUs can have a small/fast memory built in for data that is accessed frequently



source: visuals, estimates

Performance Measurements

Metrics

- throughput
- average or median latency
- "tail" latency
 - for example, 99th percentile, 99.9th percentile, etc.
(abbreviated p99 or p99.9)

Which metrics do we expect caching to help with the most?

Outline

Challenge: Latency

Cache Hierarchy

- CPU, RAM, SSD, Disk, Network
- Tradeoffs

Policy: what data should be cached?

- manual
- expiration
- eviction policies: random, FIFO, LRU

Practice

Cache Hierarchy

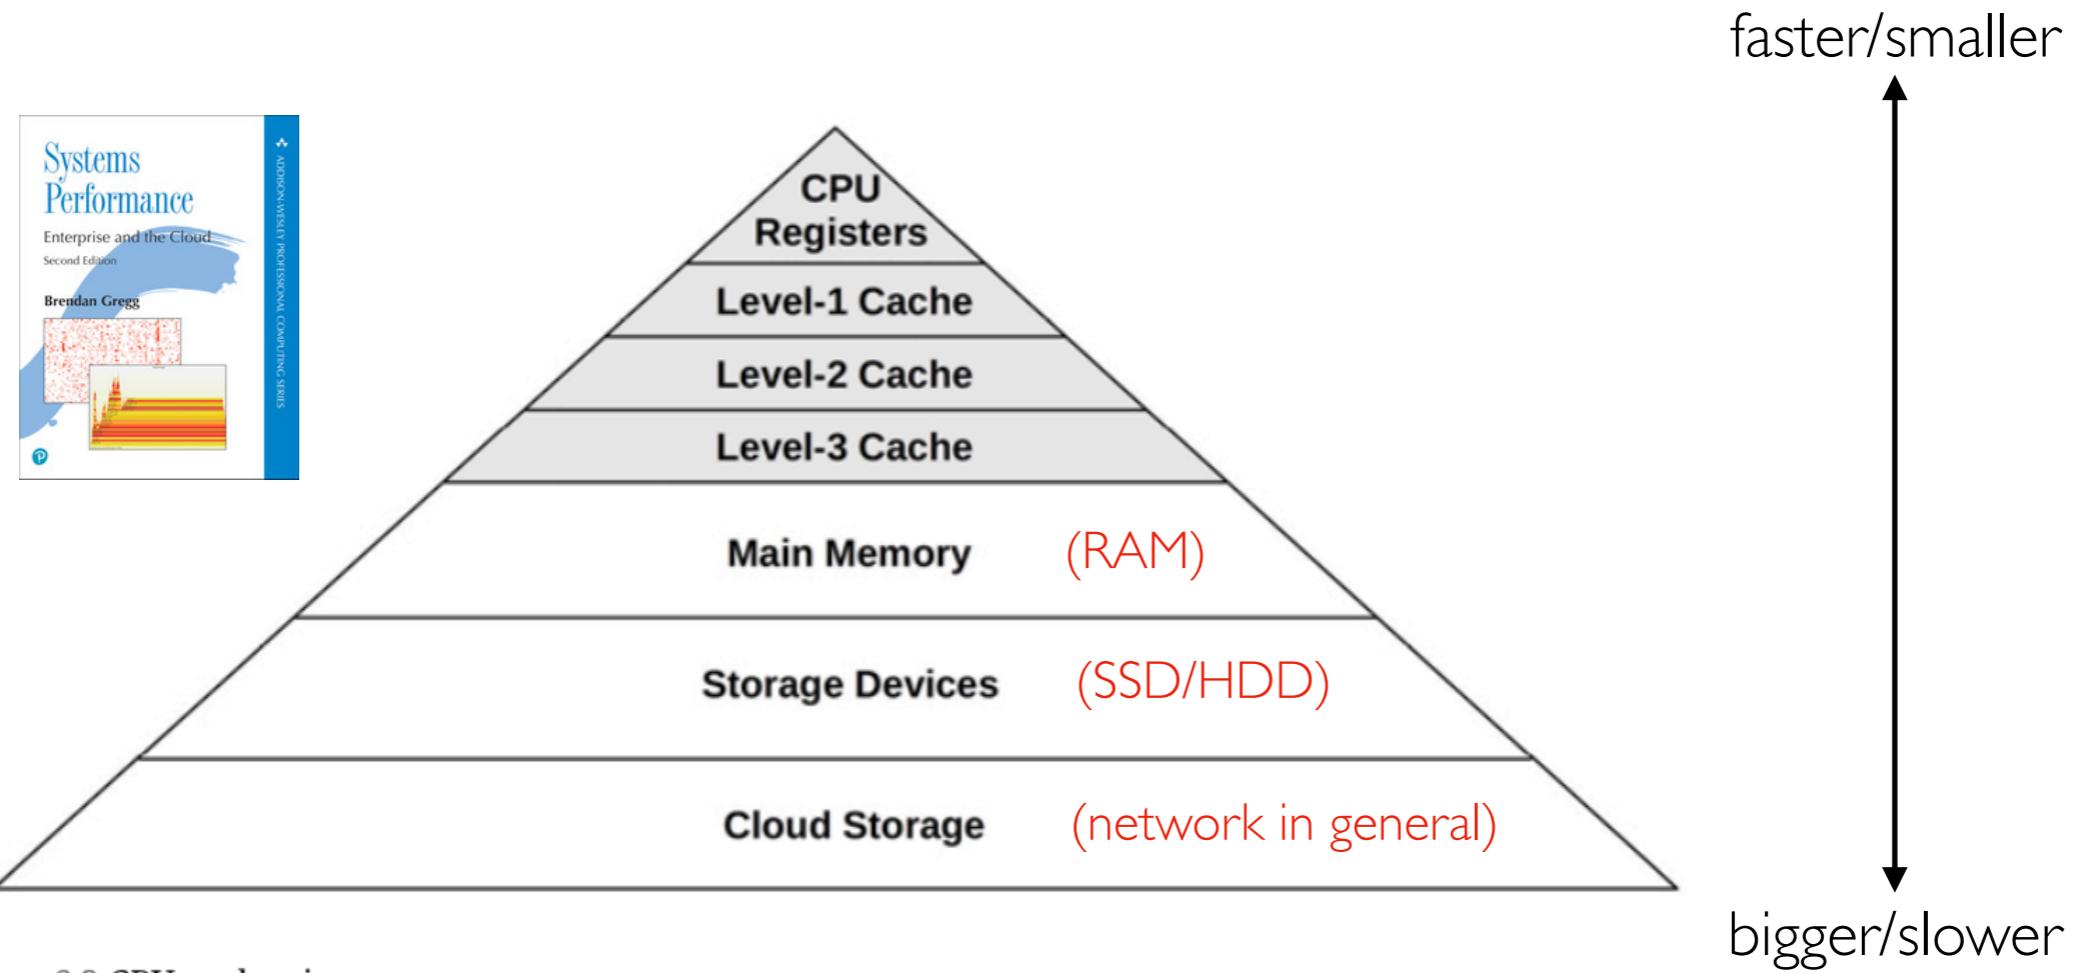


Figure 6.2 CPU cache sizes

Example: Intel Xeon Platinum 9282 (2019)

- L1: 64 KB
- L2: 1 MB
- L3: 77 MB

Cache Hierarchy

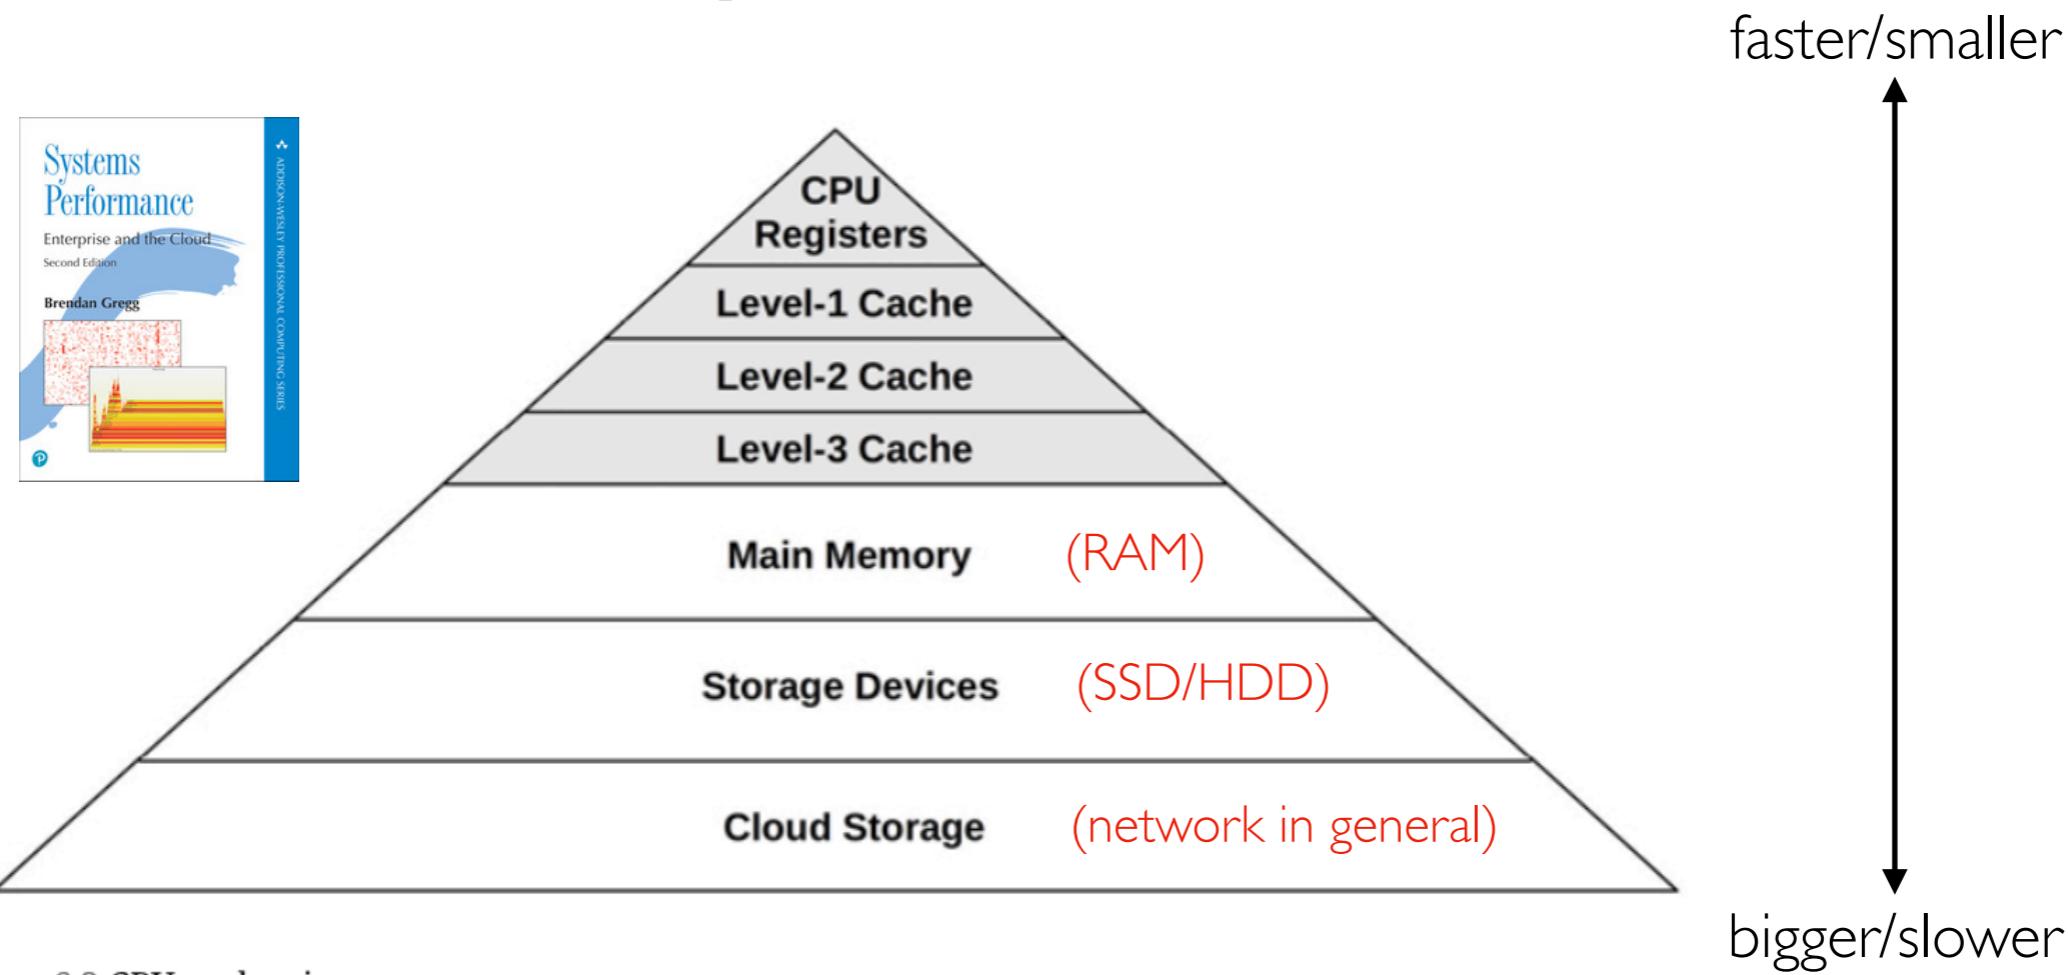
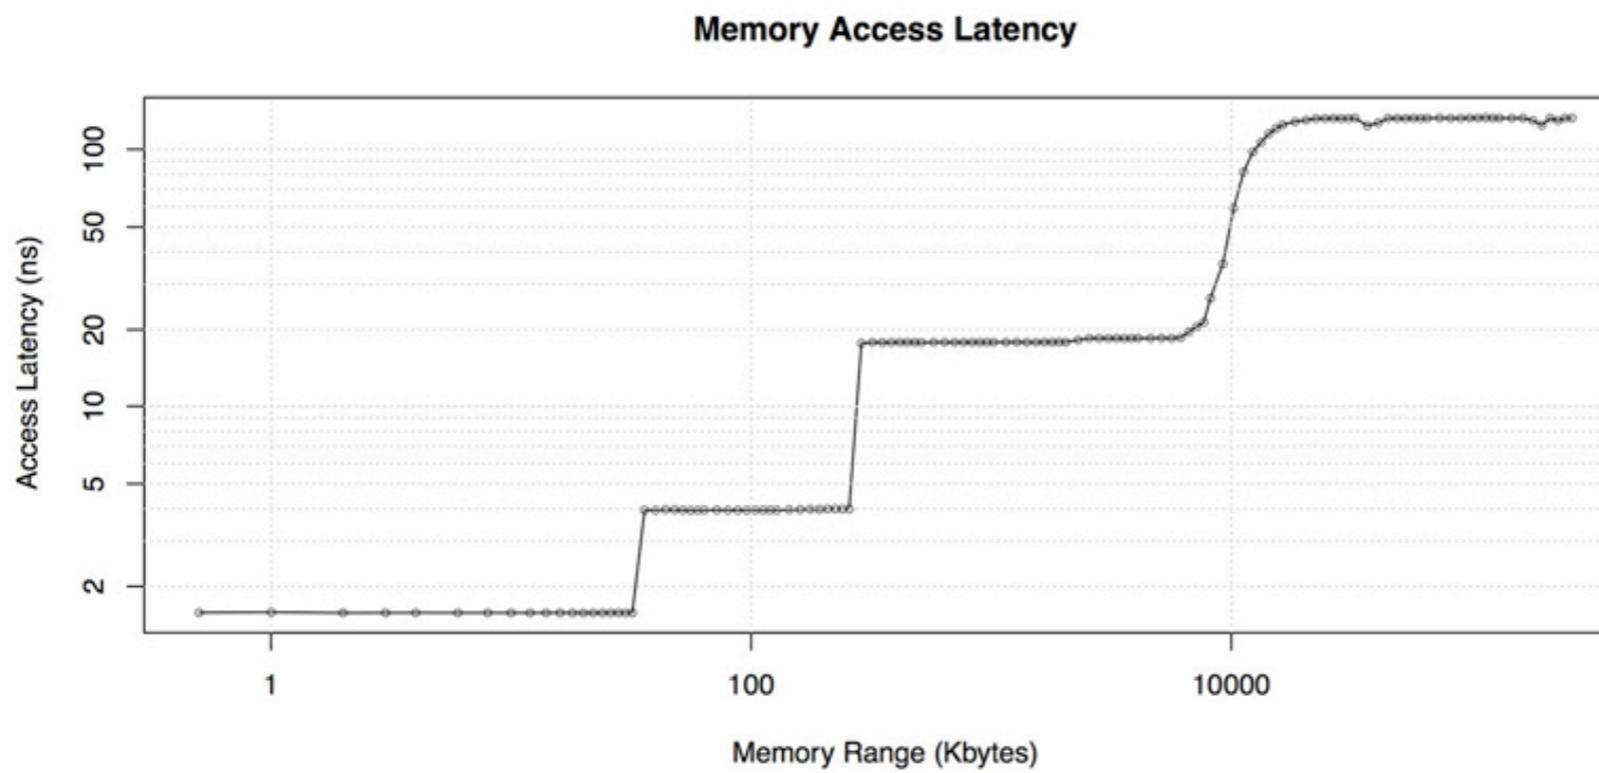


Figure 6.2 CPU cache sizes



about **how big** is the L3 cache?
what is the **latency** for an L3 cache access?

Resource Tradeoffs

Operating system caches file data in RAM

- use **memory**
- avoid **storage** reads

Browser caches recently visited page as file

- uses **storage** space/reads
- avoid **network** transfers

Python dictionary caches return values in a dict (key=args, val{return})

- uses **memory** space
- avoid repeated **compute**

```
cache = {}  
def f(x):  
    if not x in cache:  
        cache[x] = g(x)  
    return cache[x]
```

Outline

Challenge: Latency

Cache Hierarchy

- CPU, RAM, SSD, Disk, Network
- Tradeoffs

Policy: what data should be cached?

- manual
- expiration
- eviction policies: random, FIFO, LRU

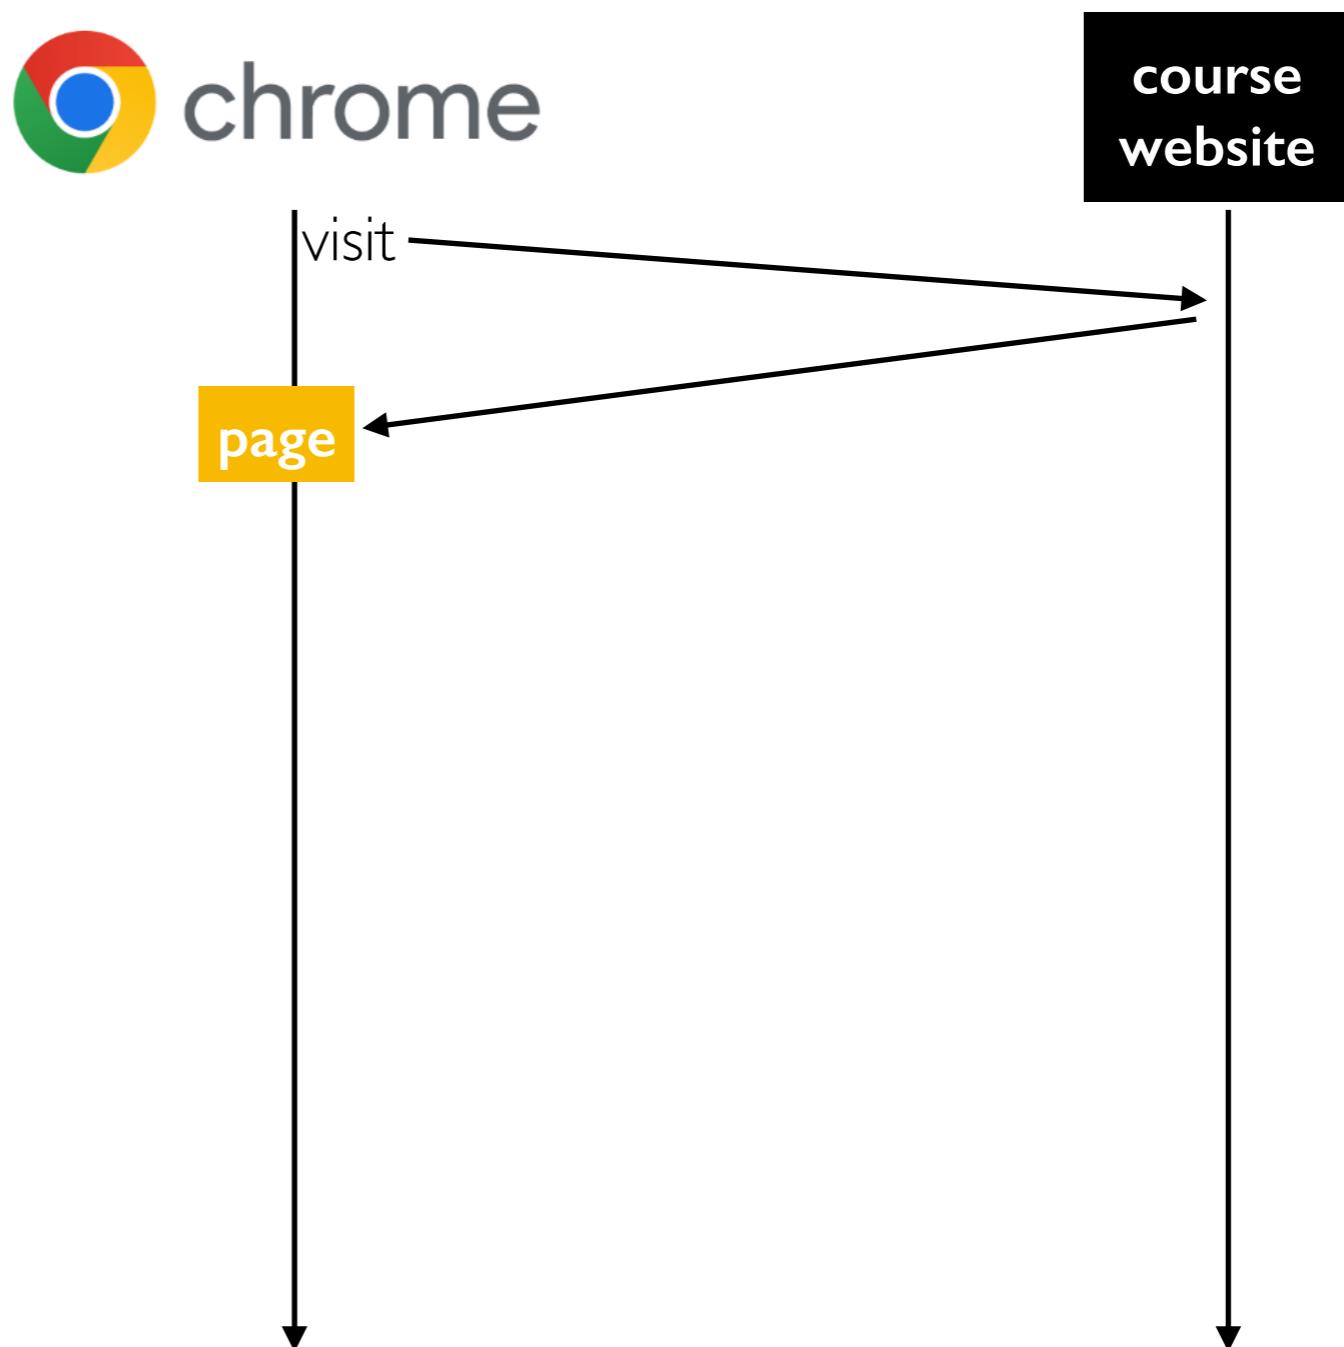
Practice

Manual Caching: Spark Example

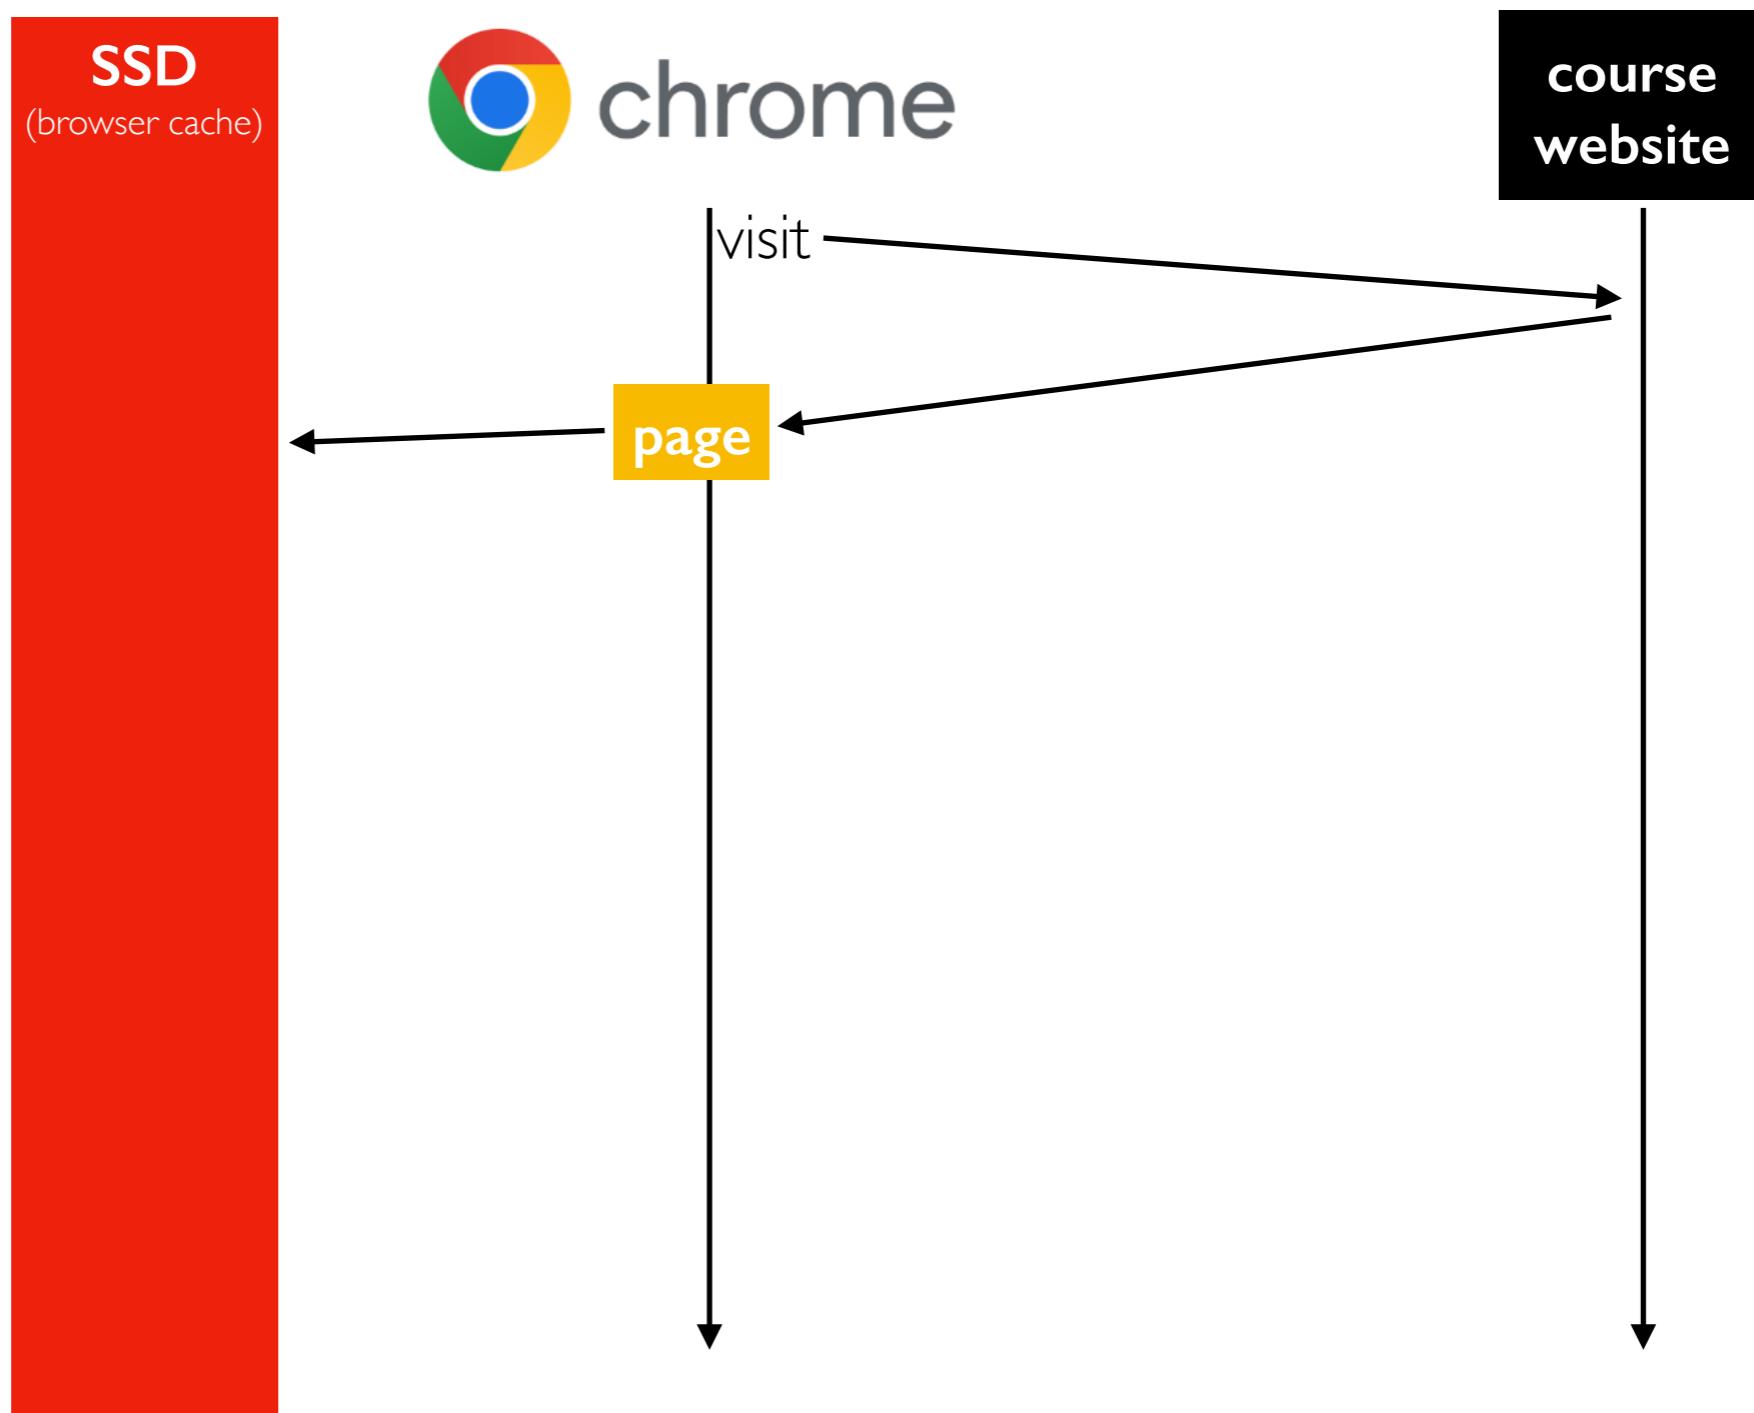
```
spark_df = ??? # not usually in memory  
  
spark_df.cache() # put it in memory  
  
# use spark_df for a lot of calculations  
  
spark_df.unpersist() # free up memory
```

we'll be spending lots of time on Spark later in the semester

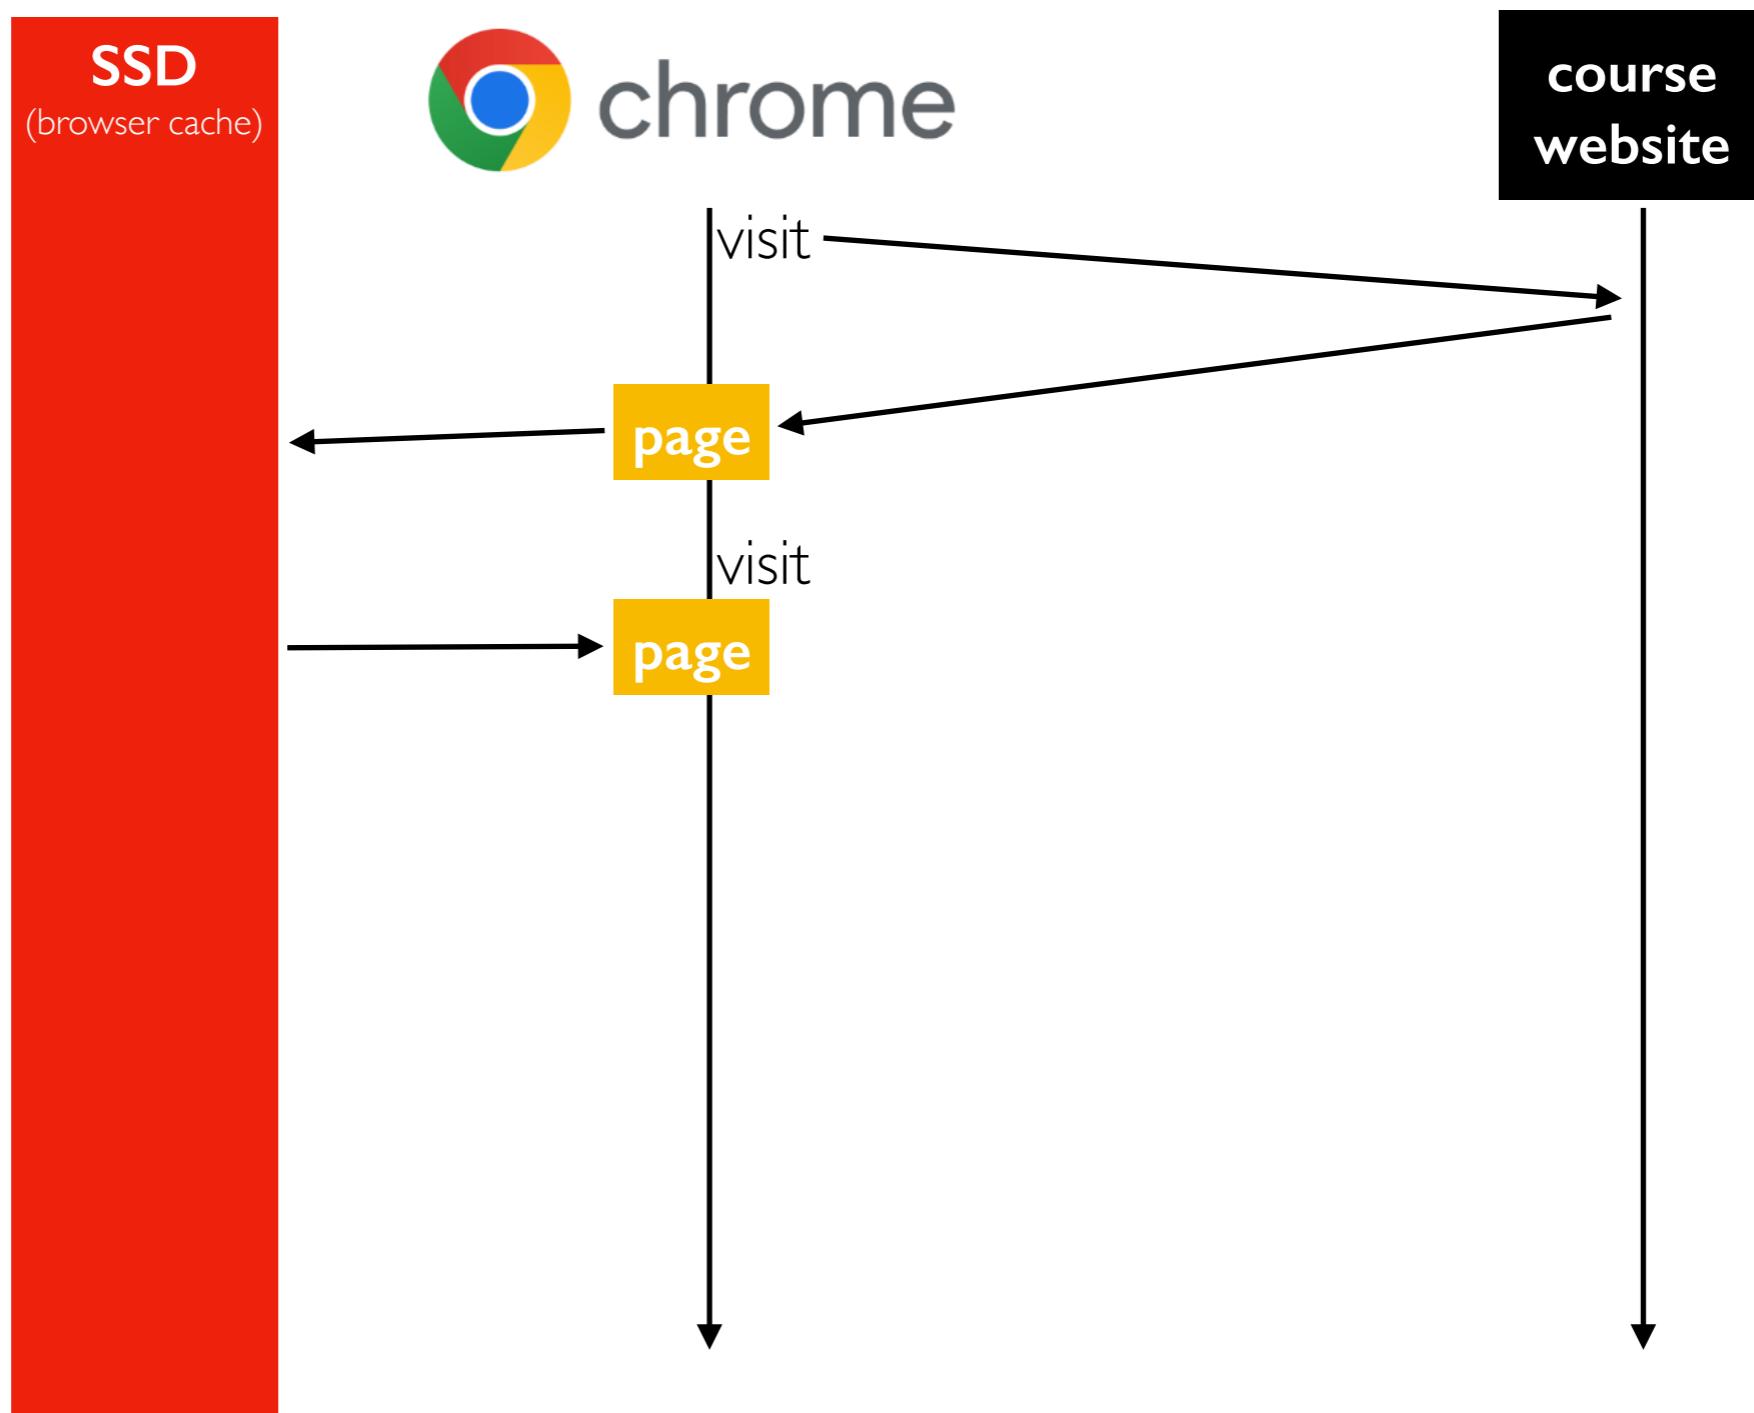
Expiration: Browser Example



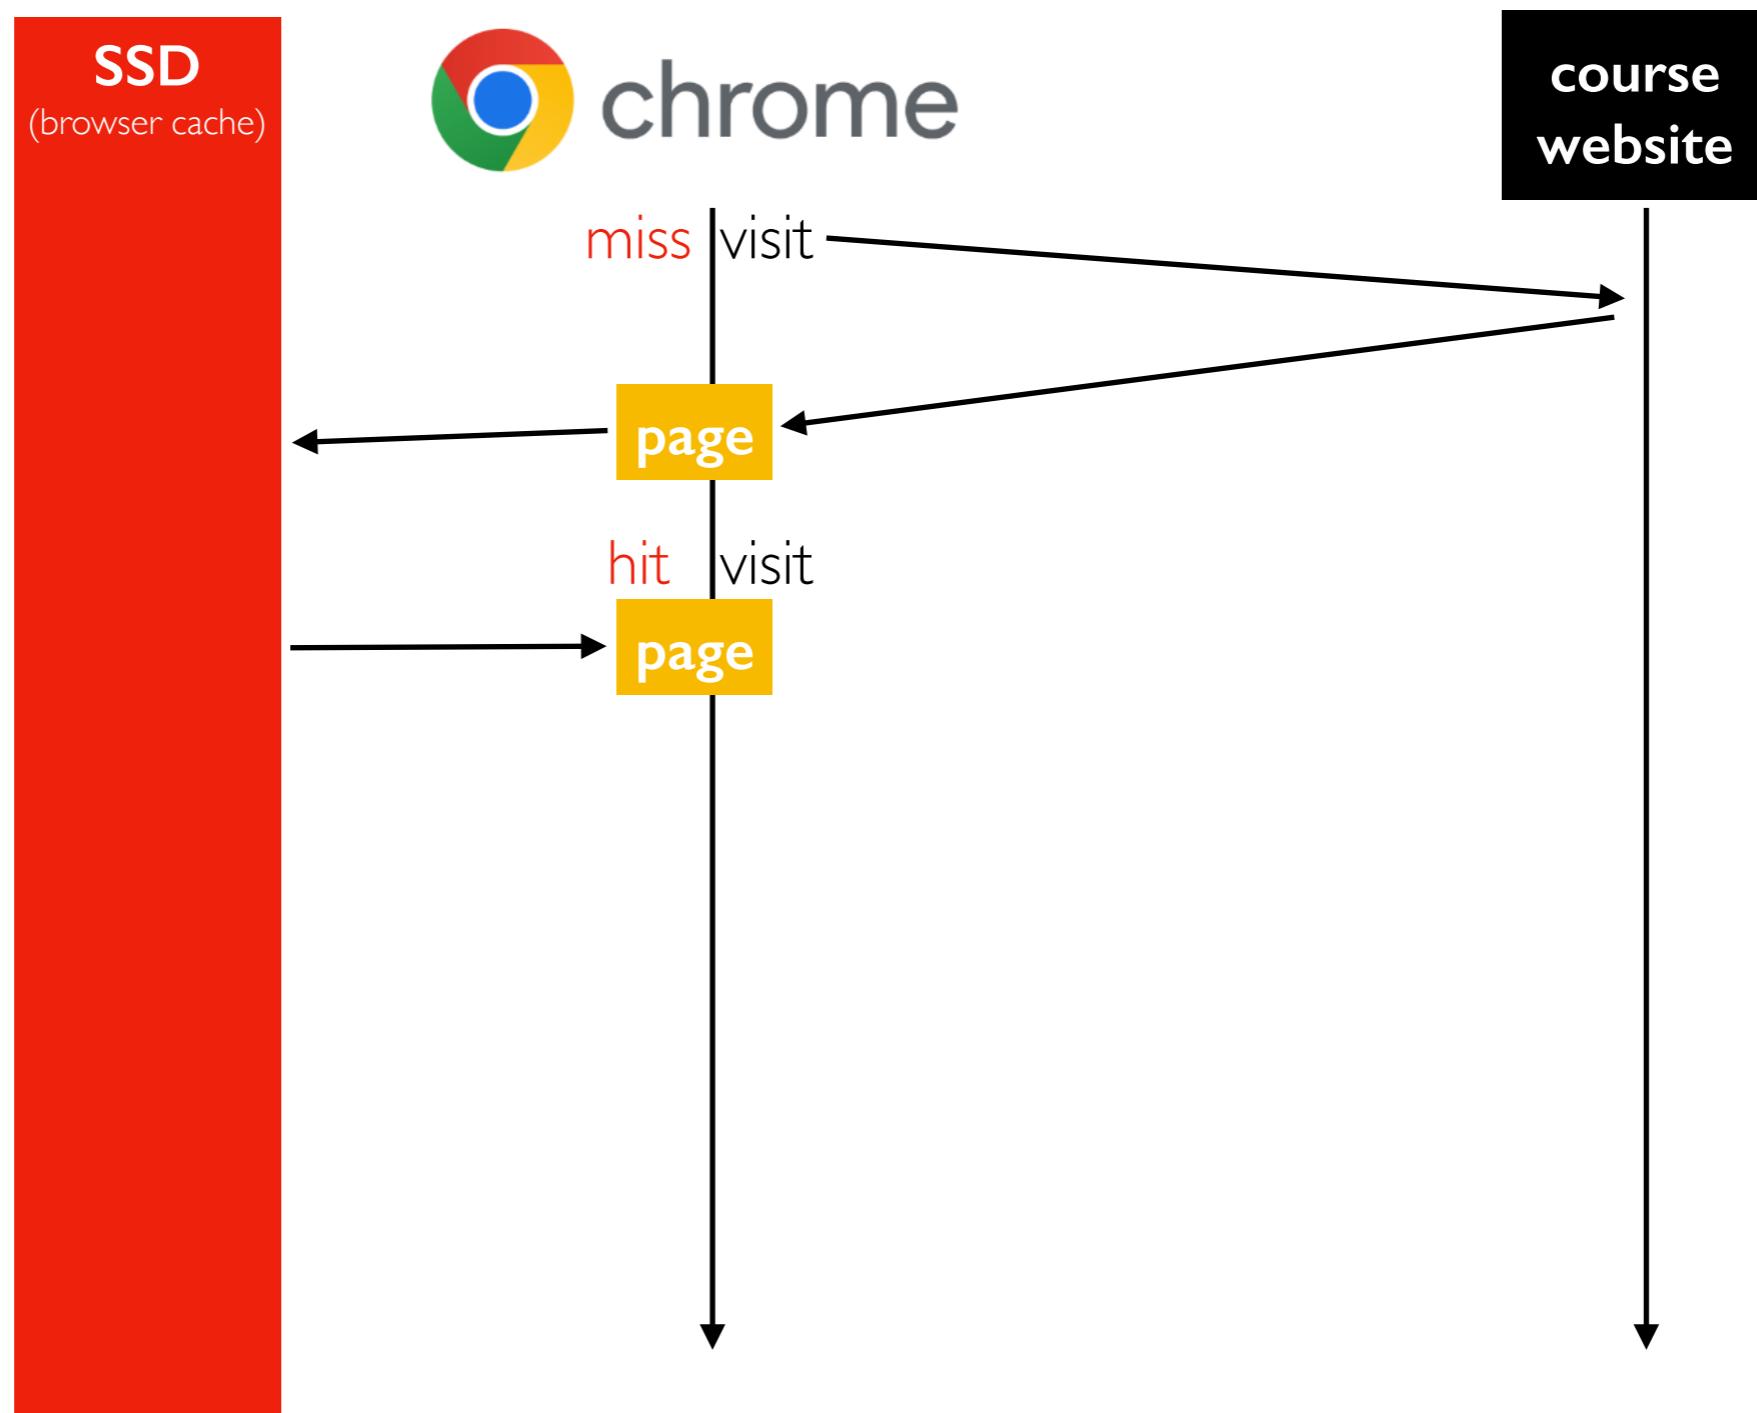
Expiration: Browser Example



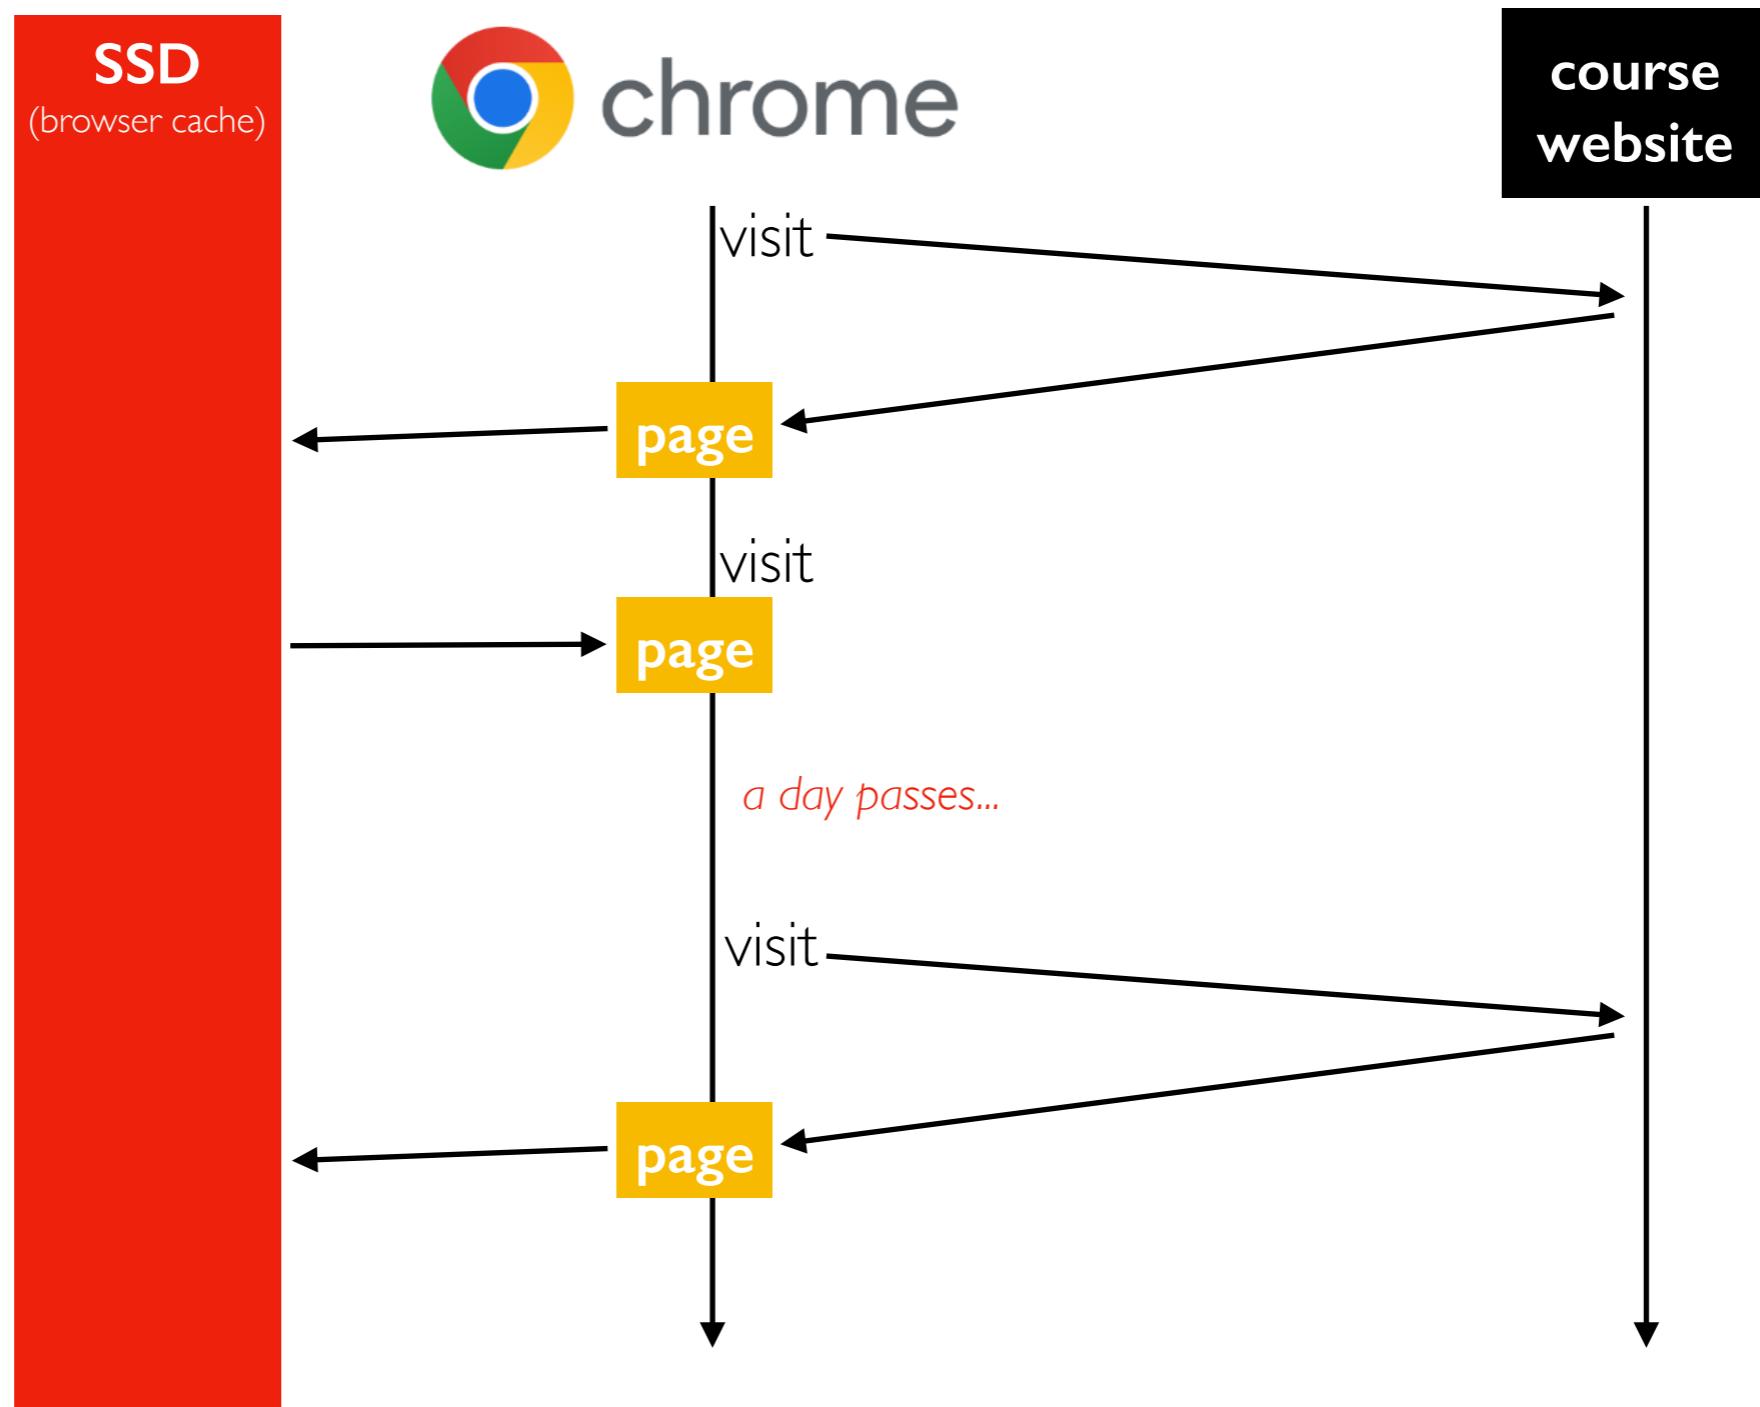
Expiration: Browser Example



Expiration: Browser Example



Expiration: Browser Example



stale data (past expiration) is deleted (or re-validated).
SSD is large so **freshness** is a more important factor than **space**.

Outline

Challenge: Latency

Cache Hierarchy

- CPU, RAM, SSD, Disk, Network
- Tradeoffs

Policy: what data should be cached?

- manual
- expiration
- eviction policies: random, FIFO, LRU

Practice

Cache Policies

When to **load** data to a cache?

- usually whenever we read something, add it
- an exception: programmer knows it will never be read again
 - for example, F_NOCACHE option in Linux. Example program:
 - read file.txt (caching off)
 - compress it
 - write file.txt.gz
 - delete file.txt

When to **evict** data to a cache? Several policies

- random
 - select any entry at random as victim for eviction
- **FIFO** (first in, first out)
 - evict whichever entry has been in the cache the longest
- **LRU** (least recently used)
 - evict whichever entry has been used the least recently

Worksheet

Outline

Challenge: Latency

Cache Hierarchy

- CPU, RAM, SSD, Disk, Network
- Tradeoffs

Policy: what data should be cached?

- manual
- expiration
- eviction policies: random, FIFO, LRU

Practice