

[544] Processes and Threads

Tyler Caraza-Harter

Learning Objectives

- describe the interactions between schedulers, CPUs, threads, and address spaces
- decide for a given scenario whether to organize code as single-threaded, multi-threaded, or multi-process
- trace through different interleavings to identify race conditions

Motivation

Modern CPUs have many cores (maybe dozens)

Trend: **more** cores rather than **faster** cores

Problem: a simple Python program can use at most ONE core
(less if it accesses files or the Internet)

Understanding threads and processes will:

- let us write programs that fully utilize CPU resources
- decide the structure of our concurrent program (threads or processes) depending on the situation

Outline

Review: Virtual Address Spaces

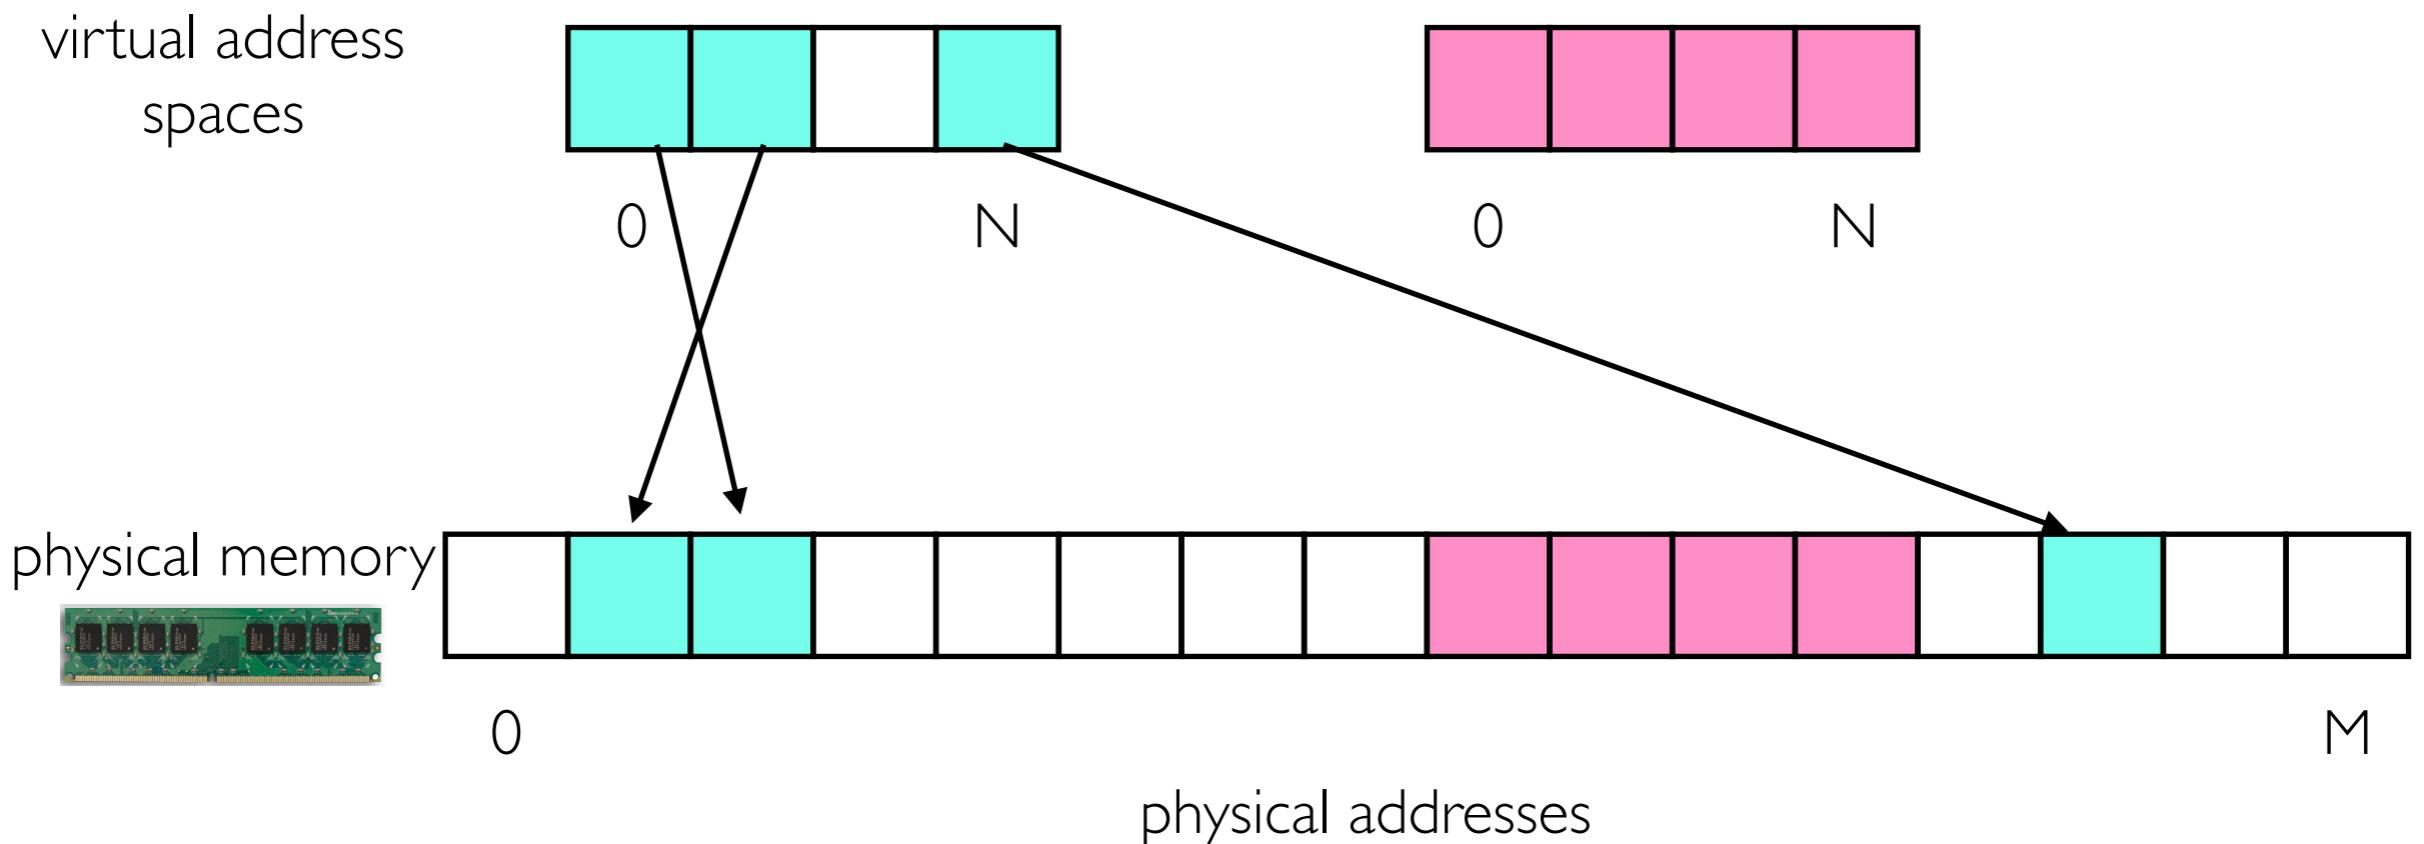
Threads

Demos and Worksheet

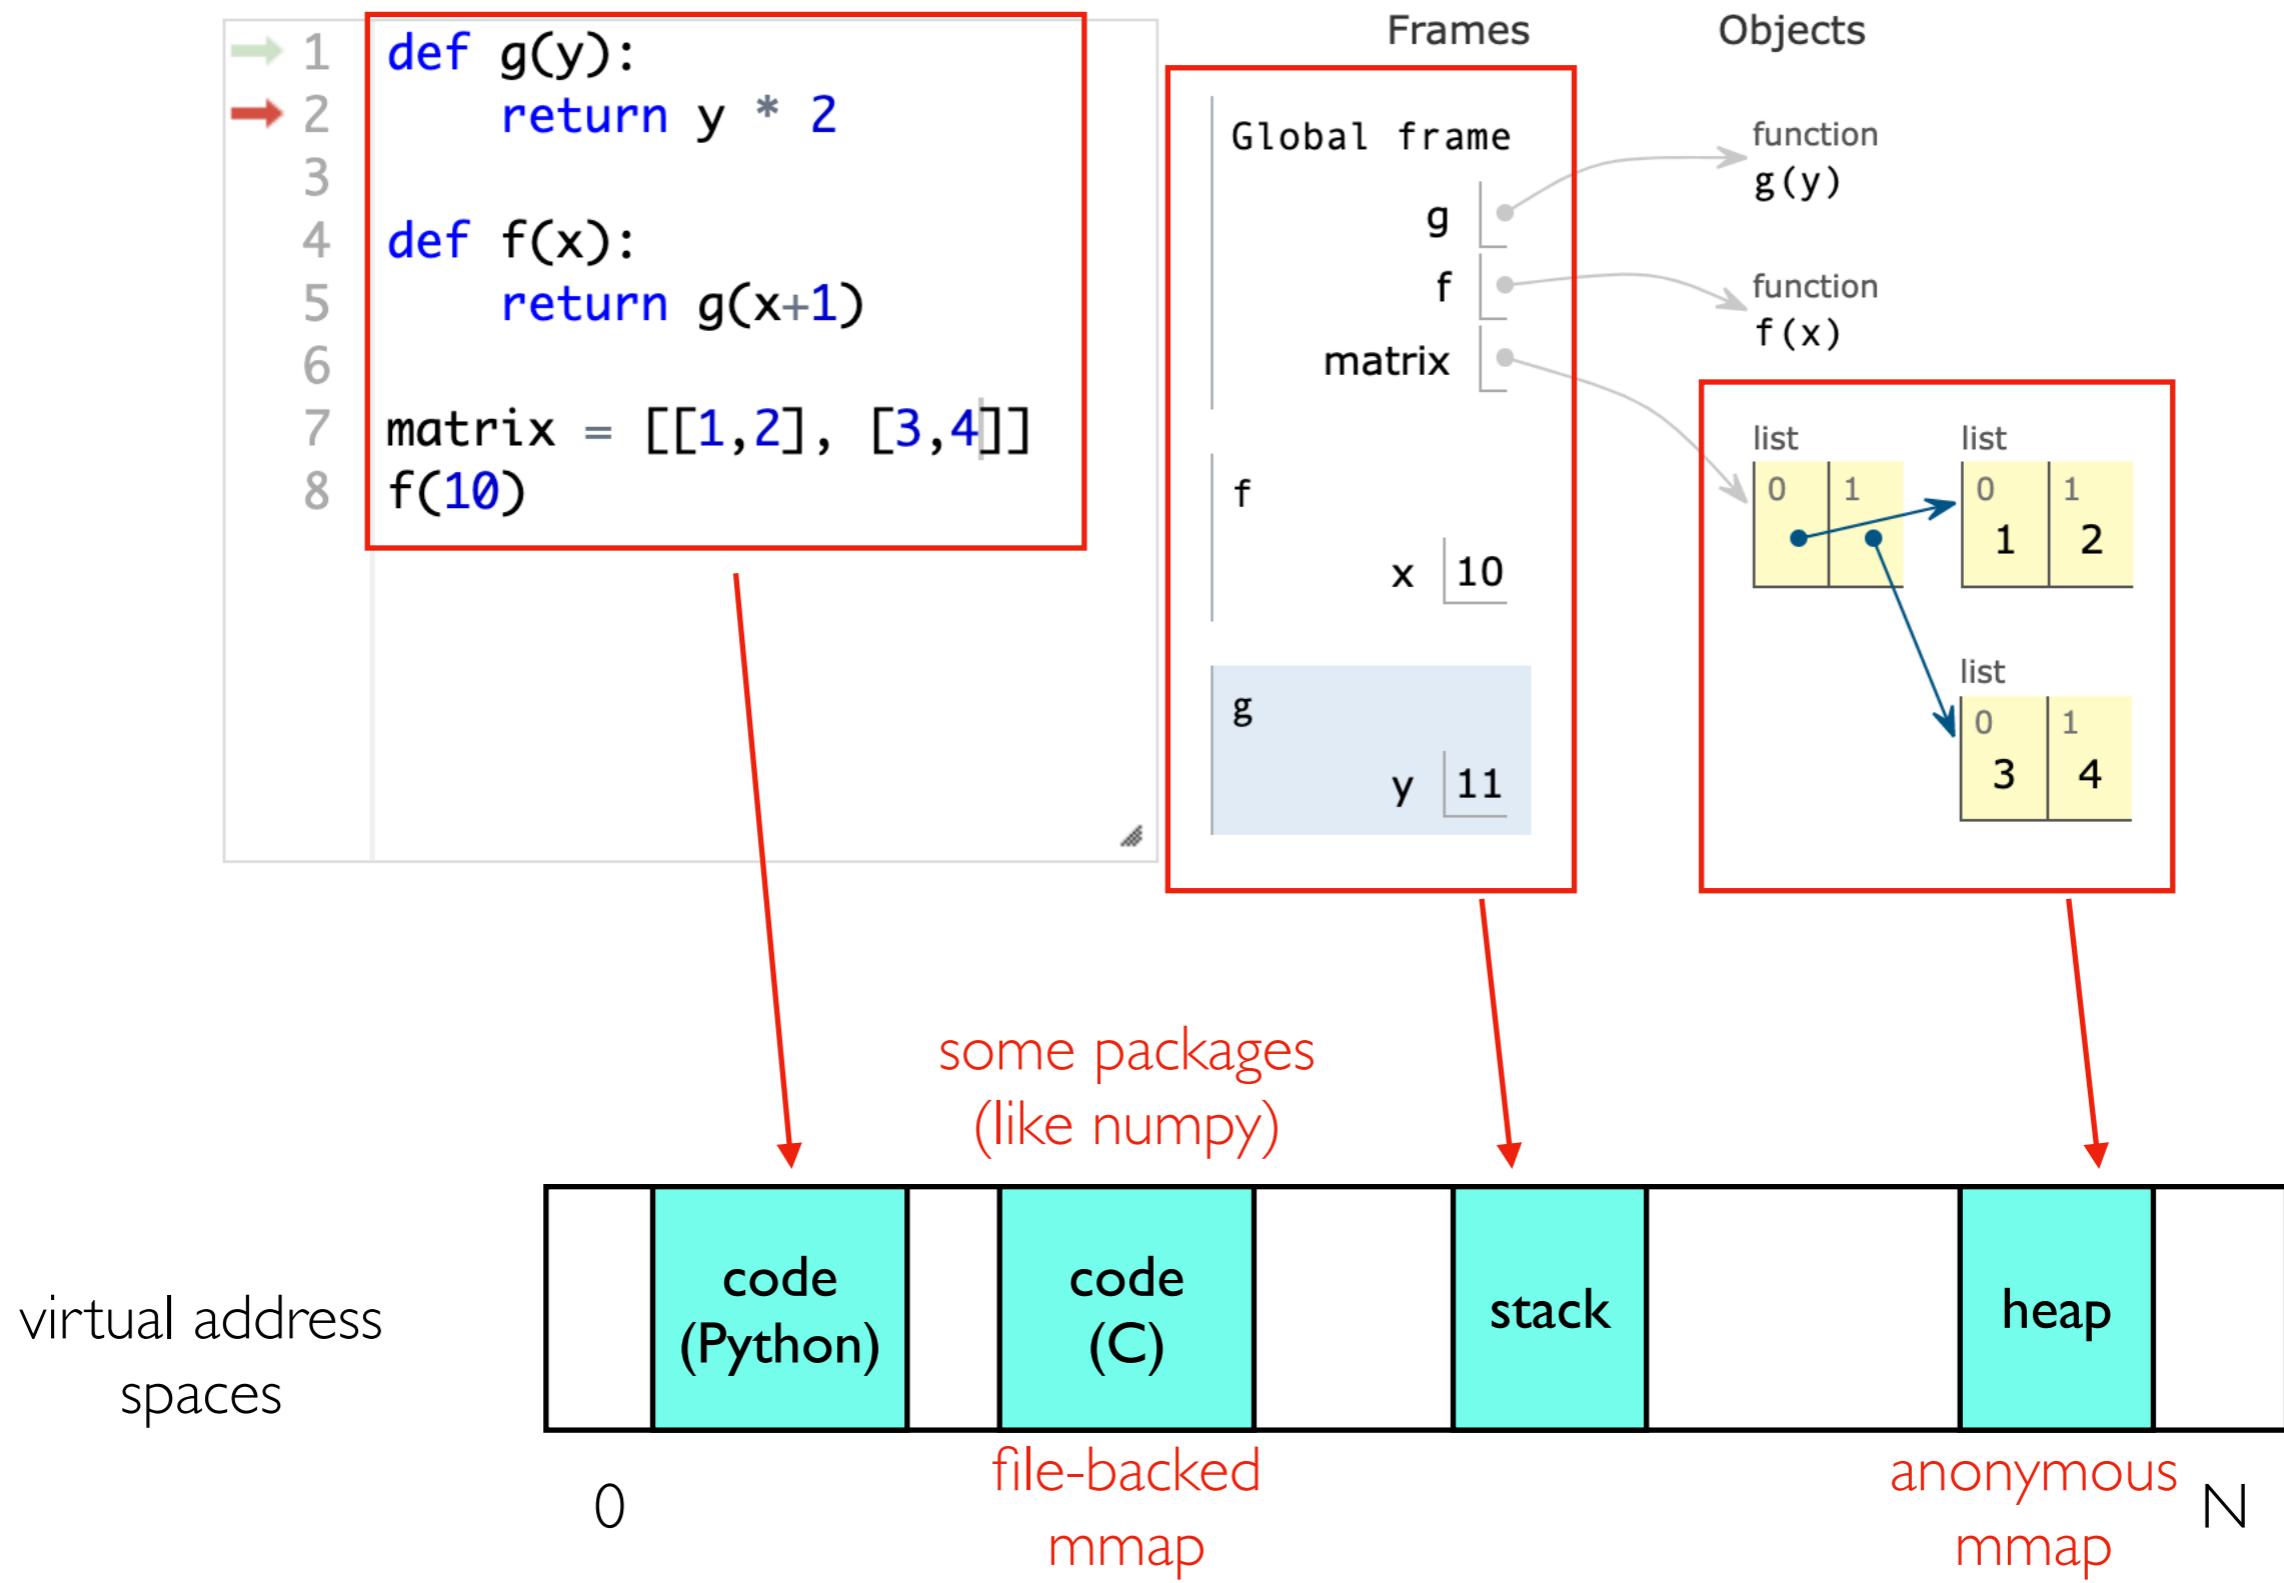
Processes and Address Spaces

Address spaces

- A **process** is a running **program**
- Each process has it's own **virtual address space**
- The same virtual address generally refers to different memory in different processes
- Regular processes cannot directly access **physical memory** or other addr spaces
- Address spaces can have holes (N is usually MUCH bigger than M)
- Physical memory for a process need not be contiguous

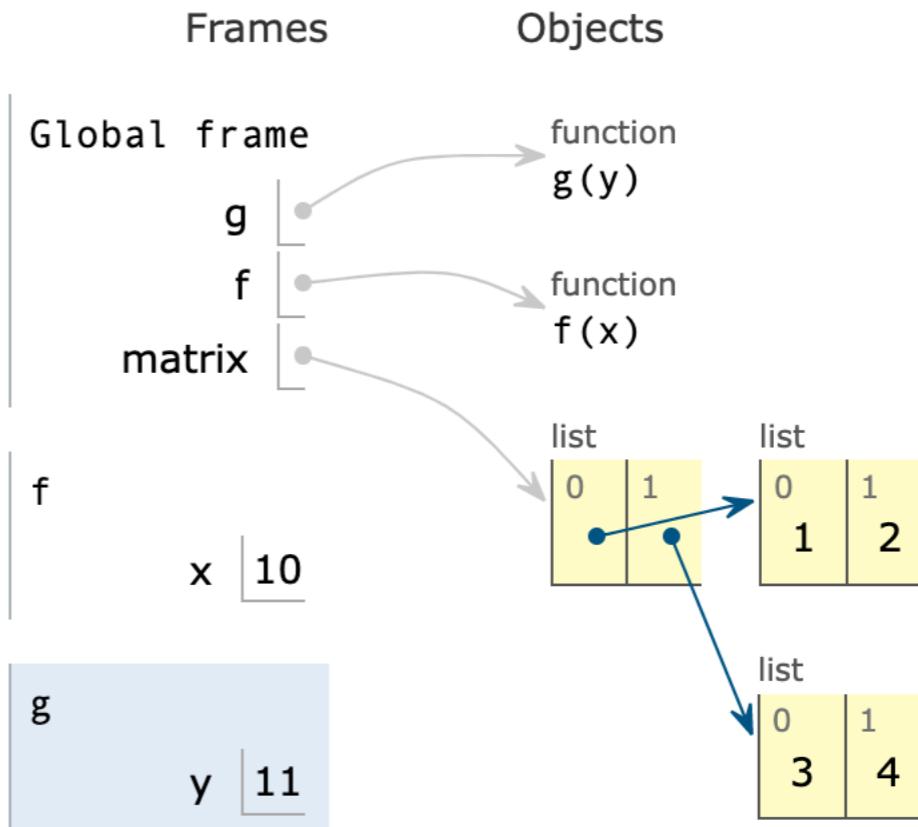


What goes in an address space?

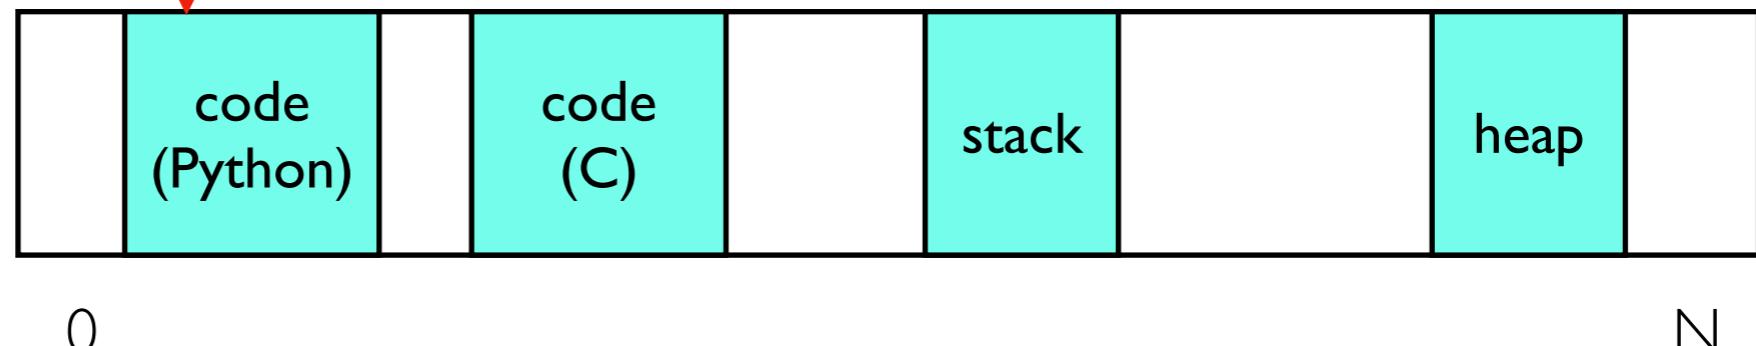


How does code execute?

```
1 def g(y):  
2     return y * 2  
3  
4 def f(x):  
5     return g(x+1)  
6  
7 matrix = [[1,2], [3,4]]  
8 f(10)
```



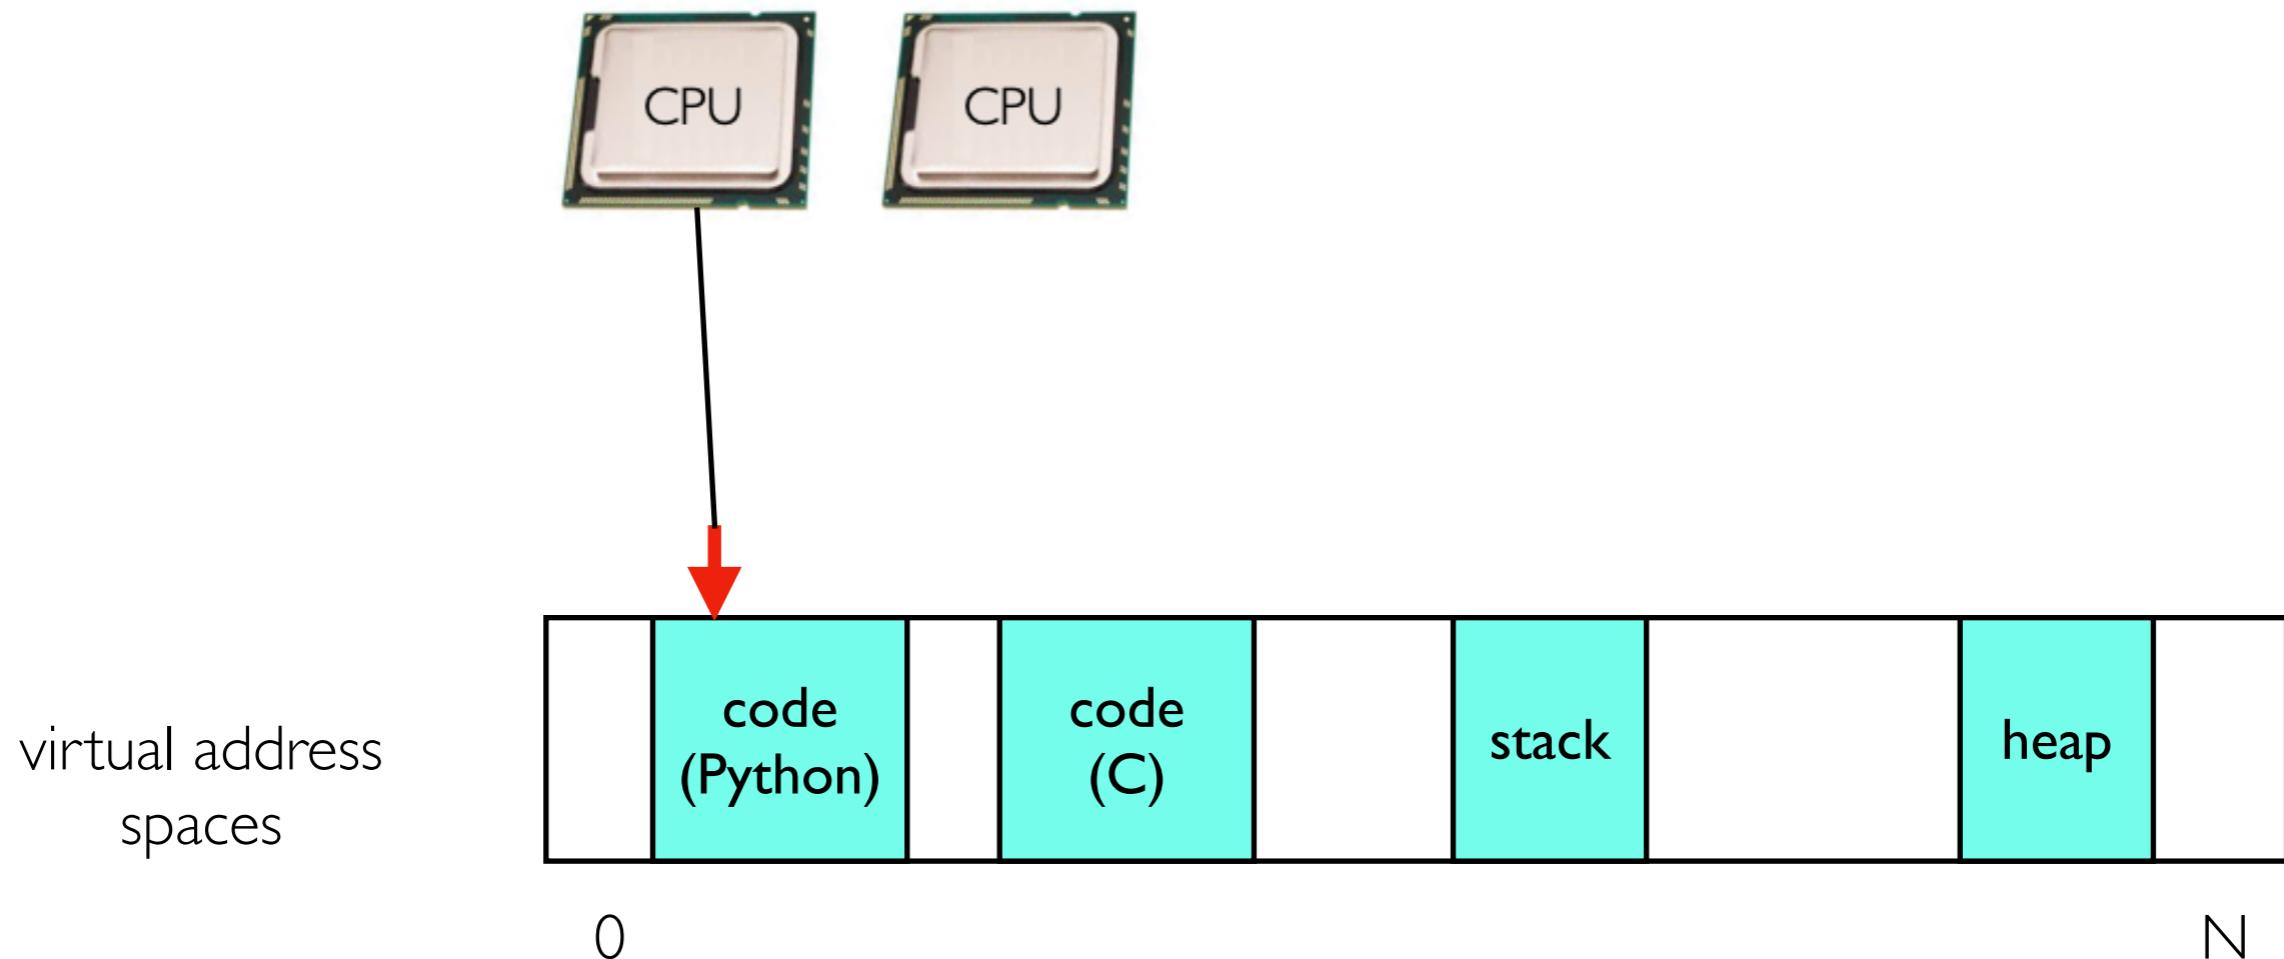
The diagram illustrates the mapping of virtual address spaces to physical memory regions. On the left, the text "virtual address spaces" is displayed. On the right, a large blue rectangle represents a memory region, divided into two vertical sections: a white section on the left and a blue section on the right. The blue section contains the text "code (Python)". A red arrow points from the text "instruction pointer" at the top to the boundary between the white and blue sections, indicating that the instruction pointer points to the start of the code region.



How does code execute?

CPUs

- CPUs are attached to at most one **instruction pointer** at any given time
- they run code by executing instructions and advancing the instruction pointer
- **Note:** interpreter left out for simplicity (CPU points to interpreter code, which points to Python bytecode)



Outline

Review: Virtual Address Spaces

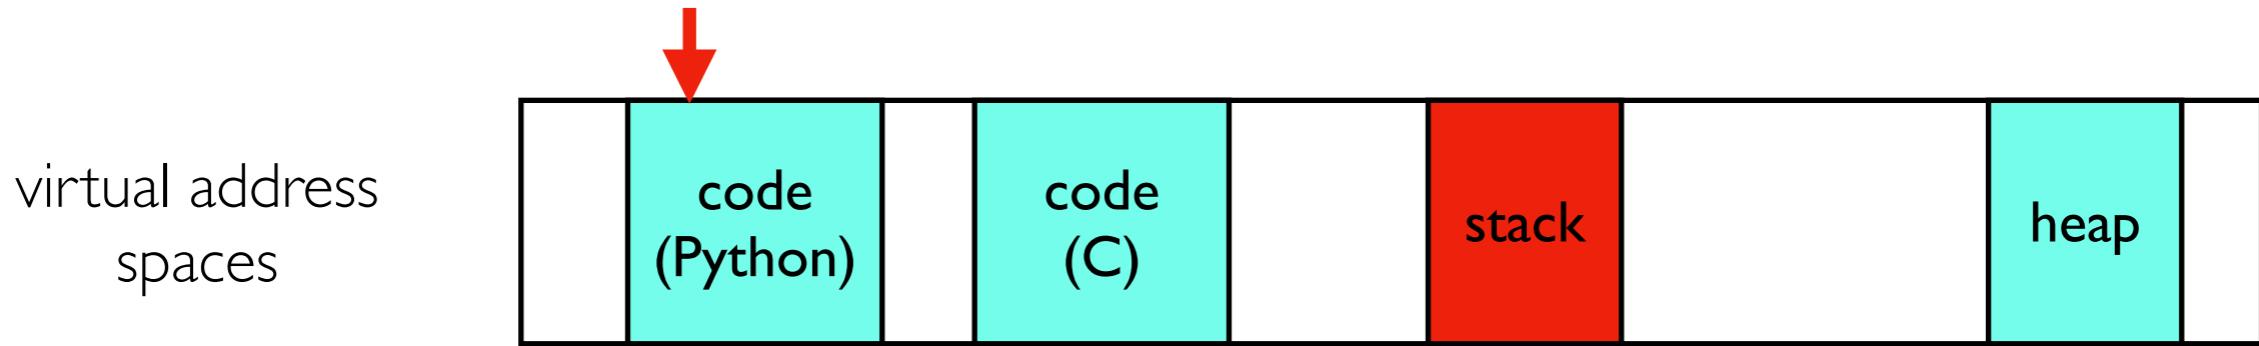
Threads

Demos and Worksheet

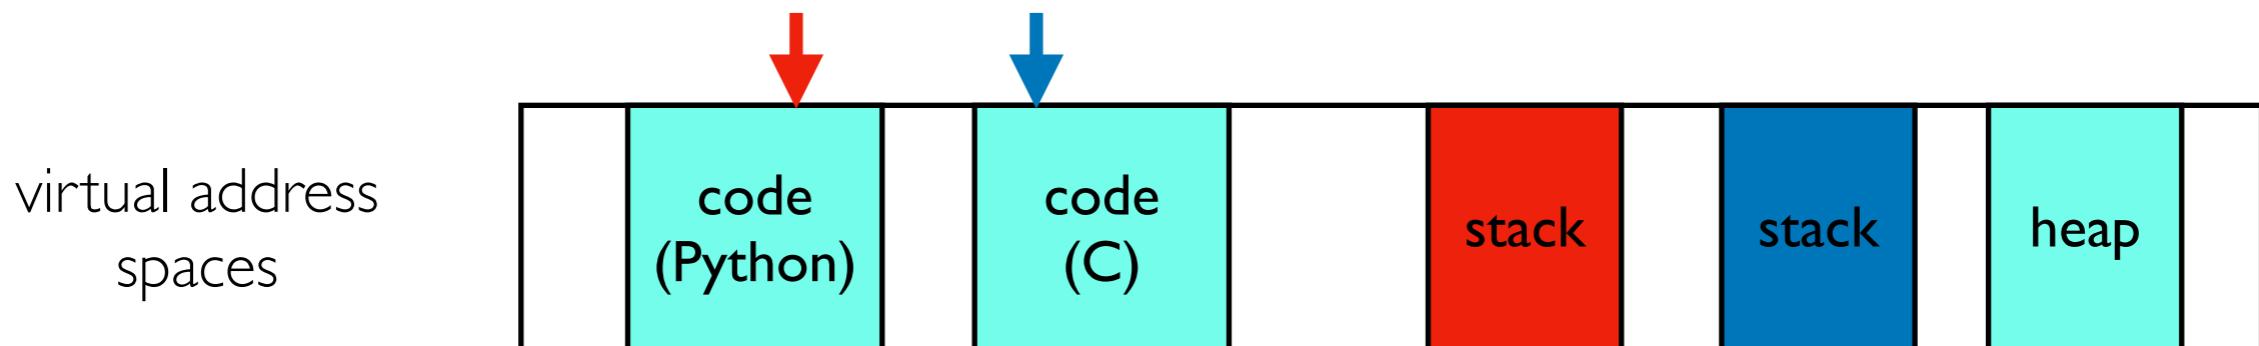
Threads

Threads have their own **instruction pointers** and **stacks**, but share the **heap**.

Single-threaded process:



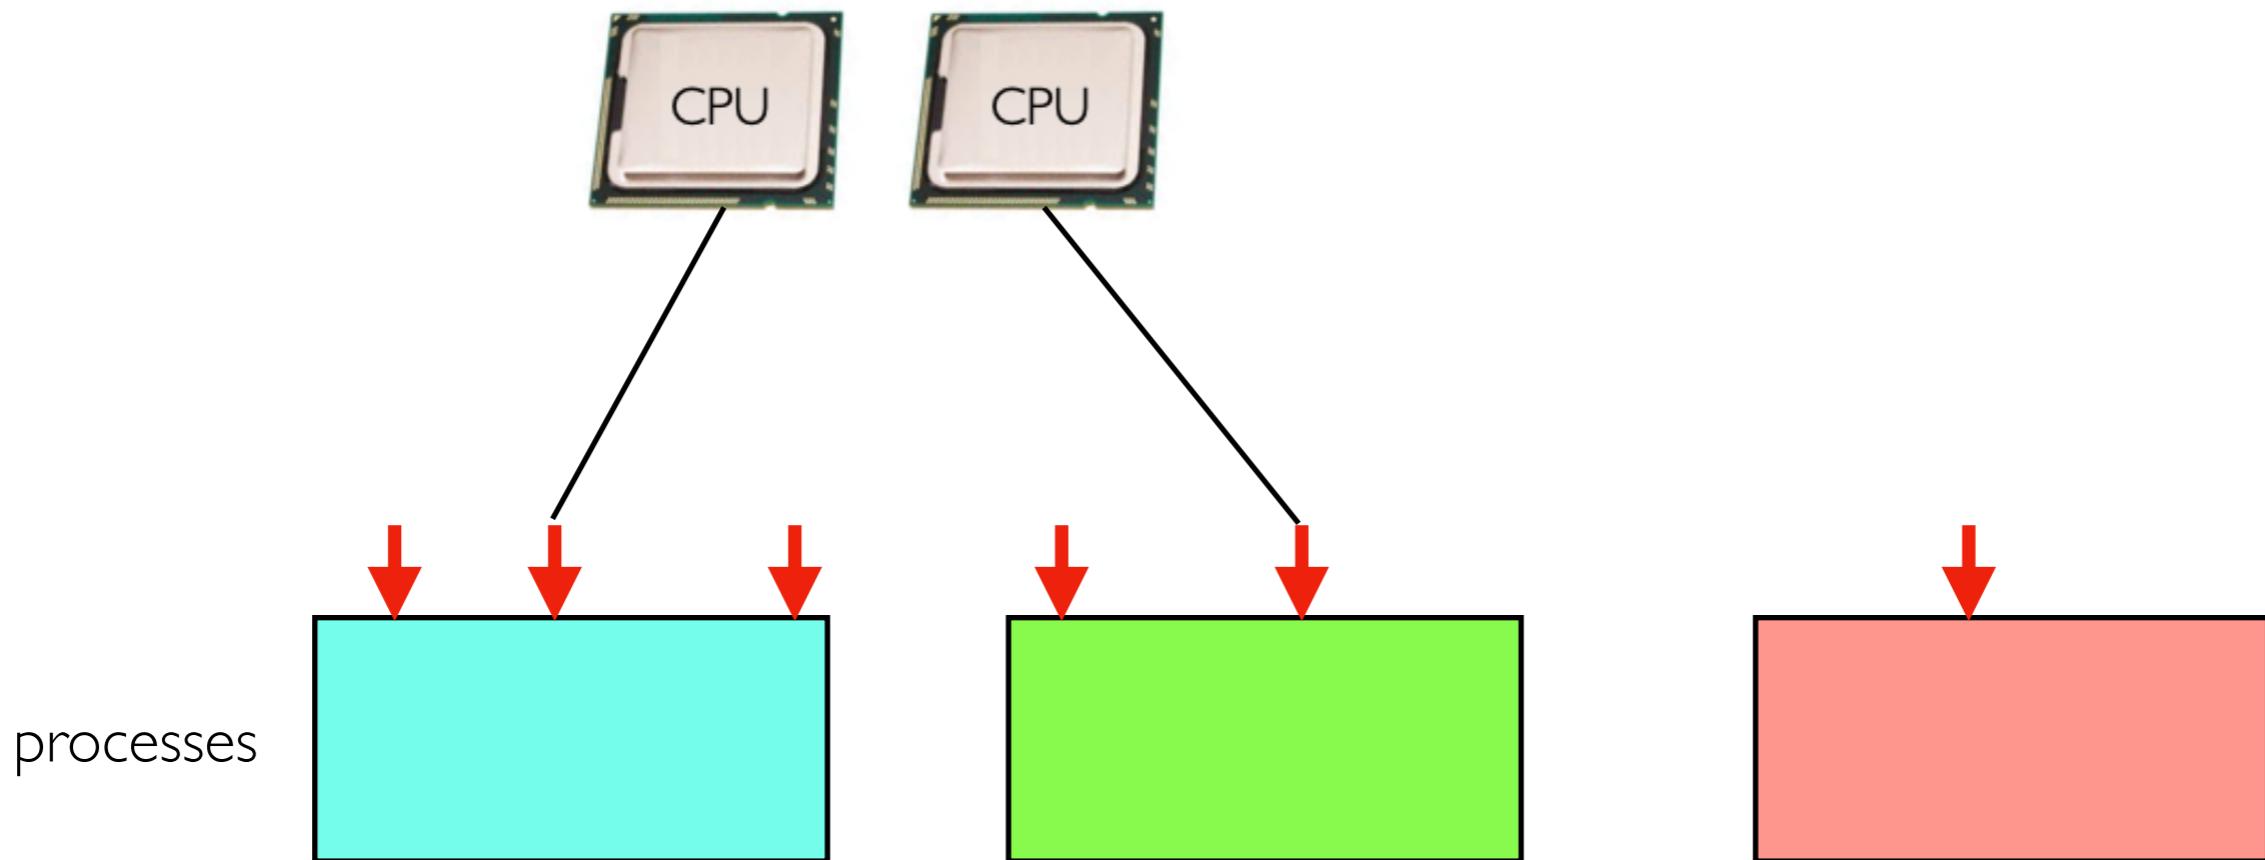
Multi-threaded process:



Context Switch

Schedulers

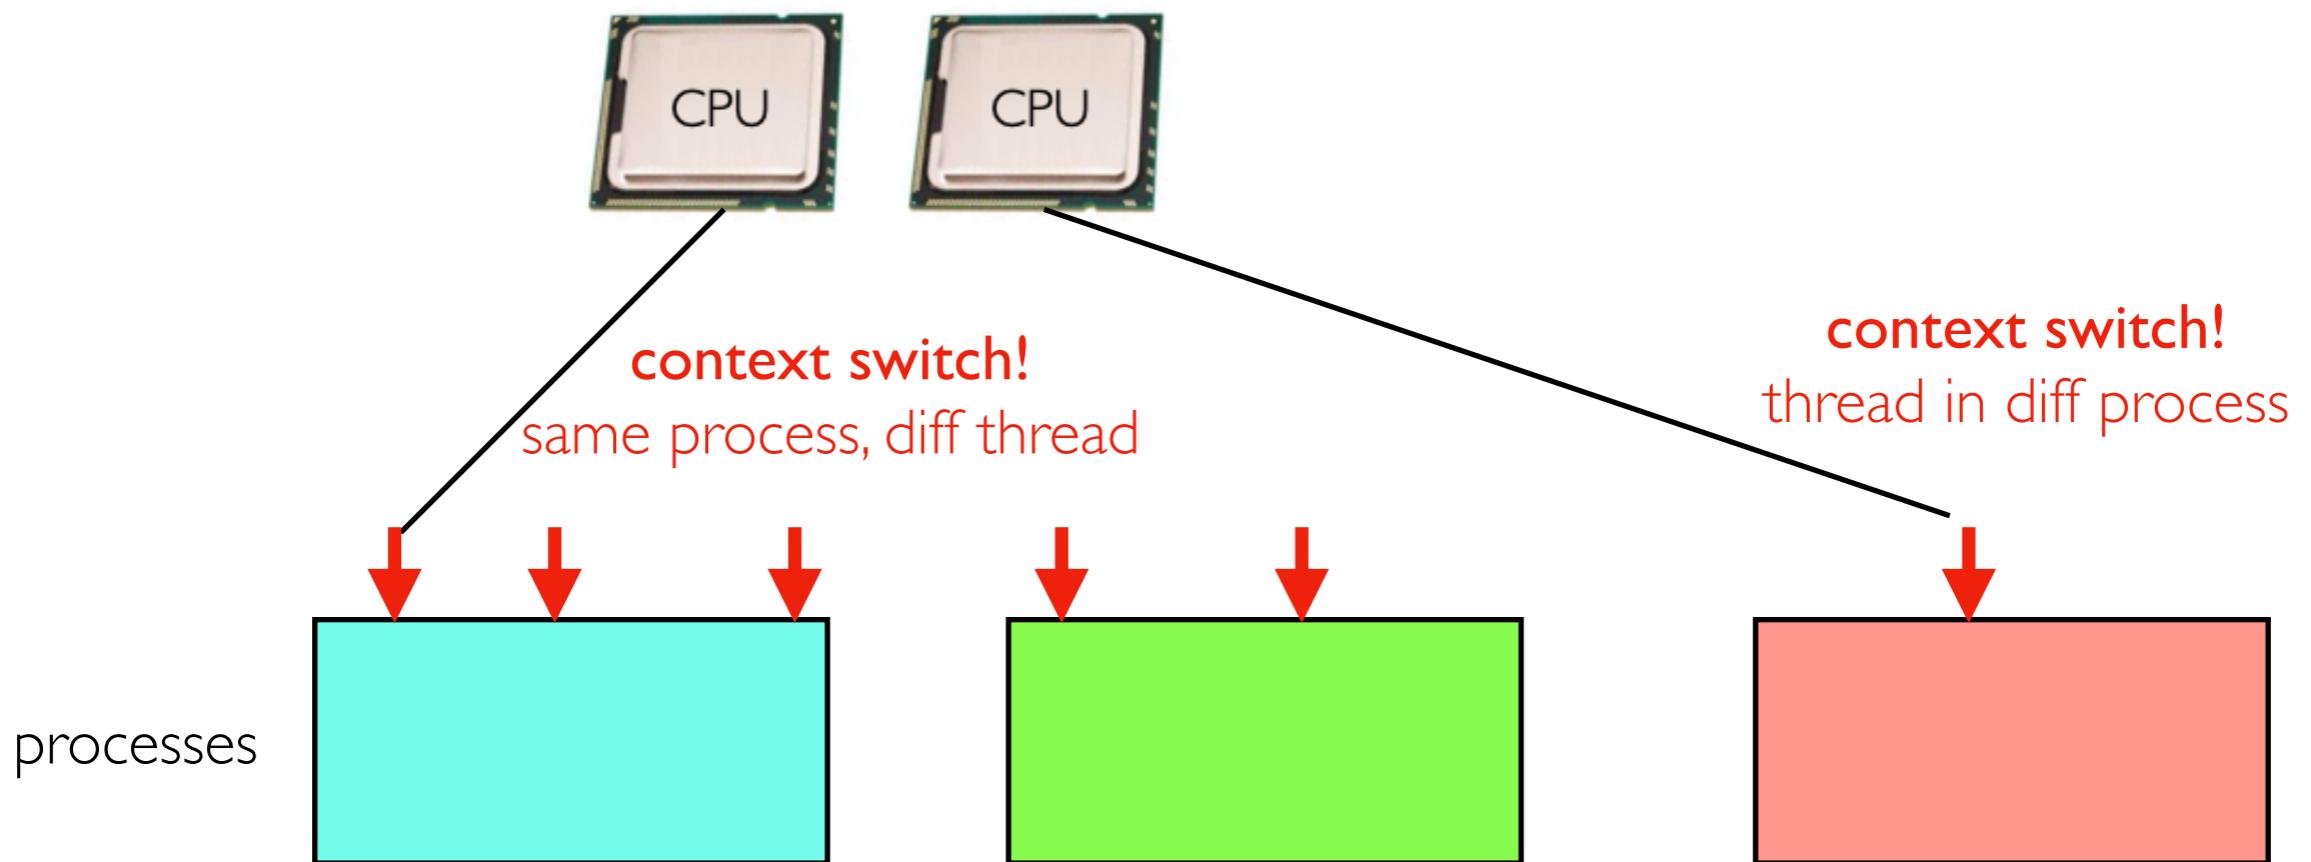
- CPU **scheduler** is an important sub system in an **operating system**
- schedulers decide when to **context switch** between threads
- context switch: change which thread a CPU is running



Context Switch

Schedulers

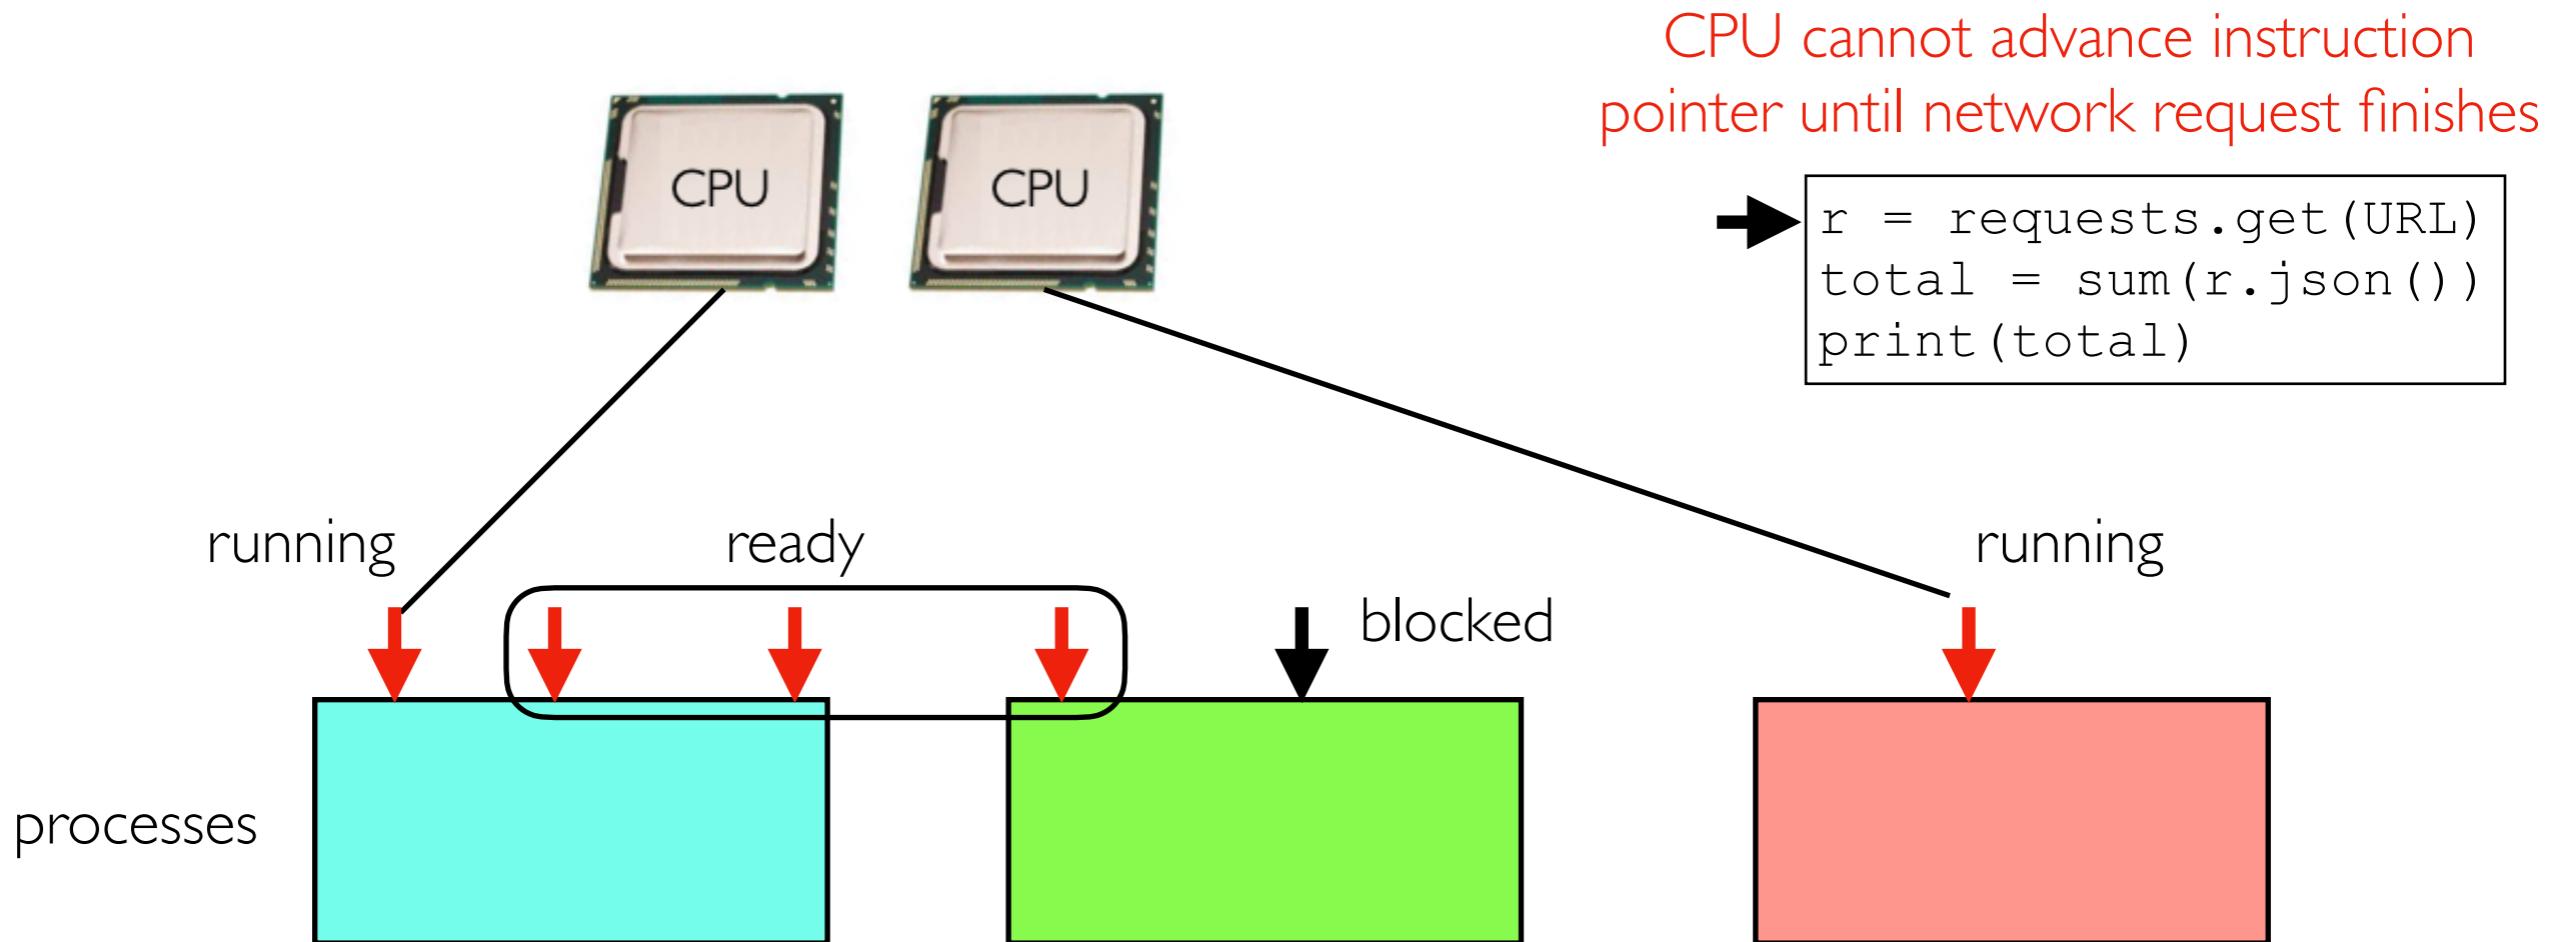
- CPU **scheduler** is an important sub system in an **operating system**
- schedulers decide when to **context switch** between threads
- context switch: change which thread a CPU is running



Scheduling Restrictions: Blocked Threads

Threads can be in one of three states

- **running**: CPU is executing it
- **blocked**: waiting on something other than CPU (network, input, disk, etc)
- **ready**: scheduler can choose to context switch to it



Efficient Use of Compute Resources

Wasted cores: (1) not enough threads (2) blocked threads

For 100% CPU utilization (difficult goal)

- need at least one ready/running thread for each CPU core
- generally need more threads than cores (threads are often blocked)
- **threads could be in one process (or many)**

Multi-threaded applications

- good when multiple threads need to access frequently modified data structures
- new kinds of bugs possible (race conditions, deadlock)

Multi-process applications (<https://docs.python.org/3/library/multiprocessing.html>)

- easier to program (or just manually launch several processes in background)
- better at keeping multiple cores busy simultaneously (Python specific)

Both approaches work well for dealing with blocked threads

Coding Demos, Worksheet

Thread operations

- `t = threading.Thread()`
- `t.start(target=????, args=[????])`
- `t.join()`
- `t.get_native_id()`