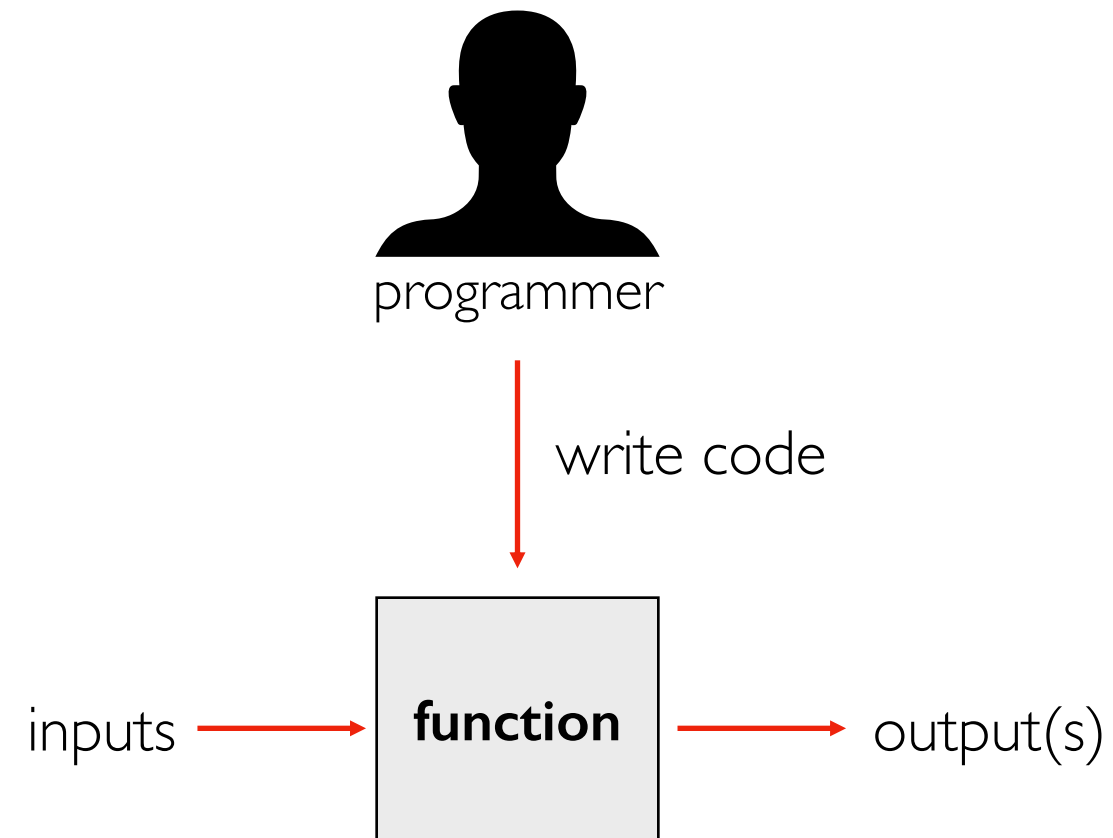


# [320] Machine Learning: Intro

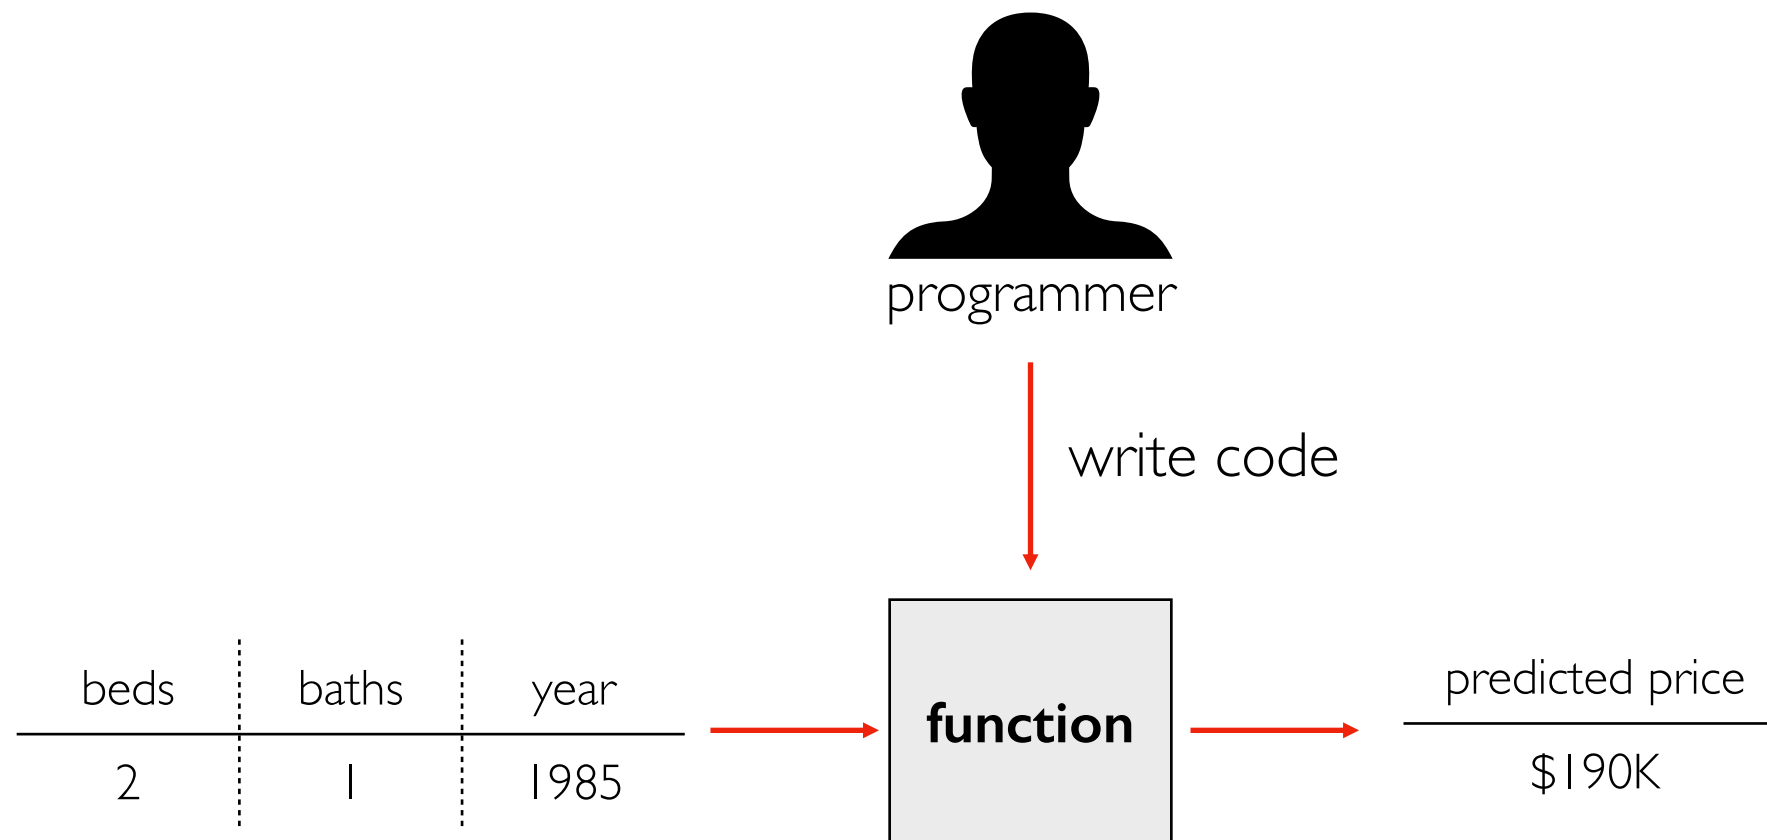
Meenakshi Syamkumar

Functions/Models

# How do we make functions?

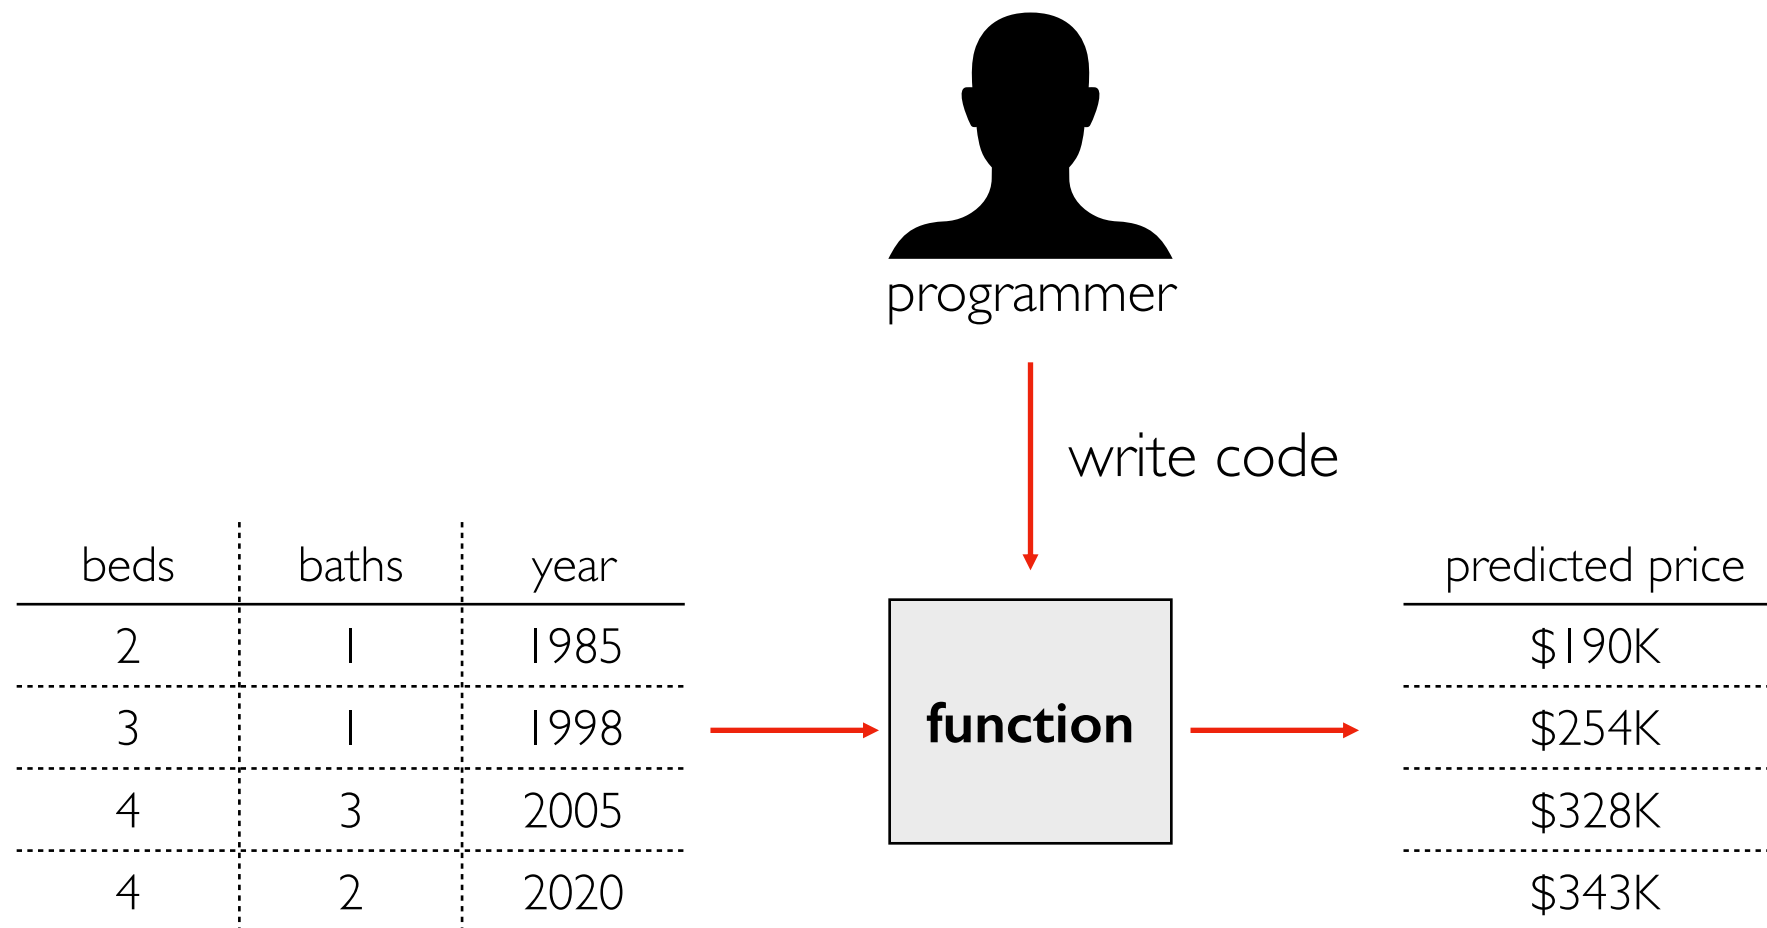


# How do we make functions?



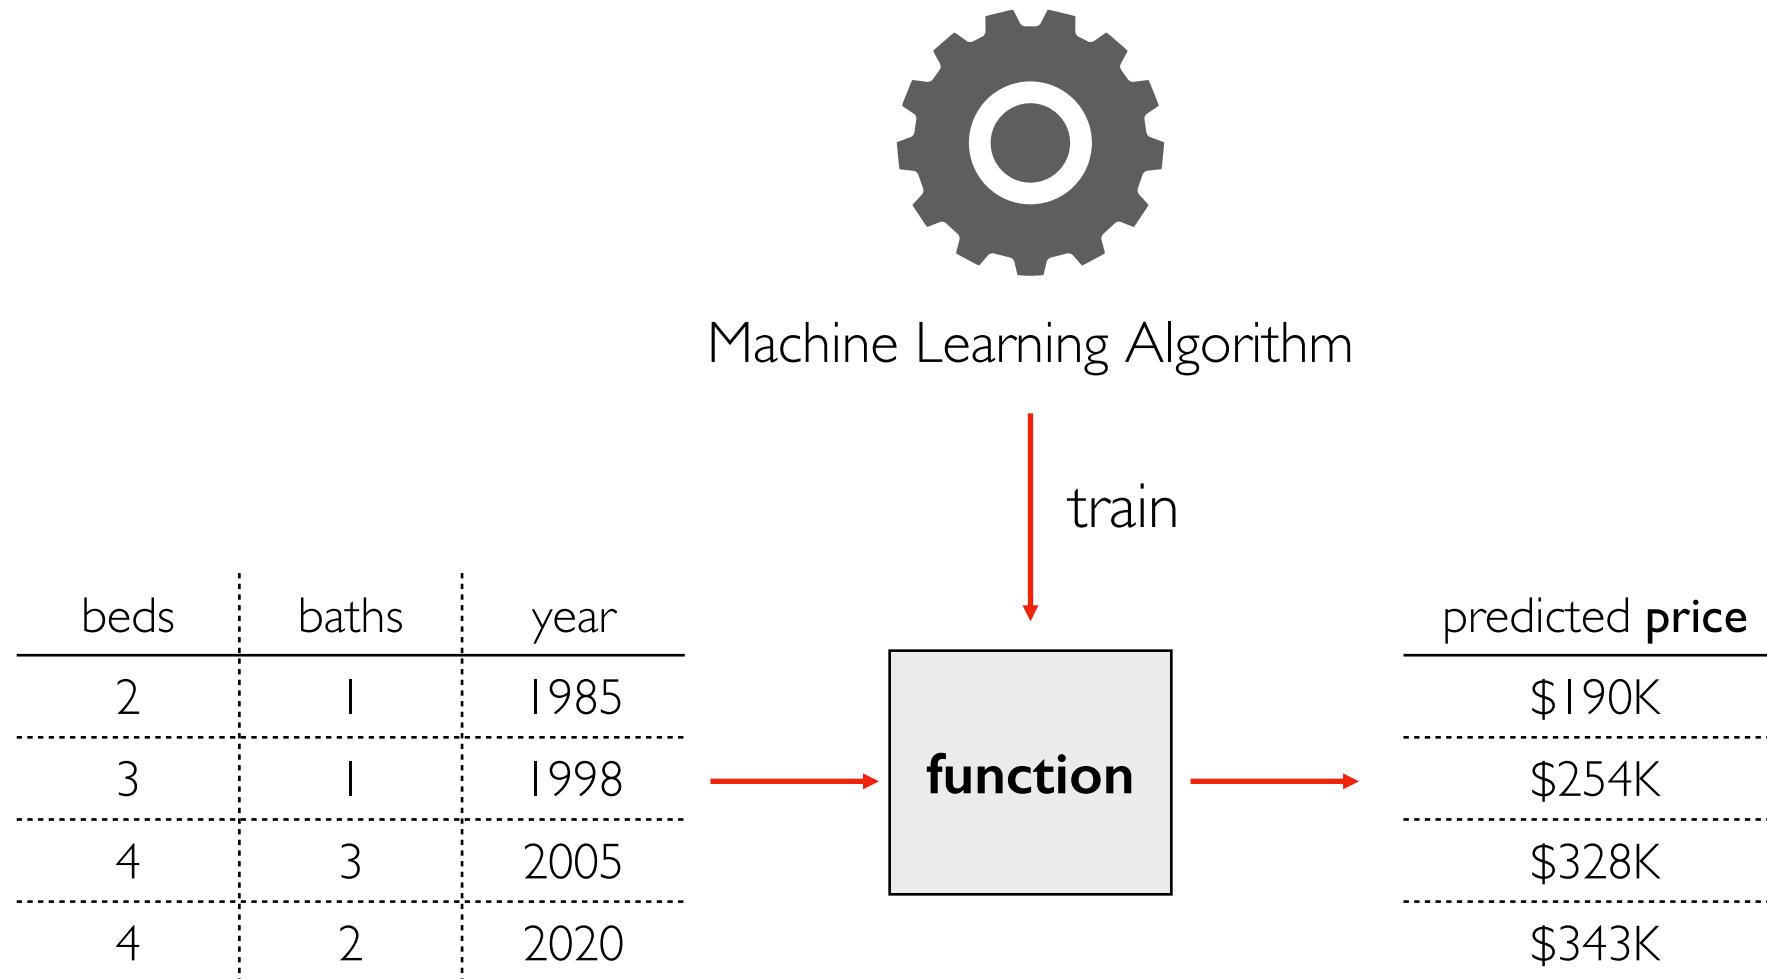
many functions are **models** that can be used to predict

# How do we make functions?



many functions are **models** that can be used to predict

# How do we make functions?




many functions are **models** that can be used to predict

# How do we make functions?


training data

beds	baths	year	price
1	1	1980	\$140K
3	1	1990	\$240K
3	4	2004	\$295K
4	3	2018	\$350K



live data

beds	baths	year
2	1	1985
3	1	1998
4	3	2005
4	2	2020



many functions are **models** that can be used to predict

# How do we make functions?


training data

beds	baths	year	price
1	1	1980	\$140K
3	1	1990	\$240K
3	4	2004	\$295K
4	3	2018	\$350K



live data

beds	baths	year
2	1	1985
3	1	1998
4	3	2005
4	2	2020



predicted price
\$190K
\$254K
\$328K
\$343K

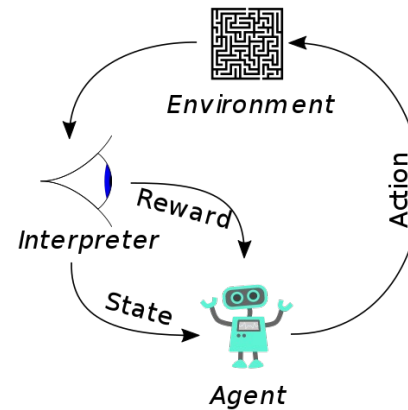
this is an example of a **regression** model, which is a type of **supervised machine learning**, which is one of the 3 main categories of ML



# Machine Learning

## Reinforcement Learning

*not covered in CS 320*



[https://en.wikipedia.org/wiki/Reinforcement\\_learning](https://en.wikipedia.org/wiki/Reinforcement_learning)

## Supervised Machine Learning

*data is labeled, we know what we want to predict*

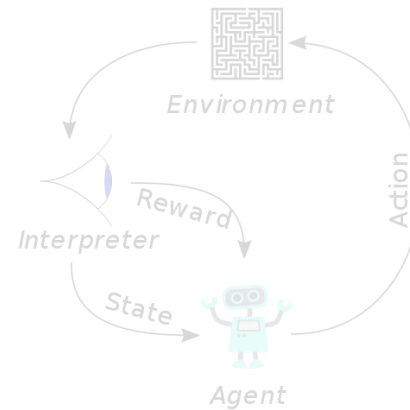
## Unsupervised Machine Learning

*data is unlabeled, we're just looking for patterns*

# Machine Learning

## Reinforcement Learning

*not covered in CS 320*



[https://en.wikipedia.org/wiki/Reinforcement\\_learning](https://en.wikipedia.org/wiki/Reinforcement_learning)

## Supervised Machine Learning

*data is labeled, we know what we want to predict*

### Regression

*predict a quantity*

### Classification

*predict a category*

## Unsupervised Machine Learning

*data is unlabeled, we're just looking for patterns*

### Clustering

*place rows in groups*

### Decomposition

*represent rows as combos of "component" rows*

# I. Regression (Supervised)

features

	<b>x0</b>	<b>x1</b>	<b>x2</b>	<b>x3</b>	<b>x4</b>	<b>y (label)</b>
<b>0</b>	37	25	40	70	68	5
<b>1</b>	50	13	7	67	79	25
<b>2</b>	56	12	5	15	90	44
<b>3</b>	89	70	85	49	68	72
<b>4</b>	36	93	52	33	14	59
<b>5</b>	53	5	67	99	55	????
<b>6</b>	47	31	9	56	27	????
<b>7</b>	50	3	20	24	63	????
<b>8</b>	36	32	66	70	7	????
<b>9</b>	27	33	16	21	9	????

problem: can we predict an unknown **quantity** based on **features**?

# I. Regression (Supervised)

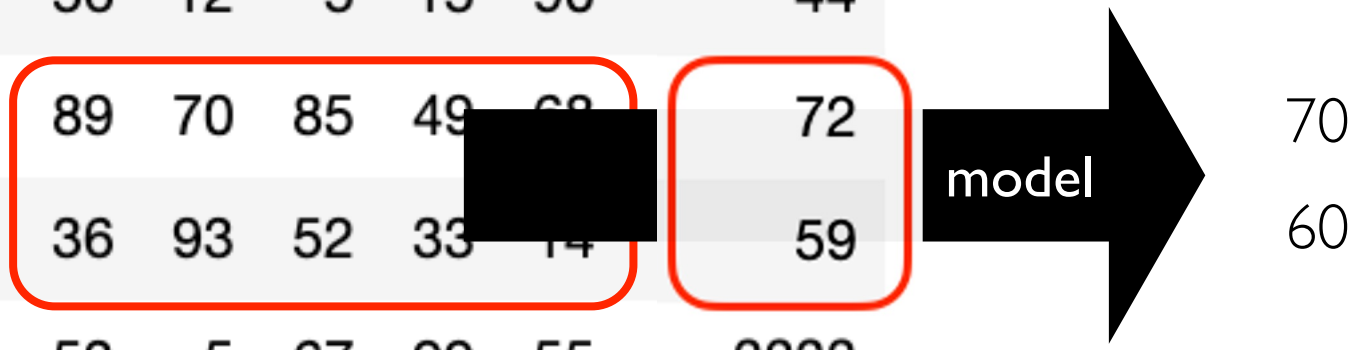
	x0	x1	x2	x3	x4	y (label)
0	37	25	40	70	68	5
1	50	13	7	67		25
2	56	12	5	15	90	44
3	89	70	85	49	68	72
4	36	93	52	33	14	59
5	53	5	67	99	55	????
6	47	31	9	56	27	????
7	50	3	20	24	63	????
8	36	32	66	70	7	????
9	27	33	16	21	9	????

model

**train:** fit a model to the relationship between some label (y) and feature (x's) values

# I. Regression (Supervised)

	x0	x1	x2	x3	x4	y (label)
0	37	25	40	70	68	5
1	50	13	7	67	79	25
2	56	12	5	15	90	44
3	89	70	85	49	68	72
4	36	93	52	33	14	59
5	53	5	67	99	55	????
6	47	31	9	56	27	????
7	50	3	20	24	63	????
8	36	32	66	70	7	????
9	27	33	16	21	9	????

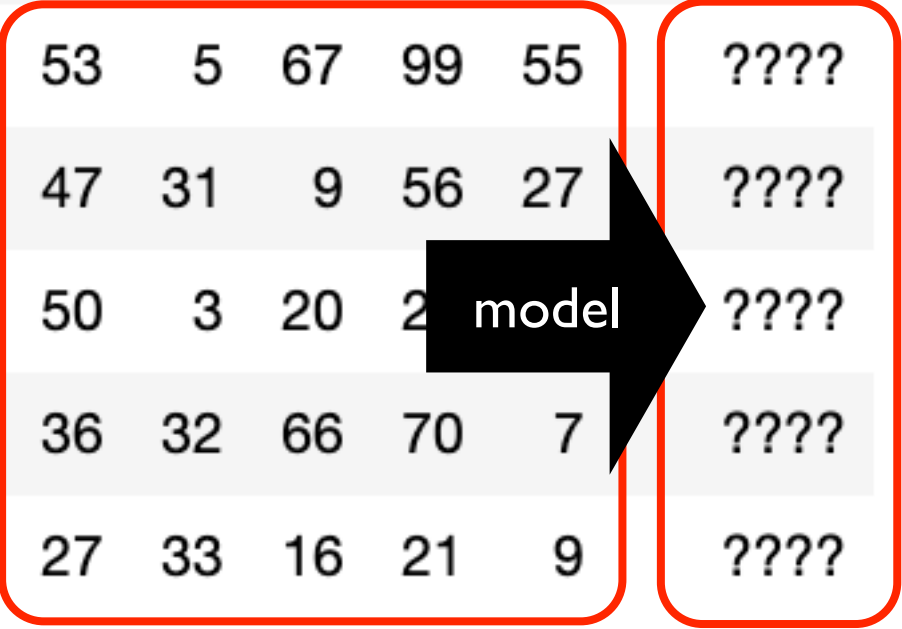


70  
60

**test:** make some predictions for known rows -- how close are we?

# I. Regression (Supervised)

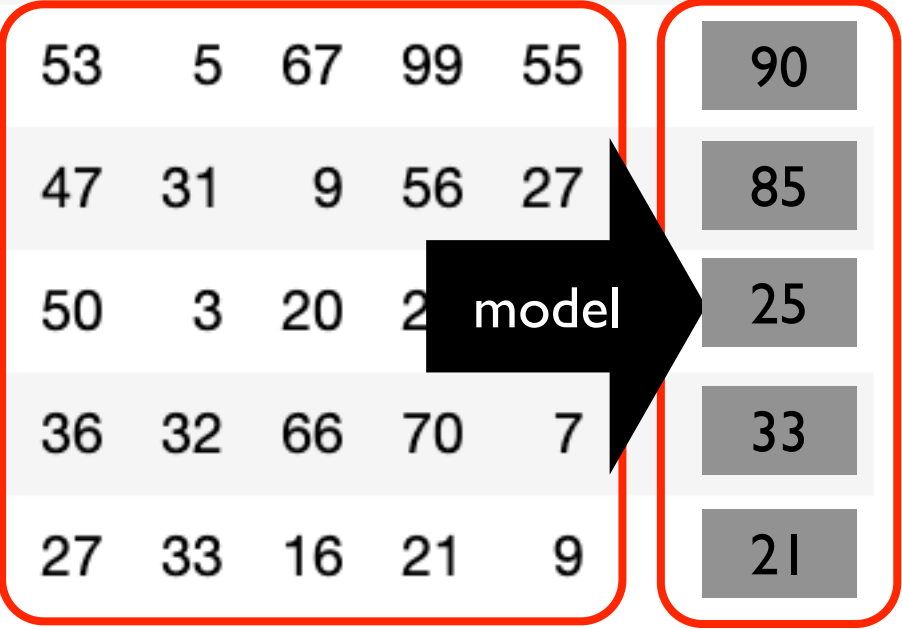
	<b>x0</b>	<b>x1</b>	<b>x2</b>	<b>x3</b>	<b>x4</b>	<b>y (label)</b>
<b>0</b>	37	25	40	70	68	5
<b>1</b>	50	13	7	67	79	25
<b>2</b>	56	12	5	15	90	44
<b>3</b>	89	70	85	49	68	72
<b>4</b>	36	93	52	33	14	59
<b>5</b>	53	5	67	99	55	????
<b>6</b>	47	31	9	56	27	????
<b>7</b>	50	3	20	2		????
<b>8</b>	36	32	66	70	7	????
<b>9</b>	27	33	16	21	9	????



**predict:** estimate for actual unknowns

# I. Regression (Supervised)


	x0	x1	x2	x3	x4	y (label)
0	37	25	40	70	68	5
1	50	13	7	67	79	25
2	56	12	5	15	90	44
3	89	70	85	49	68	72
4	36	93	52	33	14	59
5	53	5	67	99	55	90
6	47	31	9	56	27	85
7	50	3	20	2		25
8	36	32	66	70	7	33
9	27	33	16	21	9	21



**predict:** estimate for actual unknowns

# I. Regression (Supervised)

	x0	x1	x2	x3	x4	y (label)
0	37	25	40	70	68	5
1	50	13	7	67	79	25
2	56	12	5	15	90	44
3	89	70	85	49	68	72
4	36	93	52	33	14	59
5	53	5	67	99	55	90
6	47	31	9	56	27	85
7	50	3	20	24	63	25
8	36	32	66	70	7	33
9	27	33	16	21	9	21



**interpret:** what can we learn by looking directly at the model?



# I. Regression (Supervised)

category features      quantitative label

	x0	x1	x2	x3	x4	y (label)
0	37	green	40	triangle	68	5
1	50	green	7	circle	79	25
2	56	red	5	circle	90	44
3	89	blue	85	triangle	68	72
4	36	blue	52	square	14	59
5	53	green	67	triangle	55	????
6	47	blue	9	triangle	27	????
7	50	blue	20	circle	63	????
8	36	green	66	circle	7	????
9	27	red	16	circle	9	????

a problem with some **category** features is still a regression as long as the lable is **quantitative**

## 2. Classification (Supervised)

categorical  
label

	x0	x1	x2	x3	x4	y (label)
0	37	green	40	triangle	68	orange
1	50	green	7	circle	79	pear
2	56	red	5	circle	90	pear
3	89	blue	85	triangle	68	apple
4	36	blue	52	square	14	pear
5	53	green	67	triangle	55	????
6	47	blue	9	triangle	27	????
7	50	blue	20	circle	63	????
8	36	green	66	circle	7	????
9	27	red	16	circle	9	????

problem: can we predict an unknown **category**?

### 3. Clustering (Unsupervised)

no  
label!



	<b>x0</b>	<b>x1</b>	<b>x2</b>	<b>x3</b>	<b>x4</b>
<b>0</b>	37	25	40	70	68
<b>1</b>	50	13	7	67	79
<b>2</b>	56	12	5	15	90
<b>3</b>	89	70	85	49	68
<b>4</b>	36	93	52	33	14
<b>5</b>	53	5	67	99	55
<b>6</b>	47	31	9	56	27
<b>7</b>	50	3	20	24	63
<b>8</b>	36	32	66	70	7
<b>9</b>	27	33	16	21	9

problem: can we organize data into groups of similar rows?

### 3. Clustering (Unsupervised)

the algorithm  
decides groups

	x0	x1	x2	x3	x4	group
0	37	25	40	70	68	1
1	50	13	7	67	79	0
2	56	12	5	15	90	0
3	89	70	85	49	68	1
4	36	93	52	33	14	2
5	53	5	67	99	55	0
6	47	31	9	56	27	1
7	50	3	20	24	63	1
8	36	32	66	70	7	2
9	27	33	16	21	9	0

group 0

group 1

group 2

there is no official grouping to check the model against,  
but a good grouping places similar rows together

## 4. Decomposition (Unsupervised)

	<b>x0</b>	<b>x1</b>	<b>x2</b>	<b>x3</b>	<b>x4</b>
<b>0</b>	-11	-7	3	20	20
<b>1</b>	2	-19	-30	17	31
<b>2</b>	8	-20	-32	-35	42
<b>3</b>	41	38	48	-1	20
<b>4</b>	-12	61	15	-17	-34
<b>5</b>	5	-27	30	49	7
<b>6</b>	-1	-1	-28	6	-21
<b>7</b>	2	-29	-17	-26	15
<b>8</b>	-12	0	29	20	-41
<b>9</b>	-21	1	-21	-29	-39

## 4. Decomposition (Unsupervised)

original data

	x0	x1	x2	x3	x4
0	-11	-7	3	20	20
1	2	-19	-30	17	31
2	8	-20	-32	-35	42
3	41	38	48	-1	20
4	-12	61	15	-17	-34
5	5	-27	30	49	7
6	-1	-1	-28	6	-21
7	2	-29	-17	-26	15
8	-12	0	29	20	-41
9	-21	1	-21	-29	-39

components

	x0	x1	x2	x3	x4
0	-0.0	0.6	0.5	0.1	-0.6
1	0.3	-0.2	0.5	0.6	0.5
2	0.4	0.5	0.1	-0.6	0.5

-11

21

-8

## 4. Decomposition (Unsupervised)

original data

	x0	x1	x2	x3	x4
0	-11	-7	3	20	20
1	2	-19	-30	17	31
2	8	-20	-32	-35	42
3	41	38	48	-1	20
4	-12	61	15	-17	-34
5	5	-27	30	49	7
6	-1	-1	-28	6	-21
7	2	-29	-17	-26	15
8	-12	0	29	20	-41
9	-21	1	-21	-29	-39

components

	x0	x1	x2	x3	x4
0	-0.0	0.6	0.5	0.1	-0.6
1	0.3	-0.2	0.5	0.6	0.5
2	0.4	0.5	0.1	-0.6	0.5

-11

21

-8

weights

	pc0	pc1	pc2
0	-11	21	-8
1	-43	12	-6
2	-58	-14	30
3	36	41	53
4	...	...	...

...

## 4. Decomposition (Unsupervised)

original data

	x0	x1	x2	x3	x4
0	-11	-7	3	20	20
1	2	-19	-30	17	31
2	8	-20	-32	-35	42
3	41	38	48	-1	20
4	-12	61	15	-17	-34
5	5	-27	30	49	7
6	-1	-1	-28	6	-21
7	2	-29	-17	-26	15
8	-12	0	29	20	-41
9	-21	1	-21	-29	-39

components

	x0	x1	x2	x3	x4
0	-0.0	0.6	0.5	0.1	-0.6
1	0.3	-0.2	0.5	0.6	0.5
2	0.4	0.5	0.1	-0.6	0.5

-43

12

-6

weights

	pc0	pc1	pc2
0	-11	21	-8
1	-43	12	-6
2	-58	-14	30
3	36	41	53
4	...	...	...

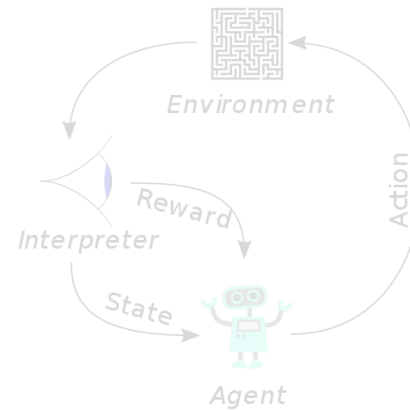
...



# Machine Learning

## Reinforcement Learning

*not covered in CS 320*



[https://en.wikipedia.org/wiki/Reinforcement\\_learning](https://en.wikipedia.org/wiki/Reinforcement_learning)

## Supervised Machine Learning

*data is labeled, we know what we want to predict*

### Regression

*predict a quantity*

### Classification

*predict a category*

## Unsupervised Machine Learning

*data is unlabeled, we're just looking for patterns*

### Clustering

*place rows in groups*

### Decomposition

*represent rows as combos of "component" rows*

this semester, we'll learn at least one technique in each of these four categories

# 1. Regression (Supervised)

+

# 2. Classification (Supervised)

```
linear_model.LogisticRegression([penalty, ...]) classification!  
linear_model.LogisticRegressionCV(*[, Cs, ...])  
linear_model.PassiveAggressiveClassifier(*  
linear_model.Perceptron(*[, penalty, alpha, ...])  
linear_model.RidgeClassifier([alpha, ...])  
linear_model.RidgeClassifierCV([alphas, ...])  
linear_model.SGDClassifier([loss, penalty, ...])
```

```
linear_model.LinearRegression(*[, ...])  
linear_model.Ridge([alpha, fit_intercept, ...])  
linear_model.RidgeCV([alphas, ...])  
linear_model.SGDRegressor([loss, penalty, ...])
```

```
svm.LinearSVC([penalty, loss, dual, tol, C, ...])  
svm.LinearSVR(*[, epsilon, tol, C, loss, ...])
```

```
tree.DecisionTreeClassifier  
tree.DecisionTreeRegressor  
tree.ExtraTreeClassifier  
tree.ExtraTreeRegressor
```

```
neighbors.KNeighborsClassifier([...])  
neighbors.KNeighborsRegressor([n_neighbors, ...])
```

# 3. Clustering (Unsupervised)

```
cluster.AffinityPropagation(*[, damping, ...])  
cluster.AgglomerativeClustering([...])  
cluster.Birch(*[, threshold, ...])  
cluster.DBSCAN([eps, min_samples, metric, ...])  
cluster.FeatureAgglomeration([n_clusters, ...])  
cluster.KMeans([n_clusters, init, n_init, ...])  
cluster.MinibatchKMeans([n_clusters, init, ...])  
cluster.MeanShift(*[, bandwidth, seeds, ...])  
cluster.OPTICS(*[, min_samples, max_eps, ...])  
cluster.SpectralClustering([n_clusters, ...])  
cluster.SpectralBiclustering([n_clusters, ...])  
cluster.SpectralCoclustering([n_clusters, ...])
```

# 4. Decomposition (Unsupervised)

```
decomposition.DictionaryLearning([...])  
decomposition.FactorAnalysis([n_components, ...])  
decomposition.FastICA([n_components, ...])  
decomposition.IncrementalPCA([n_components, ...])  
decomposition.KernelPCA([n_components, ...])  
decomposition.LatentDirichletAllocation([...])  
decomposition.MinibatchDictionaryLearning([...])  
decomposition.MinibatchSparsePCA([...])  
decomposition.NMF([n_components, init, ...])  
decomposition.PCA([n_components, copy, ...])  
decomposition.SparsePCA([n_components, ...])  
decomposition.SparseCoder(dictionary, *[, ...])  
decomposition.TruncatedSVD([n_components, ...])
```

# Foundations: Modules and Math

# Important Packages

We'll be learning the following to do ML and related calculations efficiently:

1

numpy

2

pytorch


3

scikit-learn


```
pip3 install numpy scikit-learn
```

```
pip3 install torch torchvision
```


# Linear Algebra



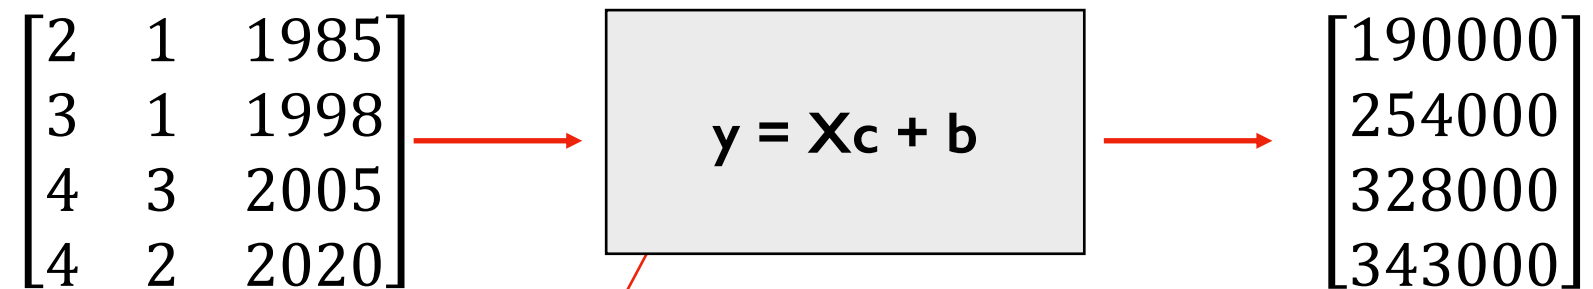
with matrices...



# Linear Algebra



with matrices...

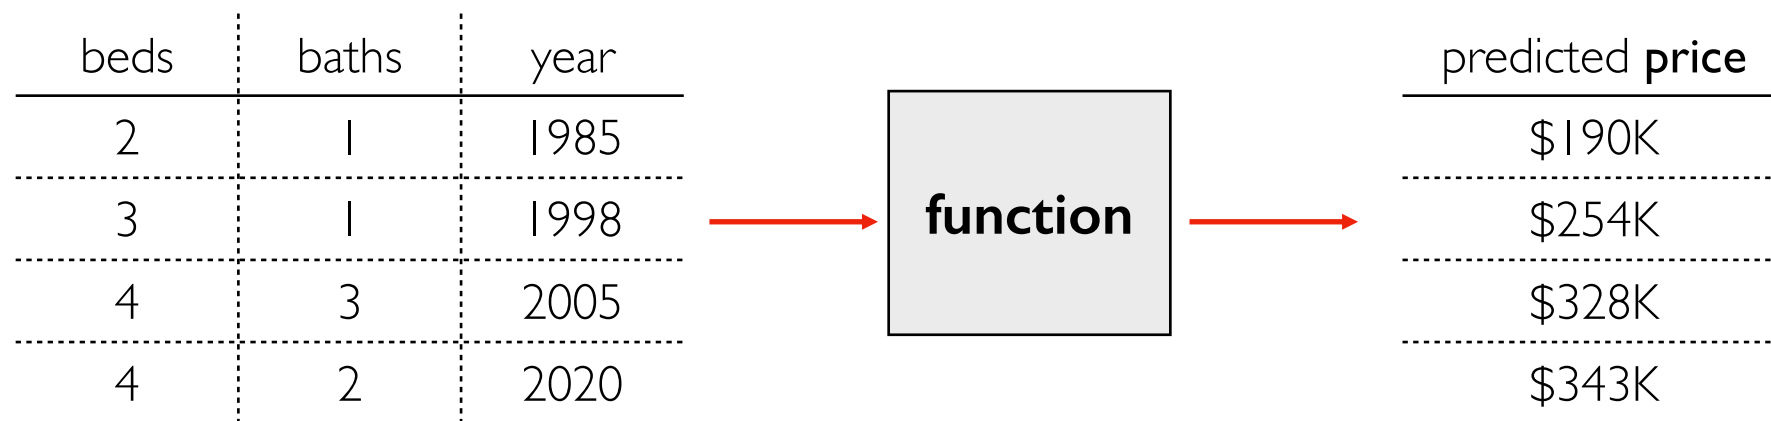


$$\begin{bmatrix} 2 & 1 & 1985 \\ 3 & 1 & 1998 \\ 4 & 3 & 2005 \\ 4 & 2 & 2020 \end{bmatrix} \times \begin{bmatrix} 41.46 \\ 10.36 \\ 1.70 \end{bmatrix} + -3277.31$$

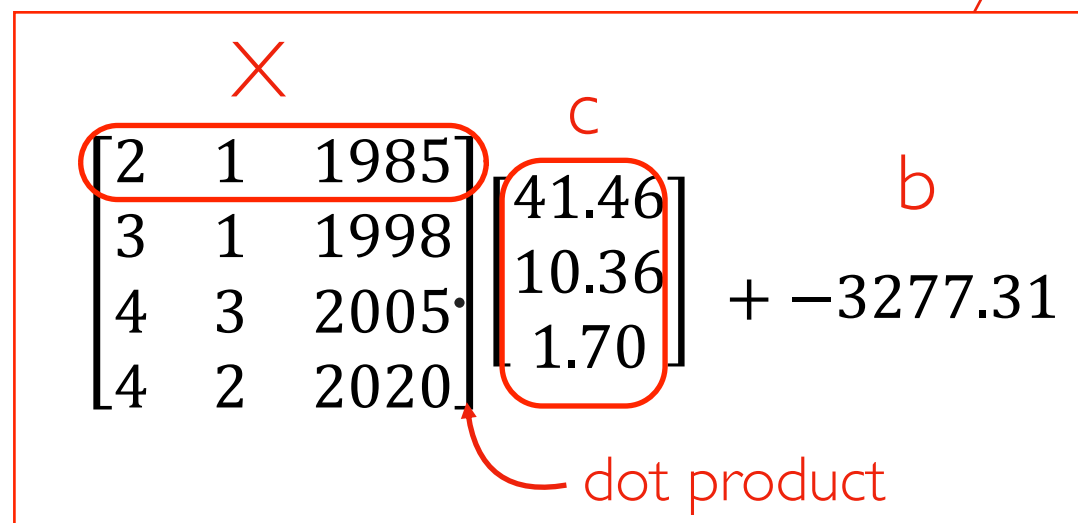
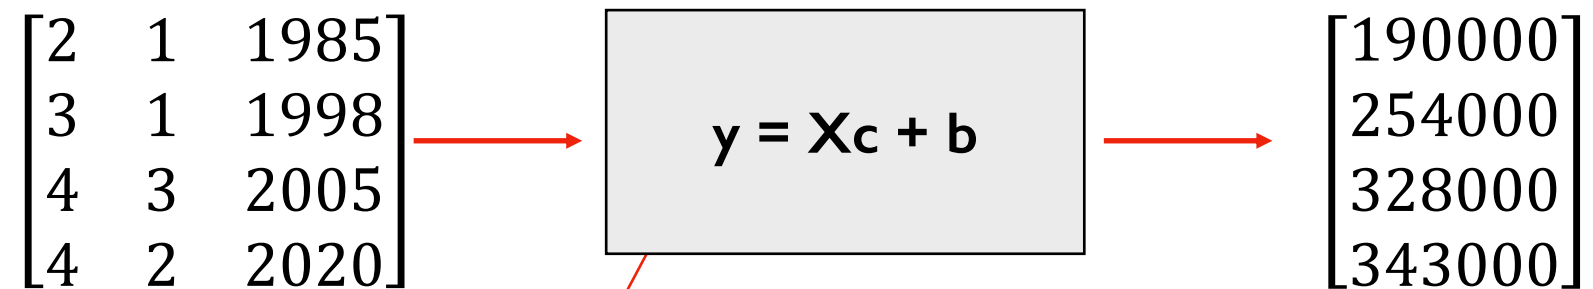
dot product\*

\*dot product is usually between vectors by definition, but numpy uses np.dot for matrix multiplication

# Linear Algebra

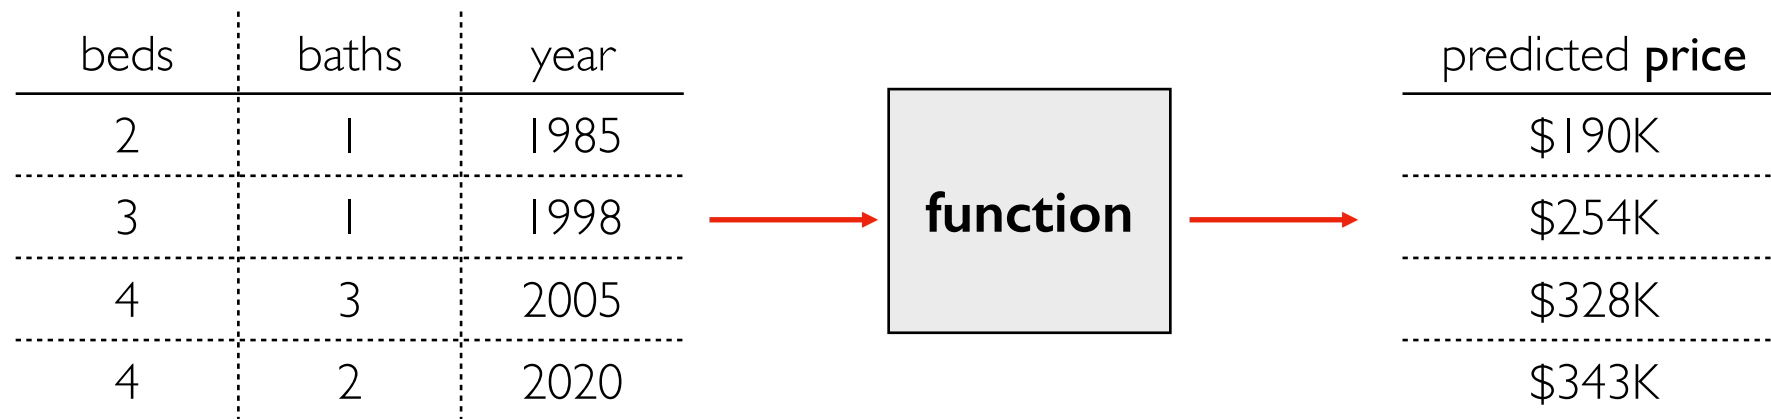


with matrices...

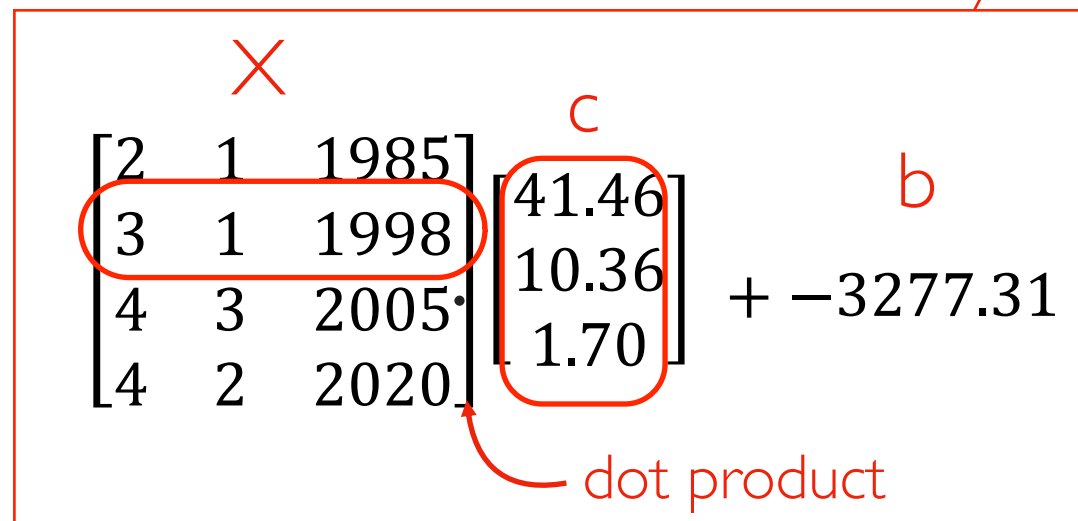



$$\begin{array}{ccccccc}
 \times & c & \times & c & \times & c & b \\
 2 * 41.46 + 1 * 10.36 + 1985 * 1.7 - 3277.31 & & & & & & \\
 & & & & & & = 190000 \\
 & & & & & & y
 \end{array}$$

# Linear Algebra



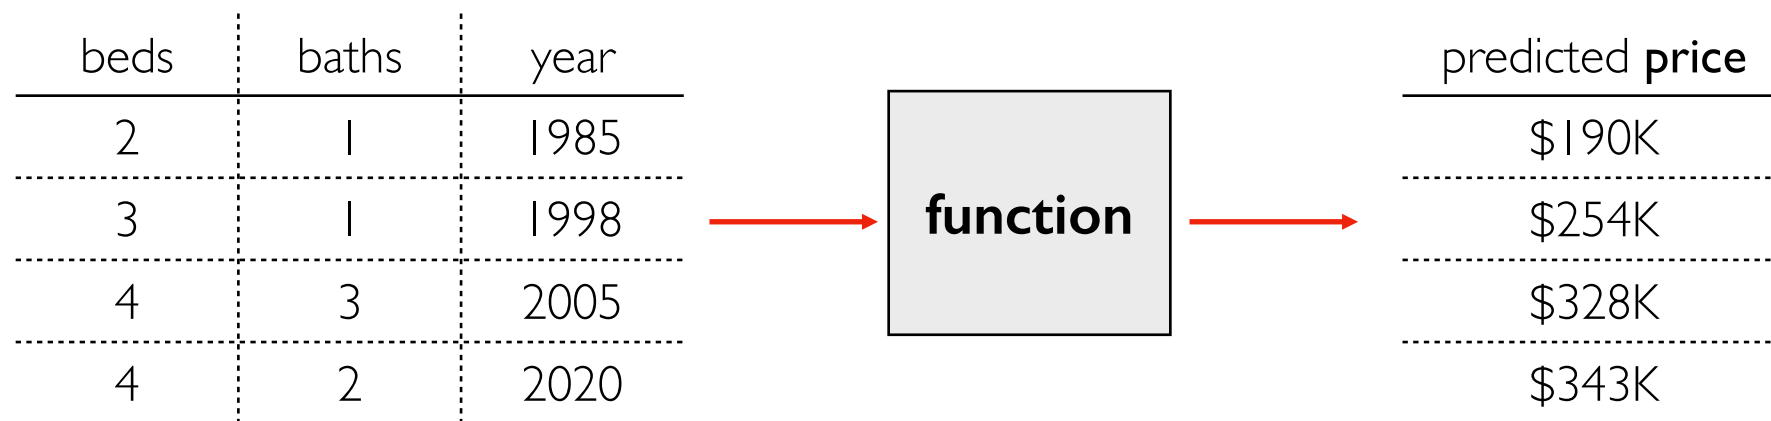
with matrices...



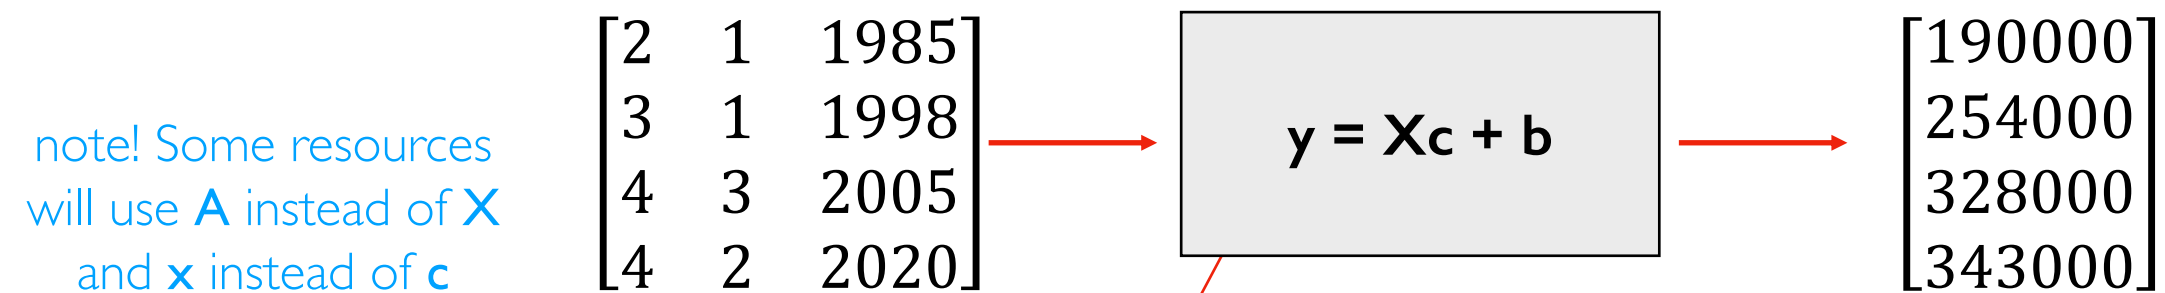
$$\begin{array}{ccccccc}
 \times & c & \times & c & \times & c & b \\
 3 * 41.46 + 1 * 10.36 + 1998 * 1.7 - 3277.31 & & & & & & \\
 & & & & & & = 254000 \\
 & & & & & & y
 \end{array}$$



# Linear Algebra



with matrices...



$$\begin{matrix} X & c & b \\ \begin{bmatrix} 2 & 1 & 1985 \\ 3 & 1 & 1998 \\ 4 & 3 & 2005 \\ 4 & 2 & 2020 \end{bmatrix} & \begin{bmatrix} 41.46 \\ 10.36 \\ 1.70 \end{bmatrix} & + -3277.31 \end{matrix}$$

dot product

```
import numpy as np
X = df.values
y = np.dot(X, c) + b
```

# Linear Algebra

$$y = x^{**2} \quad \text{not linear}$$

$$y = x_0*4 + x_1*(-1) + x_2*0.5 + \dots + x_{10}*3 \quad \text{linear}$$

with matrices...

$$\begin{bmatrix} 2 & 1 & 1985 \\ 3 & 1 & 1998 \\ 4 & 3 & 2005 \\ 4 & 2 & 2020 \end{bmatrix} \longrightarrow \boxed{y = Xc + b} \longrightarrow \begin{bmatrix} 190000 \\ 254000 \\ 328000 \\ 343000 \end{bmatrix}$$

$\times$

$$\begin{bmatrix} 2 & 1 & 1985 \\ 3 & 1 & 1998 \\ 4 & 3 & 2005 \\ 4 & 2 & 2020 \end{bmatrix} \begin{matrix} c \\ \begin{bmatrix} 41.46 \\ 10.36 \\ 1.70 \end{bmatrix} \end{matrix} + -3277.31$$


dot product

```
import numpy as np
X = df.values
y = np.dot(X, c) + b
```

# Calculus: Minimizing Something

training data

beds	baths	year	price
1	1	1980	\$140K
3	1	1990	\$240K
3	4	2004	\$295K



train

$$y = Xc + b$$

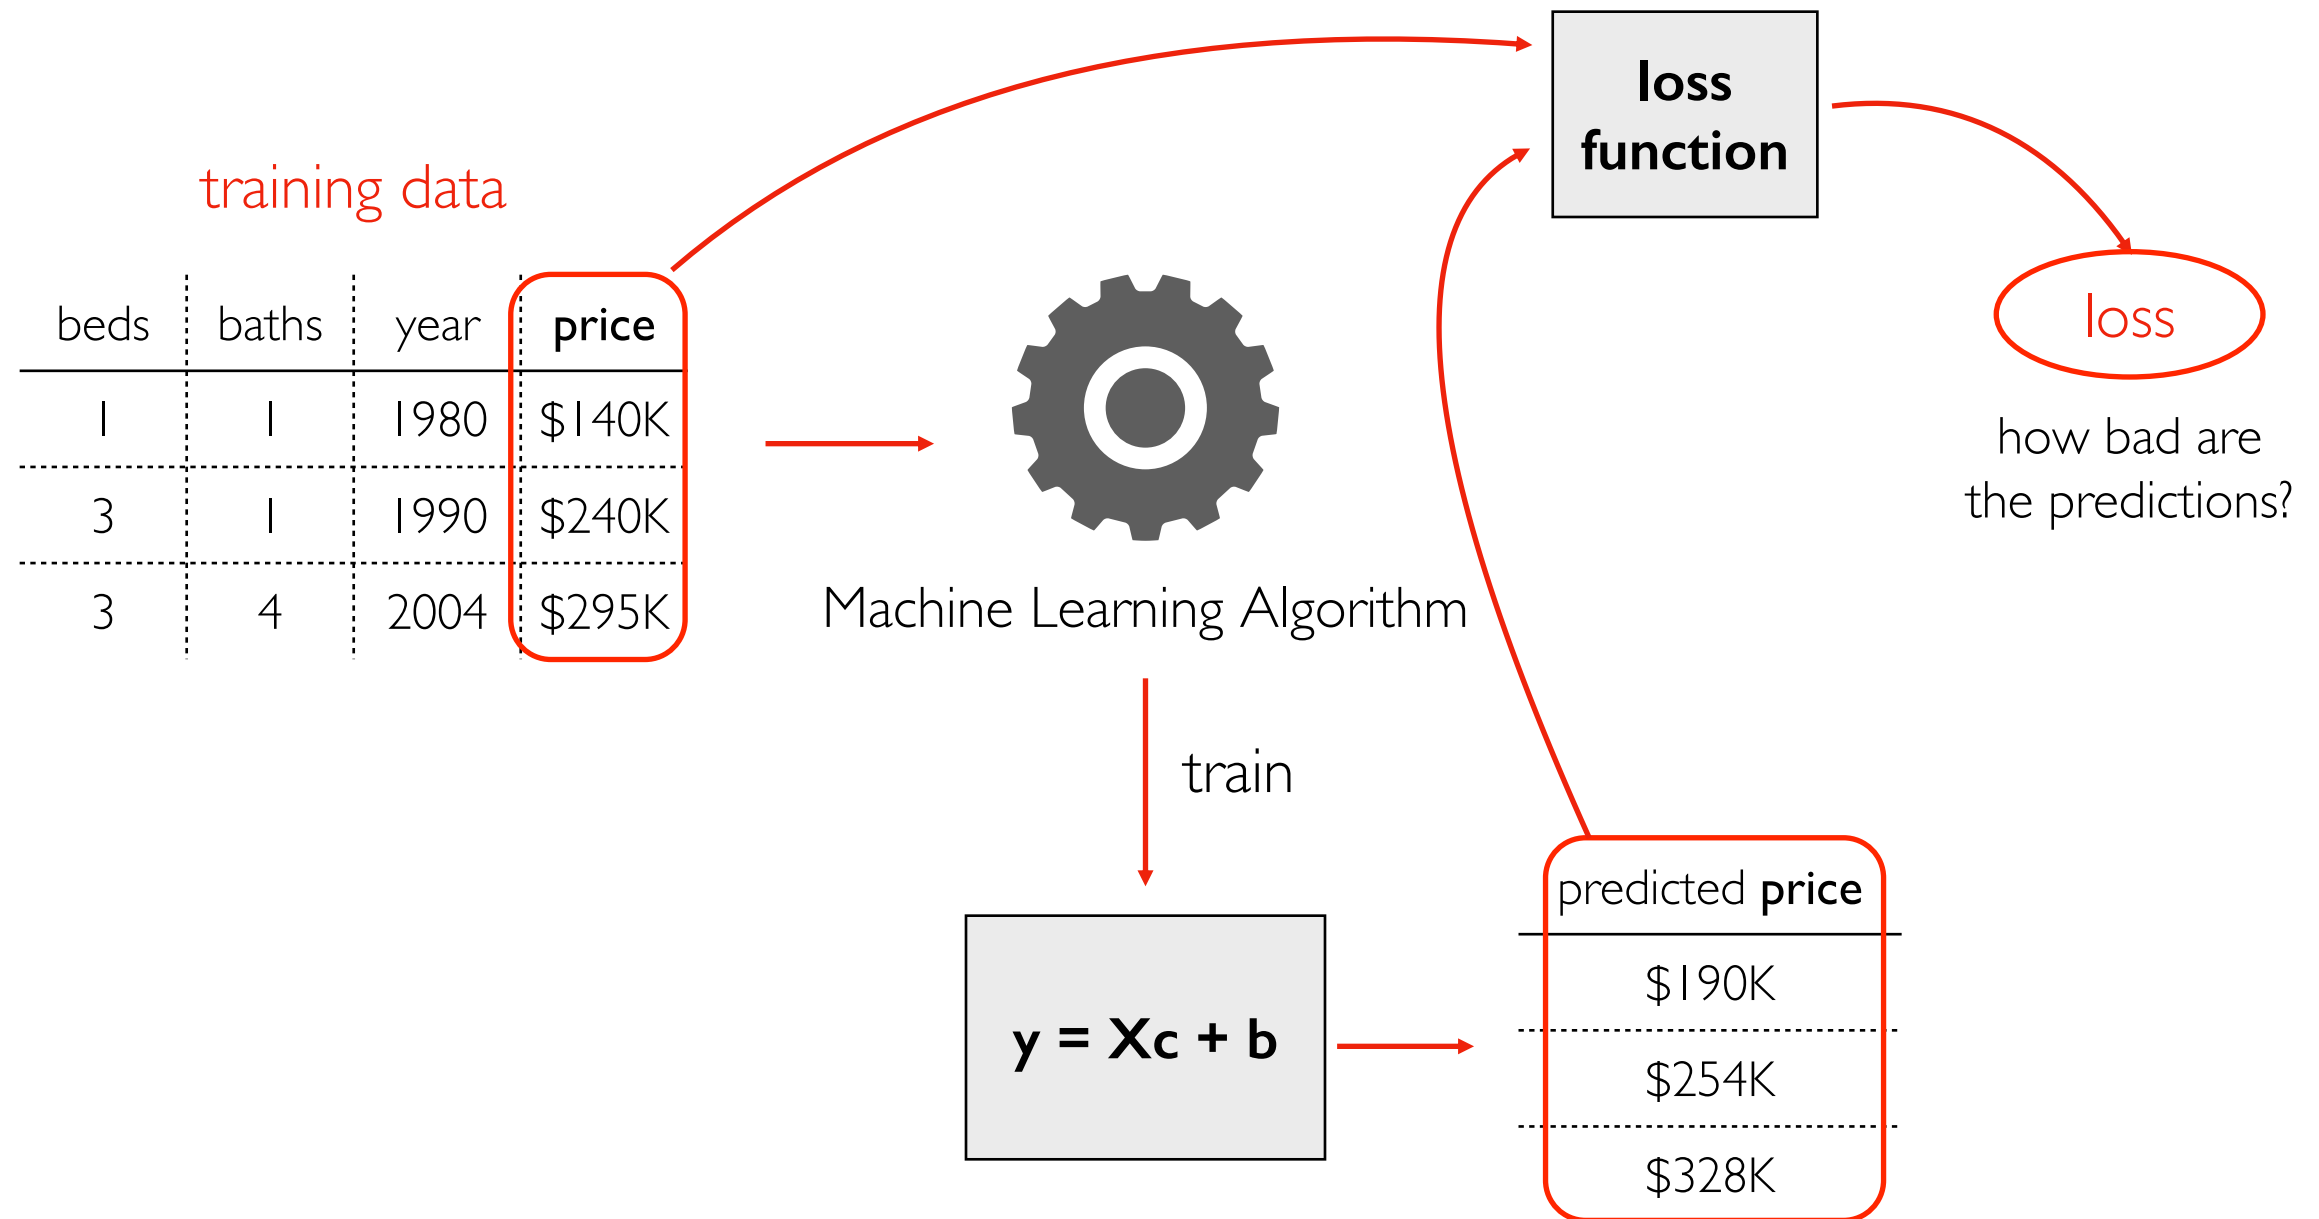
predicted **price**

\$190K

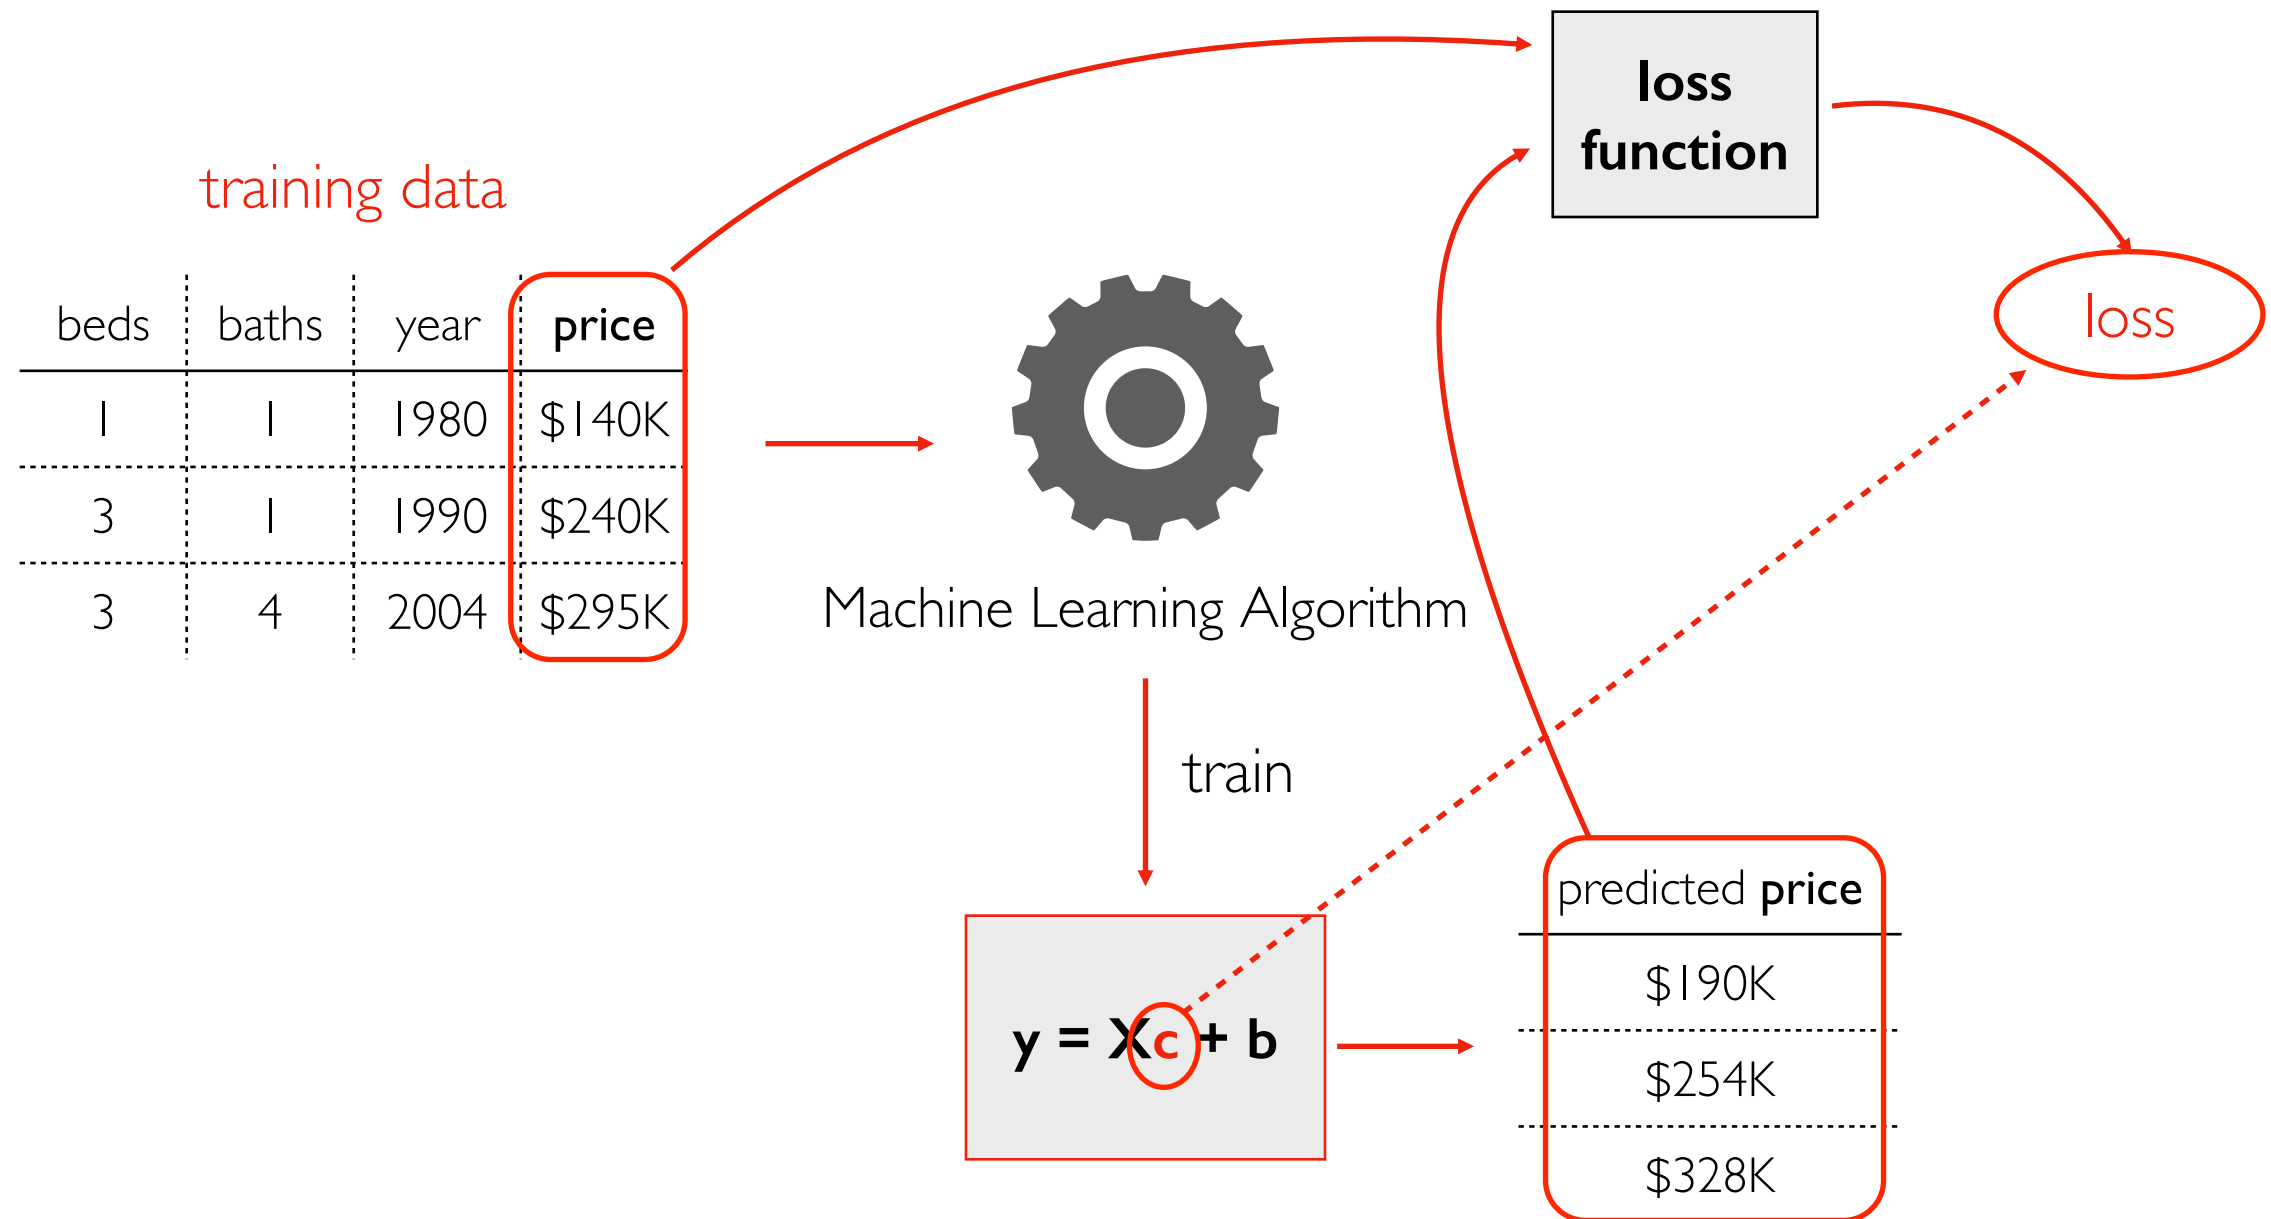
\$254K

\$328K

# Calculus: Minimizing Something



# Calculus: Minimizing Something

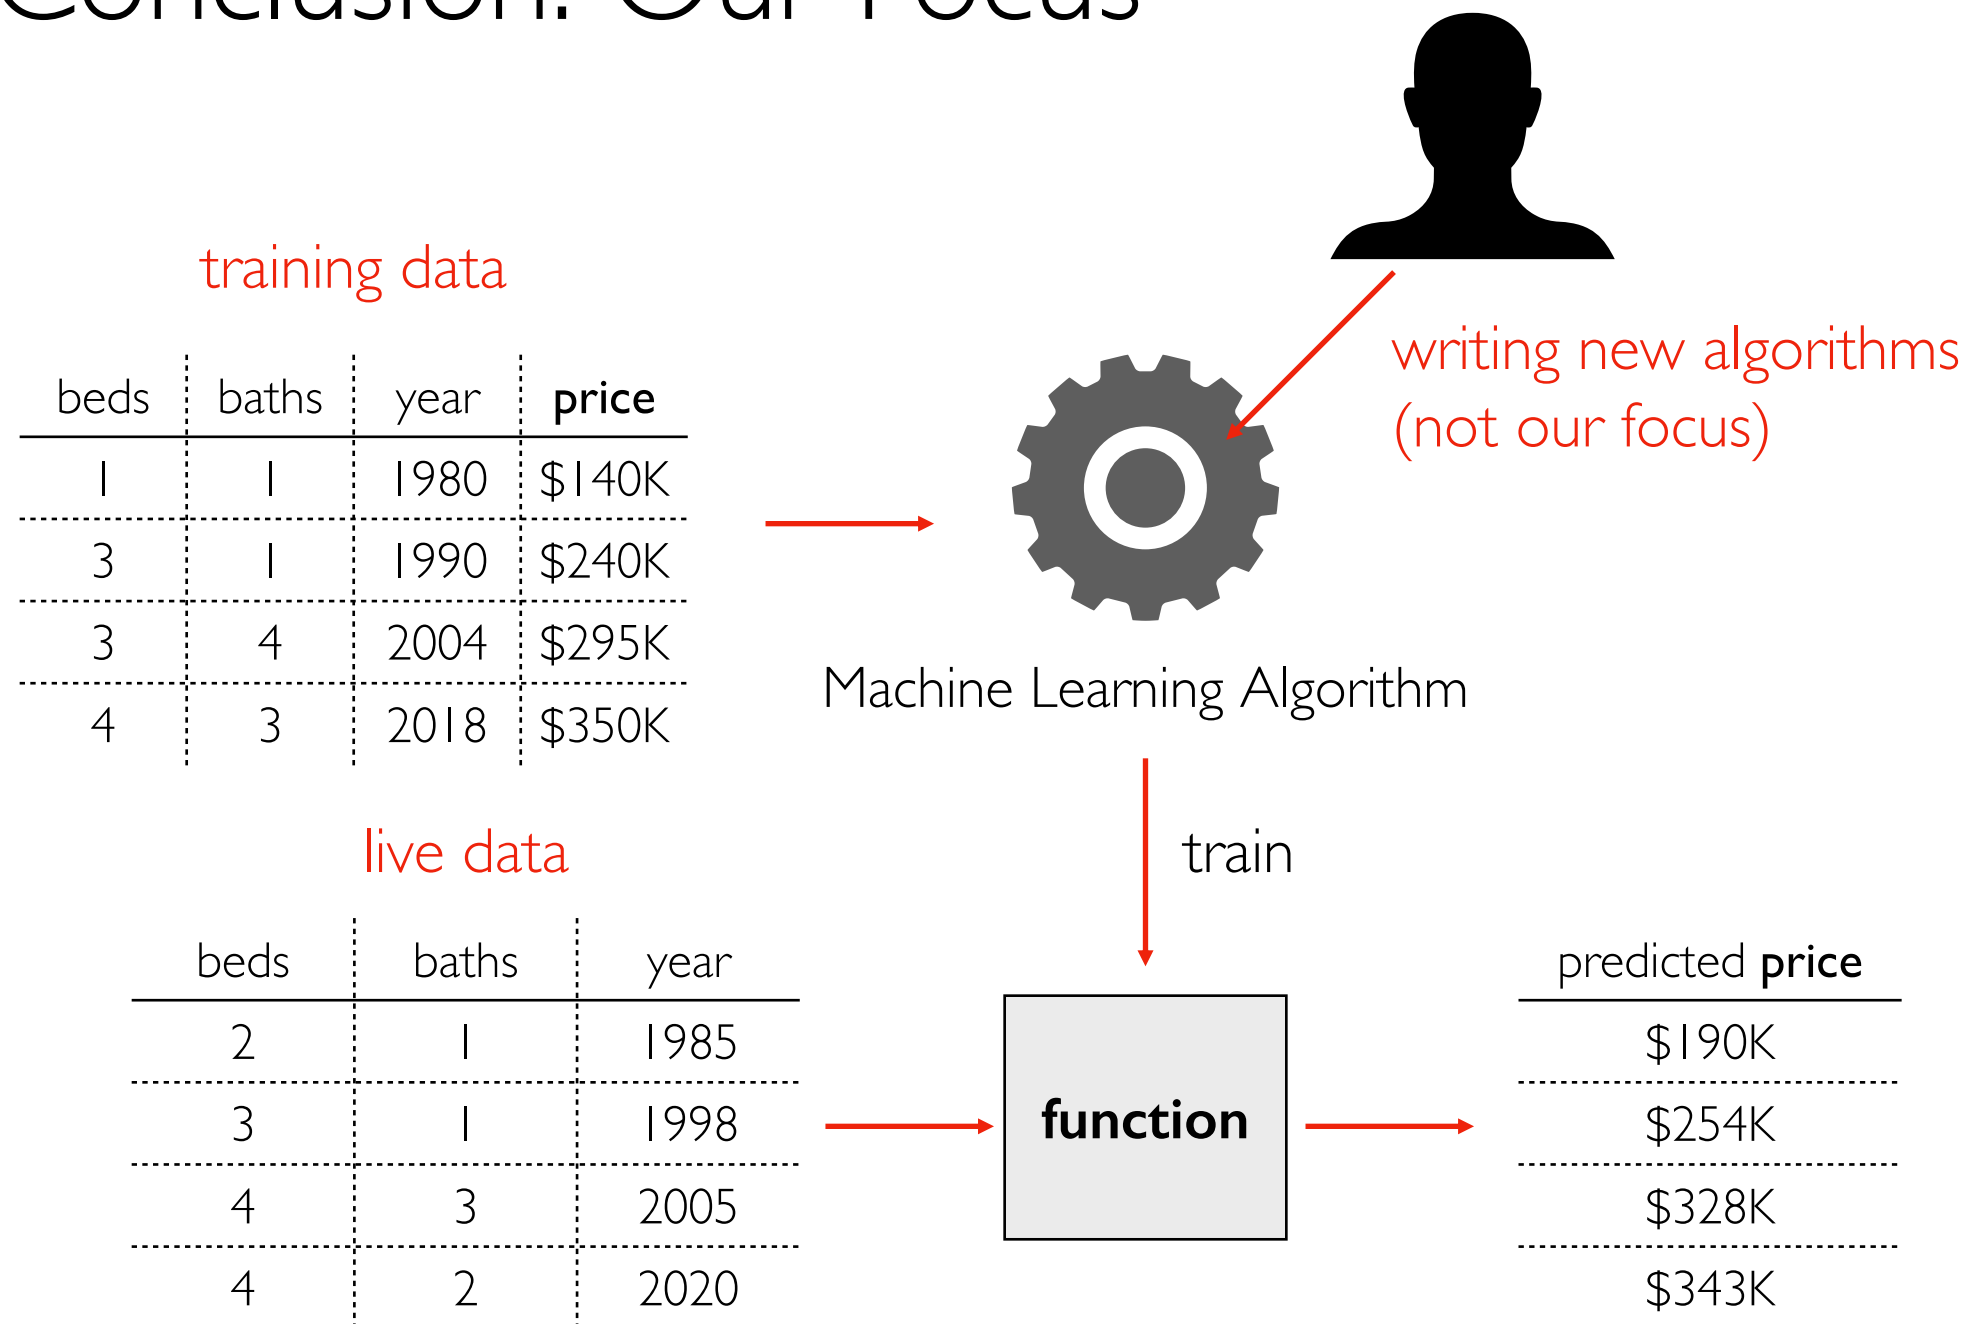


how do we optimize  $\mathbf{c}$  to minimize **loss**?  
Important concepts: derivative, gradient

(pytorch can do this)

Conclusion: Developers vs. Users

# Conclusion: Our Focus

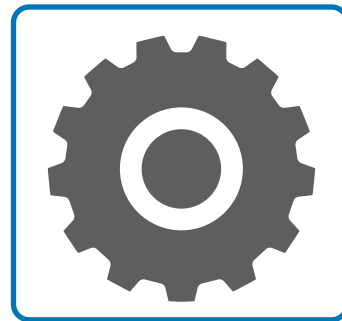


# Conclusion: Our Focus

how can we clean this up?

training data

beds	baths	year	price
1	1	1980	\$140K
3	1	1990	\$240K
3	4	2004	\$295K
4	3	2018	\$350K



which algorithm (from sklearn?) should we pick, and how should we configure it?

Machine Learning Algorithm

is it working well?  
(evaluation)

train

live data

beds	baths	year
2	1	1985
3	1	1998
4	3	2005
4	2	2020

function

predicted price

\$190K

\$254K

\$328K

\$343K